

Alan Ritter

Associate Professor, School of Interactive Computing
College of Computing, Georgia Institute of Technology

CONTACT	alan.ritter@cc.gatech.edu Website Google Scholar
RESEARCH INTERESTS	My research in natural language processing and machine learning focuses on how language models acquire knowledge, reason, and learn through interaction. My group develops methods for learning from expert solutions and improving through simulated experience, with applications to planning and decision-making. We also study privacy and safety in human–AI interaction.
PROFESSIONAL APPOINTMENTS	Georgia Institute of Technology, Atlanta, GA USA Associate Professor, School of Interactive Computing, College of Computing 2020–present The Ohio State University, Columbus, OH USA Associate Professor, Computer Science and Engineering 2020 Assistant Professor, Computer Science and Engineering 2014–2020 Carnegie Mellon University, Pittsburgh, PA USA Visiting Faculty, Language Technologies Institute 2019 Postdoctoral Fellow, Machine Learning Department 2013–2014 Microsoft Research, Redmond, WA USA Consultant, Natural Language Processing Group 2014–2016 Intern, Natural Language Processing Group (2009); Knowledge Tools Group (2008) Naval Research Laboratory, Marine Meteorology Division, Monterey, CA USA Intern, Data Assimilation Group 2006
EDUCATION	University of Washington, Seattle, WA USA Ph.D. in Computer Science & Engineering, 2013 Western Washington University, Bellingham, WA USA B.S./M.S., Computer Science, 2006
HONORS AND RECOGNITION	Microsoft Research Fellowship — Fellowship Advisor (2026) Spotlight for <i>Language Models can Self-Improve at State-Value Estimation for Better Search</i> at NeurIPS (2025) Google Academic Research Award (2025) SAC Highlight Award for <i>Balancing the Budget: Understanding Trade-offs Between Supervised and Preference-Based Finetuning</i> at ACL (2025) Sony Faculty Innovation Award (2024) Social Impact Award for <i>Having Beer after Prayer? Measuring Cultural Bias in Large Language Models</i> at ACL (2024) Reproduction Award for <i>Do CoNLL-2003 Named Entity Taggers Still Work Well in 2023?</i> at ACL (2023) NSF CAREER Award (2019) Lumley Research Award, College of Engineering, Ohio State University (2019) Amazon Research Award (2018) Outstanding Student Organization Advisor, College of Engineering, Ohio State University (2018) Best Student Paper Award for <i>Learning to Generalize for Complex Selection Tasks</i> at the ACM Conference on Intelligent User Interfaces (2009) NDSEG Fellowship (2008)

GRANTS AS PI

NSF Future CoRe, (\$997,565), III: Robust and Responsible Reasoning for Privacy-Preserving Image Geolocation, 2026–2029. Co-PIs: James Hays and Wei Xu.

Google, (\$50,000 share of \$100,000 total), AI Tools to Help Users Make Informed Decisions about Online Information Sharing, with Sauvik Das (Carnegie Mellon University), 2025–2026.

Microsoft Research, (\$47,000), LLM-Assisted Bayesian Reasoning for Collaborative Organizational Decision-Making, 2026.

Ford Motor Company, (\$231,938), Reliable Repair Chatbot, 2024–2026.

Sony Group, (\$100,000), Adapting LLM-based Multimodal Dialog Across Languages and Cultures, 2024–2025. Co-PI: Wei Xu.

NSF, (\$500,000), CAREER: Large-Scale Learning for Information Extraction, 2019–August 2027.

IARPA HIATUS (SRI International subcontract), (\$491,467), Stable, Interpretable, and General Neural Attribution for Trust of Users and Robust Effacement (SIGNATURE), 2022–2024. Co-PI: Wei Xu.

DARPA AMD (SRI International subcontract), (\$800,000), Logically Adaptive Molecular Design Automation (LAMDA), 2019–2023.

IARPA BETTER (Brown University subcontract), (\$850,000), Task and User-Aware Representation Learning for Fine-Grained Cross-Lingual Information Retrieval, 2019–2023. Co-PI: Wei Xu.

NSF, (\$151,000), CRII: Learning to Extract Events from Knowledge Base Revisions, 2015–2018.

Amazon Research Award, (\$50,000), Integrating Data-Driven Conversation in Open-Domain Question Answering, 2018.

DARPA SocialSim (Leidos subcontract), (\$600,000), Modeling the Spread of Information through Social and Knowledge Graphs, 2017–2021.

IARPA CAUSE (Leidos subcontract), (\$507,000), Extracting a Realtime Cybersecurity Knowledge Graph from Text, 2016–2019.

CMU Software Engineering Institute, (\$40,000), Events, relationships, and script learning for situational awareness, January–September 2017.

GRANTS AS CO-PI
OR SENIOR
PERSONNEL

Georgia Tech Seed Grant (PI: Wei Xu), jointly funded by the School of Civil and Environmental Engineering and the College of Computing (\$20,000 total; \$10,000 from each unit), AI-Accelerated Materials Discovery with Lab-in-the-Loop for Resource Recovery and Sustainability, beginning July 2026. Role: Co-PI; with Yongsheng Chen.

Sony Group (PI: Wei Xu), (\$100,000), Multi-Agent Reinforcement Learning for Culture, Knowledge, and Context Adaptive Translation, 2026–2027. Role: Co-PI.

NSF (PI: Sonia Chernova), (\$19,995,808), AI Institute for Collaborative Assistance and Responsive Interaction for Networked Groups (AI-CARING). Role: Senior Personnel.

NSF (PI: Betsy DiSalvo), (\$399,634), The Cost of AI: A Comparative Study of Machine Learning Training Methods, 2024–present. Role: Co-PI.

NIH (PI: Po-Yin Yen), (\$100,000), Supporting systematic review production with article similarity network visualization, 2016–2018. Role: Co-PI.

Criteo Faculty Research Award (PI: Wei Xu), (\$20,000), 2018. Role: Co-PI.

Battelle (PI: Marie de Marneffe), (\$23,000), Trustworthiness in Social Media, 2017–2019. Role: Co-PI.

PUBLICATIONS

2026

Soft Token Alignment for Cross-Lingual Reasoning
 Jiayi He, Jungsoo Park, Wei Xu, **Alan Ritter**
 NeurIPS 2026

Synthetic Worlds for Temporal Evaluation and Knowledge Updating in LLMs
Jonathan Zheng, Zirui Shao, **Alan Ritter**, Wei Xu
NeurIPS 2026

Investigating and Alleviating Harm Amplification in LLM Interactions
Ruohao Guo, Wei Xu, **Alan Ritter**
EMNLP 2026

Who Pays More for Safety? Measuring the Disparate Cost of Safety Alignment across Languages
Chanwoong Yoon, Jungsoo Park, **Alan Ritter**
EMNLP 2026

Camellia: Benchmarking Cultural Biases in LLMs for Asian Languages
Tarek Naous, Anagha Savit, Carlos Rafael Catalan, Geyang Guo, Jaehyeok Lee, Kyungdon Lee, Lheane Marie Dizon, Mengyu Ye, Neel Kothari, Sahajpreet Singh, Sarah Masud, Tanish Patwa, Trung Thanh Tran, Zohaib Khan, **Alan Ritter**, Tanmoy Chakraborty, Yuki Arase, Keisuke Sakaguchi, JinYeong Bak, Wei Xu
EMNLP 2026

LaRA: Layer-wise Representation Analysis for Detecting Data Contamination in RL Post-Training
Minju Gwak, Minseo Kwak, Dongseok Lee, Guijin Son, **Alan Ritter**, Jaehyung Kim
Findings of EMNLP 2026

CoCoA: Context-Conditional Cultural Alignment for Large Language Models
Kyungdon Lee, Wei Xu, **Alan Ritter**, Dong-Ho Lee, JinYeong Bak
Findings of EMNLP 2026

Semantic Differentiation for Tackling Challenges in Watermarking Low-Entropy Constrained Generation Outputs
Nghia T. Le, **Alan Ritter**, Kartik Goyal
COLM 2026

Making Expert Reasoning Learnable with Self-Distillation
Ethan Mendes, Jungsoo Park, **Alan Ritter**
ICML 2026

Safe and Scalable Web Agent Learning via Recreated Websites
Hyunjoo Chae, Jungsoo Park, **Alan Ritter**
ICML 2026

Learning to Route Languages for Multilingual Preference Optimization
Geyang Guo, Hiromi Wakaki, Yuki Mitsufuji, **Alan Ritter**, Wei Xu
ICML 2026

GeoRC: A Benchmark for Geolocation Reasoning Chains
Mohit Talreja, Joshua Diao, Jim Thannikary James, Radu Casapu, Tejas Santanam, Ethan Mendes, **Alan Ritter**, Wei Xu, James Hays
ACL 2026

Do Vision-Language Models Respect Contextual Integrity in Location Disclosure?
Ruixin Yang, Ethan Mendes, Arthur Wang, James Hays, Sauvik Das, Wei Xu, **Alan Ritter**
ICLR 2026

Tree-based Dialogue Reinforced Policy Optimization for Red-Teaming Attacks
Ruohao Guo, Afshin Oroojlooy, Roshan Sridhar, Miguel Ballesteros, **Alan Ritter**, Dan Roth
ICLR 2026

Supporting Informed Self-Disclosure: Design Recommendations for Presenting AI-Estimates of Privacy Risks to Users
Isadora Krsek, Meryl Ye, Wei Xu, **Alan Ritter**, Laura Dabbish, Sauvik Das
CHI 2026

Auditing Language Model Unlearning via Information Decomposition
Anmol Goel, **Alan Ritter**, Iryna Gurevych
EACL 2026

2025

Language Models can Self-Improve at State-Value Estimation for Better Search
Ethan Mendes and **Alan Ritter**
NeurIPS 2025
Spotlight

Probabilistic Reasoning with LLMs for Privacy Risk Estimation
Jonathan Zheng, **Alan Ritter**, Sauvik Das, Wei Xu
NeurIPS 2025

CARE: Assessing the Impact of Multilingual Human Preference Learning on Cultural Awareness
Geyang Guo, Tarek Naous, Hiromi Wakaki, Yukiko Nishimura, Yuki Mitsufuji, **Alan Ritter**, Wei Xu
EMNLP 2025

USimBench: Are User Simulators Reliable Proxies for Multi-Turn Evaluation of AI Assistants?
Yao Dou, Michel Galley, Baolin Peng, Chris Kedzie, Weixin Cai, **Alan Ritter**, Chris Quirk, Wei Xu, Jianfeng Gao
EMNLP 2025

How to Protect Yourself from 5G Radiation? Investigating LLM Responses to Implicit Misinformation
Ruohao Guo, Wei Xu, **Alan Ritter**
EMNLP 2025

What are Foundation Models Cooking in the Post-Soviet World?
Anton Lavrouk, Tarek Naous, **Alan Ritter**, Wei Xu
EMNLP 2025

Can LLMs Help Uncover Insights about LLMs? A Large-Scale, Evolving Literature Analysis of Frontier LLMs
Jungsoo Park, Junmo Kang, Gabriel Stanovsky, **Alan Ritter**
ACL 2025

Balancing the Budget: Understanding Trade-offs Between Supervised and Preference-Based Fine-tuning
Mohit Raghavendra, Junmo Kang, **Alan Ritter**
ACL 2025
SAC Highlight Award

Translation and Fusion Improves Cross-lingual Information Extraction
Yang Chen, Vedaant Shah, **Alan Ritter**
ACL 2025

Self-MoE: Towards Compositional Large Language Models with Self-Specialized Experts
Junmo Kang, Leonid Karlinsky, Hongyin Luo, Zhen Wang, Jacob A Hansen, James R. Glass, David Daniel Cox, Rameswar Panda, Rogerio Feris, **Alan Ritter**
ICLR 2025

CROSSNEWS: A Cross-Genre Authorship Verification and Attribution Benchmark
Marcus Ma, Duong Minh Le, Junmo Kang, Yao Dou, John Cadigan, Dayne Freitag, **Alan Ritter**,
Wei Xu
AAAI 2025

Measuring, Modeling, and Helping People Account for Privacy Risks in Online Self-Disclosures with
AI
Isadora Krsek, Anubha Kabra, Yao Dou, Tarek Naous, Laura A Dabbish, **Alan Ritter**, Wei Xu,
Sauvik Das
CSCW 2025

2024

Granular Privacy Control for Geolocation with Vision Language Models
Ethan Mendes, Yang Chen, James Hays, Sauvik Das, Wei Xu, **Alan Ritter**
EMNLP 2024

Schema-Driven Information Extraction from Heterogeneous Tables
Fan Bai, Junmo Kang, Gabriel Stanovsky, Dayne Freitag, Mark Dredze, **Alan Ritter**
Findings of EMNLP 2024

ChatHF: Collecting Rich Human Feedback from Real-time Conversations
Andrew Li, Zhenduo Wang, Ethan Mendes, Duong Minh Le, Wei Xu, **Alan Ritter**
EMNLP 2024 (Demo Track)

Are Language Models Robust Coreference Resolvers?
Nghia T. Le, **Alan Ritter**
COLM 2024

UniIR: Training and Benchmarking Universal Multimodal Information Retrievers
Cong Wei, Yang Chen, Haonan Chen, Hexiang Hu, Ge Zhang, Jie Fu, **Alan Ritter**, Wenhui Chen
ECCV 2024

Constrained Decoding for Cross-lingual Label Projection
Duong Minh Le, Yang Chen, **Alan Ritter**, Wei Xu
ICLR 2024

Tensor Trust: Interpretable Prompt Injection Attacks from an Online Game
Sam Toyer, Olivia Watkins, Ethan Adrian Mendes, Justin Svegliato, Luke Bailey, Tiffany Wang,
Isaac Ong, Karim Elmaaroufi, Pieter Abbeel, Trevor Darrell, **Alan Ritter**, Stuart Russell
ICLR 2024 (**Spotlight**)

Having Beer after Prayer? Measuring Cultural Bias in Large Language Models
Tarek Naous, Michael J. Ryan, **Alan Ritter**, Wei Xu
ACL 2024
Social Impact Award

Meta-Tuning LLMs to Leverage Lexical Knowledge for Generalizable Language Style Understanding
Ruohao Guo, Wei Xu, **Alan Ritter**
ACL 2024

NEO-BENCH: Evaluating Robustness of Large Language Models with Neologisms
Jonathan Zheng, **Alan Ritter**, Wei Xu
ACL 2024

Reducing Privacy Risks in Online Self-Disclosures with Language Models
Yao Dou, Isadora Krsek, Tarek Naous, Anubha Kabra, Sauvik Das, **Alan Ritter**, Wei Xu
ACL 2024

Self-Specialization: Uncovering Latent Expertise within Large Language Models
Junmo Kang, Hongyin Luo, Yada Zhu, Jacob Hansen, James Glass, David Cox, **Alan Ritter**,
Rogerio Feris, Leonid Karlinsky
Findings of ACL 2024

2023

Can Pre-trained Vision and Language Models Answer Visual Information-Seeking Questions?
Yang Chen, Hexiang Hu, Yi Luan, Haitian Sun, Soravit Changpinyo, **Alan Ritter**, Ming-Wei Chang
EMNLP 2023

Distill or Annotate? Cost-Efficient Fine-Tuning of Compact Models
Junmo Kang, Wei Xu and **Alan Ritter**
ACL 2023

Improved Instruction Ordering in Recipe-Grounded Conversation
Duong Minh Le, Ruohao Guo, Wei Xu and **Alan Ritter**
ACL 2023

Human-in-the-loop Evaluation for Early Misinformation Detection: A Case Study of COVID-19
Treatments
Ethan Mendes, Yang Chen, Wei Xu, **Alan Ritter**
ACL 2023

Do CoNLL-2003 Named Entity Taggers Still Work Well in 2023?
Shuheng Liu, **Alan Ritter**
ACL 2023

Reproduction Award

Frustratingly Easy Label Projection for Cross-lingual Transfer
Yang Chen, Chao Jiang, **Alan Ritter**, Wei Xu
Findings of ACL 2023

2022

Few-Shot Anaphora Resolution in Scientific Protocols via Mixtures of In-Context Experts
Nghia Le, Fan Bai, **Alan Ritter**
Findings of EMNLP 2022

SynKB: Semantic Search for Synthetic Procedures
Fan Bai, **Alan Ritter**, Peter Madrid, Dayne Freitag and John Niekrasz
EMNLP 2022 (Demo Track)

Extracting a Knowledge Base of COVID-19 Events from Social Media
Shi Zong, Ashutosh Baheti, Wei Xu, **Alan Ritter**
COLING 2022

2021

Pre-train or Annotate? Domain Adaptation with a Constrained Budget
Fan Bai, **Alan Ritter** and Wei Xu
EMNLP 2021

Just Say No: Analyzing the Stance of Neural Dialogue Generation in Offensive Contexts
Ashutosh Baheti, Maarten Sap, **Alan Ritter**, Mark Riedl
EMNLP 2021 [\[Covered by TNW\]](#)

Model Selection for Cross-Lingual Transfer
Yang Chen and **Alan Ritter**
EMNLP 2021

Process-Level Representation of Scientific Protocols with Interactive Annotation
Ronen Tamari, Fan Bai, **Alan Ritter** and Gabriel Stanovsky
EACL 2021

2020

An Empirical Study of Pre-trained Transformers for Arabic Information Extraction
Wuwei Lan, Yang Chen, Wei Xu, **Alan Ritter**
EMNLP 2020 (short paper)

Measuring Forecasting Skill from Text
Shi Zong, **Alan Ritter** and Eduard Hovy
ACL 2020

Fluent Response Generation for Conversational Question Answering
Ashutosh Baheti, **Alan Ritter** and Kevin Small
ACL 2020

Code and Named Entity Recognition in StackOverflow
Jeniya Tabassum, Mounica Maddela, Wei Xu and **Alan Ritter**
ACL 2020

2019

Structured Minimally Supervised Learning for Neural Relation Extraction
Fan Bai, **Alan Ritter**
NAACL 2019

Analyzing the Perceived Severity of Cybersecurity Threats Reported on Social Media
Shi Zong, **Alan Ritter**, Graham Mueller and Evan Wright
NAACL 2019 [\[Covered by WIRED\]](#)

2018

Generating More Interesting Responses in Neural Conversation Models with Distributional Constraints
Ashutosh Baheti, **Alan Ritter**, Jiwei Li, Bill Dolan
EMNLP 2018

An Annotated Corpus for Machine Reading of Instructions in Wet Lab Protocols
Chaitanya Kulkarni, Wei Xu, **Alan Ritter**, Raghu Machiraju
NAACL 2018 (short paper)

2017

Adversarial Learning for Neural Dialogue Generation
Jiwei Li, Will Monroe, Tianlin Shi, Sébastien Jean, **Alan Ritter**, Dan Jurafsky
EMNLP 2017

i have a feeling trump will win.....: Forecasting Winners and Losers from User Predictions on Twitter
Sandesh Swamy, **Alan Ritter** and Marie-Catherine de Marneffe
EMNLP 2017

Learning to Extract Events from Knowledge Base Revisions
Alexander Kononov, Benjamin Strauss, **Alan Ritter** and Brendan O'Connor
WWW 2017

2016

Deep Reinforcement Learning for Dialogue Generation
Jiwei Li, Will Monroe, **Alan Ritter**, Michel Galley, Jianfeng Gao, Dan Jurafsky
EMNLP 2016

A Minimally Supervised Method for Recognizing and Normalizing Time Expressions in Twitter
Jeniya Tabassum, **Alan Ritter**, Wei Xu
EMNLP 2016

Developing a successful SemEval task in sentiment analysis of Twitter and other social media texts
Preslav Nakov, Sara Rosenthal, Svetlana Kiritchenko, Saif M. Mohammad, Zornitsa Kozareva, **Alan Ritter**, Veselin Stoyanov, Xiaodan Zhu
Language Resources and Evaluation 2016

2015

Sense Discovery via Co-Clustering on Images and Text
Xinlei Chen, **Alan Ritter**, Abhinav Gupta and Tom Mitchell
CVPR 2015

Weakly Supervised Extraction of Computer Security Events from Twitter
Alan Ritter, Evan Wright, William Casey and Tom Mitchell
WWW 2015

Never Ending Learning
Tom Mitchell, William Cohen, Estevam Hruschka, Partha Talukdar, Justin Betteridge, Andrew Carlson, Bhavana Dalvi, Matt Gardner, Bryan Kisiel, Jayant Krishnmurthy, Ni Lao, Kathryn Mazaitis, Tahir Mohammad, Ndapa Nakashole, Emmanouil Antonios Platanios, **Alan Ritter**, Mehdi Samadi, Burr Settles, Richard Wang, Derry Wijaya, Abhinav Gupta, Xinlei Chen, Abulhair Saparov, Malcolm Greaves and Joel Welling
AAAI 2015

2014

Major Life Event Extraction from Twitter based on Congratulations/Condolences Speech Acts
Jiwei Li, **Alan Ritter**, Claire Cardie and Eduard Hovy
EMNLP 2014

Weakly Supervised User Profile Extraction from Twitter
Jiwei Li, **Alan Ritter** and Eduard Hovy
ACL 2014

Extracting Lexically Divergent Paraphrases from Twitter
Wei Xu, **Alan Ritter**, Chris Callison-Burch, William B. Dolan and Yangfeng Ji
TACL 2014

2013

Modeling Missing Data in Distant Supervision for Information Extraction
Alan Ritter, Luke Zettlemoyer, Mausam and Oren Etzioni
TACL 2013

2012

Open Domain Event Extraction from Twitter
Alan Ritter, Mausam, Oren Etzioni, Sam Clark
KDD 2012

Paraphrasing for Style
Wei Xu, **Alan Ritter**, Bill Dolan, Ralph Grishman, Colin Cherry
COLING 2012

2011

Named Entity Recognition in Tweets: An Experimental Study
Alan Ritter, Sam Clark, Mausam, Oren Etzioni
EMNLP 2011

Data-Driven Response Generation in Social Media
Alan Ritter, Colin Cherry, Bill Dolan
EMNLP 2011

2010

A Latent Dirichlet Allocation Method for Selectional Preferences
Alan Ritter, Mausam and Oren Etzioni
ACL 2010

Unsupervised Modeling of Twitter Conversations
Alan Ritter, Colin Cherry and Bill Dolan
HLT-NAACL 2010

2009

Learning to Generalize for Complex Selection Tasks
Alan Ritter and Sumit Basu
IUI 2009
Best Student Paper Award

2008

It's a Contradiction – No, It's Not: A Case Study using Functional Relations
Alan Ritter, Doug Downey, Stephen Soderland, and Oren Etzioni
EMNLP 2008

PH.D. STUDENTS	Hyunjoo Chae Geyang Guo (co-advised with Wei Xu) Ruohao Guo Ivy He (co-advised with Wei Xu) Junmo Kang (co-advised with Wei Xu) Duong Minh Le (co-advised with Wei Xu) Nghia (Max) Le (co-advised with Kartik Goyal) Ethan Mendes — Microsoft Research Fellowship (2026) Jungsoo Park Jonathan Zheng (co-advised with Wei Xu)
POSTDOCTORAL MENTORING	Zhenduo Wang (Postdoctoral Researcher, 2024–2025) Research with Ford on generating and evaluating physical safety warnings in AI assistants for repair tasks. Current: AI/ML Engineer, GrowthProtocol.ai.
PH.D. ALUMNI	Yang Chen (Ph.D., Georgia Institute of Technology, 2024; co-advised with Wei Xu) Extracting Knowledge with Multimodal and Multilingual Intelligent Systems. Current: Research Scientist, NVIDIA. Fan Bai (Ph.D., Georgia Institute of Technology, 2023; M.S., Ohio State University, 2020) Current: Bloomberg AI; previously Postdoctoral Fellow, Johns Hopkins University. Jeniya Tabassum (Ph.D., Ohio State University, 2020; co-advised with Wei Xu) Information Extraction From User Generated Noisy Texts. Current: Senior Applied Engineer, Deepgram. Shi Zong (Ph.D., Ohio State University, 2020) Extracting Events from Social Media and the Web. Current: Applied Scientist, Dialpad; previously Mitacs Elevate Postdoctoral Fellow, University of Waterloo. Xiaonan Ji (Ph.D., Ohio State University, 2018; co-advised with Po-Yin Yen) An Integrated Framework of Text and Visual Analytics to Facilitate Information Retrieval towards Biomedical Literature. Current: Motional.
M.S. ALUMNI	Mohit Raghavendra (M.S., Georgia Institute of Technology, 2025) Current: Machine Learning Researcher, Scale AI. Andrew Li (M.S., Georgia Institute of Technology, 2024) Vedaant Shah (M.S., Georgia Institute of Technology, 2024) Current: Databricks. Sandesh Swamy (M.S., Ohio State University, 2017; co-advised with Marie-Catherine de Marneffe) Forecasting Event Outcomes from User Predictions on Twitter. Current: Senior Applied Scientist, AWS AI Labs.
TEACHING	GT CS 7652: Large Language Models (Fall 2026; 66 students) GT CS 4650: Natural Language Processing (Spring 2026; 209 students) GT CS 8803: Large Language Models (Fall 2025; 39 students) GT CS 7650: Natural Language Processing (Spring 2025; 124 students) GT CS 4650: Natural Language Processing (Fall 2024; 54 students) GT CS 8803: Large Language Models (Spring 2024; 32 students) GT CS 4650: Natural Language Processing (Fall 2023; 48 students) GT CS 7650: Natural Language Processing (Spring 2023; 102 students) GT CS 4650: Natural Language Processing (Fall 2022; 61 students) GT CS 7650: Natural Language Processing (Spring 2022; 113 students)

GT CS 4641: Machine Learning (Fall 2021; 222 students)
 GT CS 7650: Natural Language Processing (Spring 2021; 91 students)
 OSU CSE 5539: Social Media and Text Analytics (Spring 2020; 18 students)
 OSU CSE 5525: Speech and Language Processing (Fall 2019; 60 students)
 OSU CSE 5525: Speech and Language Processing (Spring 2019; 33 students)
 OSU CSE 5523: Machine Learning and Statistical Pattern Recognition (Fall 2018; 35 students)
 OSU CSE 5539: Social Media and Text Analytics (Spring 2018; 26 students)
 OSU CSE 5522: Survey of Artificial Intelligence II: Advanced Techniques (Fall 2017; 36 students)
 OSU CSE 5523: Machine Learning and Statistical Pattern Recognition (Spring 2017; 38 students)
 OSU CSE 5525: Speech and Language Processing (Spring 2016; 29 students)
 OSU CSE 5539: Web Information Extraction (Fall 2015; 15 students)
 OSU CSE 5522: Survey of Artificial Intelligence II: Advanced Techniques (Spring 2015; 37 students)
 OSU CSE 5539: NLP and IE for the Social Web (Fall 2014; 14 students)

SERVICE

Program Chair

NAACL 2025

Best Paper Awards Committee

EACL 2024

Demo Chair

ACL 2023

Organizer

Workshop on Evaluation of Interactive Agents (NeurIPS 2026)
 ACL/EMNLP Workshop on Noisy User-Generated Text (WNUT 2015–present)
 Sentiment Analysis in Twitter Shared Task (SemEval 2013, 2014, 2015, 2016)
 Named Entity Recognition in Twitter Shared Task (at WNUT 2015, 2016)

Area Chair and Senior Area Chair

Area Chair, COLM 2026
 Senior Area Chair, ACL 2026
 Area Chair, COLM 2025
 Area Chair, ACL 2024
 Area Chair, EMNLP 2023
 Area Chair, ACL 2023
 Senior Area Chair, EMNLP 2022
 Senior Area Chair, ACL 2022
 Senior Area Chair, EMNLP 2021
 Area Chair, ACL 2021
 Senior Area Chair, NAACL 2021
 Senior Area Chair, EMNLP 2020
 Area Chair, AKBC 2020
 Area Chair, ACL 2020
 Area Chair, EMNLP 2019
 Area Chair, ACL 2019
 Area Chair, CoNLL 2018
 Area Chair, EMNLP 2018
 Area Chair, NAACL 2018

Area Chair, EMNLP 2016

Area Chair, NAACL 2015

Area Chair, ACL 2014

Senior Program Committee

AAAI 2022

Steering Committee

SemEval 2016, 2018, 2019, 2020

Proposal Reviewing

Panelist for U.S. National Science Foundation, Division of Information and Intelligent Systems (CISE/IIS) (twice in 2017, twice in 2018, once in 2019, once in 2023, once in 2024, twice in 2026)

Independent Reviewer for U.S. National Science Foundation, Division of Information and Intelligent Systems (CISE/IIS) (2020)

Independent Reviewer for Hong Kong Research Grants Council (2016)

Independent Reviewer for Natural Sciences and Engineering Research Council of Canada (2016)

Independent Reviewer for U.S. Department of Energy (2015)

Departmental/College Service

GT College of Computing Undergraduate Curriculum Committee, 2020–present

GT Faculty Lead, Intelligence Thread, 2020–present

GT Computational Media Undergraduate Curriculum Committee, 2020–2026

GT CoC Liaison to the Minor in Applications of Artificial Intelligence and Machine Learning, Spring 2024–Fall 2026

GT School of Interactive Computing Faculty Hiring Committee, 2020–2022, 2024–2026

OSU CSE Graduate Studies Committee (2016–2020)

OSU Artificial Intelligence Club, Founding Faculty Advisor (2017–2020)

OSU CSE Curriculum Committee (2019–2020)

OSU CSE Graduate Admissions Committee — Artificial Intelligence Group Representative (2014–2020)

OSU CSE Prospective Ph.D. Student Visit Day Organizer (Fall 2014–2016)

EXTERNAL THESIS COMMITTEES

Hassan Shavarani

Ph.D., December 2024

Structured Prediction for Compute Efficient and High Accuracy NLP (Advisor: Anoop Sarkar)
Simon Fraser University

Evangelia Spiliopoulou

Ph.D., January 2023

Modeling Event Implications via Multi-faceted Entity Representations (Advisor: Eduard Hovy)
Carnegie Mellon University

Alex Memory

Ph.D., October 2019

Collective Relational Data Integration with Diverse and Noisy Evidence (Advisor: Lise Getoor)
University of Maryland

Young-Bum Kim

Ph.D., February 2015

Natural Language Technologies for Low-Resource Languages (Advisor: Benjamin Snyder)
University of Wisconsin

INVITED TALKS

Google Research, December 2025
 Sungkyunkwan University, May 2025
 University of Tokyo, October 2024
 Shonan Seminar on Empowering Local Open Source LLMs, October 2024
 University of California San Diego, May 2024
 University of California Irvine, March 2024
 BrainLink Workshop, December 2023
 University of Chicago / TTIC, November 2023
 Allen Institute for Artificial Intelligence, October 2023
 University of Washington, October 2023
 Matching Workshop at ACL, July 2023
 University of Pennsylvania, February 2023
 Princeton University, October 2022
 Columbia University, October 2022
 New York University, October 2022
 Google New York, October 2022
 Johns Hopkins University, May 2022
 Megagon Labs, January 2022
 Twitter, September 2021
 University of Notre Dame, September 2021
 University of Maryland, April 2020
 UMass Amherst, February 2020
 NeurIPS Workshop on Conversational AI, December 2019
 Georgia Institute of Technology, December 2019
 Twitter, June 2019
 University of Michigan, April 2019
 University of Maryland, April 2019
 Rice University, March 2019
 University of Wisconsin, January 2019
 Stanford NLP Seminar, May 2018
 NeurIPS Workshop on Learning With Limited Labeled Data, December 2017
 Microsoft Research Faculty Summit, July 2017
 Midwest Machine Learning Symposium, June 2017
 University of Colorado, March 2017
 Tsinghua University, December 2016
 University of Edinburgh, July 2016
 University of North Carolina, March 2016
 Machine Learning Seminar, Purdue University, December 2015
 Statistics Department, Ohio State University, October 2015
 School of Computer Science, Carnegie Mellon University, April 2015
 CLIP, University of Maryland, March 2015
 Clunch, University of Pennsylvania, September 2014
 University of Pittsburgh, January 2014
 Machine Learning Lunch, Carnegie Mellon University, October 2013
 IEEE Intelligence and Security Informatics, June 2013
 Toyota Technological Institute at Chicago, April 2013
 Microsoft Research, April 2013
 University of Texas at Dallas, March 2013
 The Ohio State University, March 2013
 Machine Learning and Friends Lunch, University of Massachusetts Amherst, March 2013
 Human Language Technology Center of Excellence, March 2013
 Pacific Northwest National Laboratory, January 2013
 New York University, September 2012
 Wavii, April 2012

Facebook, March 2012
 Microsoft Research, September 2011
 Google Research, September 2011
 Yahoo! Labs, September 2011
 Twitter, May 2011
 Information Sciences Institute, University of Southern California, February 2011
 USC ICT workshop on Reasoning with Text, February 2011
 UW CSE Industrial Affiliates Meeting, October 2009
 UW/Microsoft Quarterly Symposium on Computational Linguistics, October 2009
 Boeing Research, February 2008

PATENTS

U.S. Patent 10,108,934 (granted October 23, 2018)
 Items Selection Via Automatic Generalization
 Work from my 2008 internship at Microsoft Research, which also resulted in a best student paper award at the ACM Conference on Intelligent User Interfaces.

PRESS COVERAGE

LLMs exhibit significant Western cultural bias, study finds, VentureBeat (March 2024)
 Coverage of *Having Beer after Prayer? Measuring Cultural Bias in Large Language Models* (ACL 2024).

GPT-3 mimics human love for ‘offensive’ Reddit comments, study finds, TNW (August 2021)
 Coverage of *Just Say No: Analyzing the Stance of Neural Dialogue Generation in Offensive Contexts* (EMNLP 2021).

Machine Learning Can Use Tweets to Spot Critical Security Flaws, WIRED (March 2019)
 Coverage of *Analyzing the Perceived Severity of Cybersecurity Threats Reported on Social Media* (NAACL 2019).