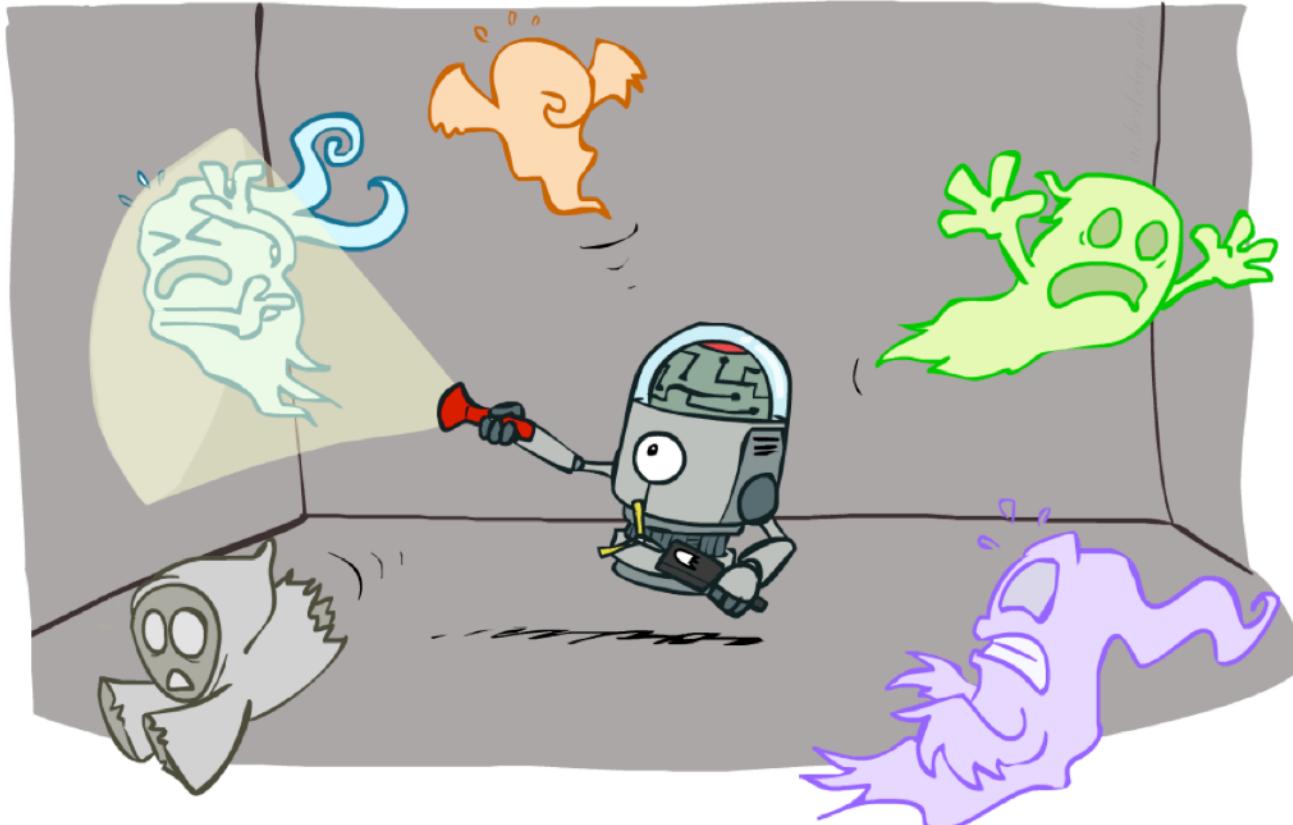


CS 5522: Artificial Intelligence II

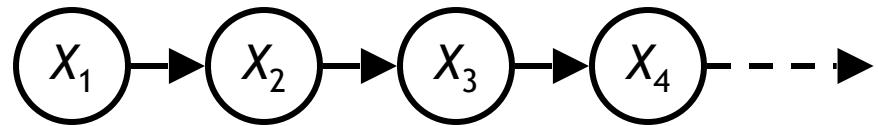
Particle Filters and Applications of HMMs



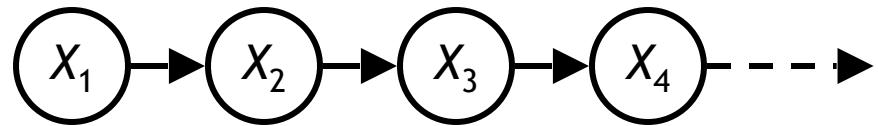
Instructor: Alan Ritter

Ohio State University

Recap: Reasoning Over Time



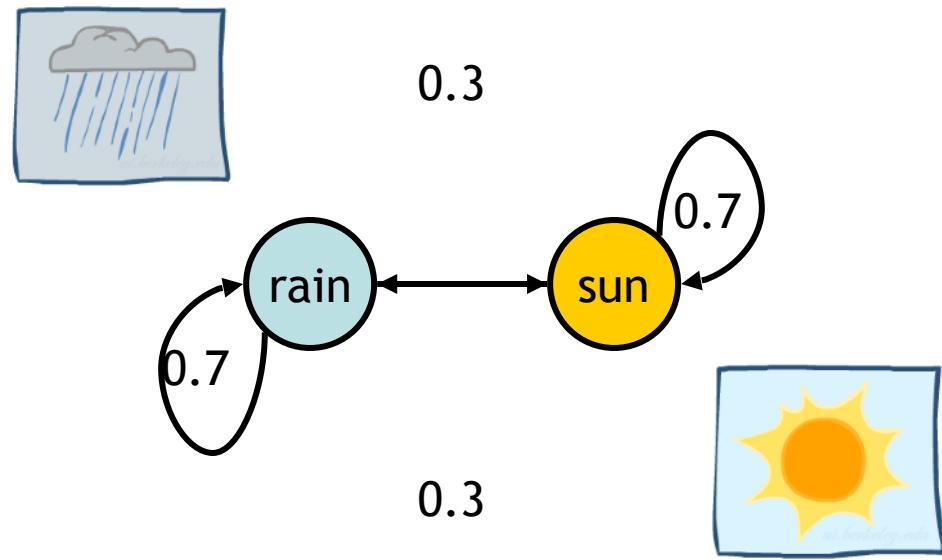
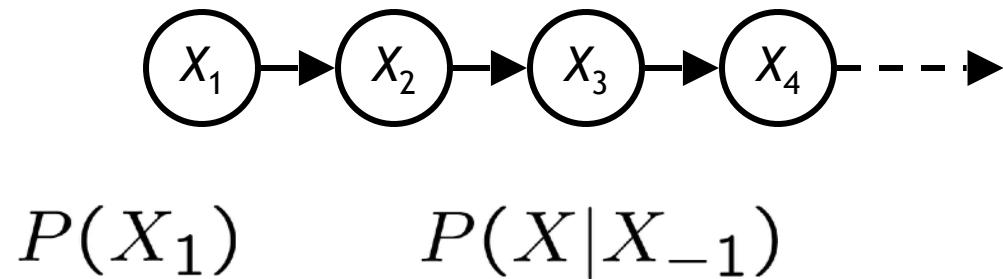
Recap: Reasoning Over Time



$$P(X_1)$$

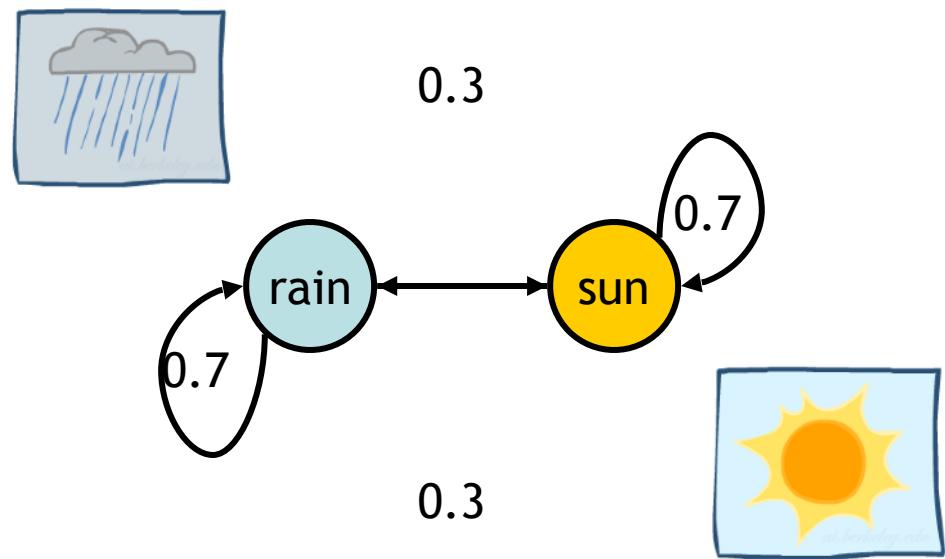
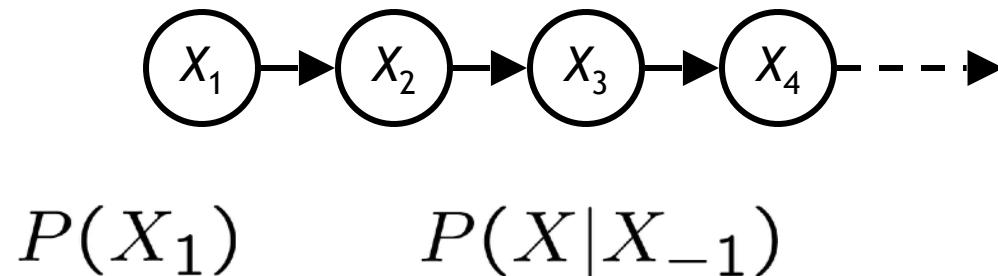
$$P(X|X_{-1})$$

Recap: Reasoning Over Time

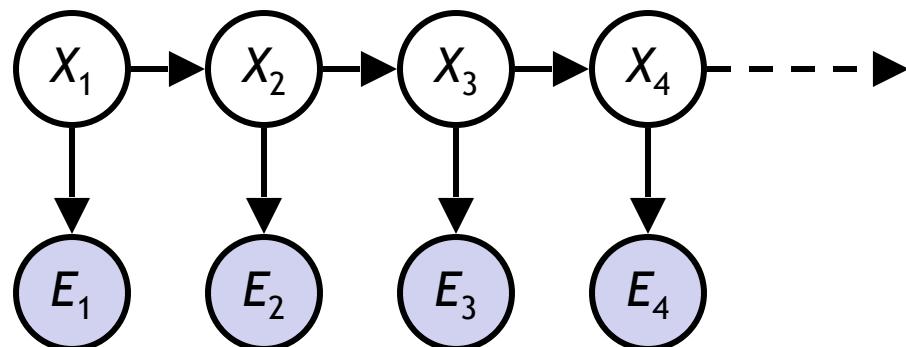


Recap: Reasoning Over Time

- Markov models

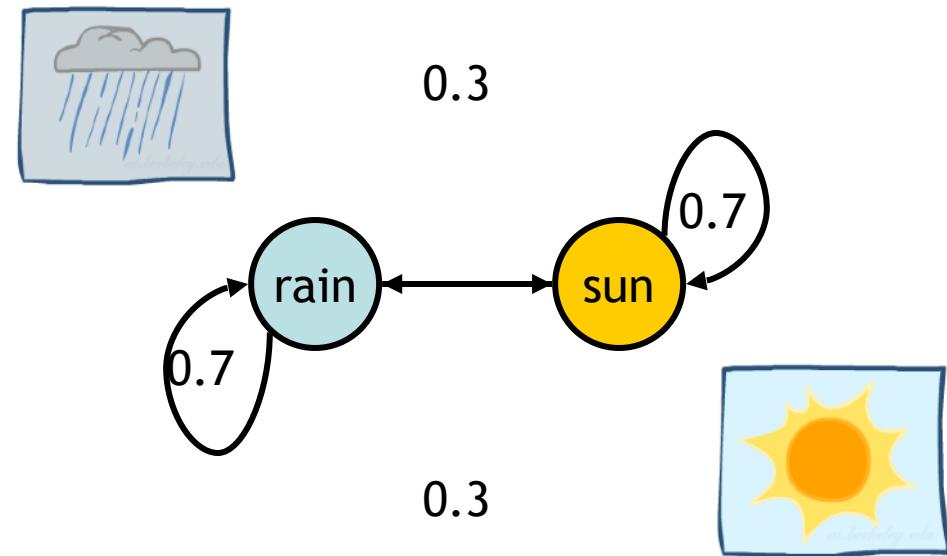
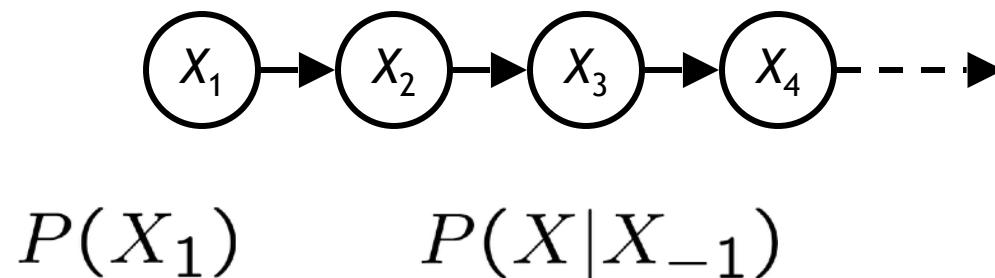


- Hidden Markov models

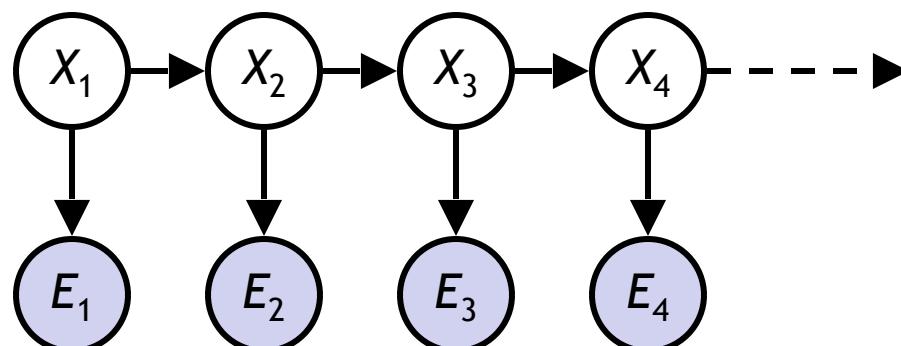


Recap: Reasoning Over Time

- Markov models

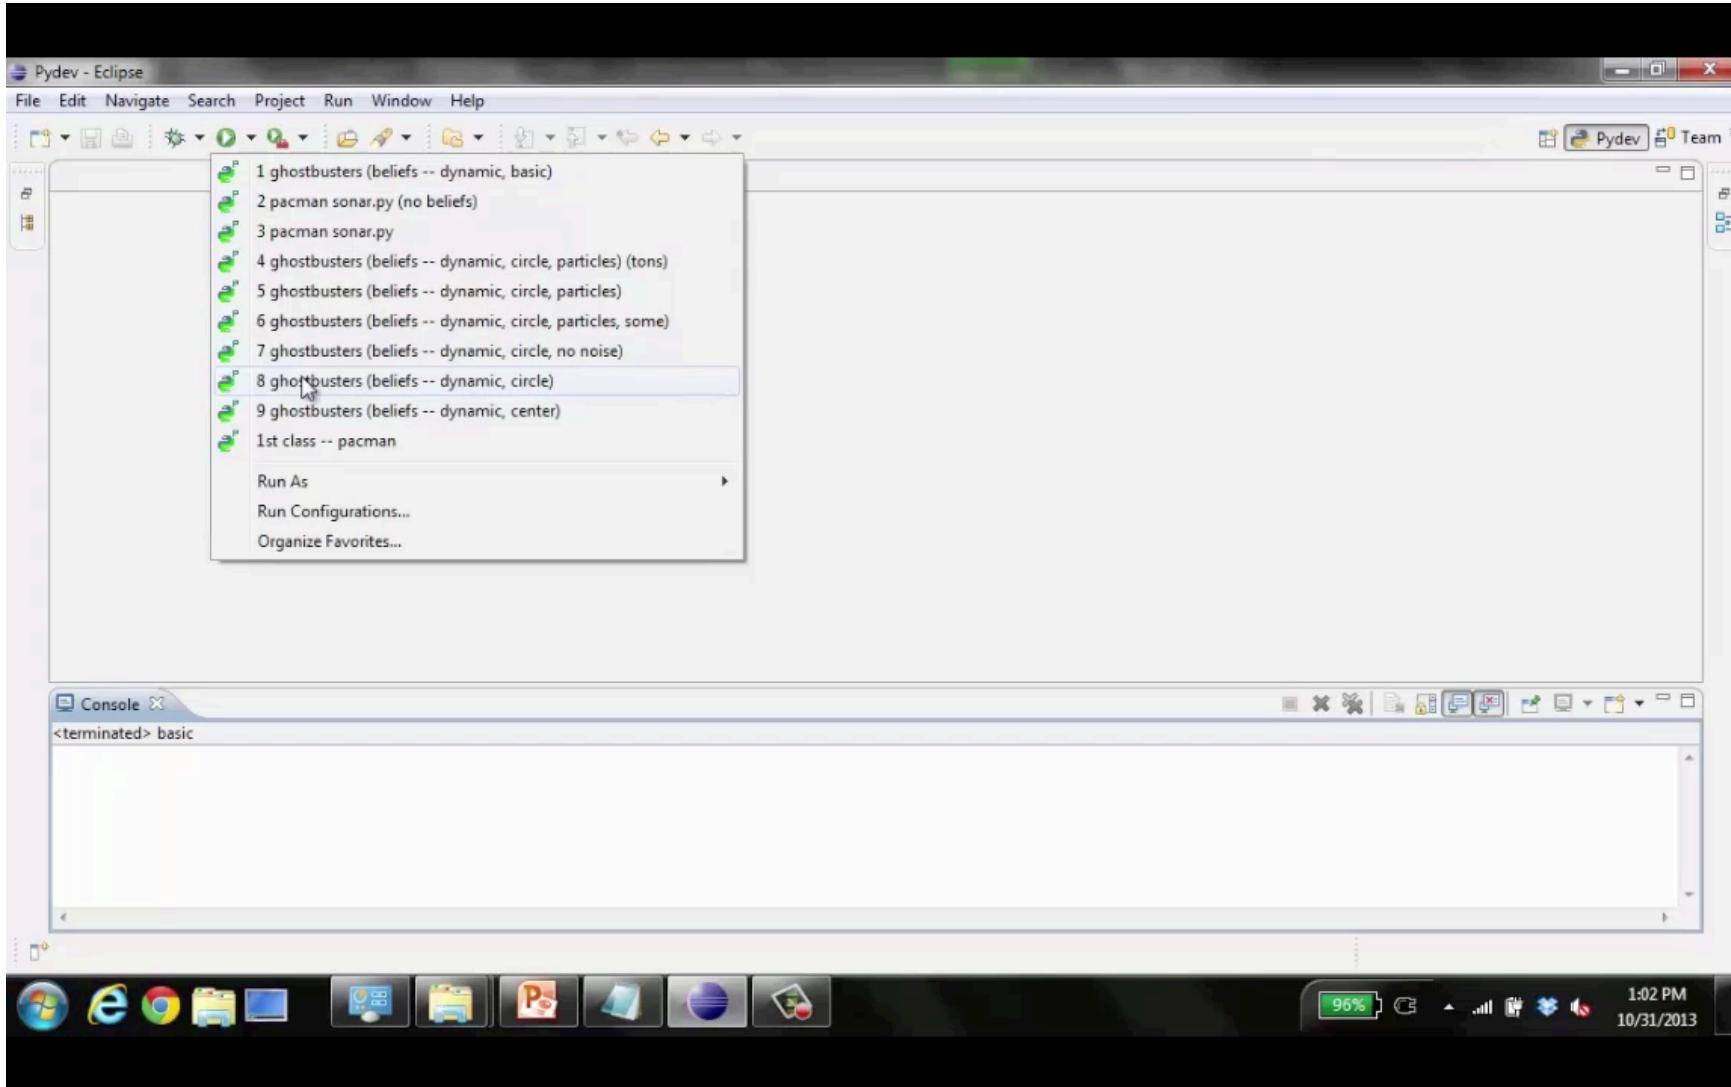


- Hidden Markov models

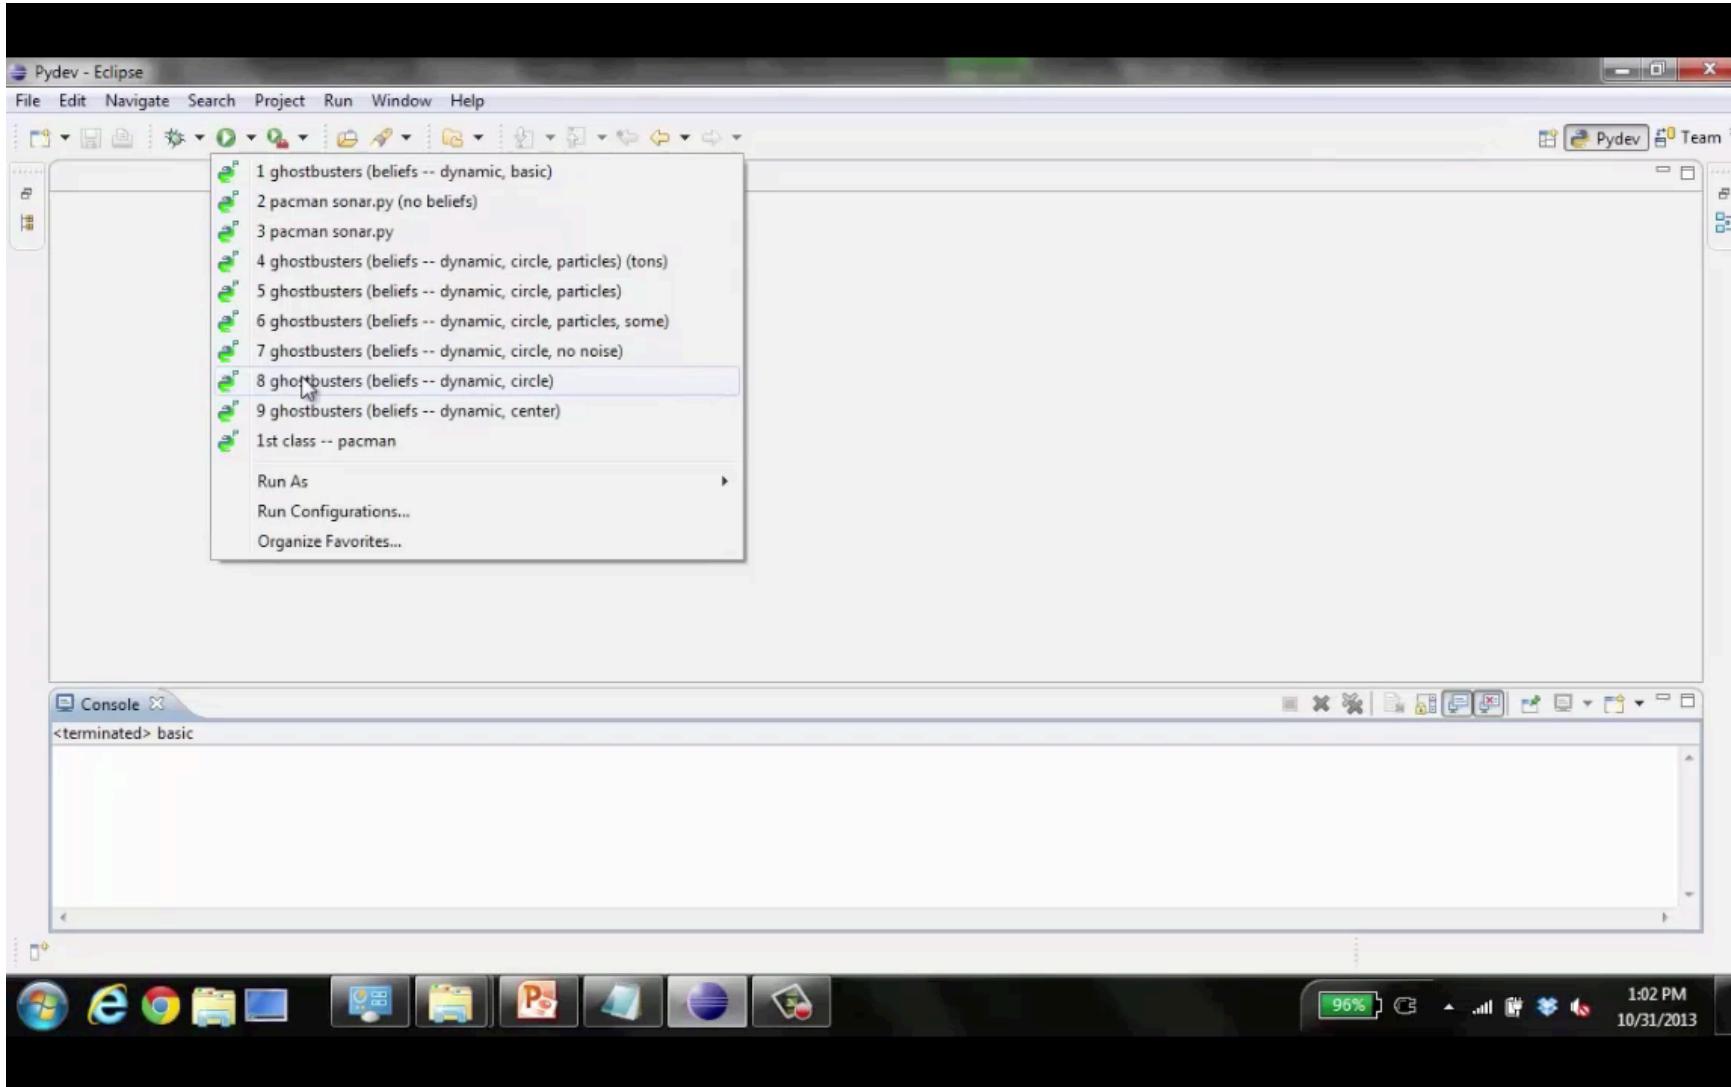


X	E	P
rain	umbrella	0.9
rain	no umbrella	0.1
sun	umbrella	0.2
sun	no umbrella	0.8

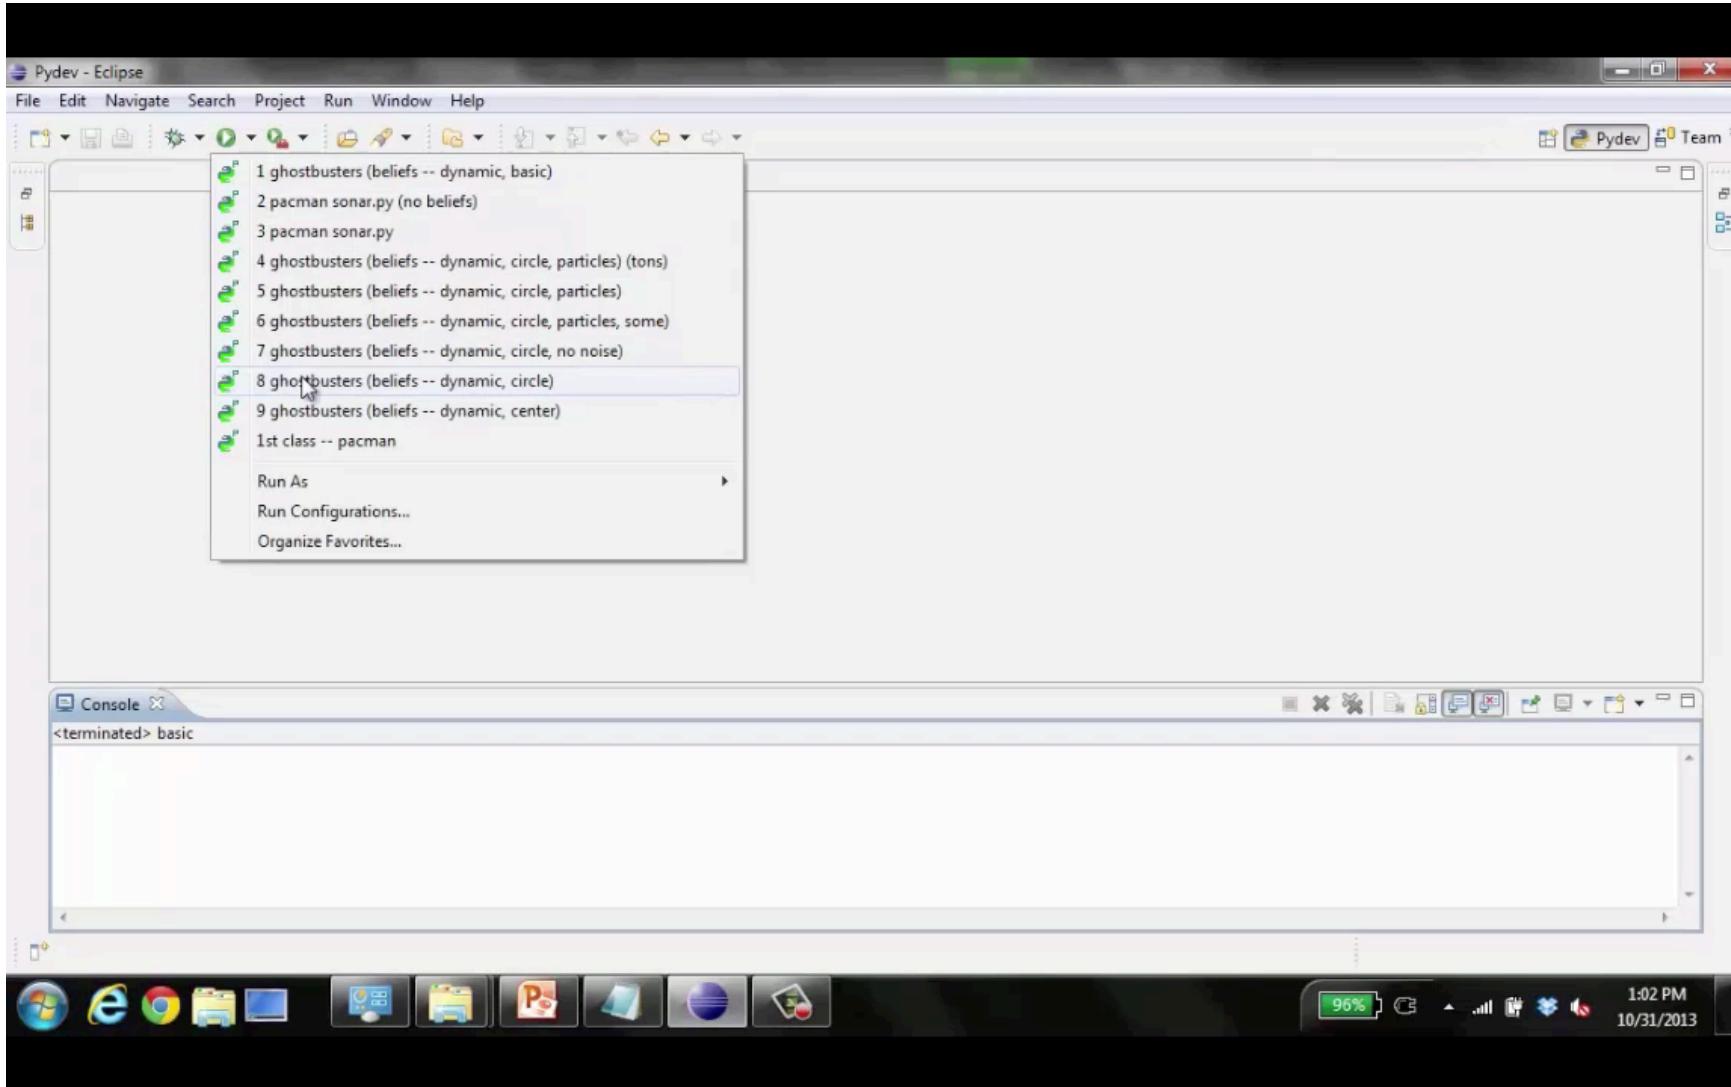
Video of Demo Ghostbusters Markov Model (Reminder)



Video of Demo Ghostbusters Markov Model (Reminder)



Video of Demo Ghostbusters Markov Model (Reminder)



Recap: Filtering

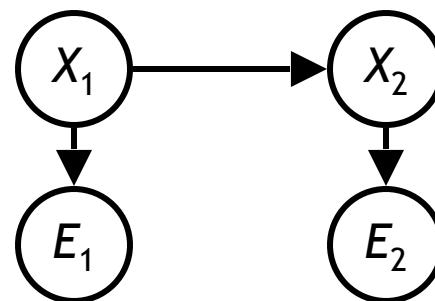
Elapse time: compute $P(X_t | e_{1:t-1})$

$$P(x_t | e_{1:t-1}) = \sum_{x_{t-1}} P(x_{t-1} | e_{1:t-1}) \cdot P(x_t | x_{t-1})$$

Observe: compute $P(X_t | e_{1:t})$

$$P(x_t | e_{1:t}) \propto P(x_t | e_{1:t-1}) \cdot P(e_t | x_t)$$

<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01
<0.01	0.76	0.06	0.06	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01



Belief: $\langle P(\text{rain}), P(\text{sun}) \rangle$

Recap: Filtering

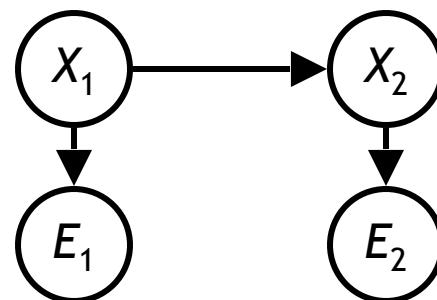
Elapse time: compute $P(X_t | e_{1:t-1})$

$$P(x_t | e_{1:t-1}) = \sum_{x_{t-1}} P(x_{t-1} | e_{1:t-1}) \cdot P(x_t | x_{t-1})$$

Observe: compute $P(X_t | e_{1:t})$

$$P(x_t | e_{1:t}) \propto P(x_t | e_{1:t-1}) \cdot P(e_t | x_t)$$

<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01
<0.01	0.76	0.06	0.06	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01



Belief: $\langle P(\text{rain}), P(\text{sun}) \rangle$

$$P(X_1) \quad \langle 0.5, 0.5 \rangle \quad \text{Prior on } X_1$$

Recap: Filtering

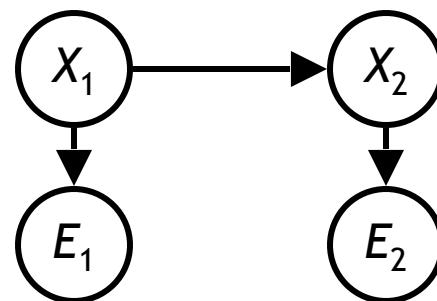
Elapse time: compute $P(X_t | e_{1:t-1})$

$$P(x_t | e_{1:t-1}) = \sum_{x_{t-1}} P(x_{t-1} | e_{1:t-1}) \cdot P(x_t | x_{t-1})$$

Observe: compute $P(X_t | e_{1:t})$

$$P(x_t | e_{1:t}) \propto P(x_t | e_{1:t-1}) \cdot P(e_t | x_t)$$

<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01
<0.01	0.76	0.06	0.06	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01



Belief: $\langle P(\text{rain}), P(\text{sun}) \rangle$

$$P(X_1) \quad <0.5, 0.5> \quad \text{Prior on } X_1$$

$$P(X_1 | E_1 = \text{umbrella}) \quad <0.82, 0.18> \quad \text{Observe}$$

Recap: Filtering

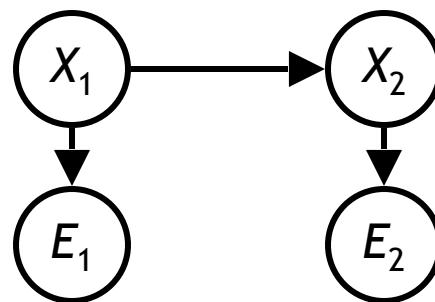
Elapse time: compute $P(X_t | e_{1:t-1})$

$$P(x_t | e_{1:t-1}) = \sum_{x_{t-1}} P(x_{t-1} | e_{1:t-1}) \cdot P(x_t | x_{t-1})$$

Observe: compute $P(X_t | e_{1:t})$

$$P(x_t | e_{1:t}) \propto P(x_t | e_{1:t-1}) \cdot P(e_t | x_t)$$

<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01
<0.01	0.76	0.06	0.06	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01



Belief: $\langle P(\text{rain}), P(\text{sun}) \rangle$

$$P(X_1) \quad \langle 0.5, 0.5 \rangle \quad \text{Prior on } X_1$$

$$P(X_1 | E_1 = \text{umbrella}) \quad \langle 0.82, 0.18 \rangle \quad \text{Observe}$$

$$P(X_2 | E_1 = \text{umbrella}) \quad \langle 0.63, 0.37 \rangle \quad \text{Elapse time}$$

Recap: Filtering

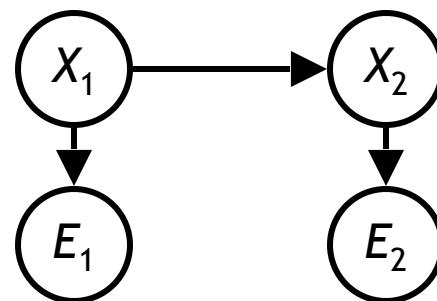
Elapse time: compute $P(X_t | e_{1:t-1})$

$$P(x_t | e_{1:t-1}) = \sum_{x_{t-1}} P(x_{t-1} | e_{1:t-1}) \cdot P(x_t | x_{t-1})$$

Observe: compute $P(X_t | e_{1:t})$

$$P(x_t | e_{1:t}) \propto P(x_t | e_{1:t-1}) \cdot P(e_t | x_t)$$

<0.01	<0.01	<0.01	<0.01	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01
<0.01	0.76	0.06	0.06	<0.01	<0.01
<0.01	<0.01	0.06	<0.01	<0.01	<0.01



Belief: $\langle P(\text{rain}), P(\text{sun}) \rangle$

$$P(X_1) \quad <0.5, 0.5> \quad \text{Prior on } X_1$$

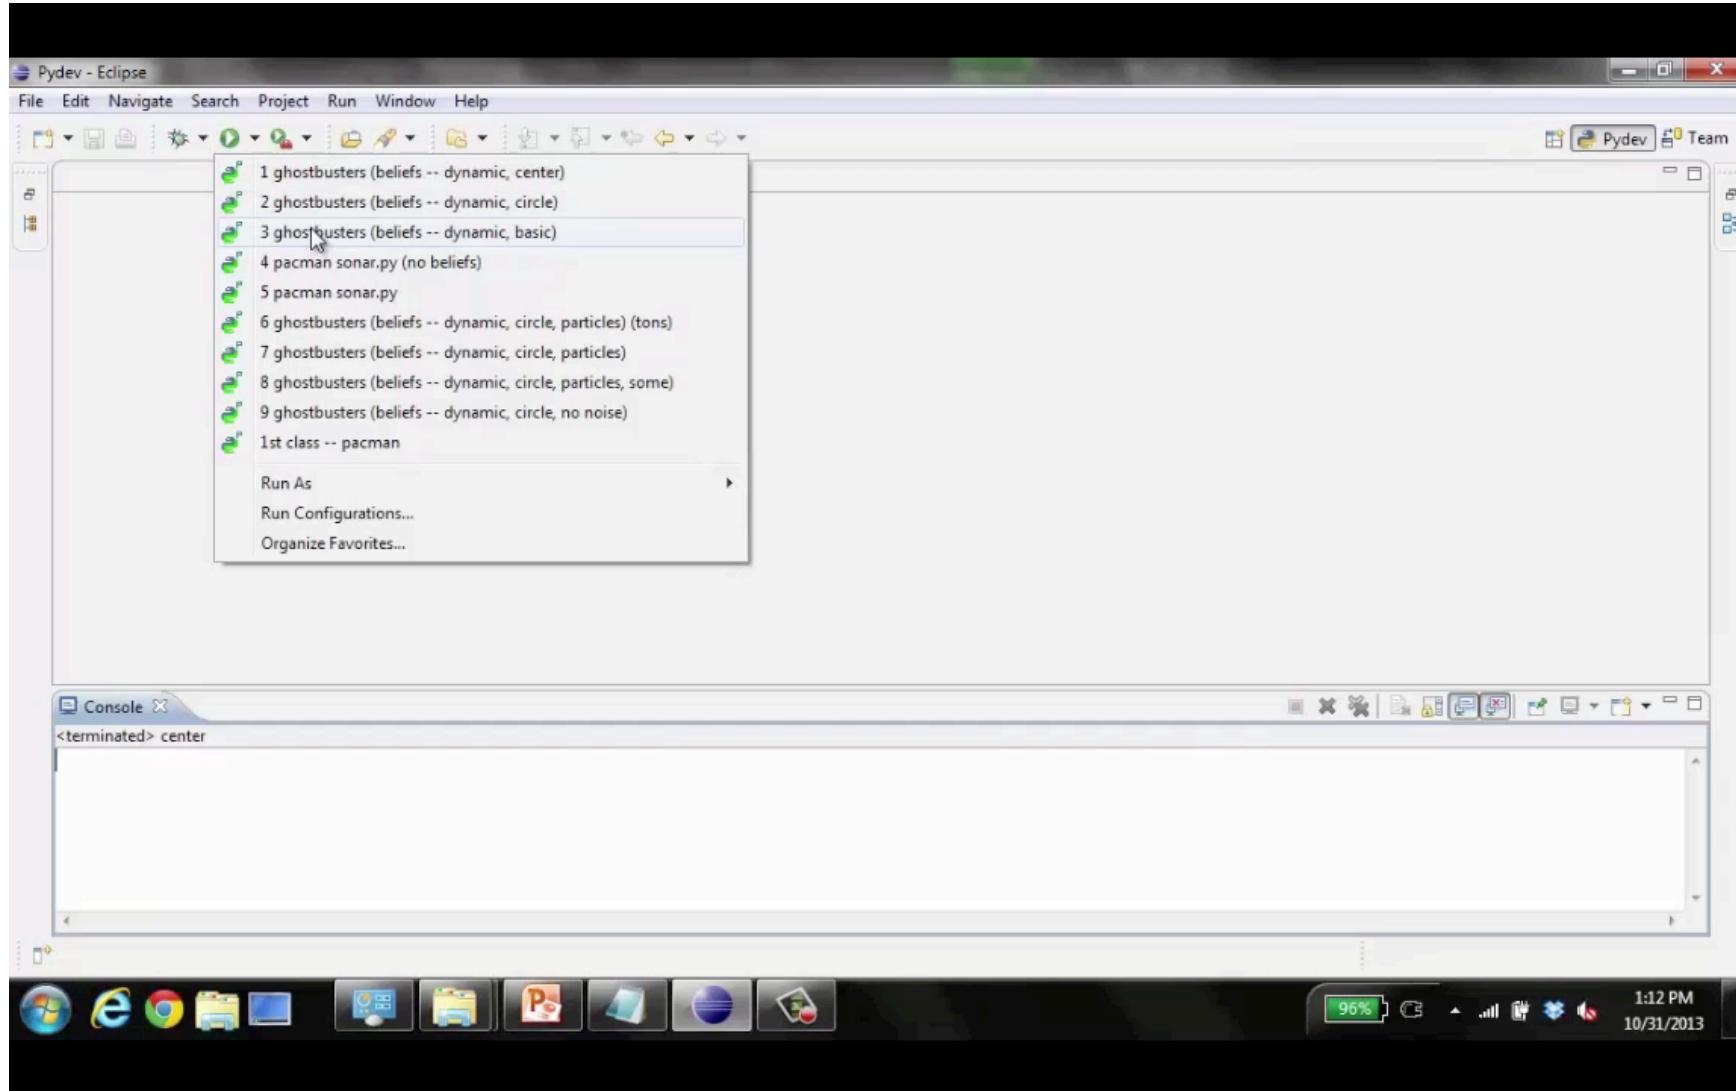
$$P(X_1 | E_1 = \text{umbrella}) \quad <0.82, 0.18> \quad \text{Observe}$$

$$P(X_2 | E_1 = \text{umbrella}) \quad <0.63, 0.37> \quad \text{Elapse time}$$

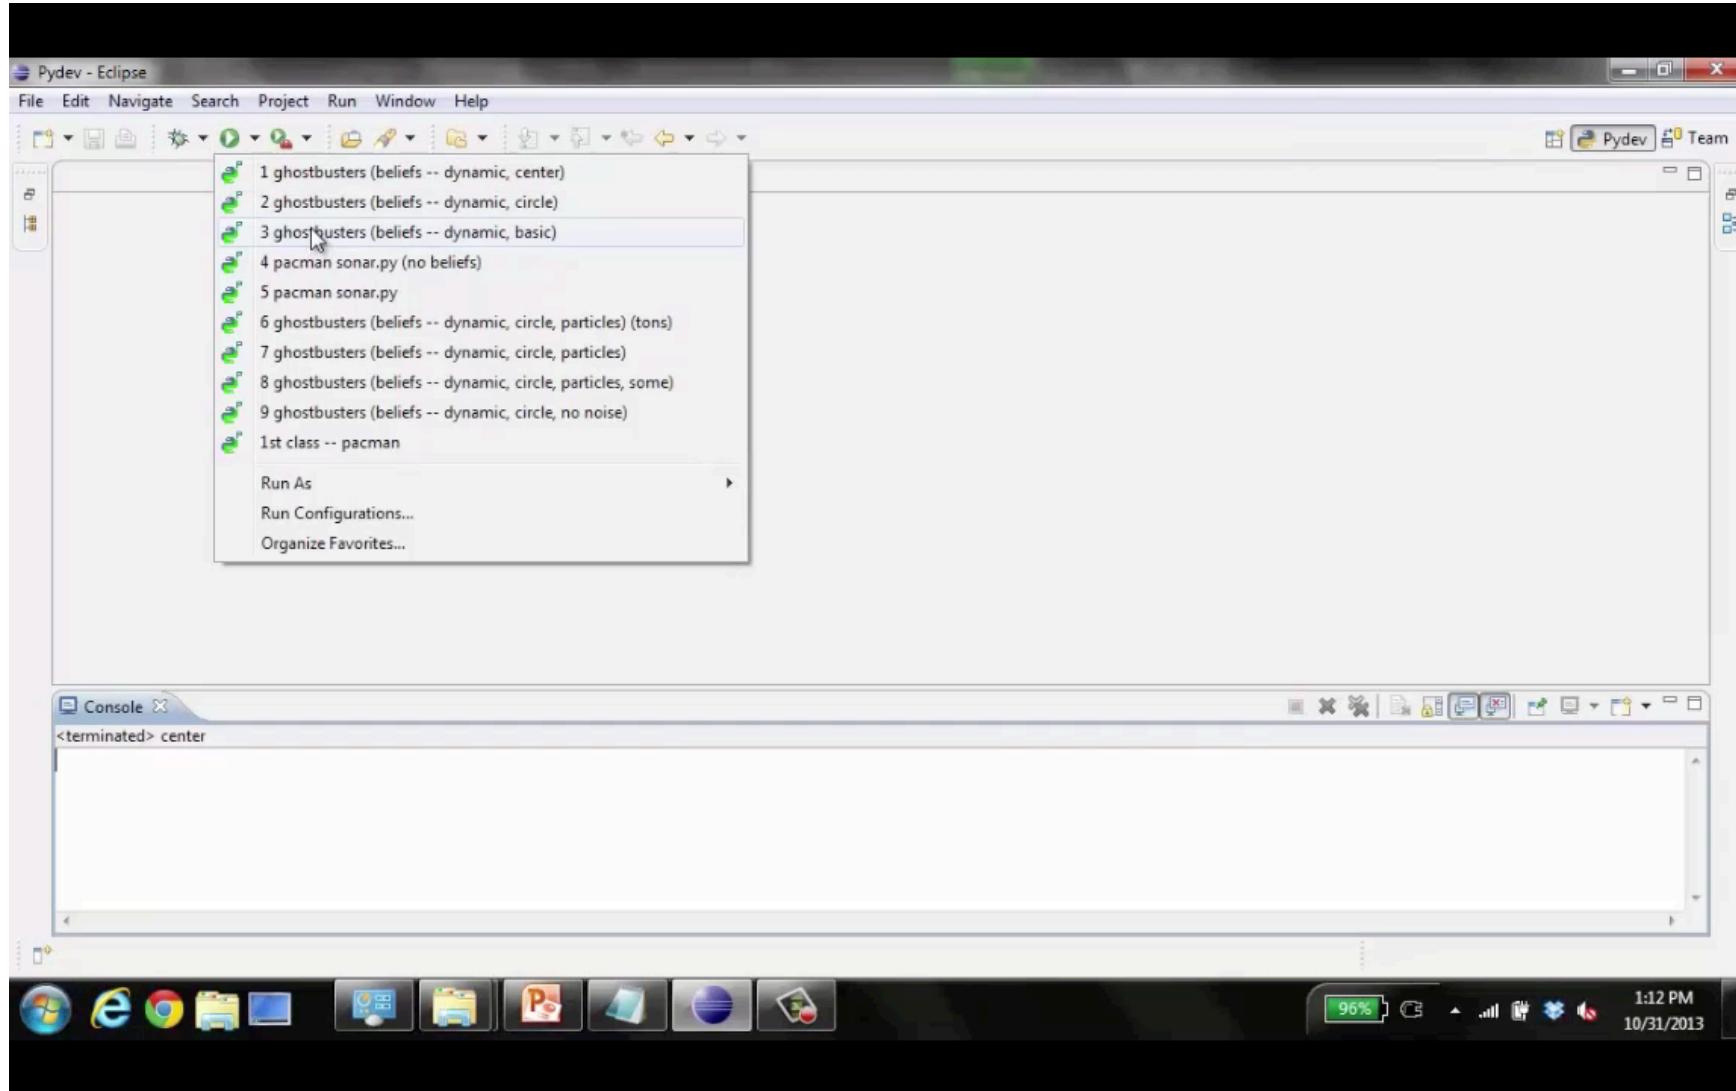
$$P(X_2 | E_1 = \text{umb}, E_2 = \text{umb}) \quad <0.88, 0.12> \quad \text{Observe}$$

[Demo: Ghostbusters Exact Filtering

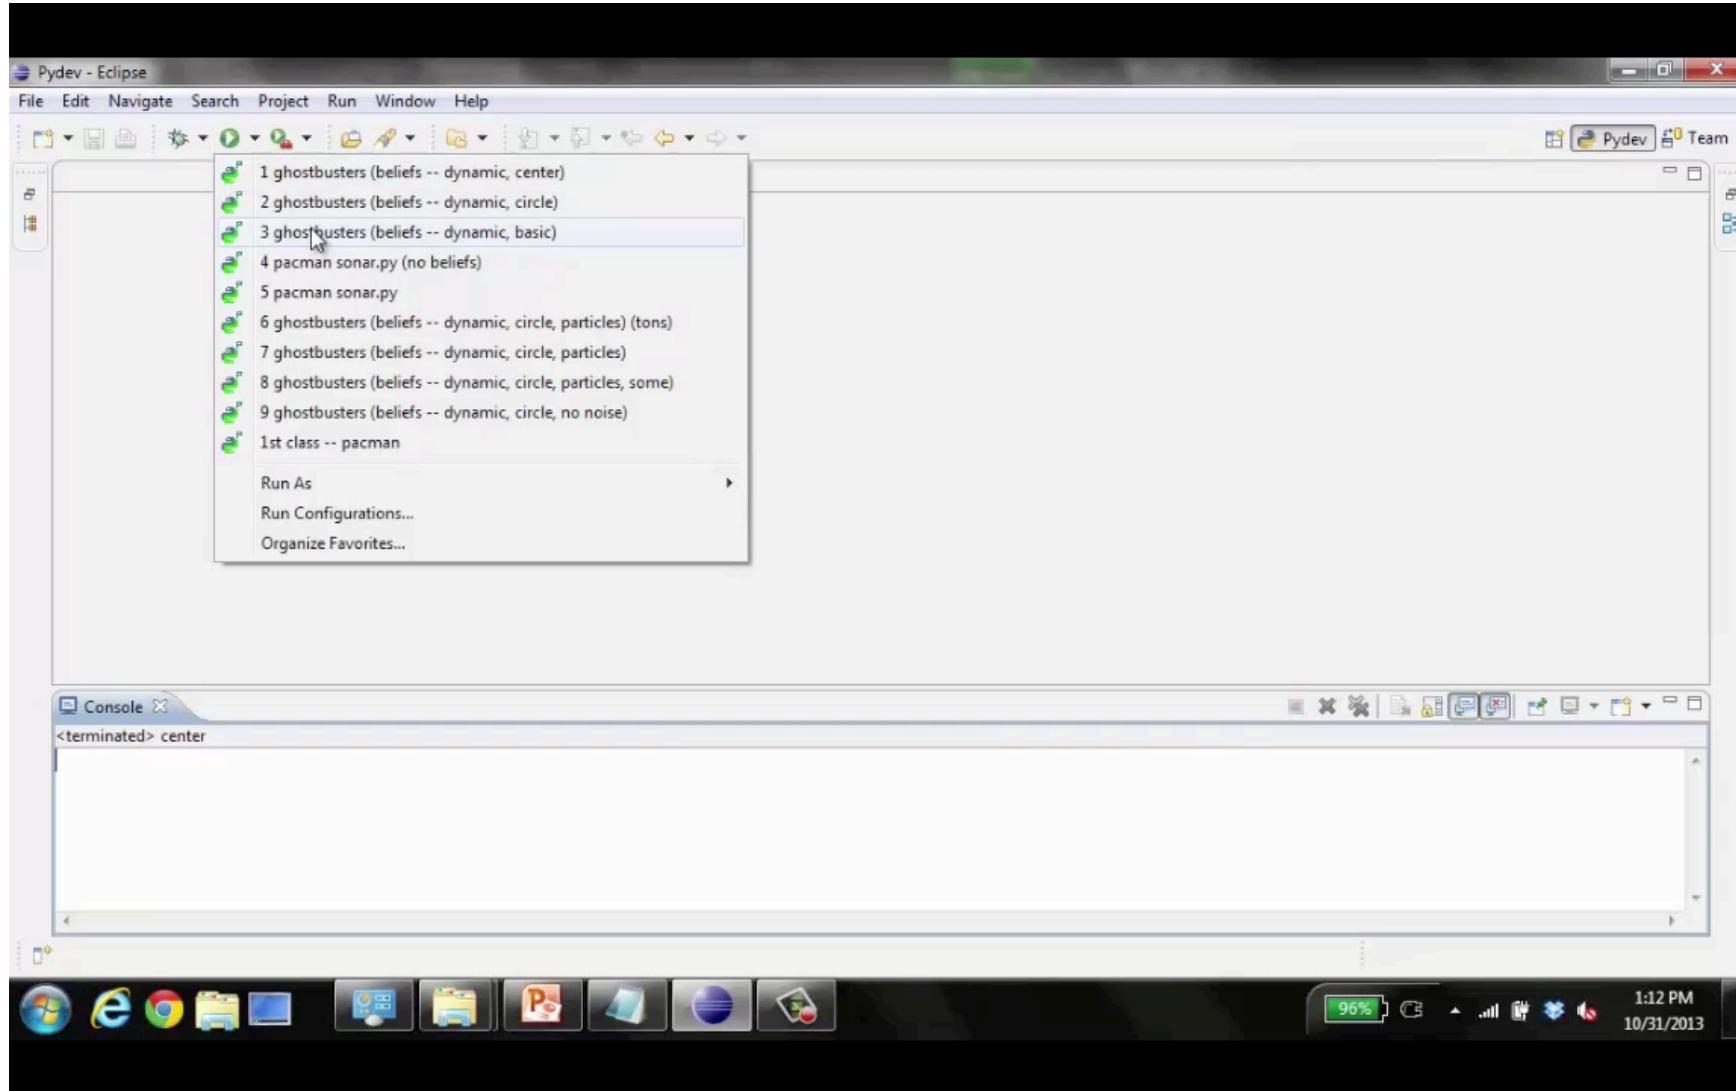
Video of Ghostbusters Exact Filtering (Reminder)



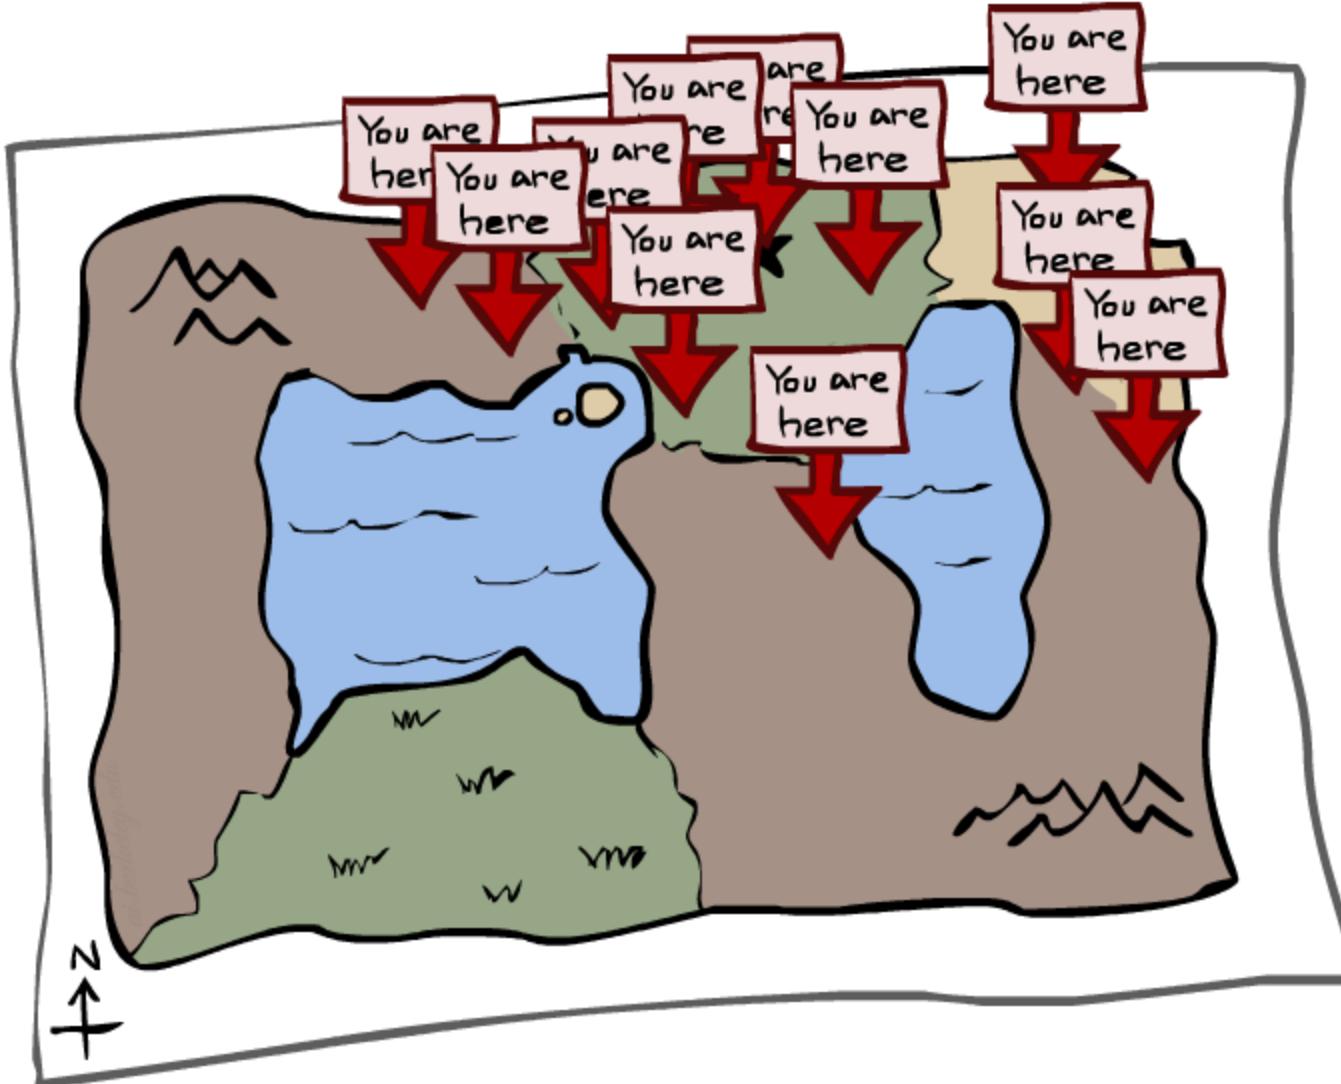
Video of Ghostbusters Exact Filtering (Reminder)



Video of Ghostbusters Exact Filtering (Reminder)



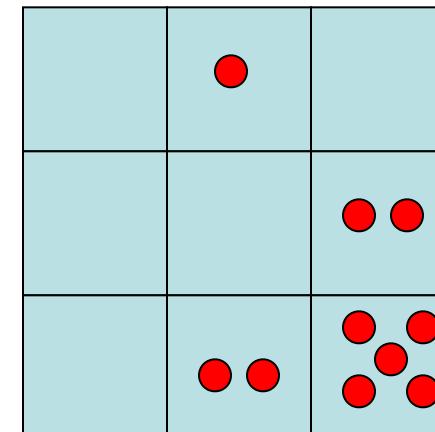
Particle Filtering



Particle Filtering

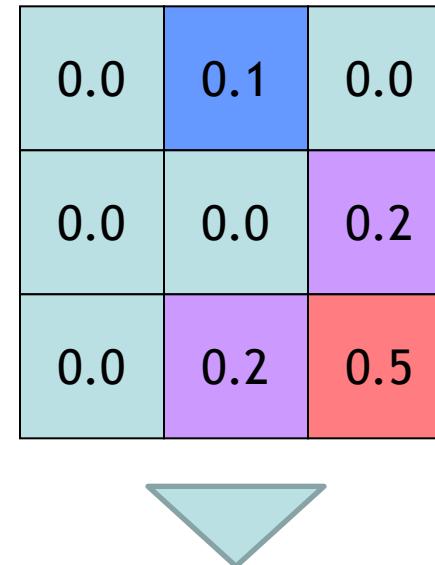
- Filtering: approximate solution
- Sometimes $|X|$ is too big to use exact inference
 - $|X|$ may be too big to even store $B(X)$
 - E.g. X is continuous

0.0	0.1	0.0
0.0	0.0	0.2
0.0	0.2	0.5



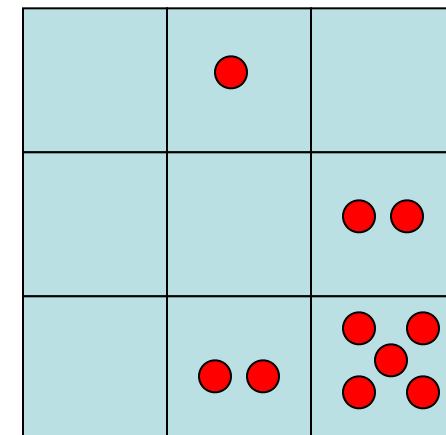
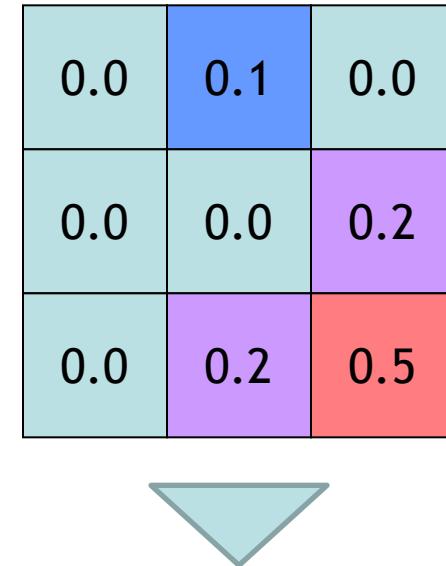
Particle Filtering

- Filtering: approximate solution
- Sometimes $|X|$ is too big to use exact inference
 - $|X|$ may be too big to even store $B(X)$
 - E.g. X is continuous
- Solution: approximate inference



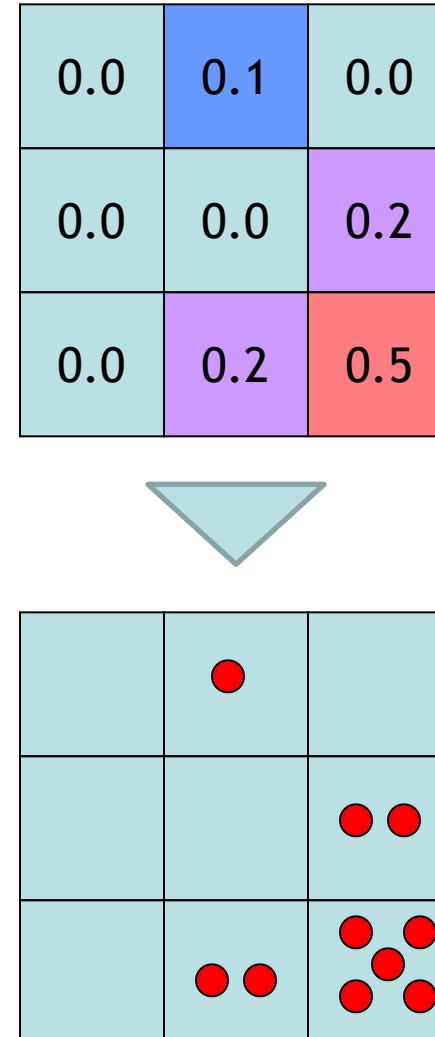
Particle Filtering

- Filtering: approximate solution
- Sometimes $|X|$ is too big to use exact inference
 - $|X|$ may be too big to even store $B(X)$
 - E.g. X is continuous
- Solution: approximate inference
 - Track samples of X , not all values
 - Samples are called particles
 - Time per step is linear in the number of samples
 - But: number needed may be large
 - In memory: list of particles, not states



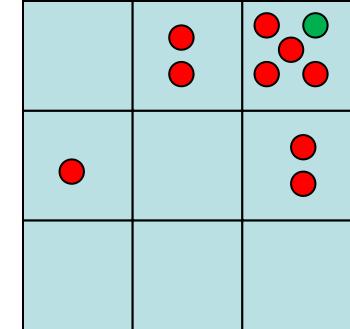
Particle Filtering

- Filtering: approximate solution
- Sometimes $|X|$ is too big to use exact inference
 - $|X|$ may be too big to even store $B(X)$
 - E.g. X is continuous
- Solution: approximate inference
 - Track samples of X , not all values
 - Samples are called particles
 - Time per step is linear in the number of samples
 - But: number needed may be large
 - In memory: list of particles, not states
- This is how robot localization works in practice
- Particle is just new name for sample



Representation: Particles

- Our representation of $P(X)$ is now a list of N particles (samples)
 - Generally, $N \ll |X|$
 - Storing map from X to counts would defeat the point
- $P(x)$ approximated by number of particles with value x
 - So, many x may have $P(x) = 0!$
 - More particles, more accuracy
- For now, all particles have a weight of 1



Particles:
(3,3)
(2,3)
(3,3)
(3,2)
(3,3)
(3,2)
(1,2)
(3,3)
(3,3)
(2,3)

Particle Filtering: Elapse Time

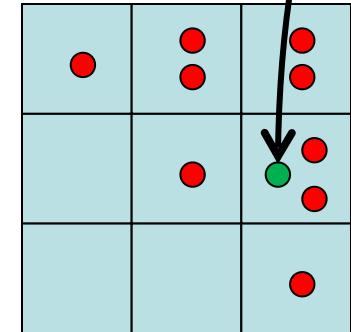
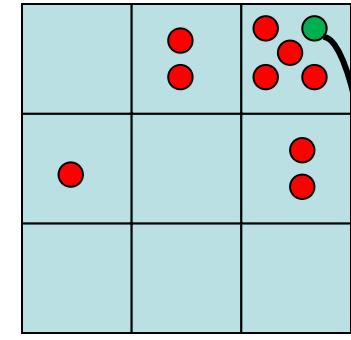
- Each particle is moved by sampling its next position from the transition model

$$x' = \text{sample}(P(X'|x))$$

- This is like prior sampling - samples' frequencies reflect the transition probabilities
- Here, most samples move clockwise, but some move in another direction or stay in place
- This captures the passage of time**
 - If enough samples, close to exact values before and after (consistent)

Particles:
(3,3)
(2,3)
(3,3)
(3,2)
(3,3)
(3,2)
(1,2)
(3,3)
(3,3)
(2,3)

Particles:
(3,2)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)



Particle Filtering: Observe

- Slightly trickier:

- Don't sample observation, fix it
- Similar to likelihood weighting, downweight samples based on the evidence

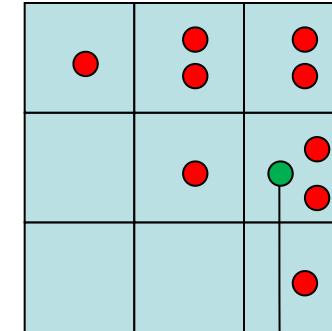
$$w(x) = P(e|x)$$

$$B(X) \propto P(e|X)B'(X)$$

- As before, the probabilities don't sum to one, since all have been downweighted (in fact they now sum to (N times) an approximation of $P(e)$)

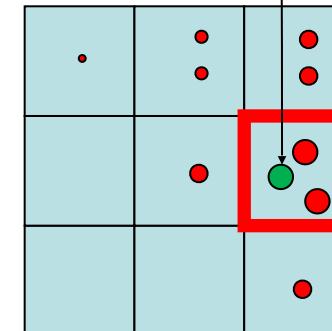
Particles:

(3,2)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)



Particles:

(3,2) w=.9
(2,3) w=.2
(3,2) w=.9
(3,1) w=.4
(3,3) w=.4
(3,2) w=.9
(1,3) w=.1
(2,3) w=.2
(3,2) w=.9
(2,2) w=.4

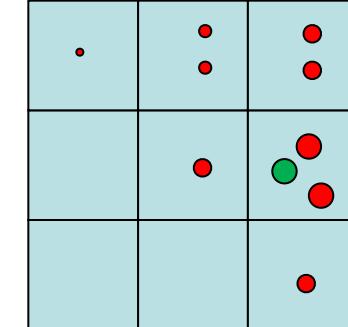


Particle Filtering: Resample

- Rather than tracking weighted samples, we resample
- N times, we choose from our weighted sample distribution (i.e. draw with replacement)
- This is equivalent to renormalizing the distribution
- Now the update is complete for this time step, continue with the next one

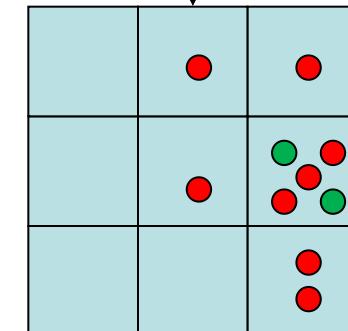
Particles:

(3,2) w=.9
(2,3) w=.2
(3,2) w=.9
(3,1) w=.4
(3,3) w=.4
(3,2) w=.9
(1,3) w=.1
(2,3) w=.2
(3,2) w=.9
(2,2) w=.4



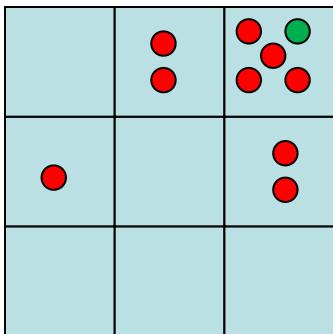
(New) Particles:

(3,2)
(2,2)
(3,2)
(2,3)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(3,2)



Recap: Particle Filtering

- Particles: track samples of states rather than an explicit distribution



Particles:

(3,3)

(2,3)

(3,3)

(3,2)

(3,3)

(3,2)

(1,2)

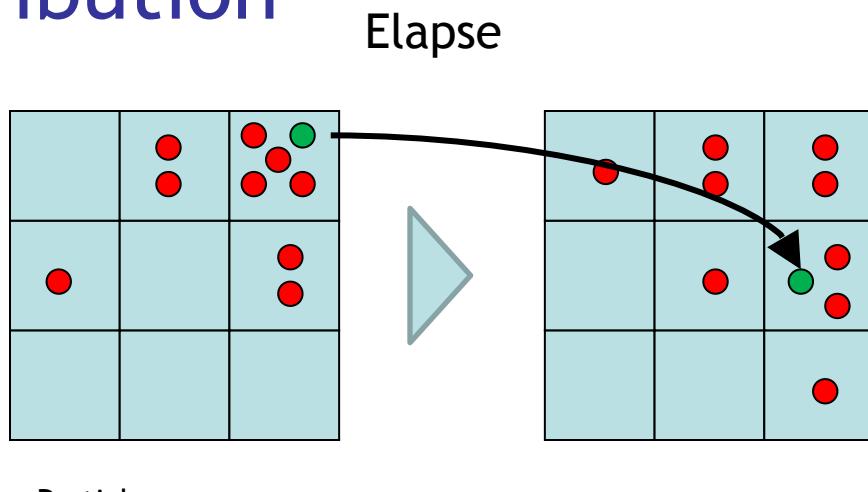
(3,3)

(3,3)

(2,3)

Recap: Particle Filtering

- Particles: track samples of states rather than an explicit distribution



Particles:

(3,3)

(2,3)

(3,3)

(3,2)

(3,3)

(3,2)

(1,2)

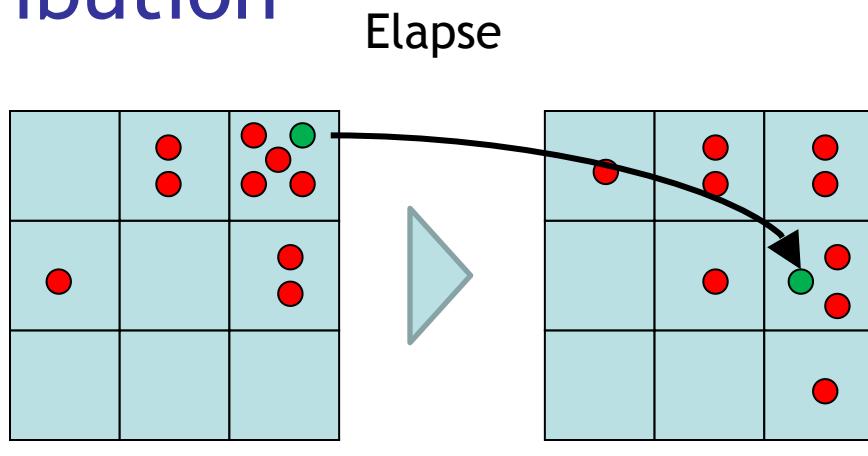
(3,3)

(3,3)

(2,3)

Recap: Particle Filtering

- Particles: track samples of states rather than an explicit distribution



Particles:

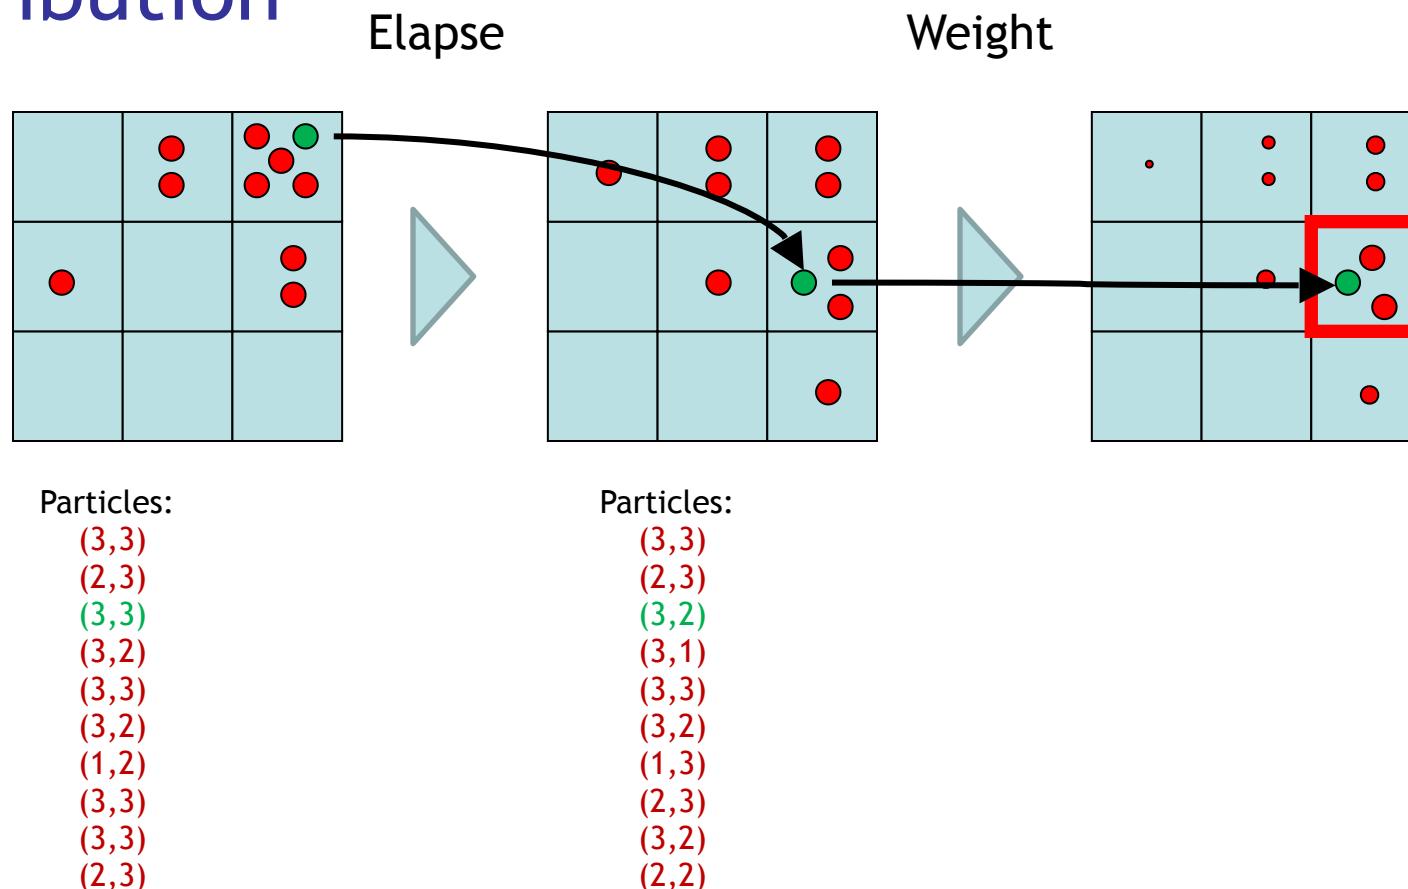
(3,3)
(2,3)
(3,3)
(3,2)
(3,3)
(3,2)
(1,2)
(3,3)
(3,3)
(2,3)

Particles:

(3,3)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)

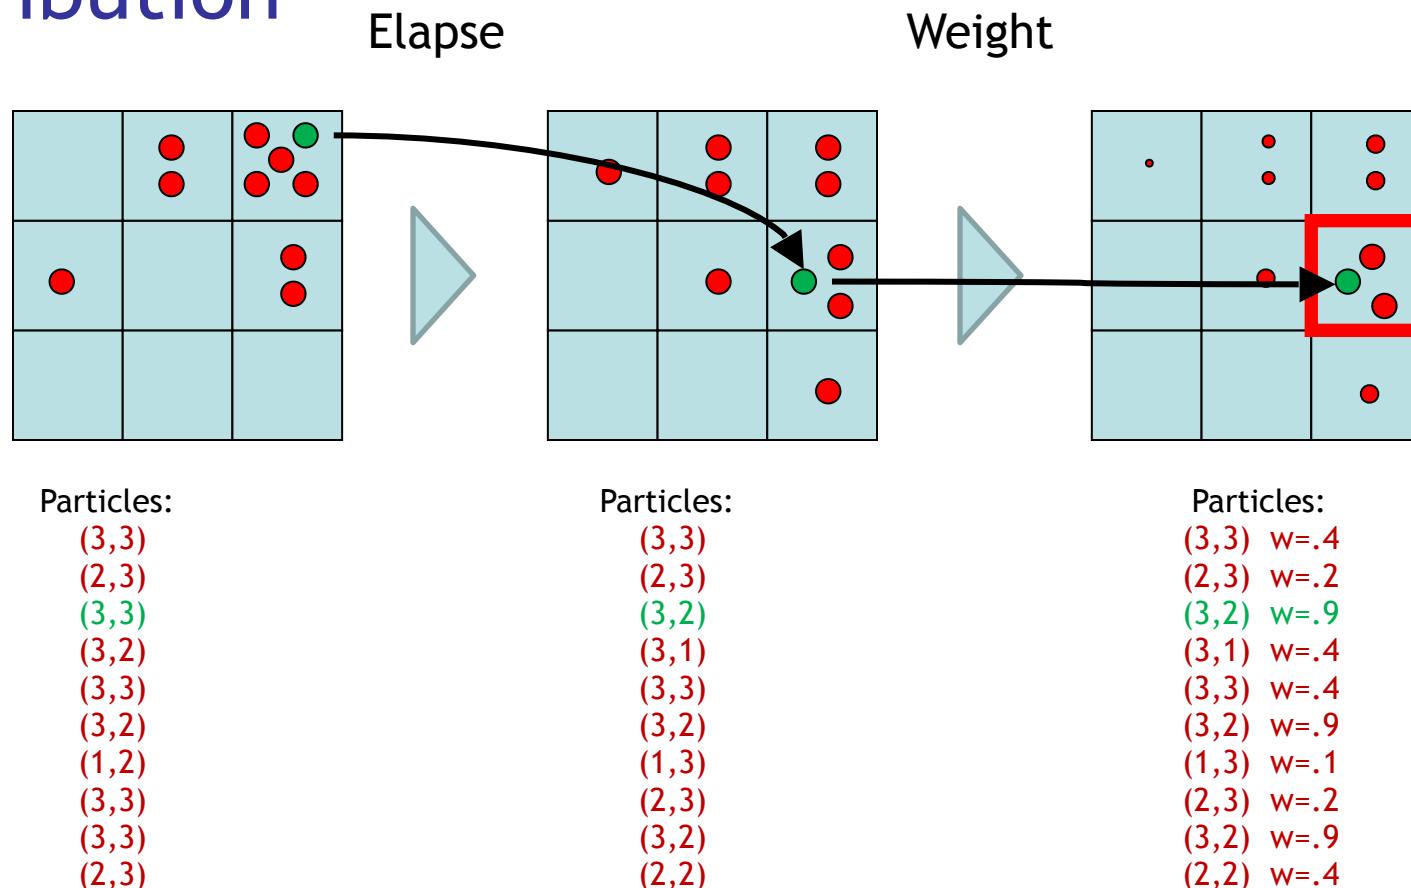
Recap: Particle Filtering

- Particles: track samples of states rather than an explicit distribution



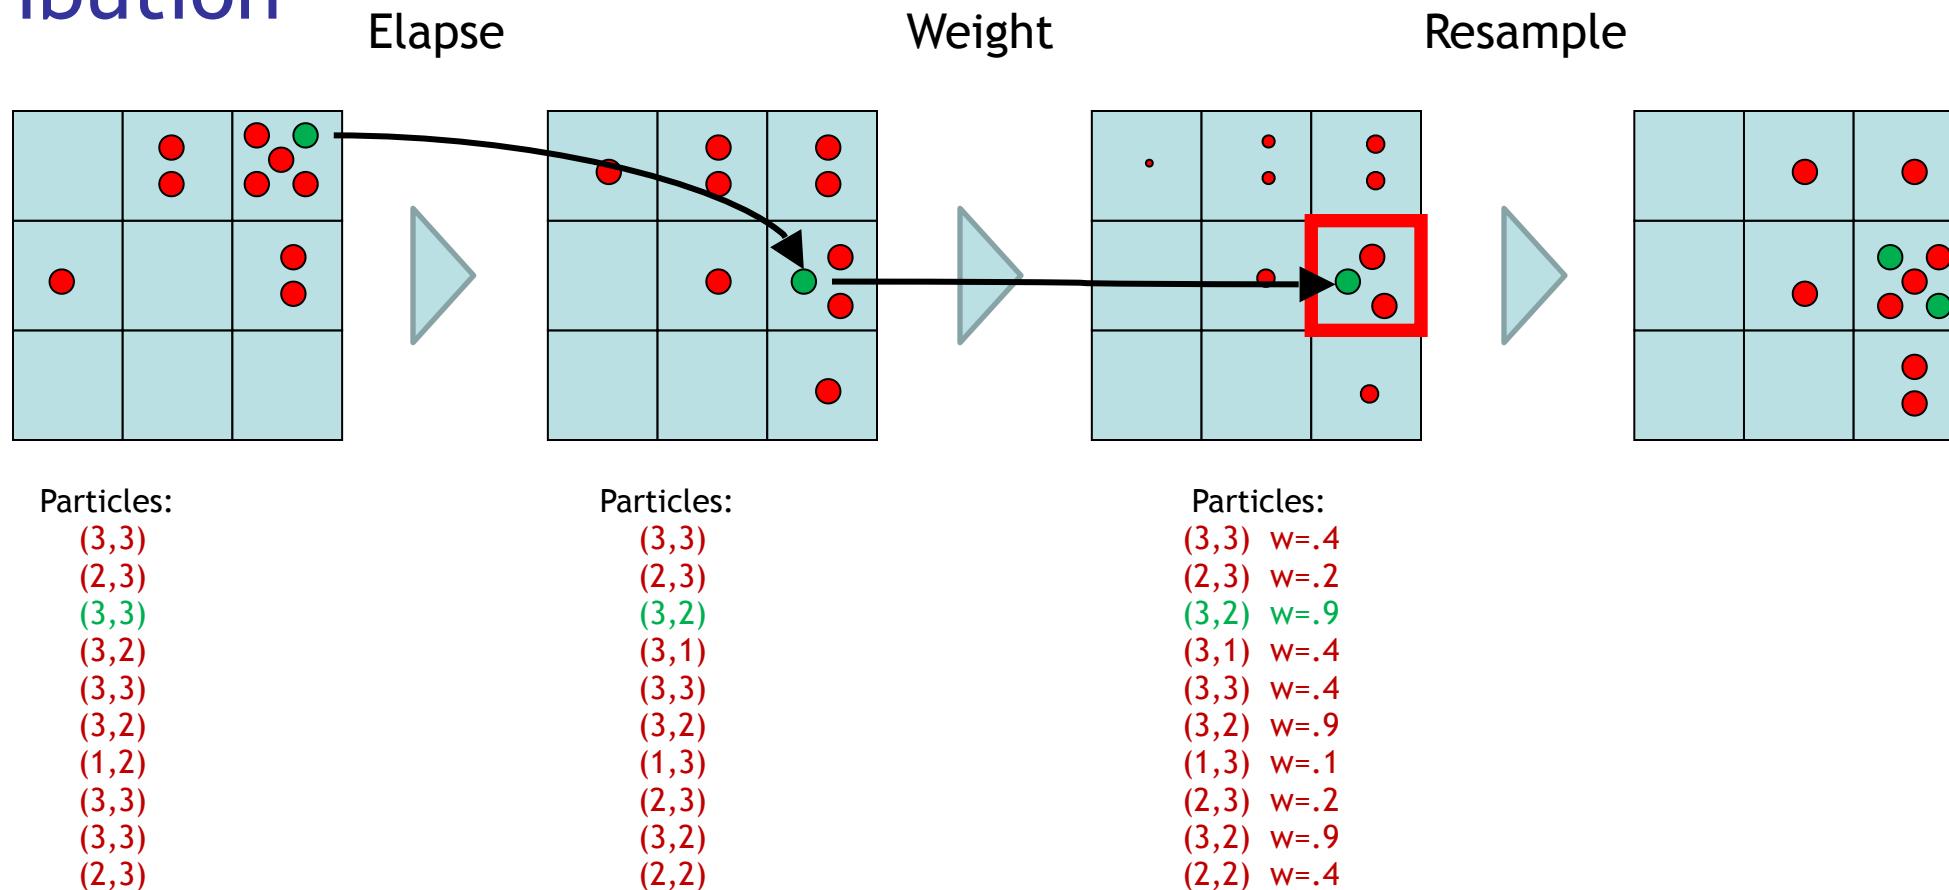
Recap: Particle Filtering

- Particles: track samples of states rather than an explicit distribution



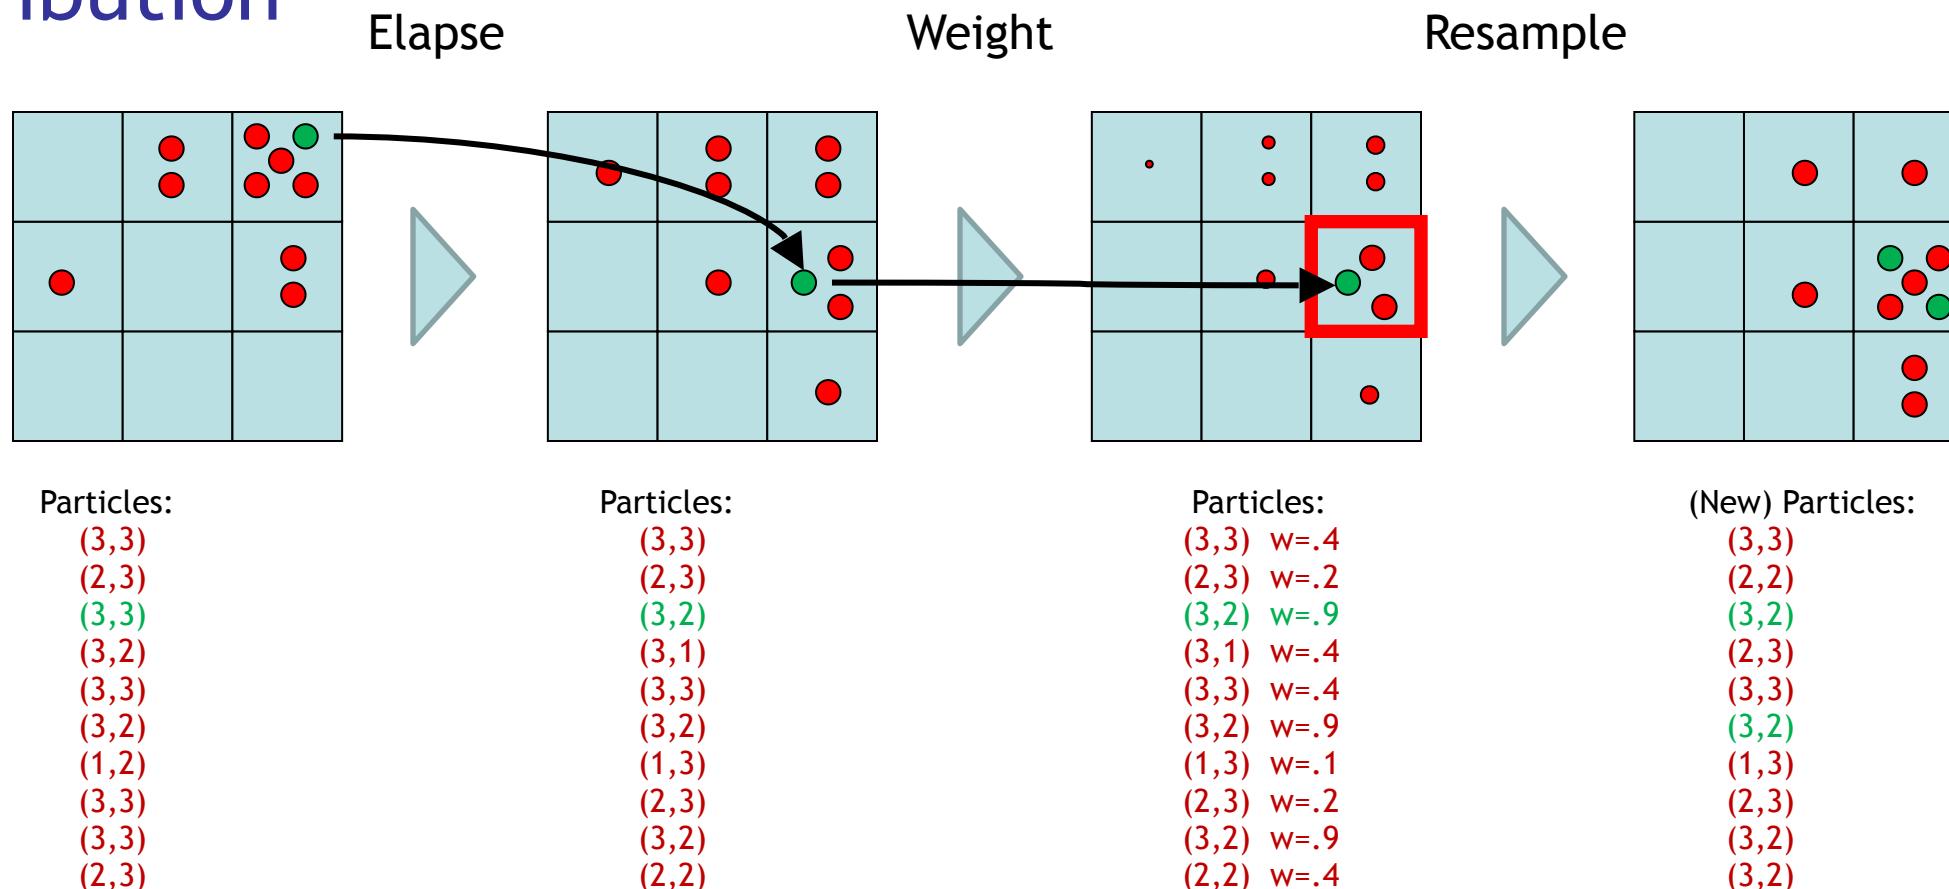
Recap: Particle Filtering

- Particles: track samples of states rather than an explicit distribution

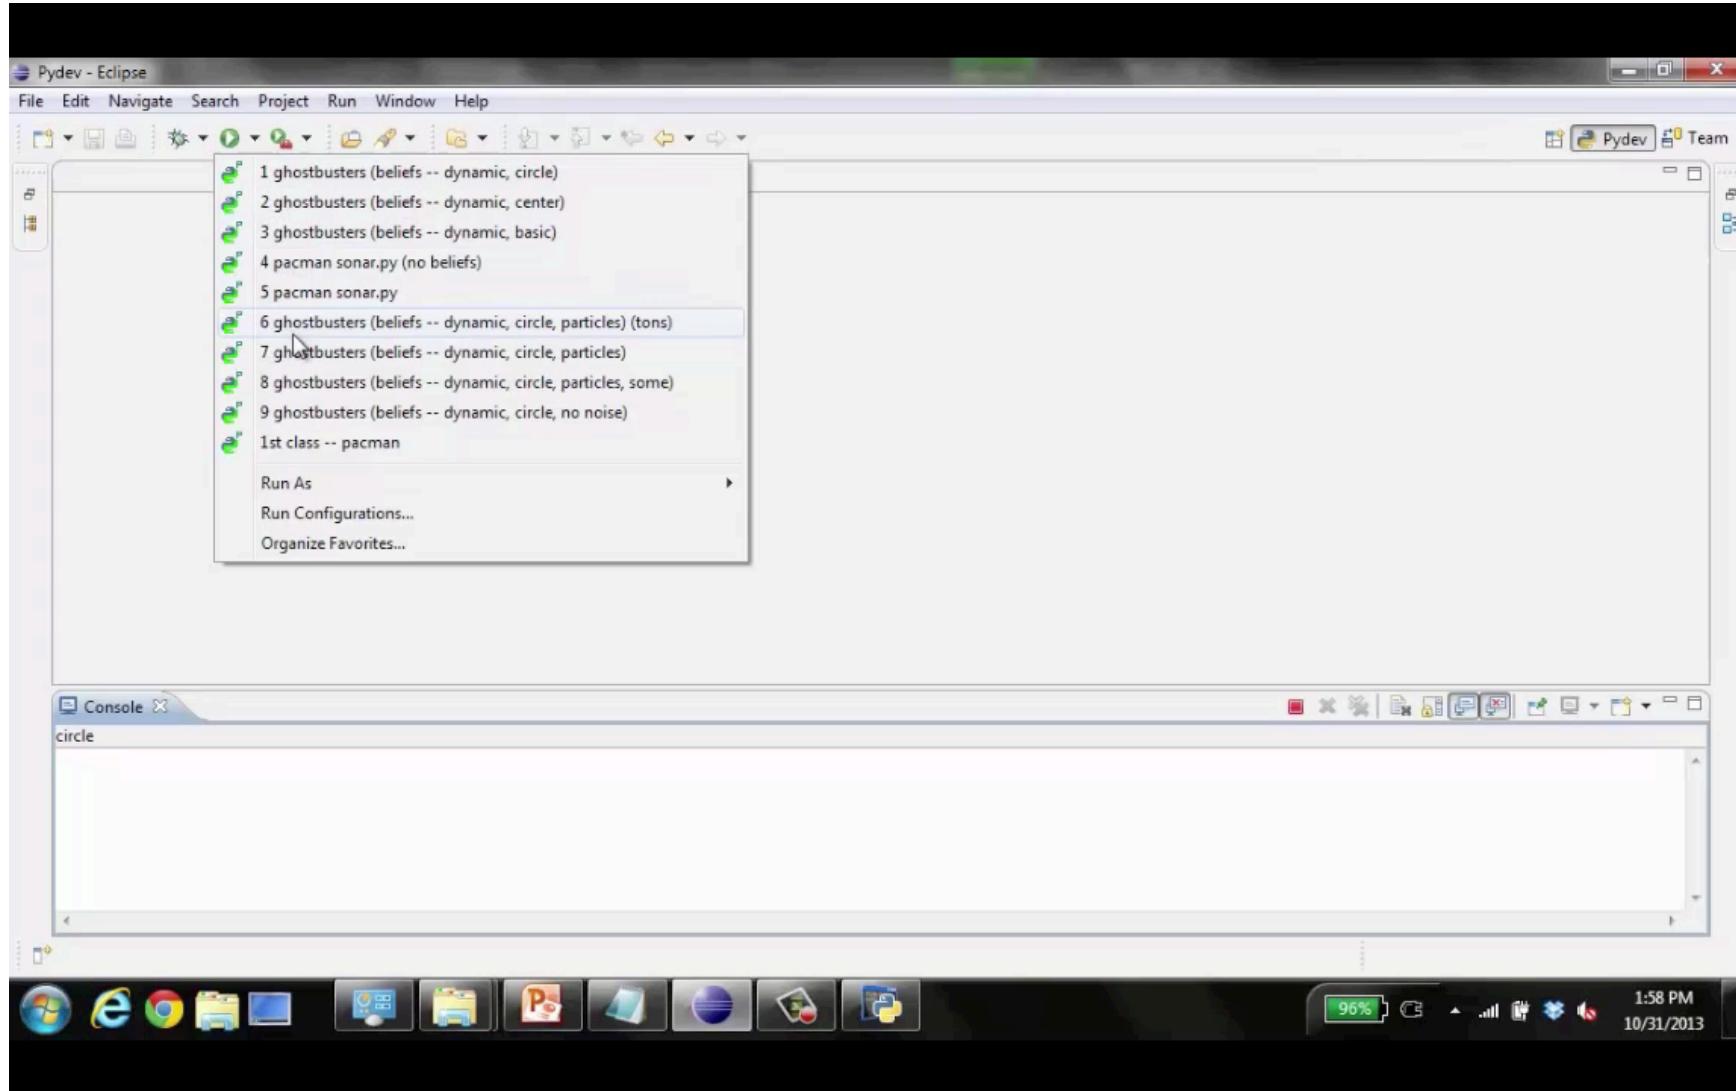


Recap: Particle Filtering

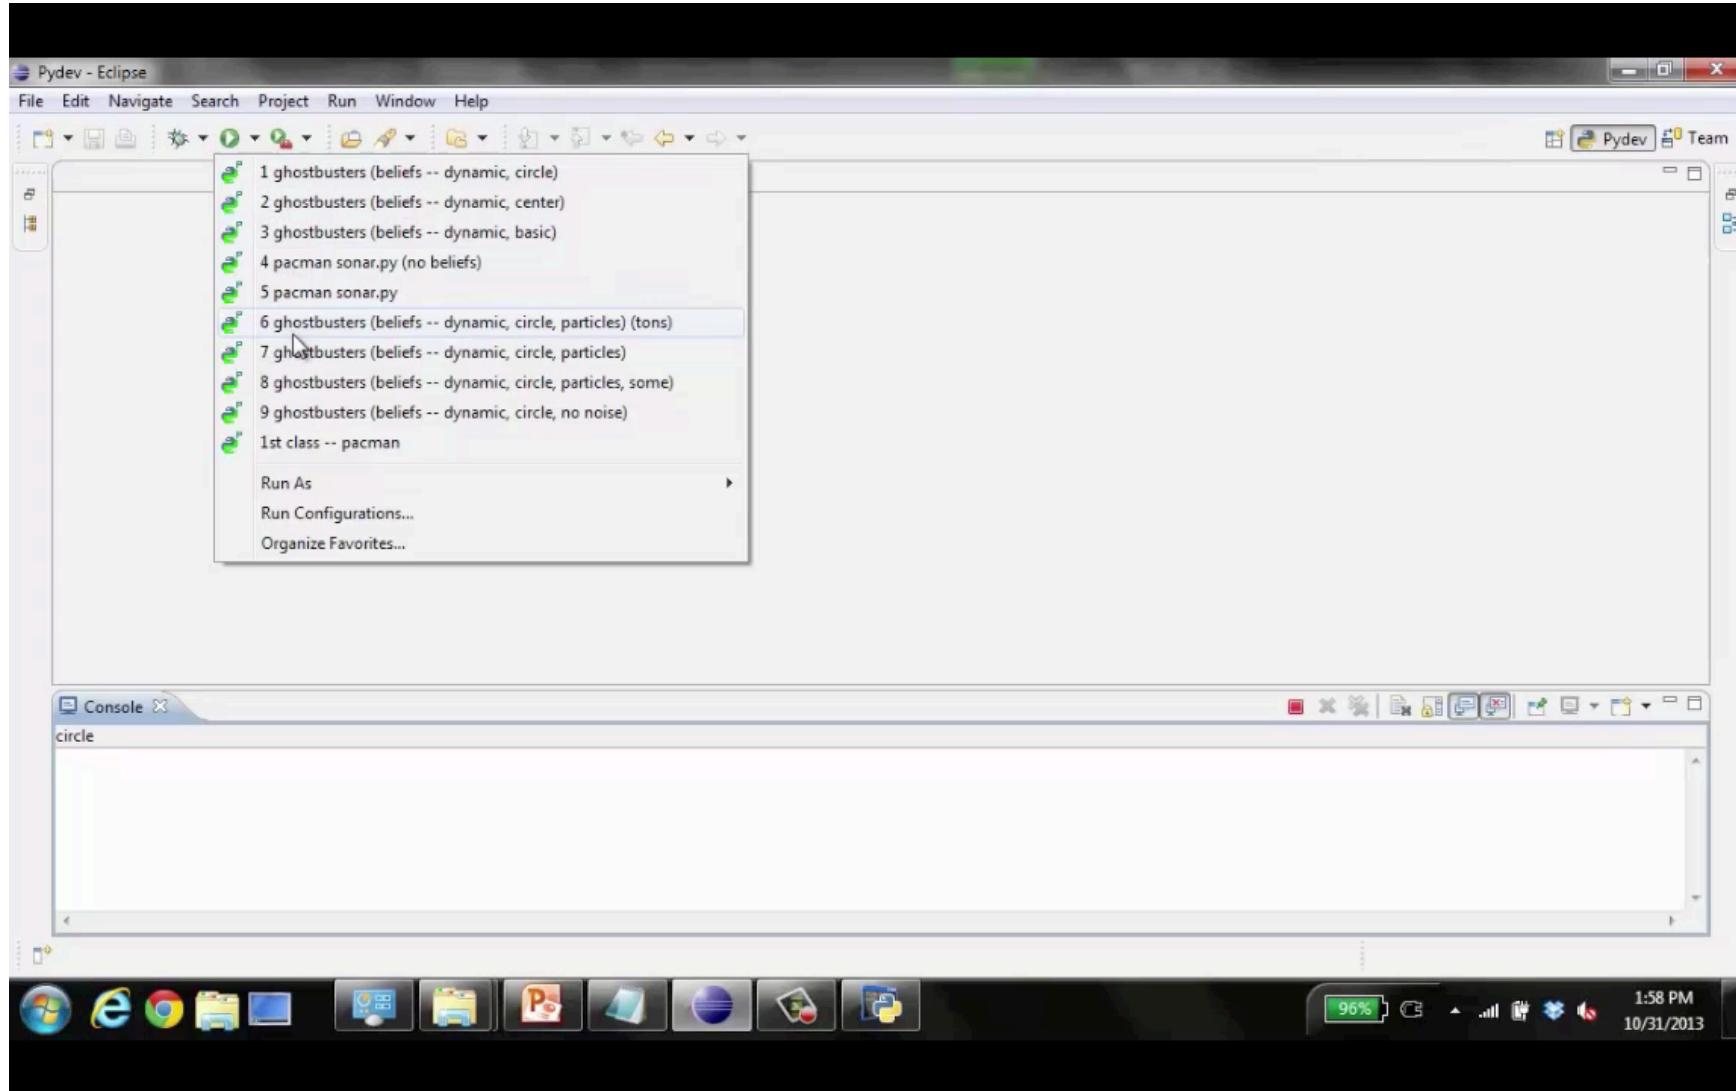
- Particles: track samples of states rather than an explicit distribution



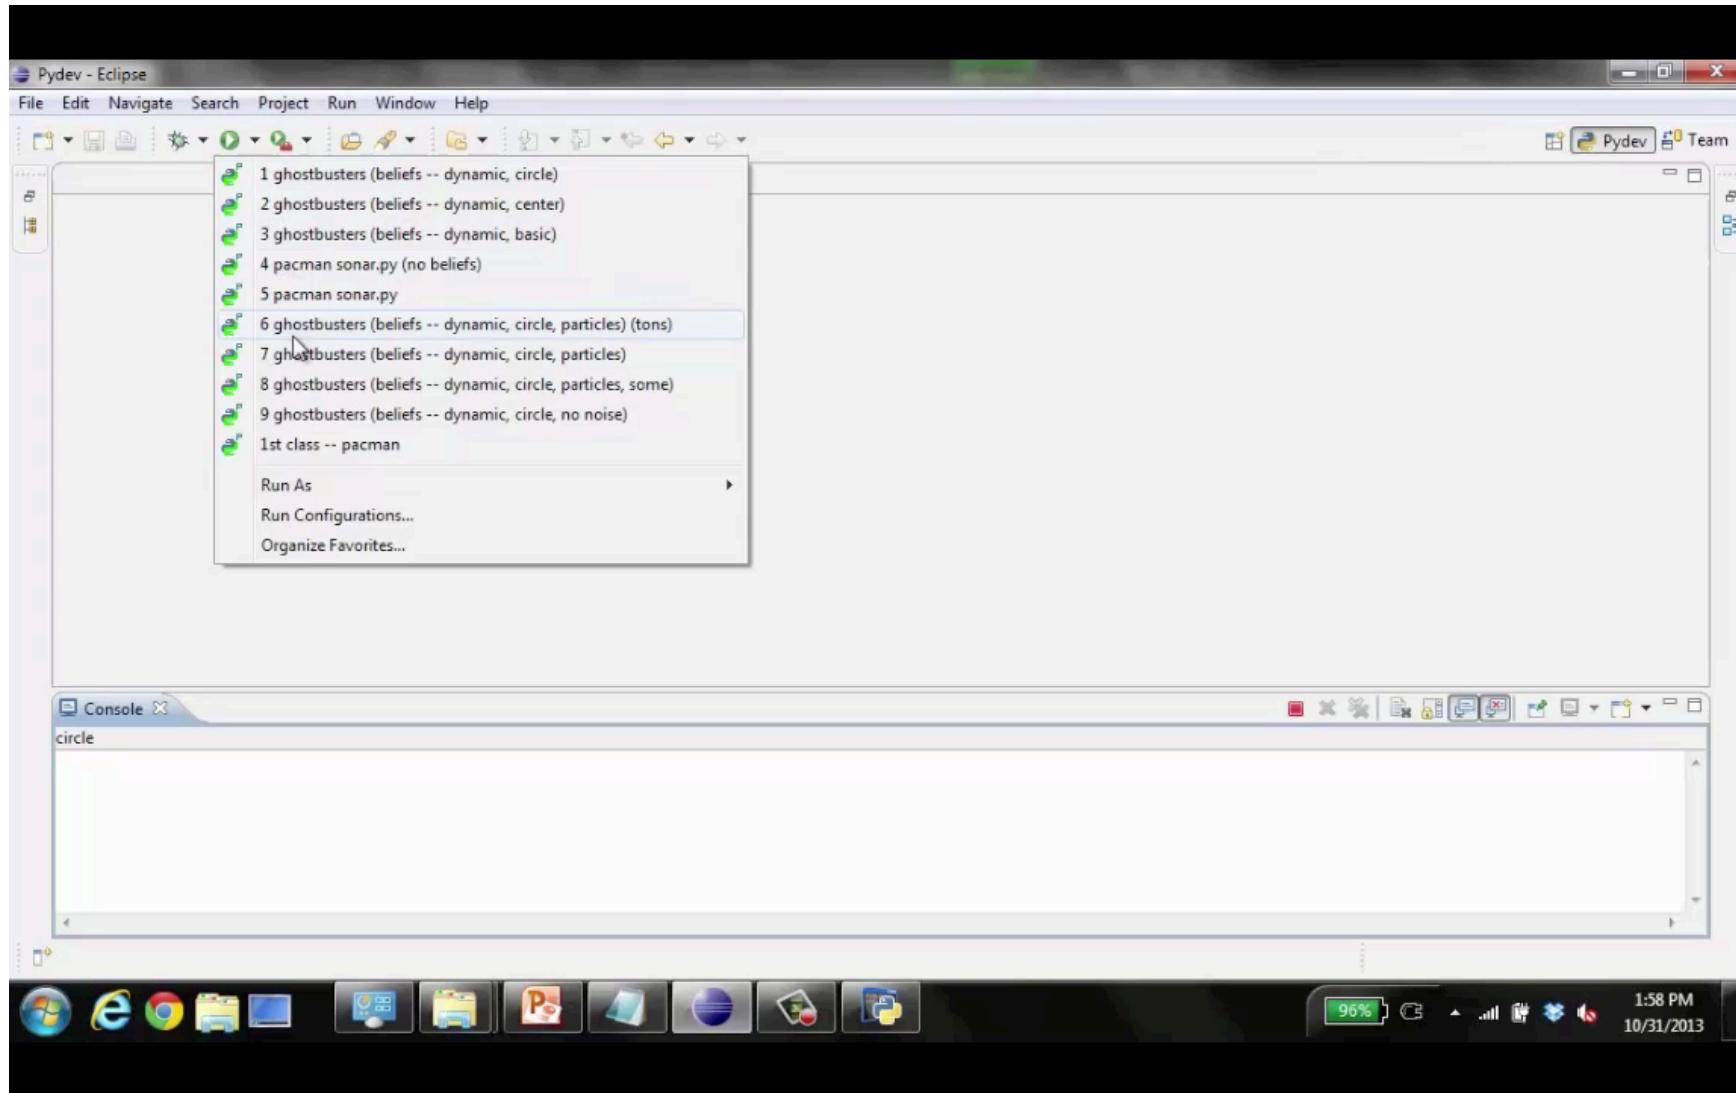
Video of Demo - Moderate Number of Particles



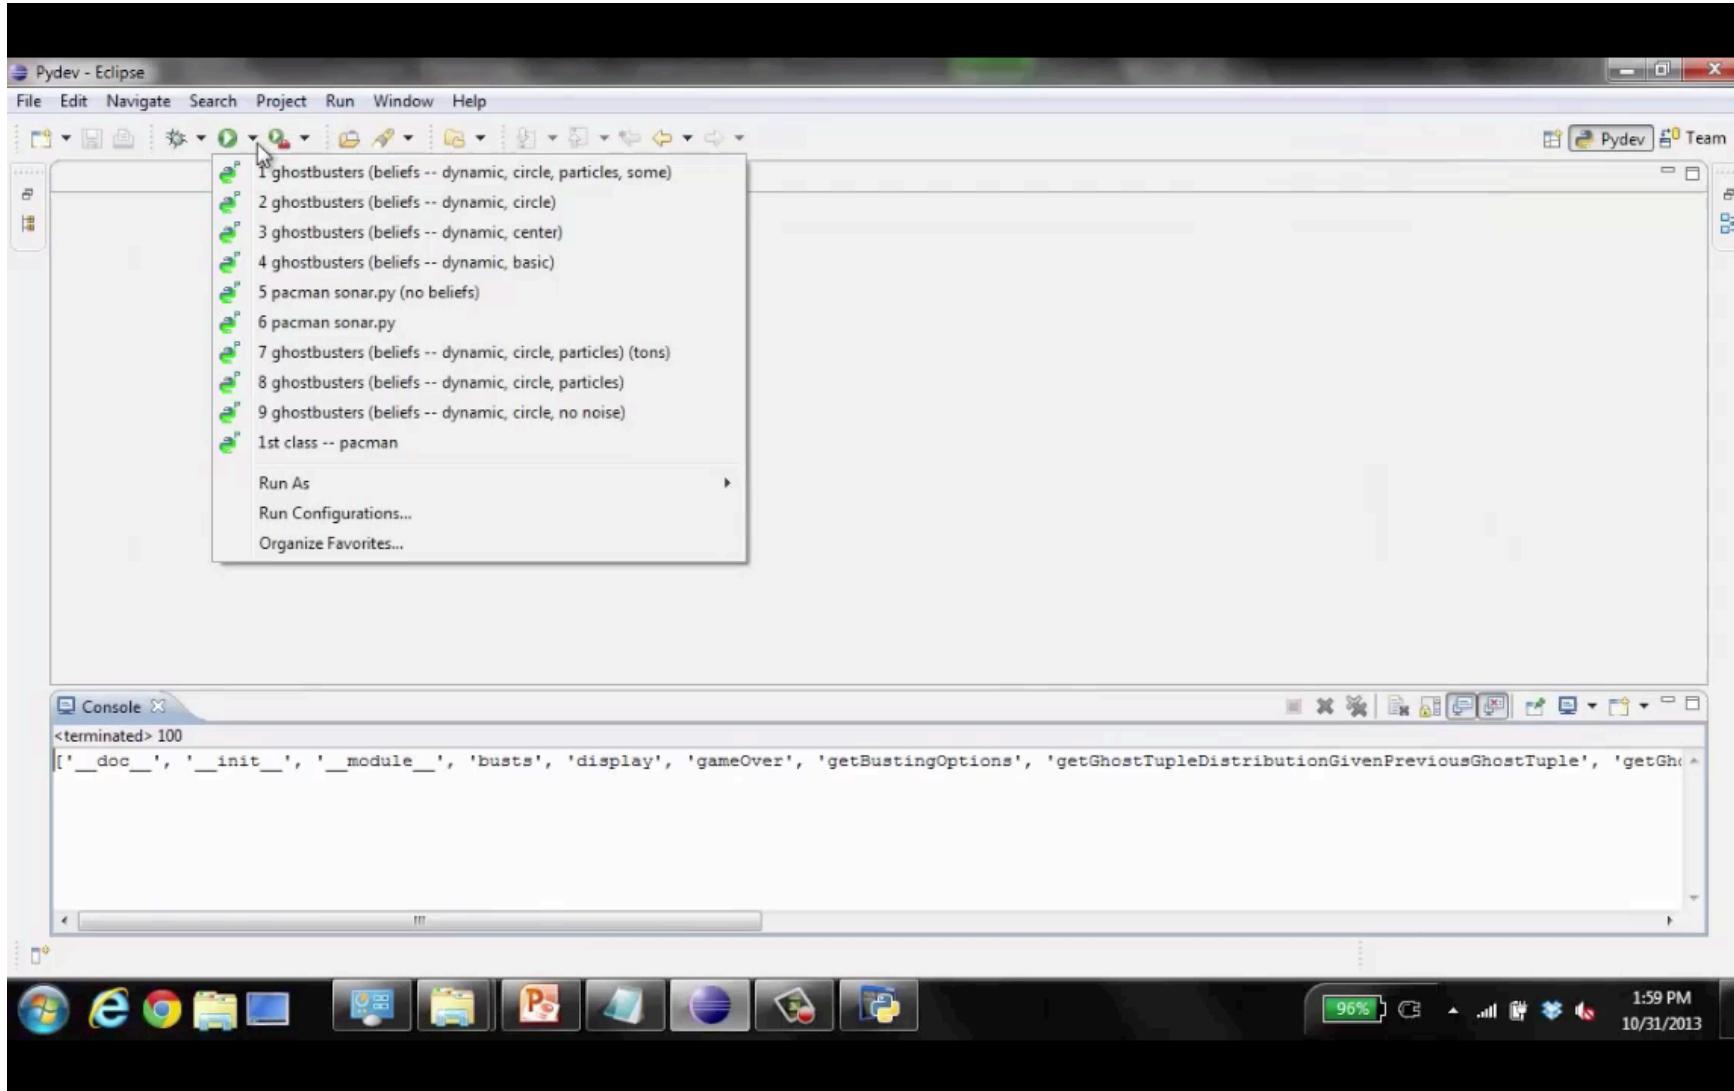
Video of Demo - Moderate Number of Particles



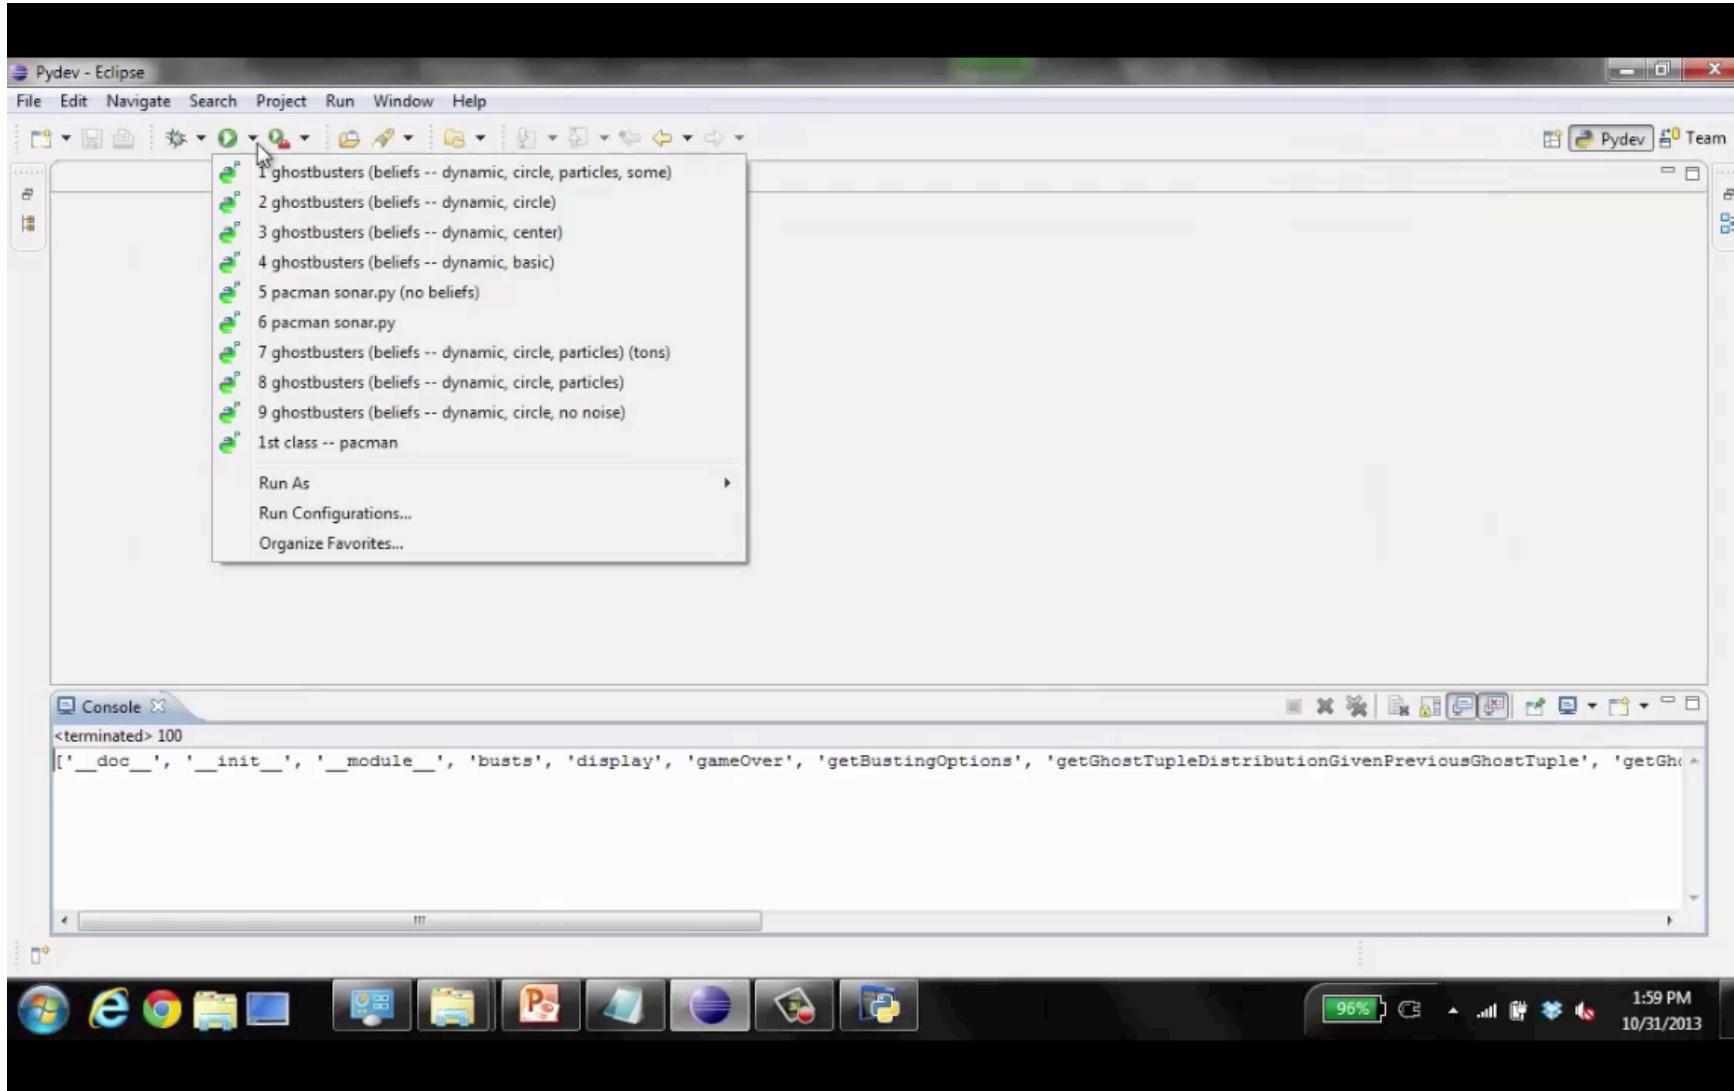
Video of Demo - Moderate Number of Particles



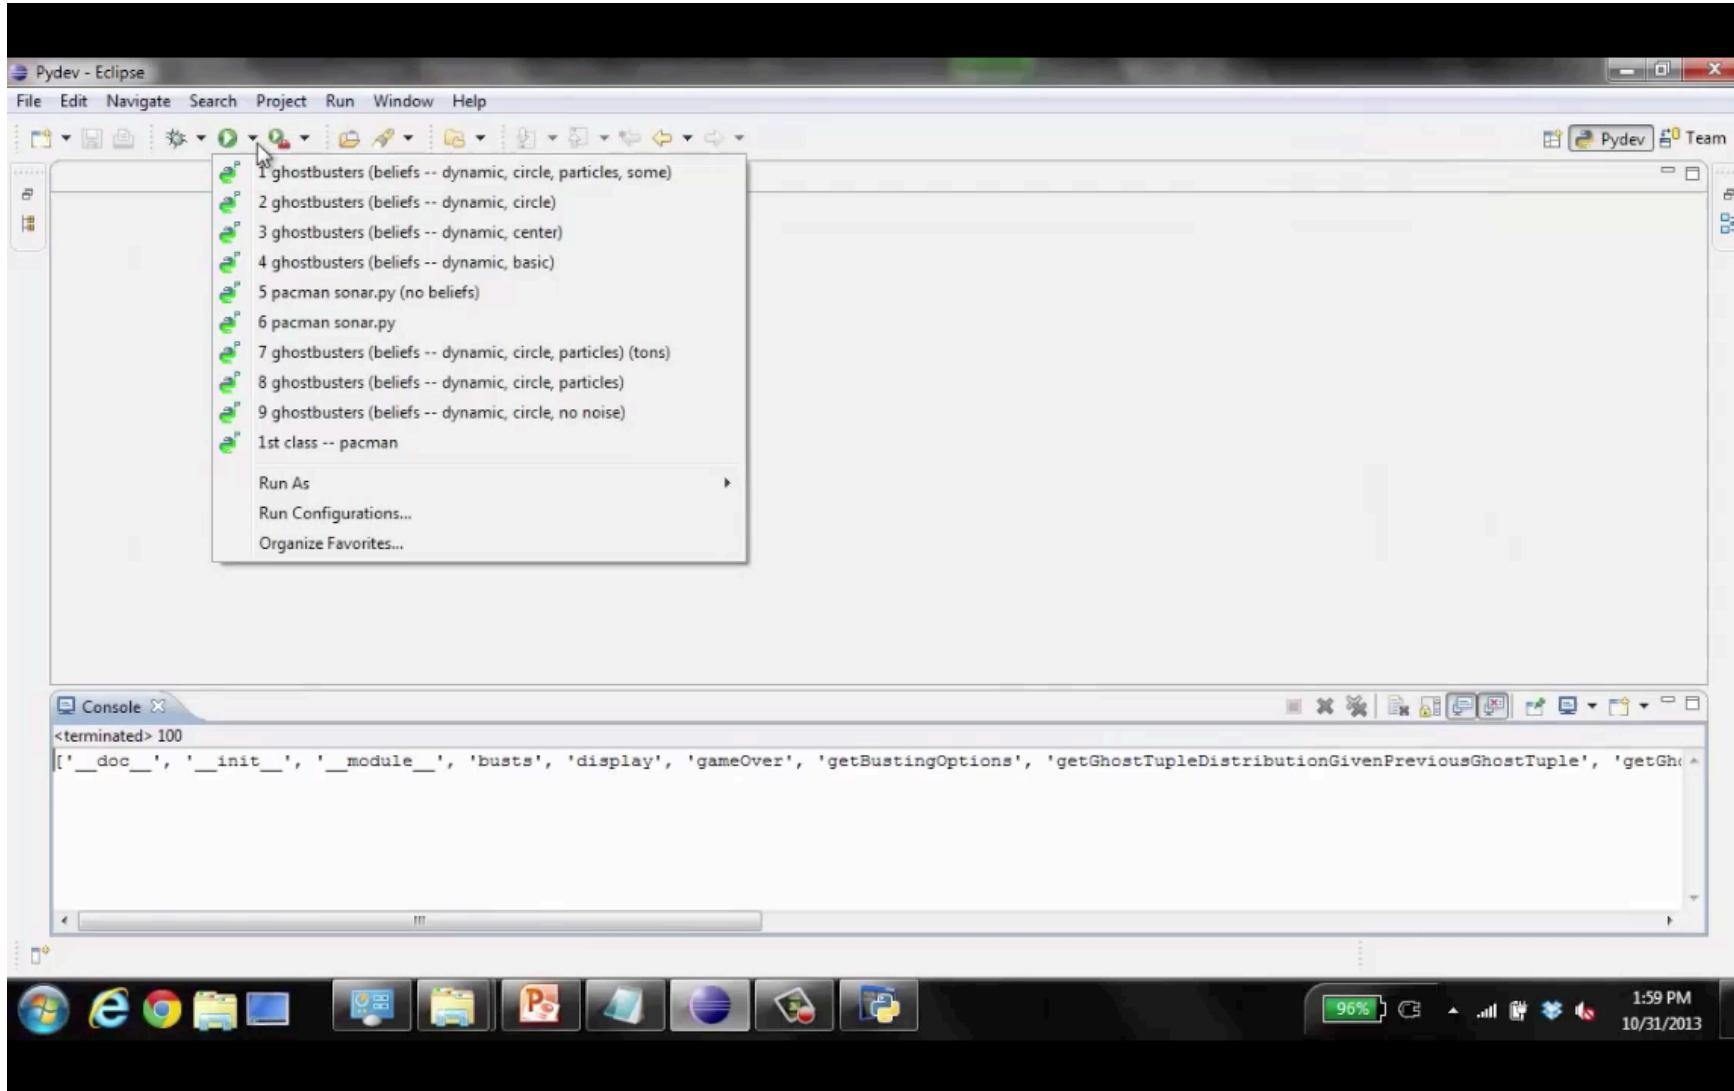
Video of Demo - One Particle



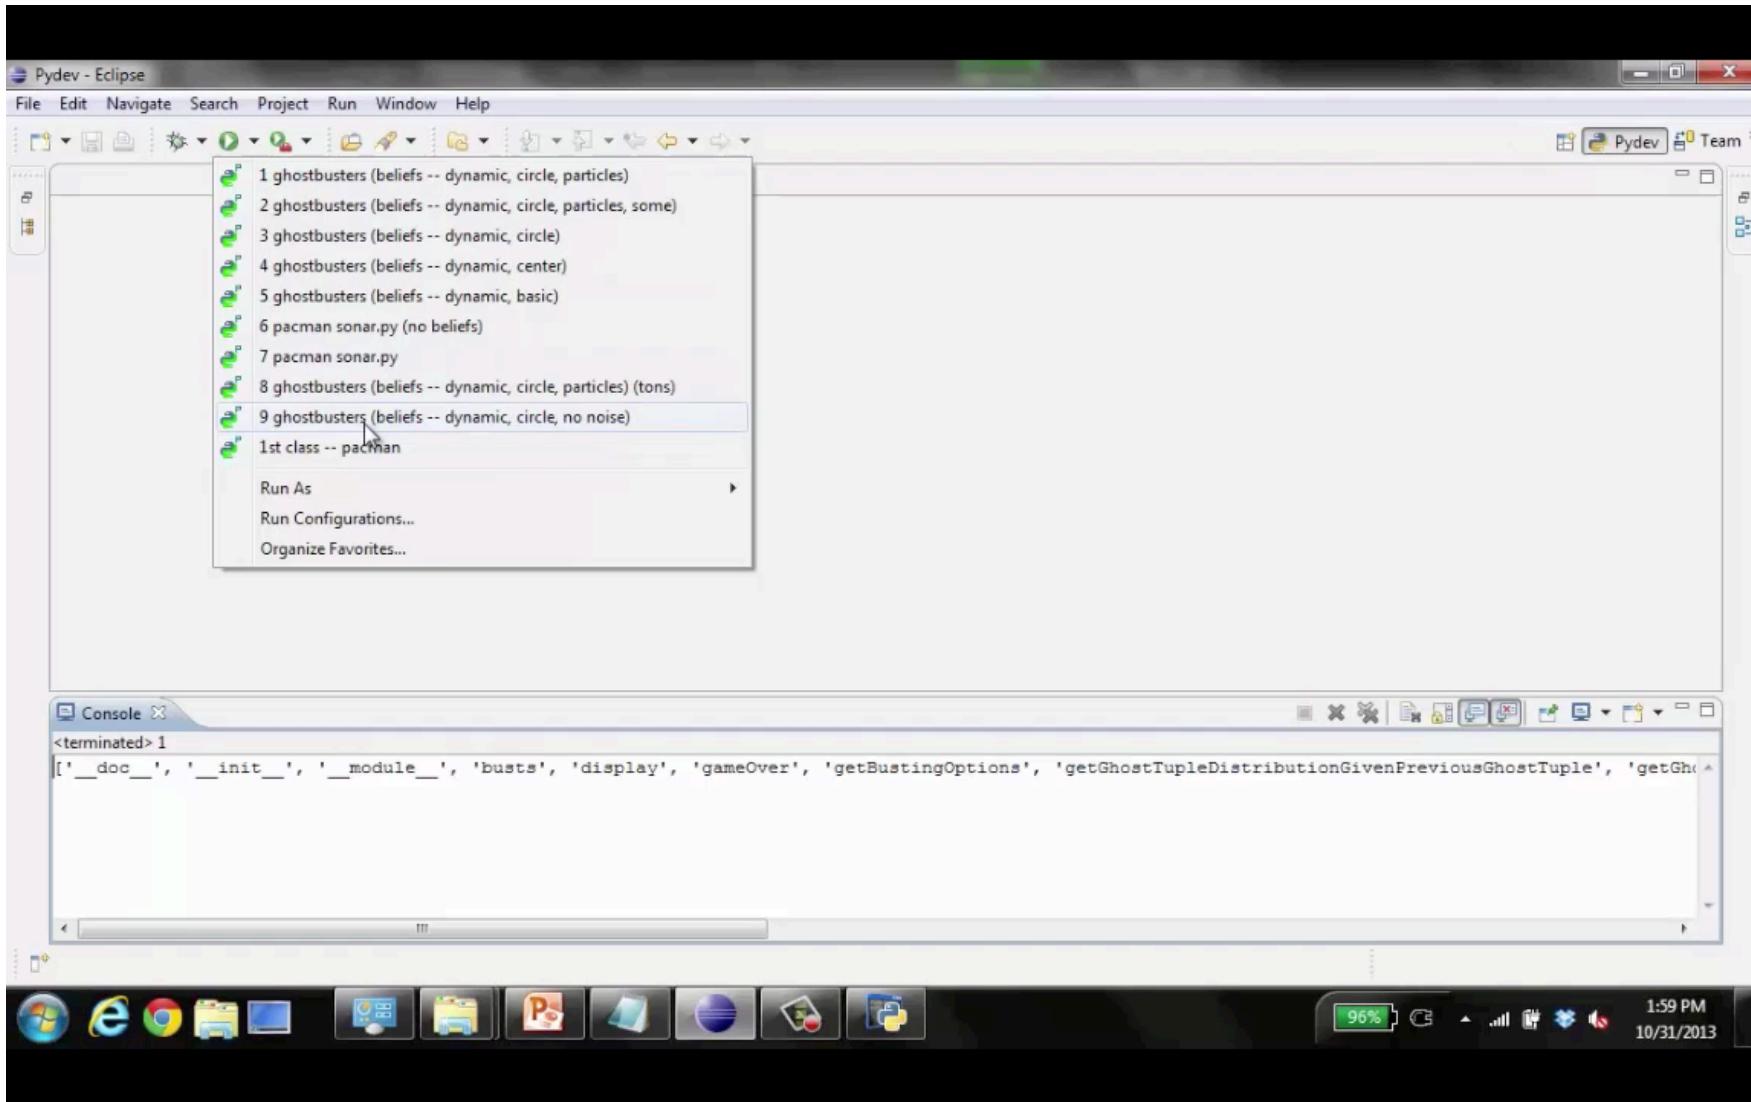
Video of Demo - One Particle



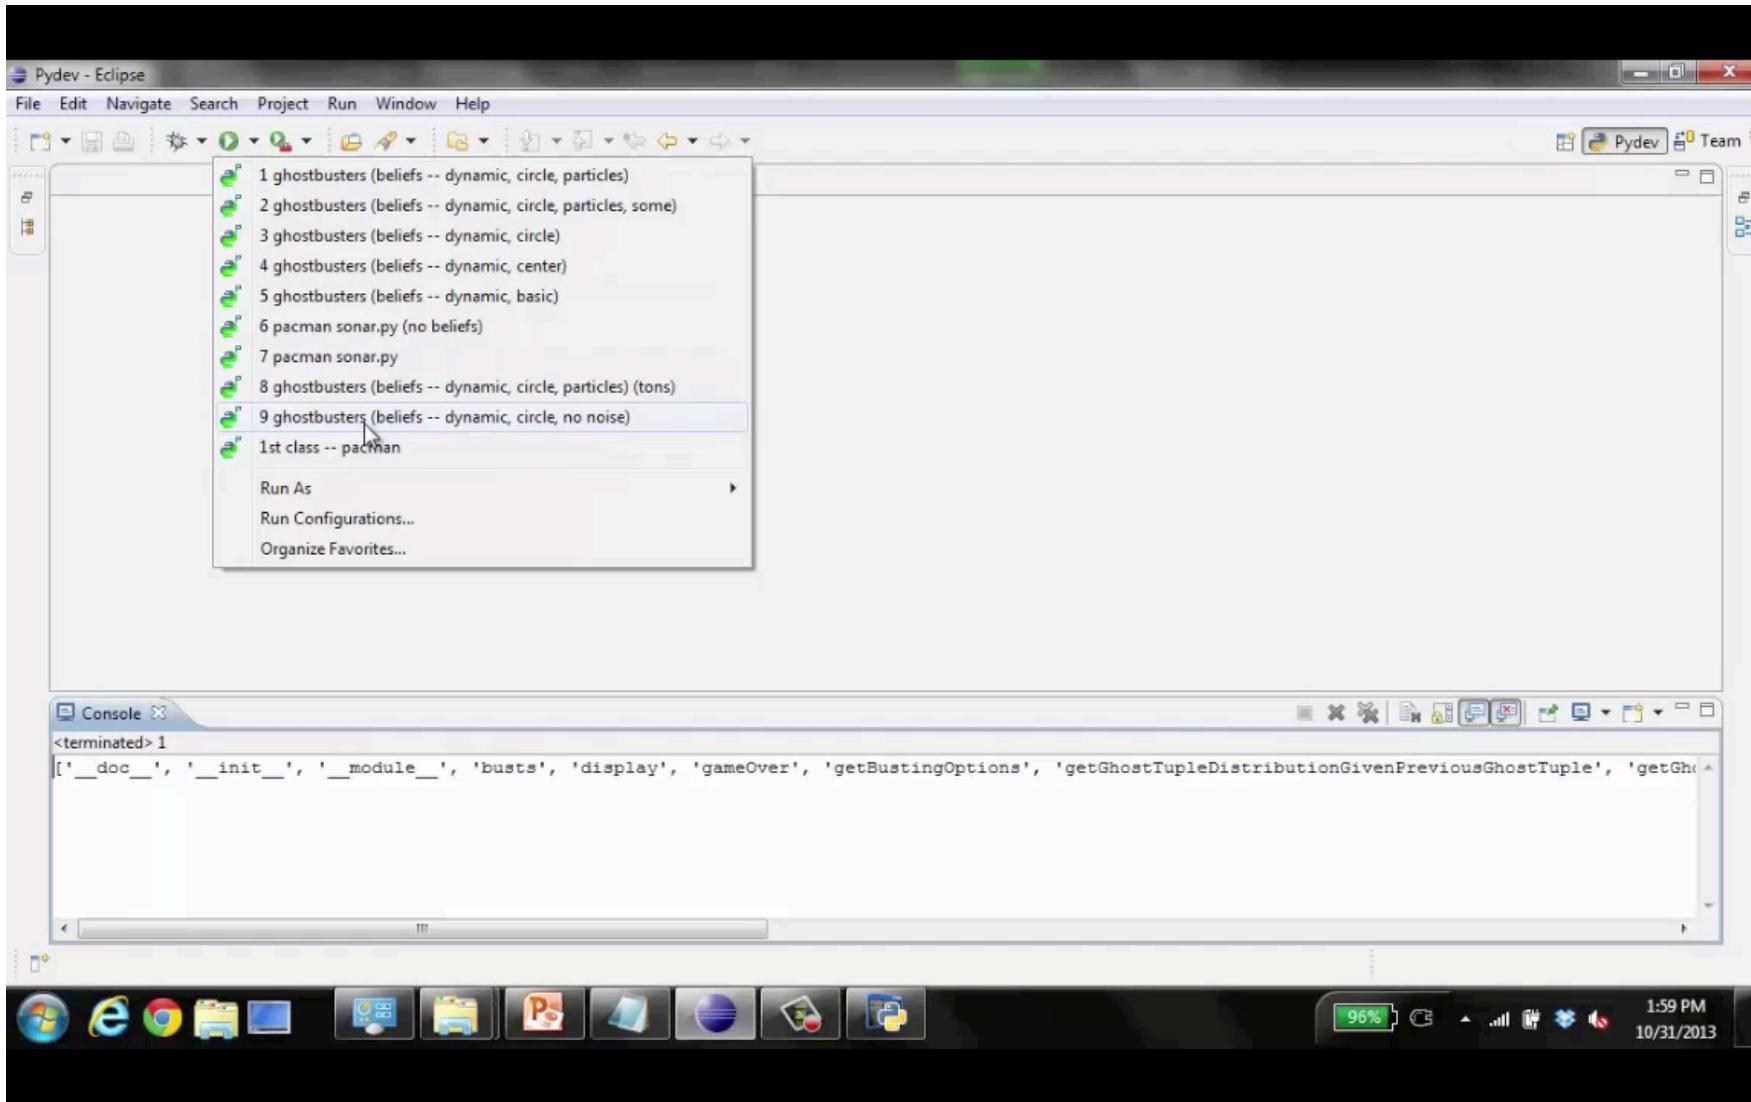
Video of Demo - One Particle



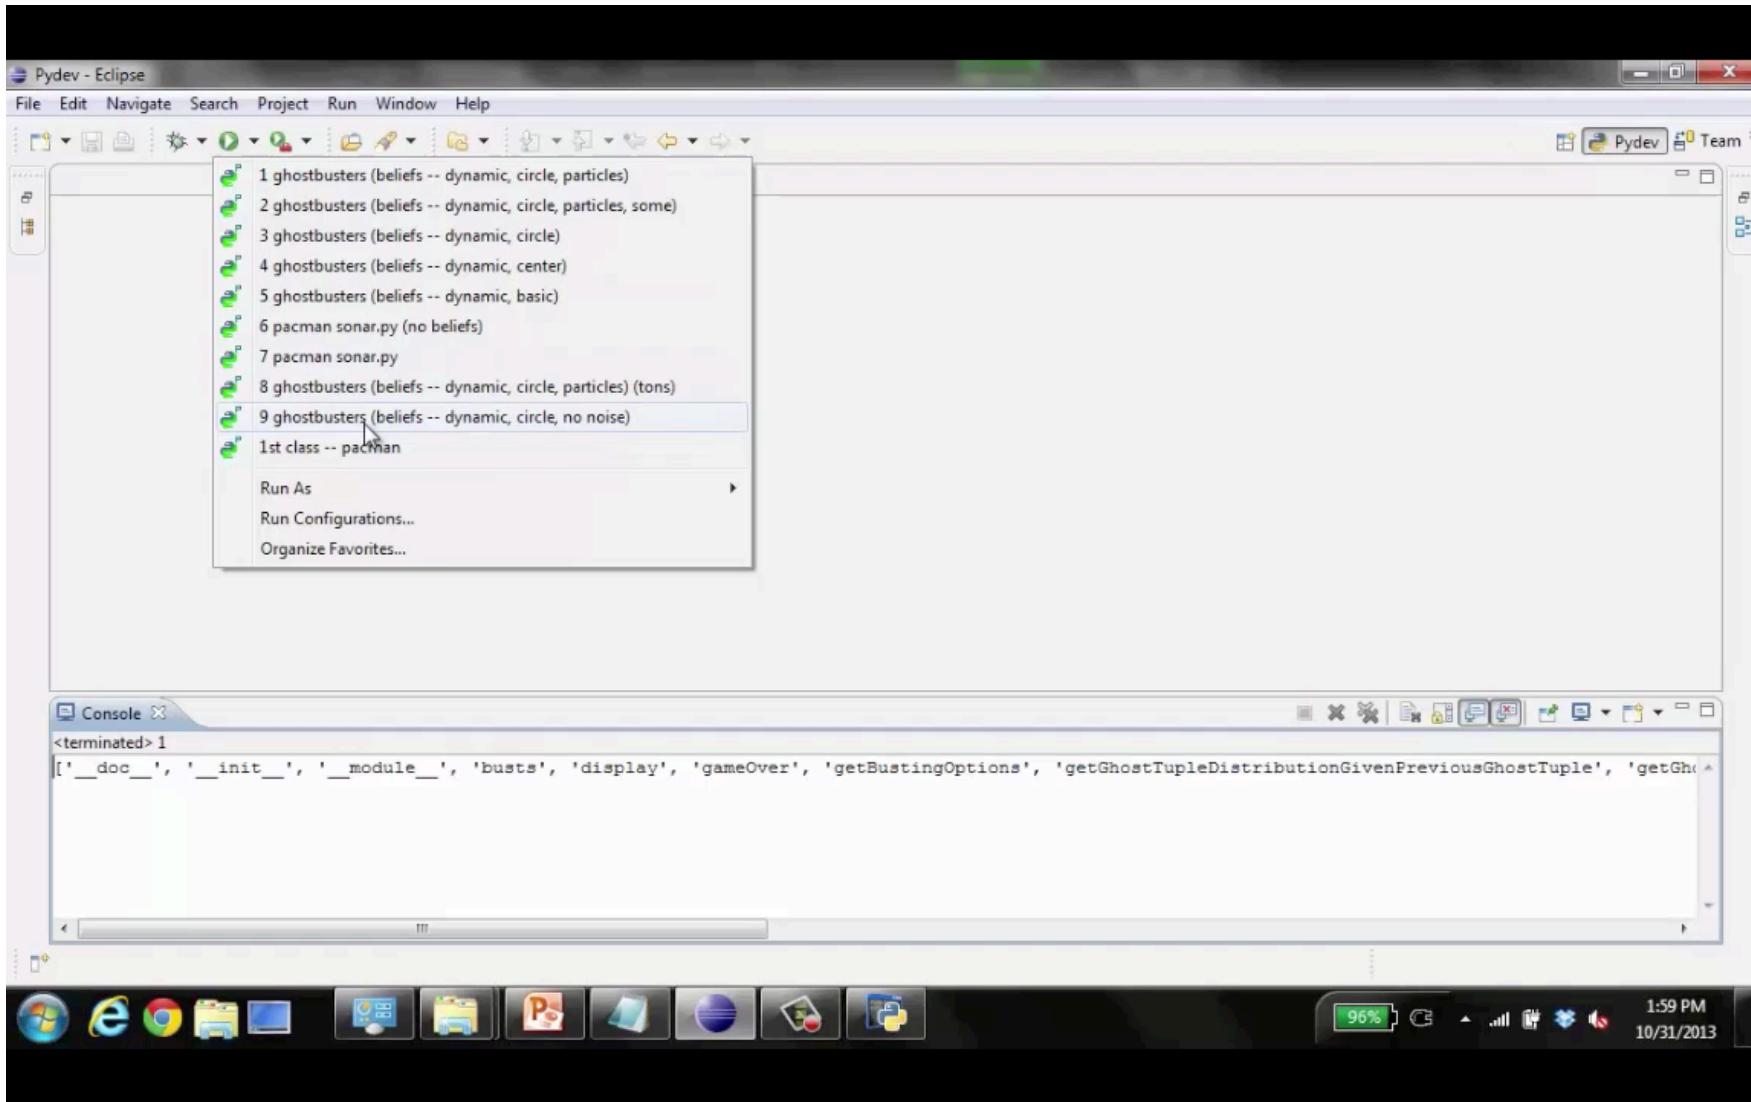
Video of Demo - Huge Number of Particles



Video of Demo - Huge Number of Particles



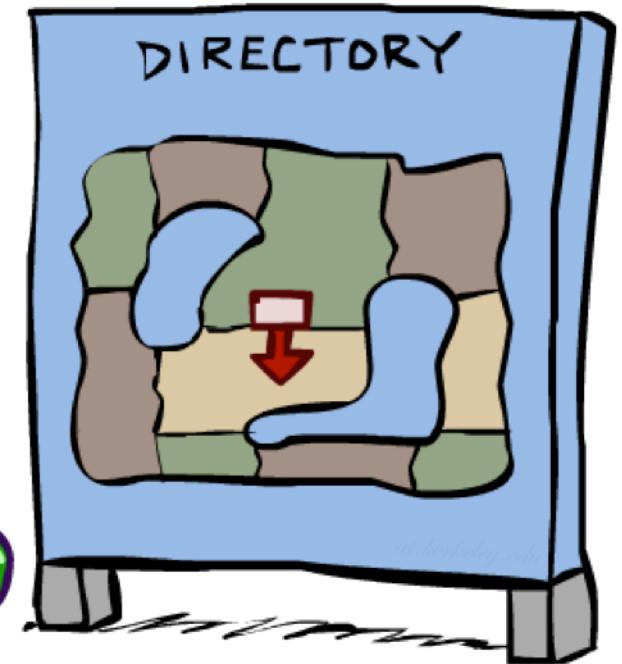
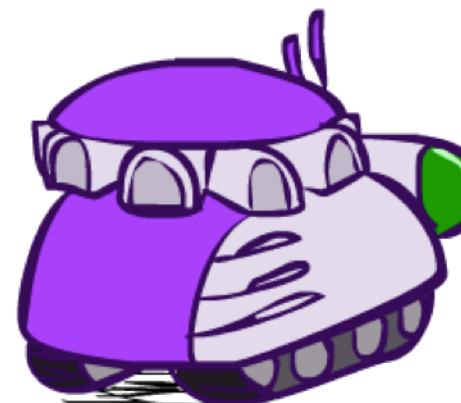
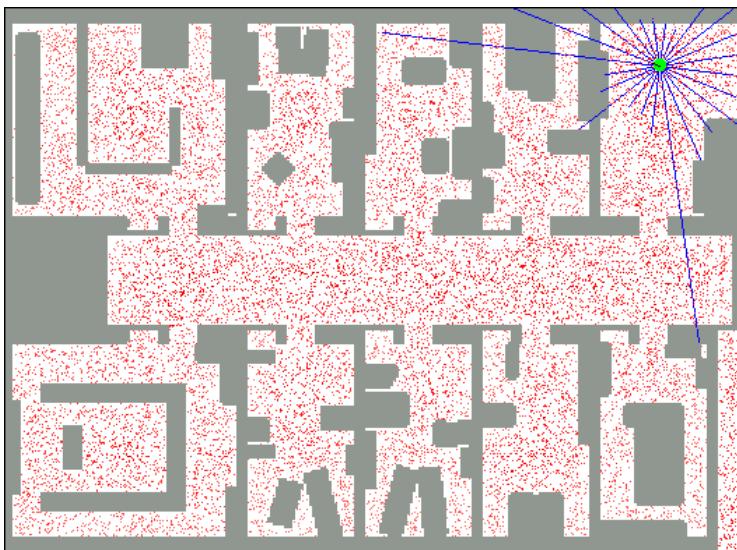
Video of Demo - Huge Number of Particles



Robot Localization

- In robot localization:

- We know the map, but not the robot's position
- Observations may be vectors of range finder readings
- State space and readings are typically continuous (works basically like a very fine grid) and so we cannot store $B(X)$
- Particle filtering is a main technique



Particle Filter Localization (Sonar)



The image shows a 3D point cloud representation of an environment. The floor and walls are grey, while obstacles are represented by red points. A small blue dot with a green crosshair indicates the estimated position of the robot. In the bottom-left corner, there is a blue rectangular box containing the number "40000".

**Global localization with
sonar sensors**

[Video: global-sonar-uw-annotated.avi]

Particle Filter Localization (Sonar)



Global localization with
sonar sensors

40000

[Video: global-sonar-uw-annotated.avi]

Particle Filter Localization (Sonar)

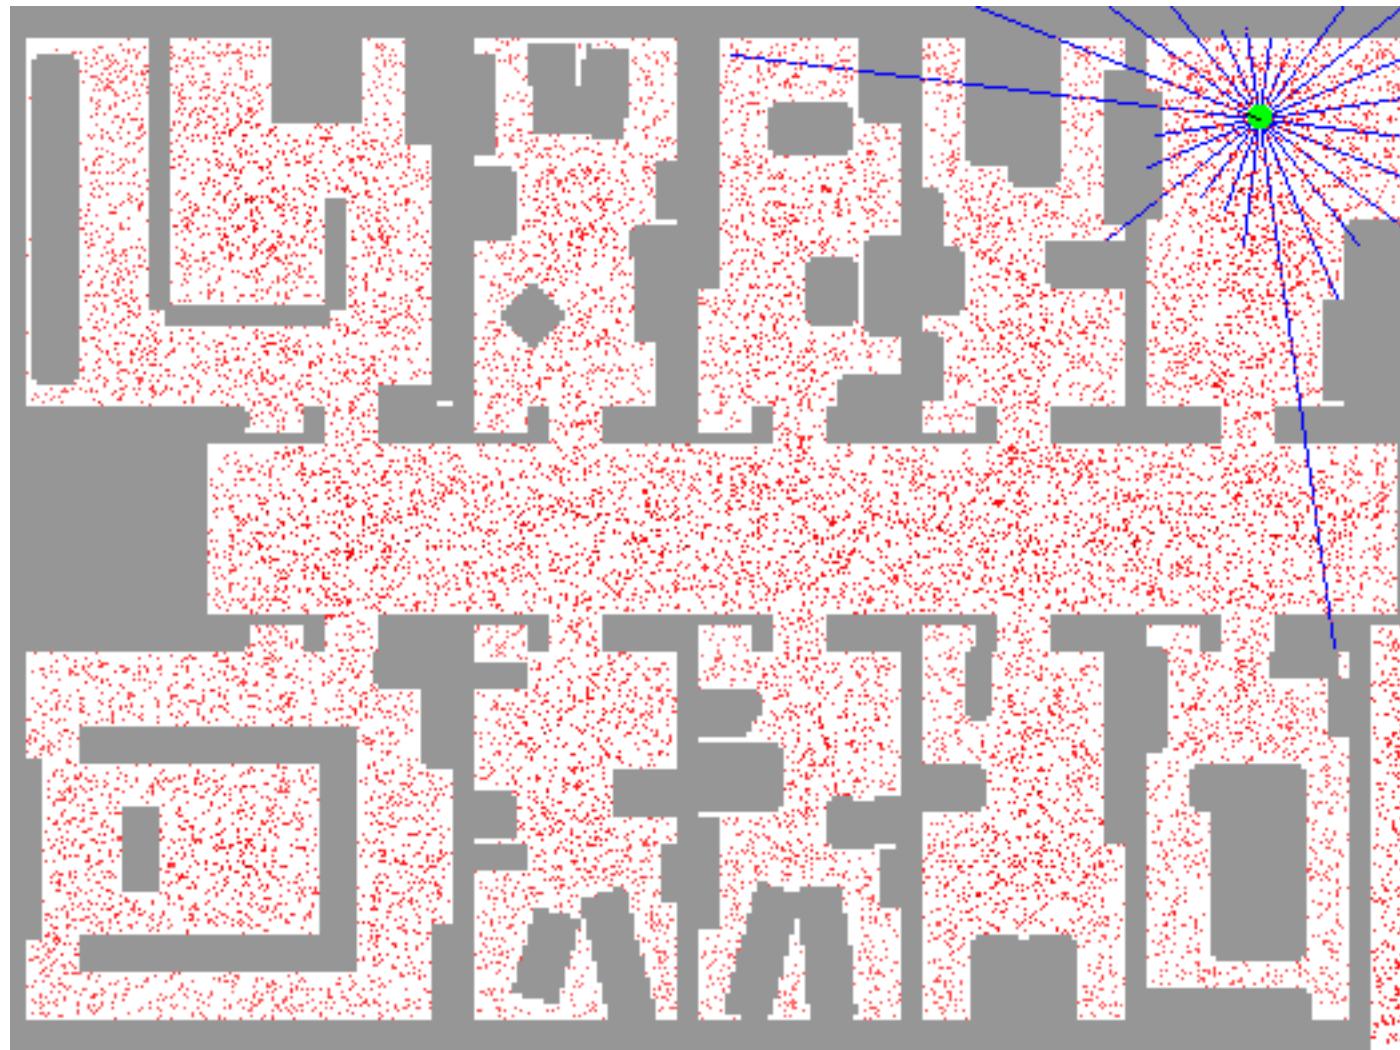


Global localization with
sonar sensors

40000

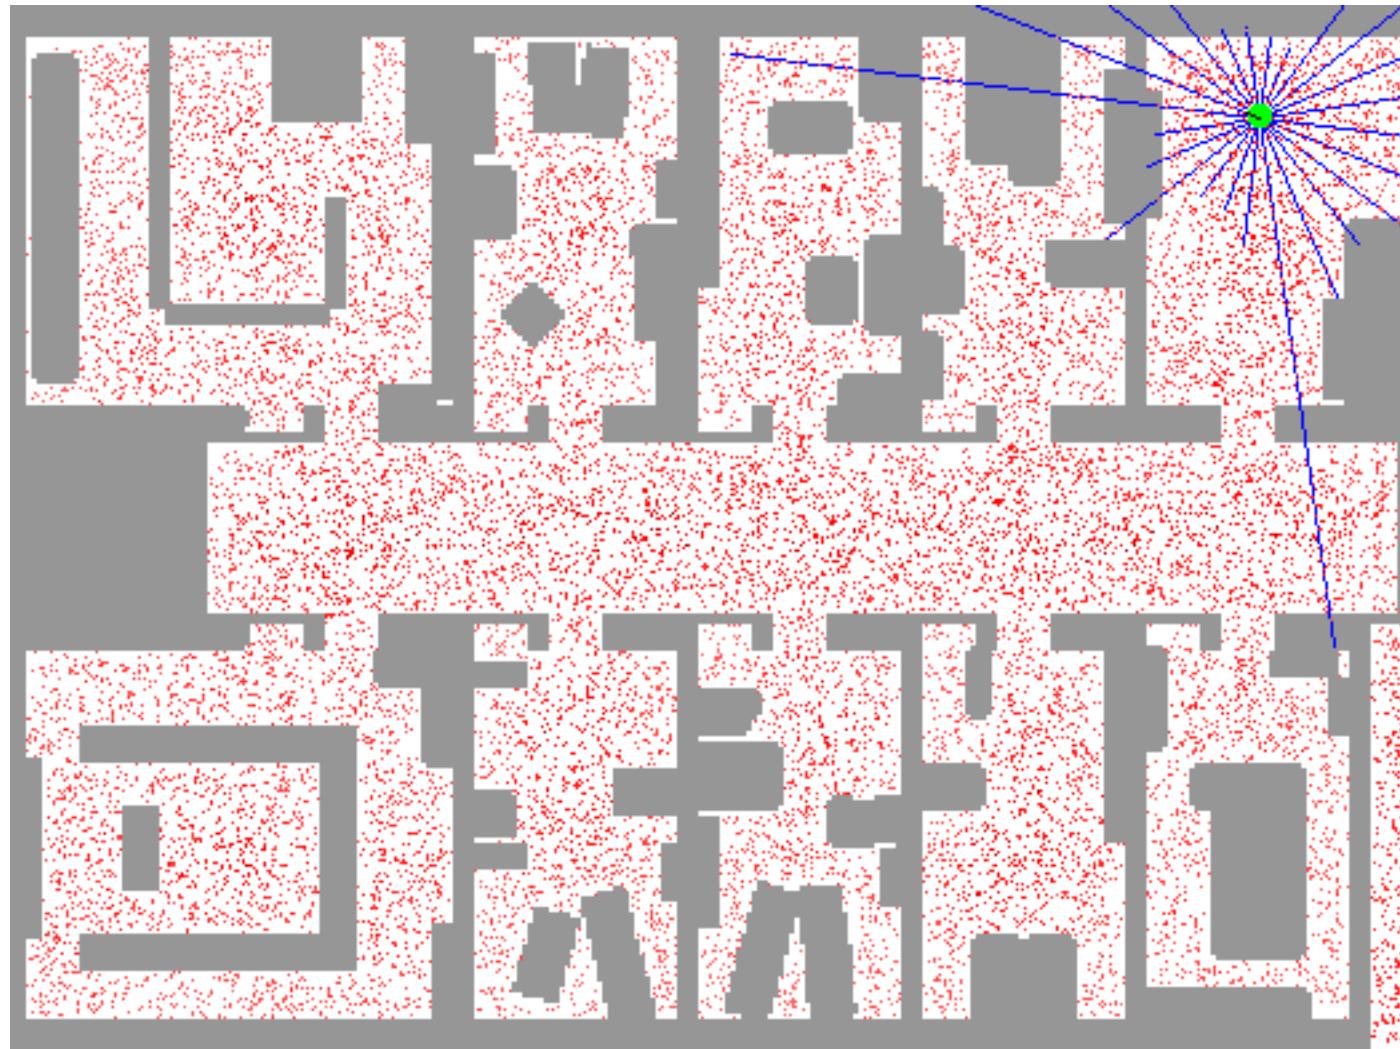
[Video: global-sonar-uw-annotated.avi]

Particle Filter Localization (Laser)



[Video: global-floor.gif]

Particle Filter Localization (Laser)

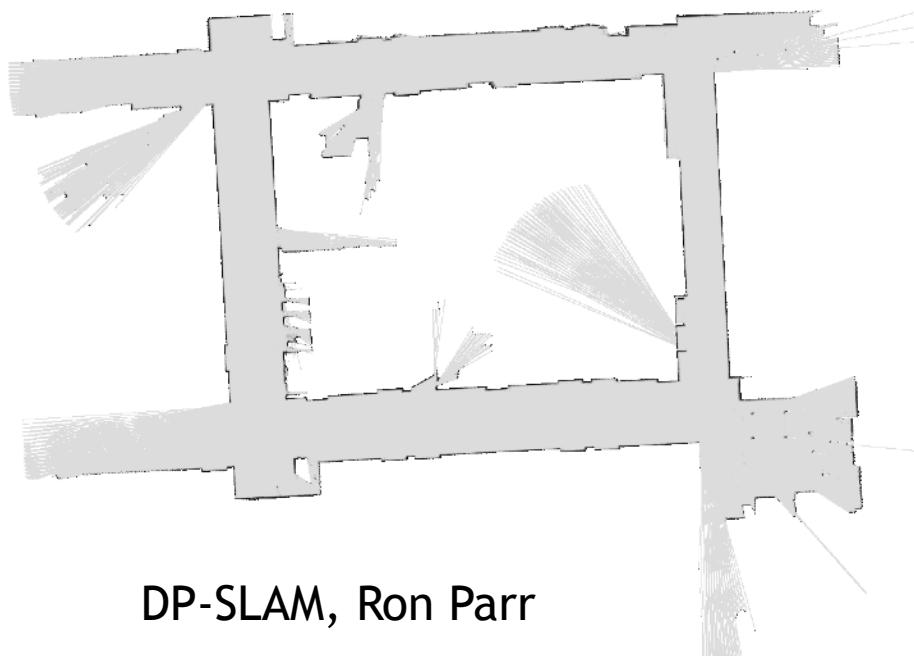


[Video: global-floor.gif]

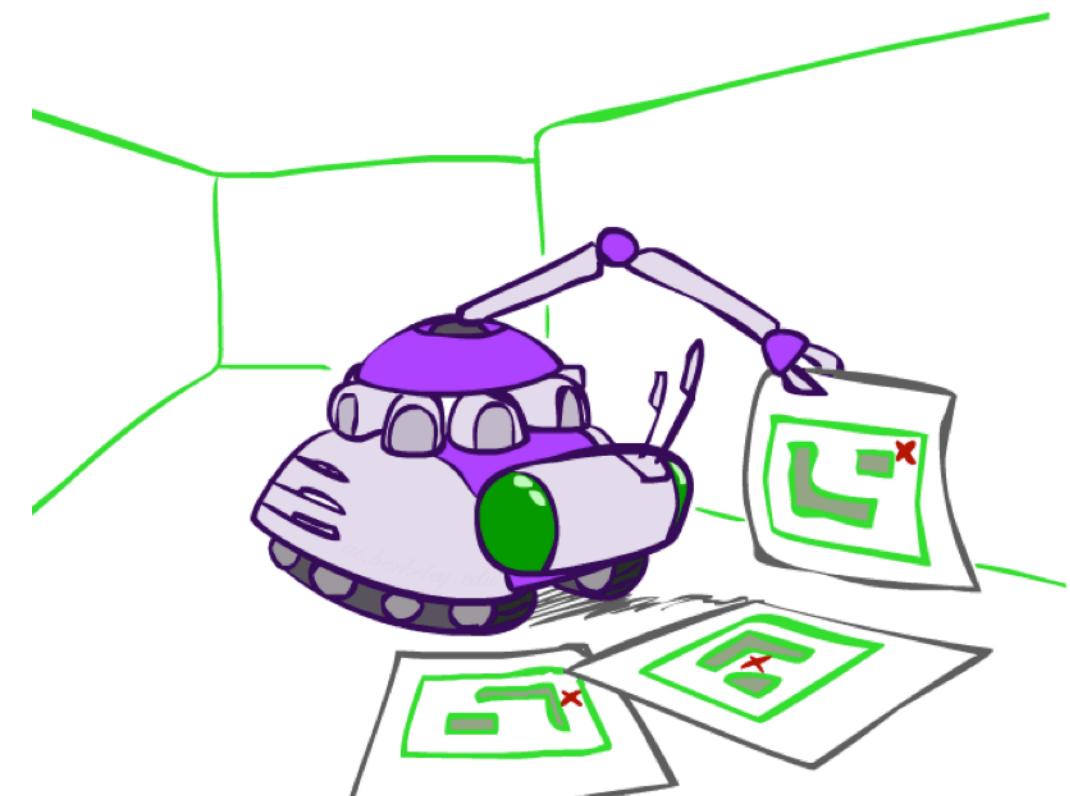
Robot Mapping

- SLAM: Simultaneous Localization And Mapping

- We do not know the map or our location
- State consists of position AND map!
- Main techniques: Kalman filtering (Gaussian HMMs) and particle methods

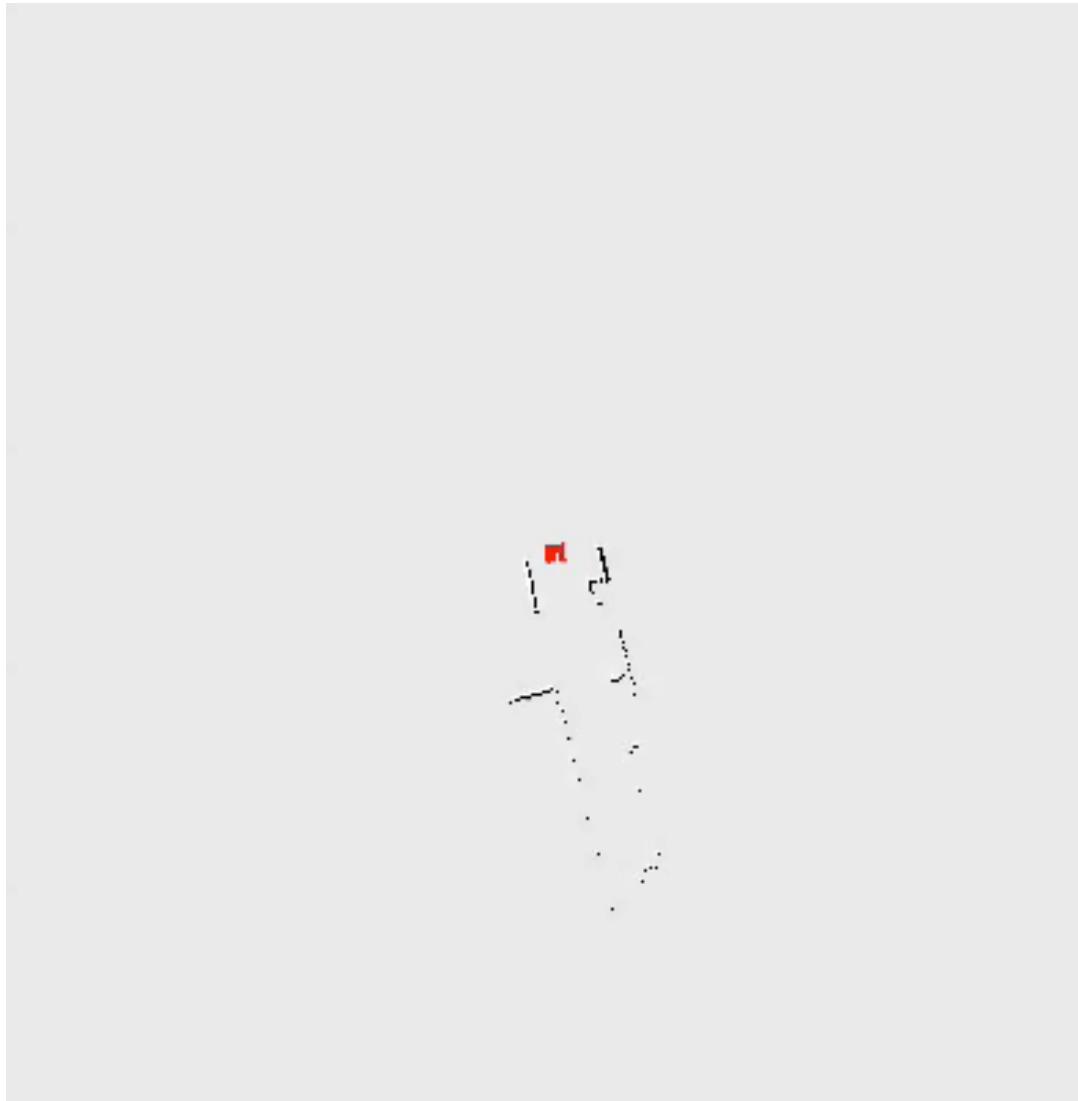


DP-SLAM, Ron Parr



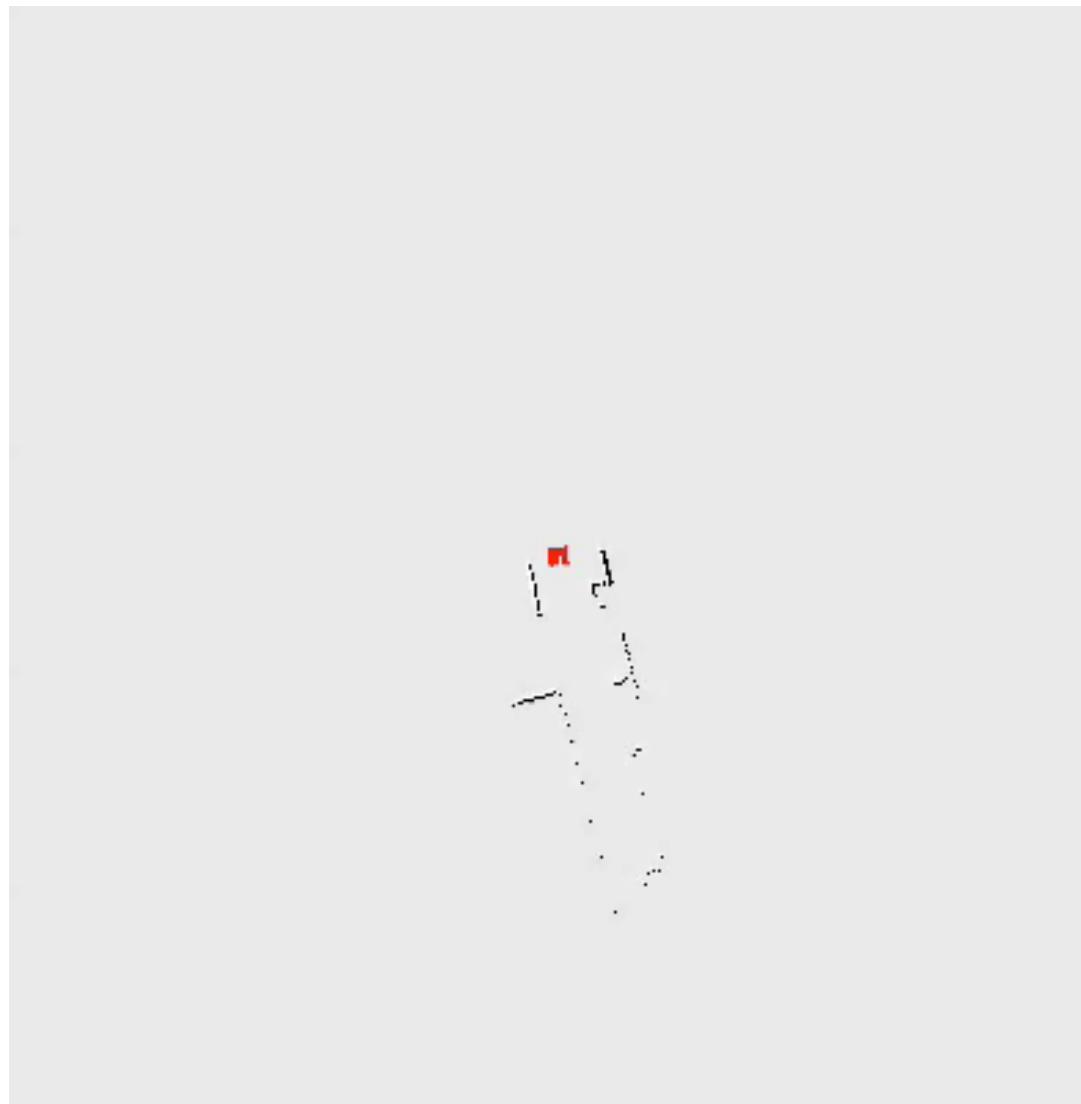
[Demo: PARTICLES-SLAM-mapping1-new.a]

Particle Filter SLAM - Video 1



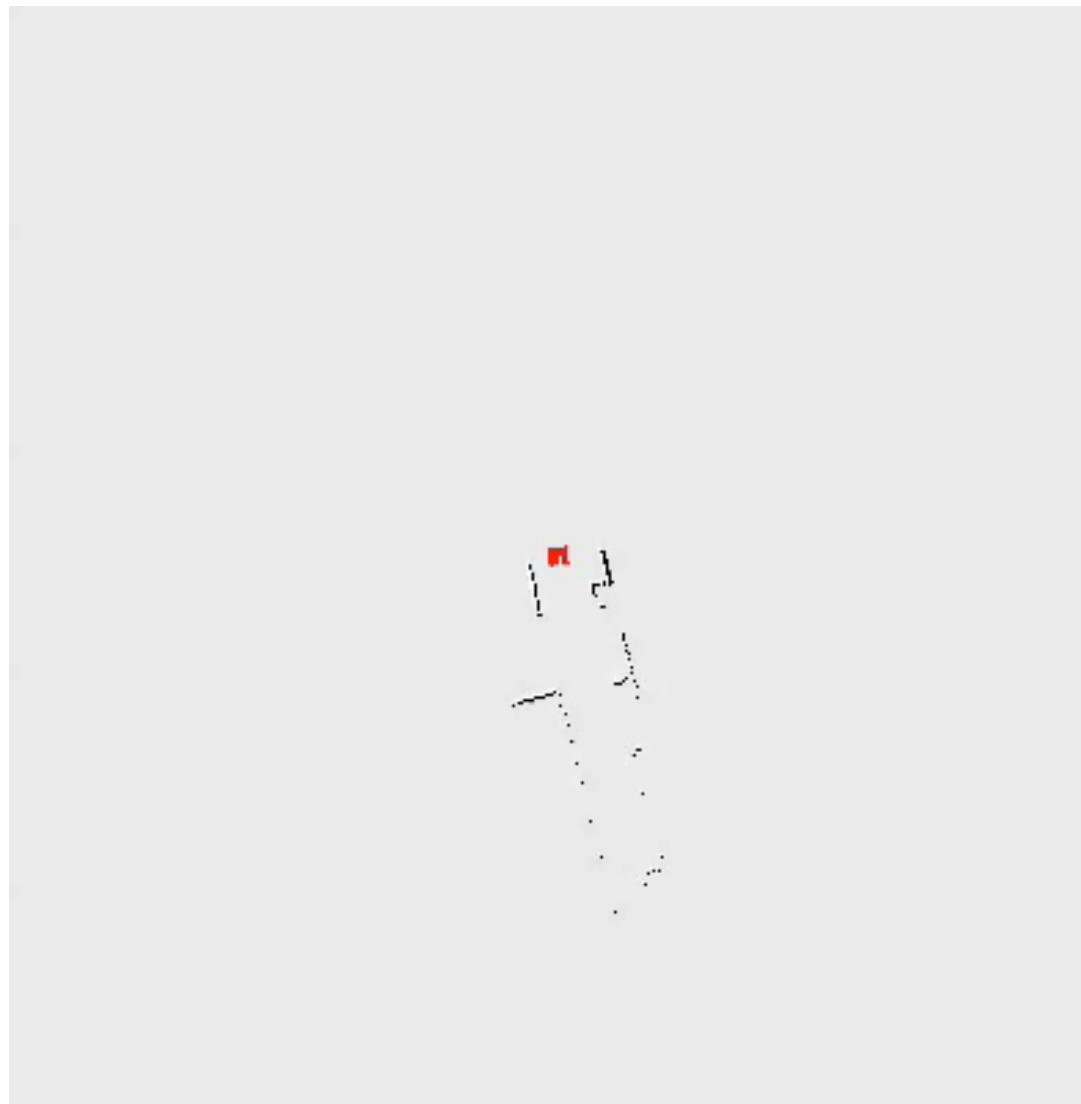
[Demo: PARTICLES-SLAM-mapping1-new.a]

Particle Filter SLAM - Video 1



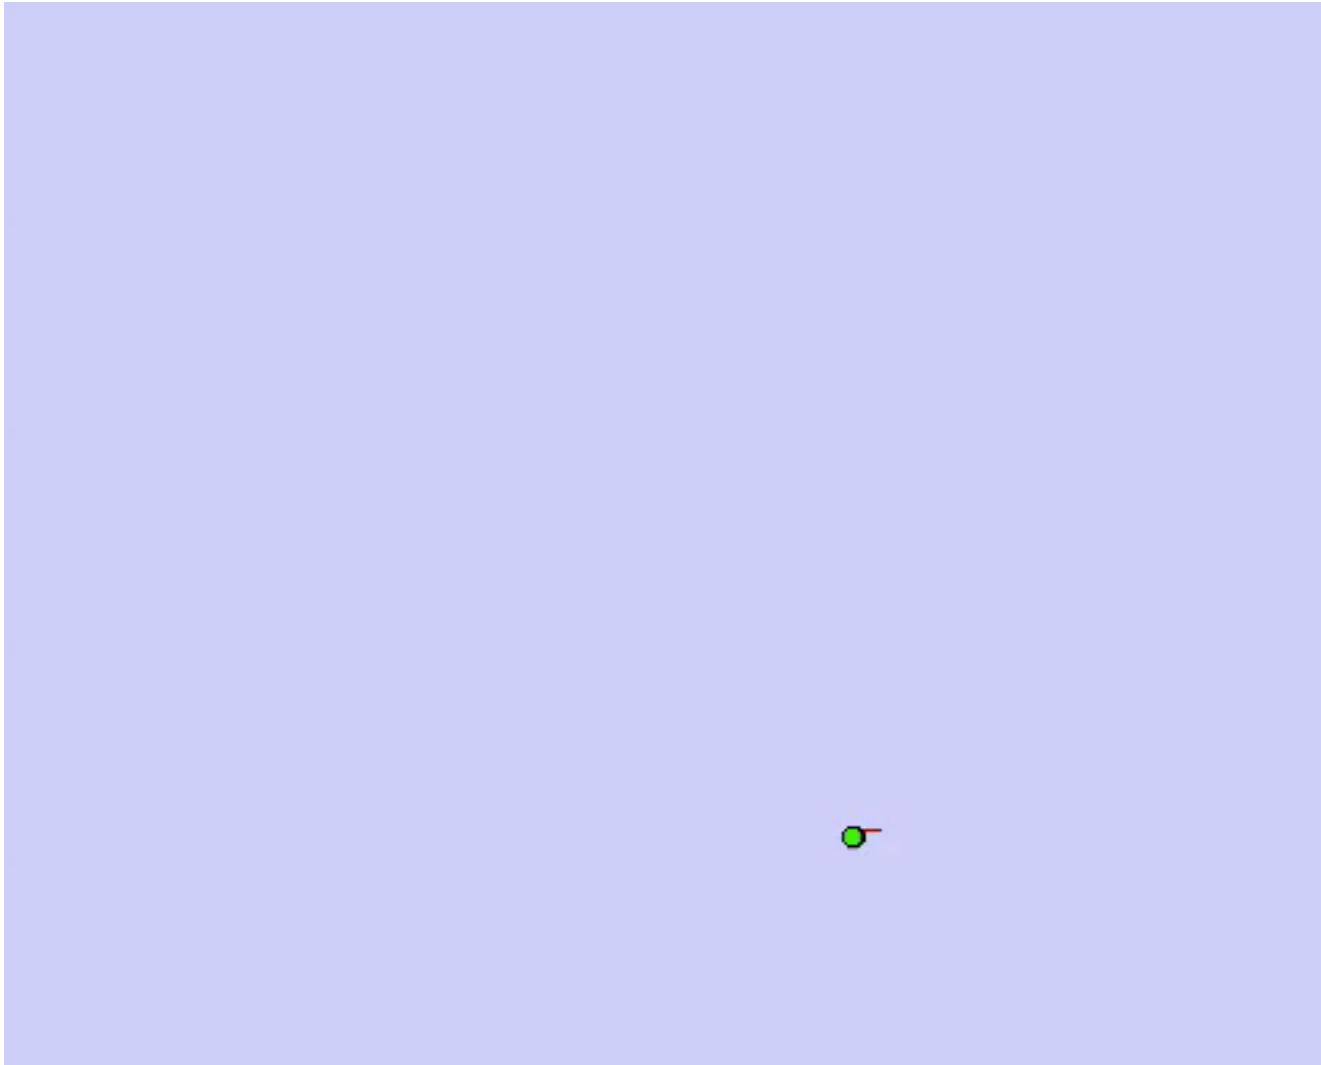
[Demo: PARTICLES-SLAM-mapping1-new.a]

Particle Filter SLAM - Video 1



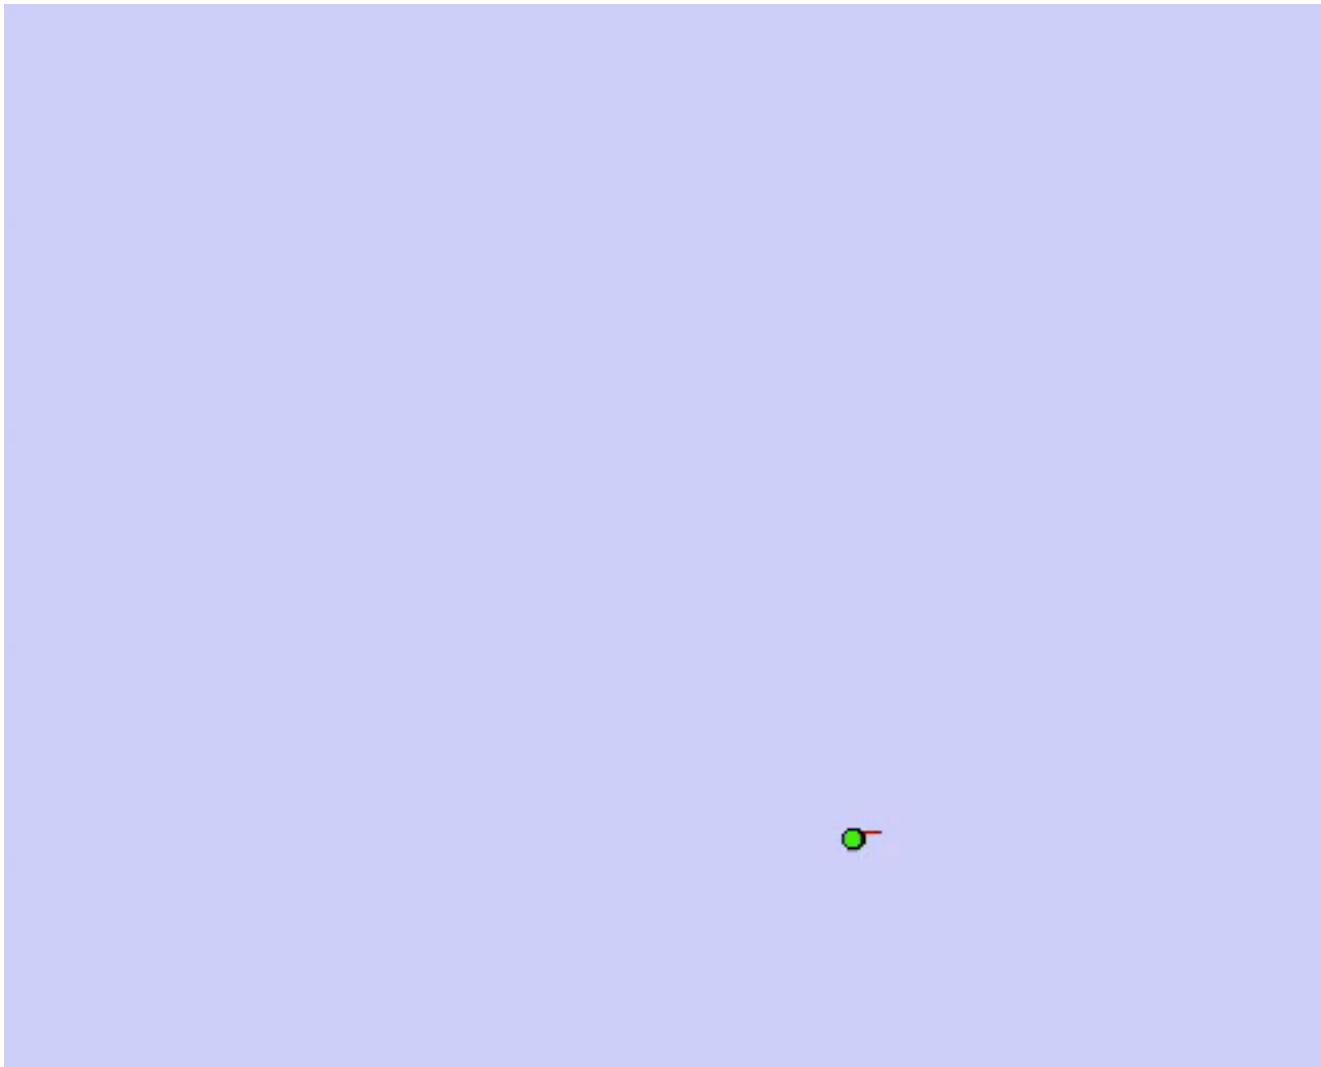
[Demo: PARTICLES-SLAM-mapping1-new.a]

Particle Filter SLAM - Video 2



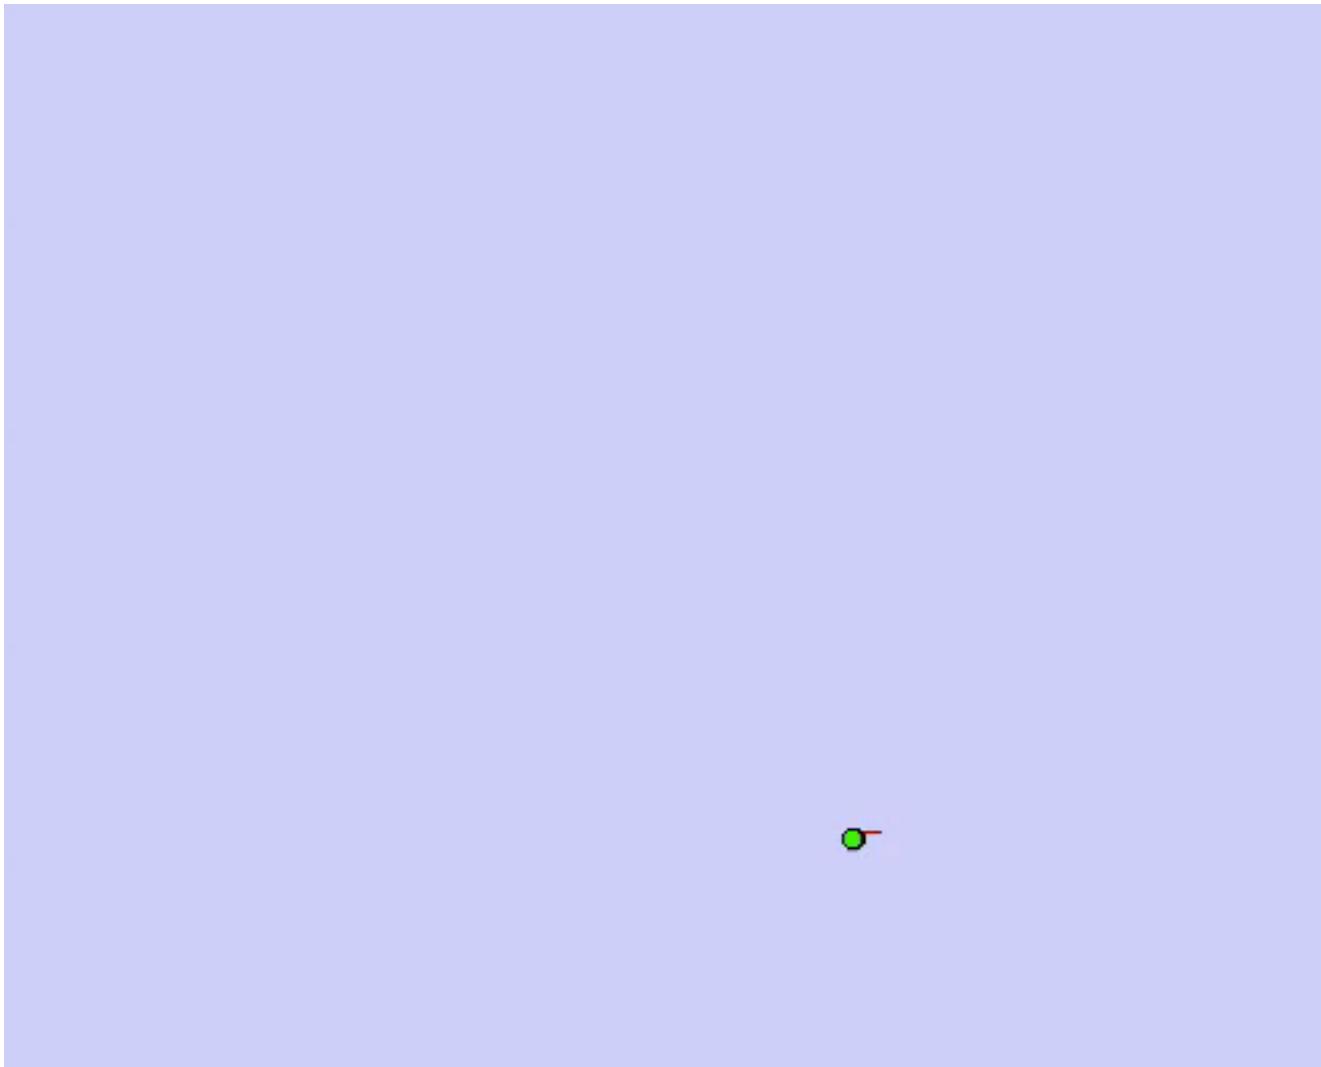
[Demo: PARTICLES-SLAM-fastslam.a]

Particle Filter SLAM - Video 2



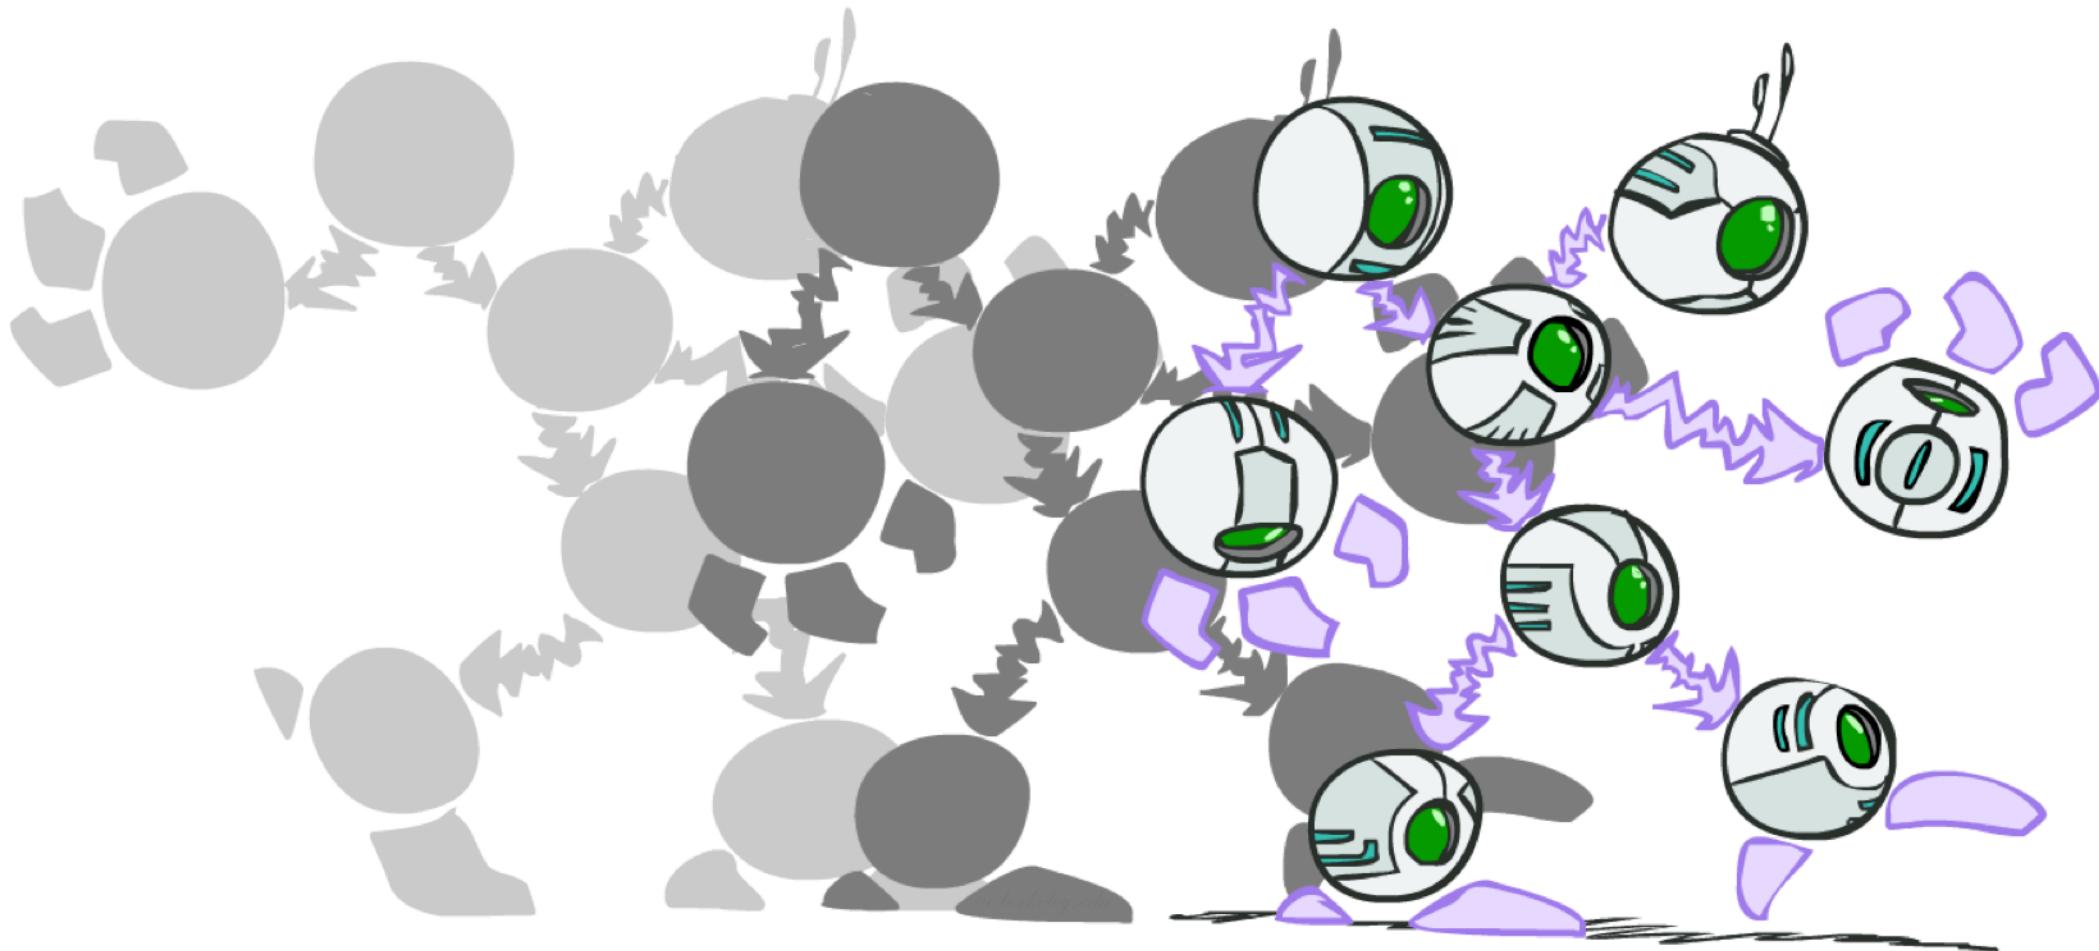
[Demo: PARTICLES-SLAM-fastslam.a]

Particle Filter SLAM - Video 2

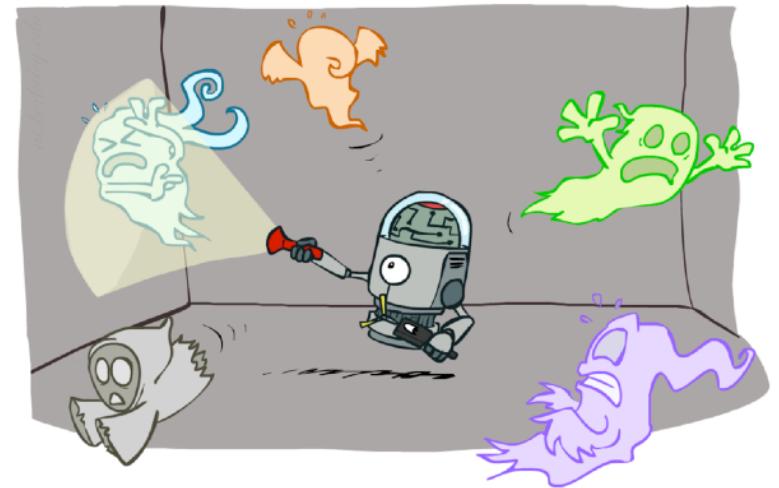


[Demo: PARTICLES-SLAM-fastslam.a]

Dynamic Bayes Nets

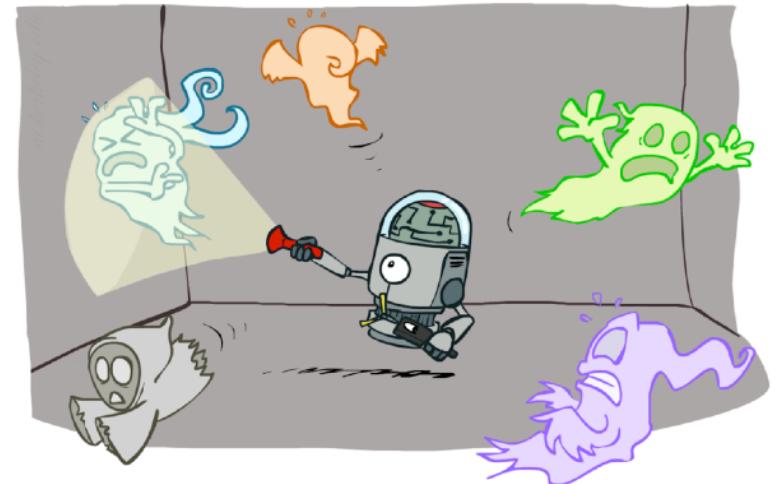
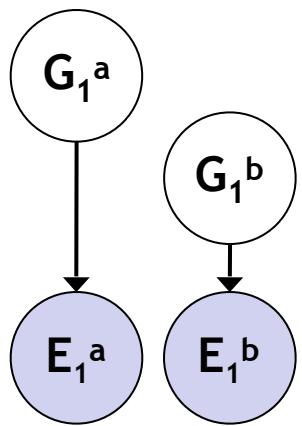


Dynamic Bayes Nets (DBNs)



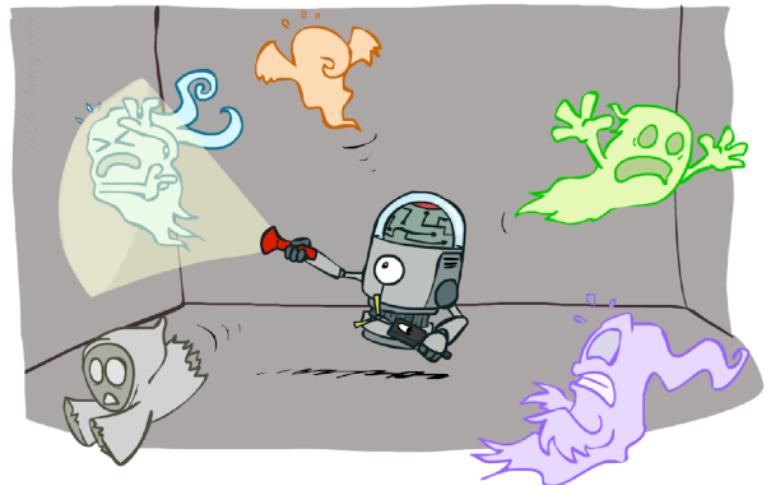
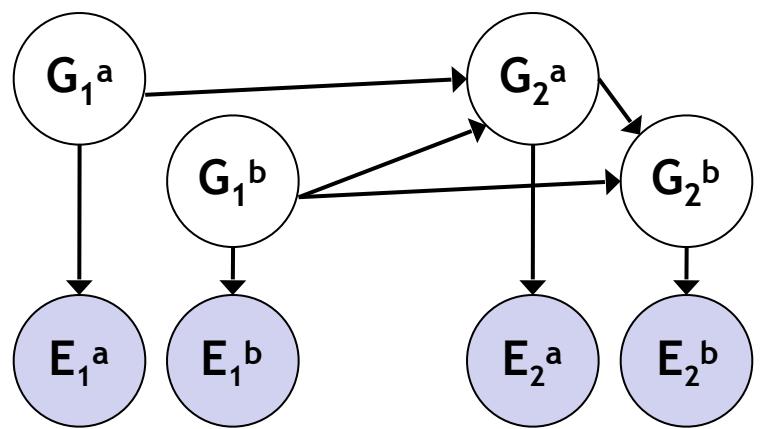
[Demo: pacman sonar ghost DBN model (L15D6)]

Dynamic Bayes Nets (DBNs)



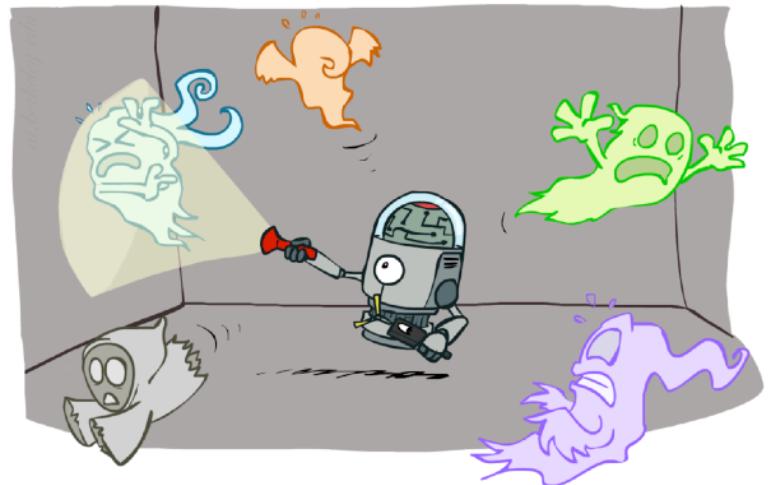
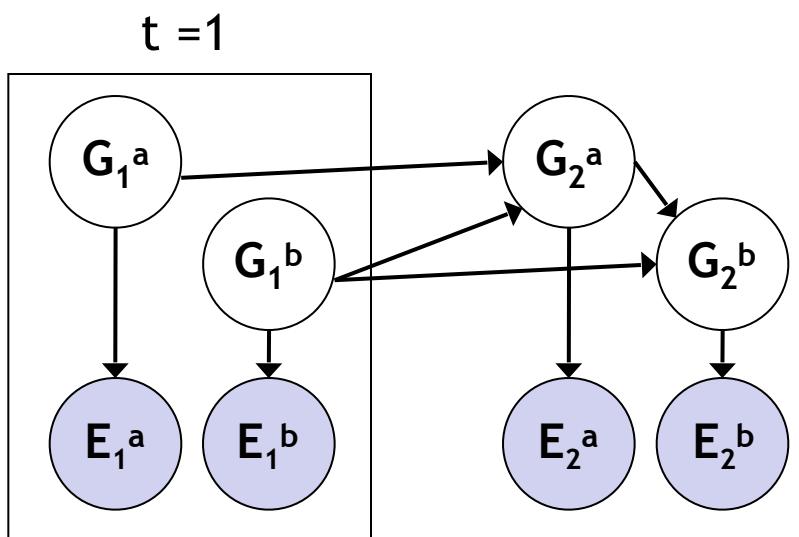
[Demo: pacman sonar ghost DBN model (L15D6)]

Dynamic Bayes Nets (DBNs)



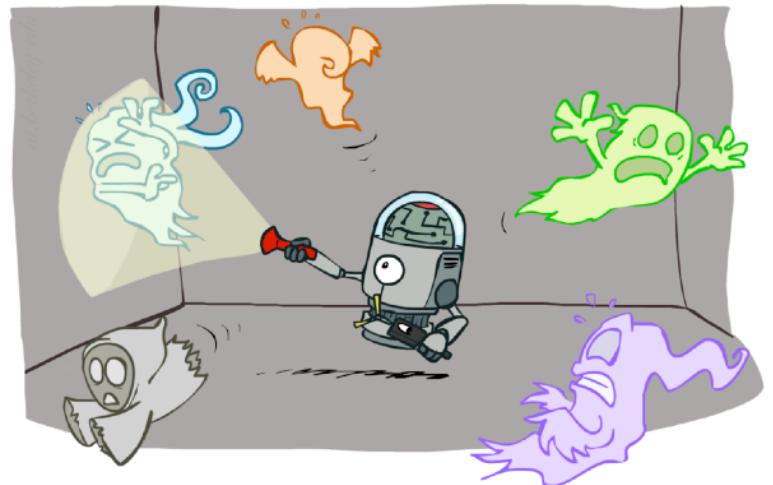
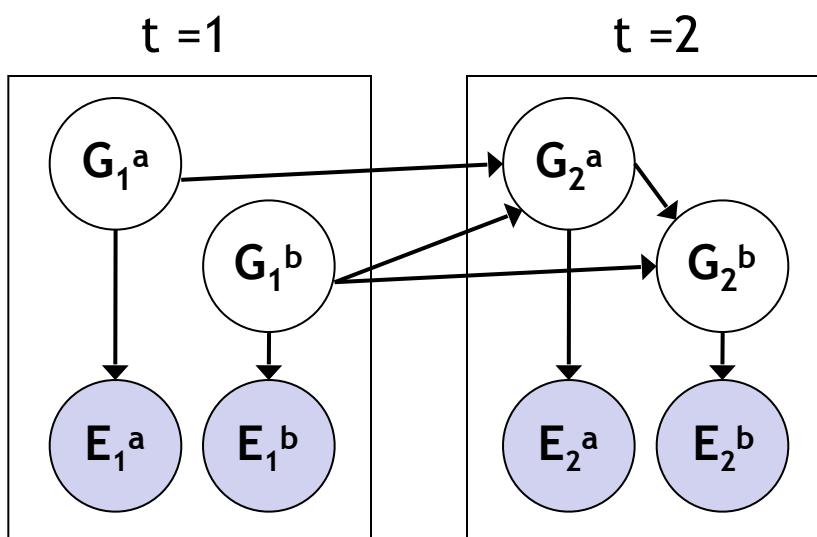
[Demo: pacman sonar ghost DBN model (L15D6)]

Dynamic Bayes Nets (DBNs)



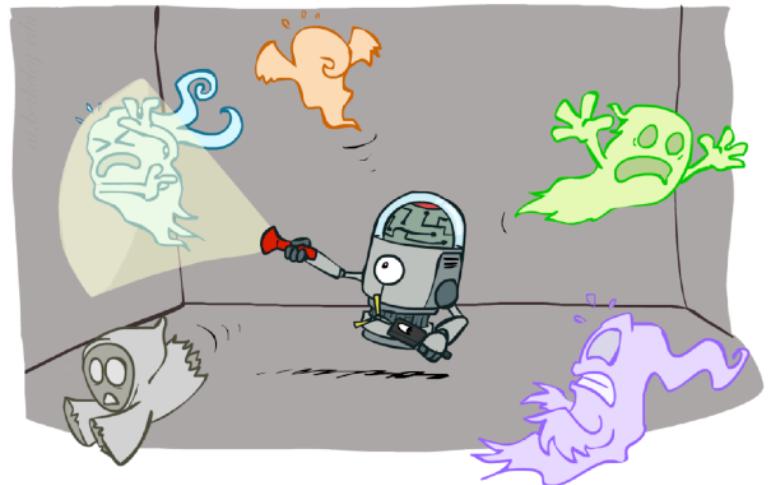
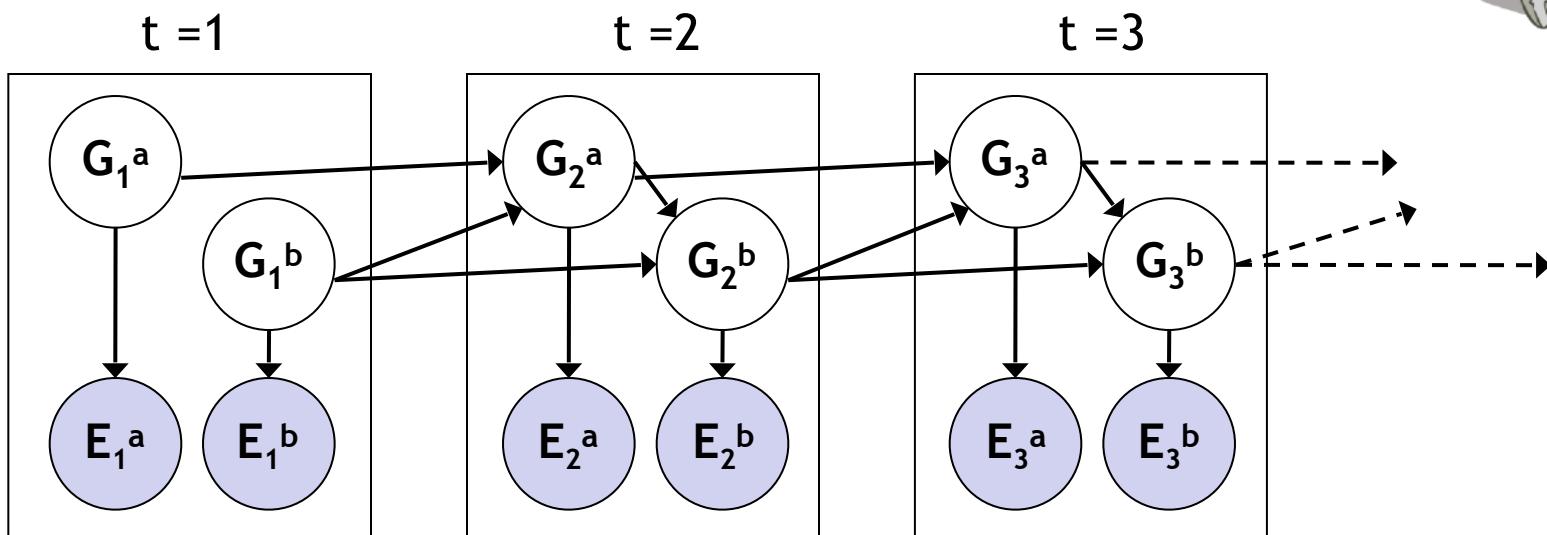
[Demo: pacman sonar ghost DBN model (L15D6)]

Dynamic Bayes Nets (DBNs)



[Demo: pacman sonar ghost DBN model (L15D6)]

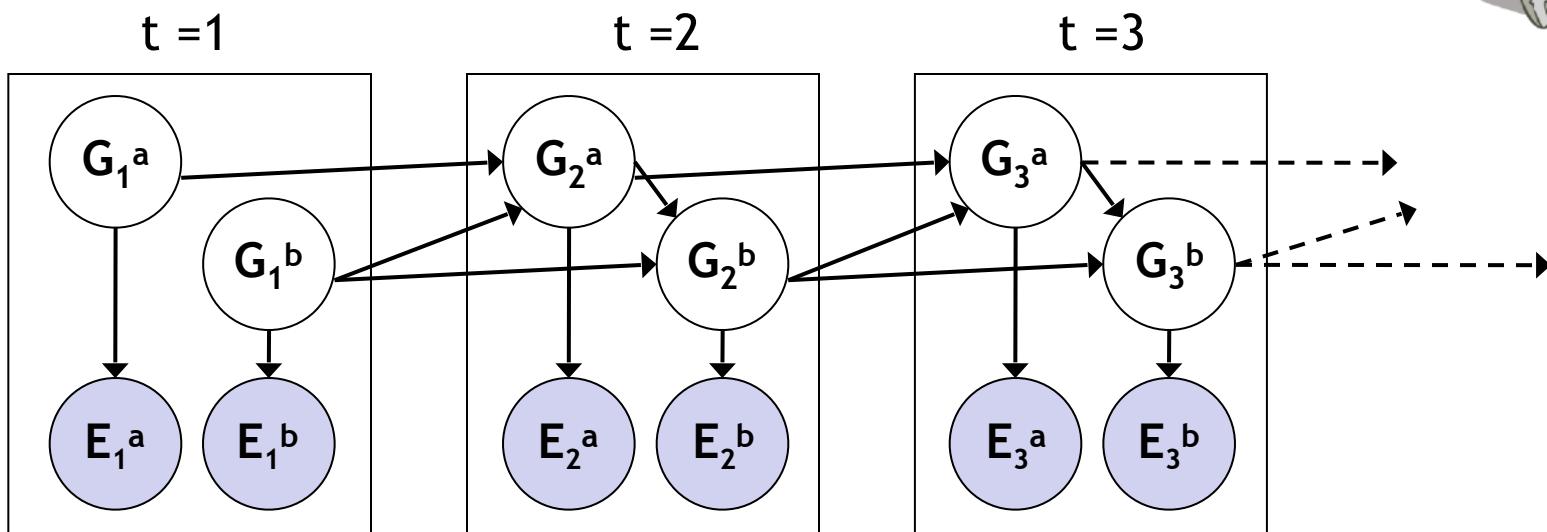
Dynamic Bayes Nets (DBNs)



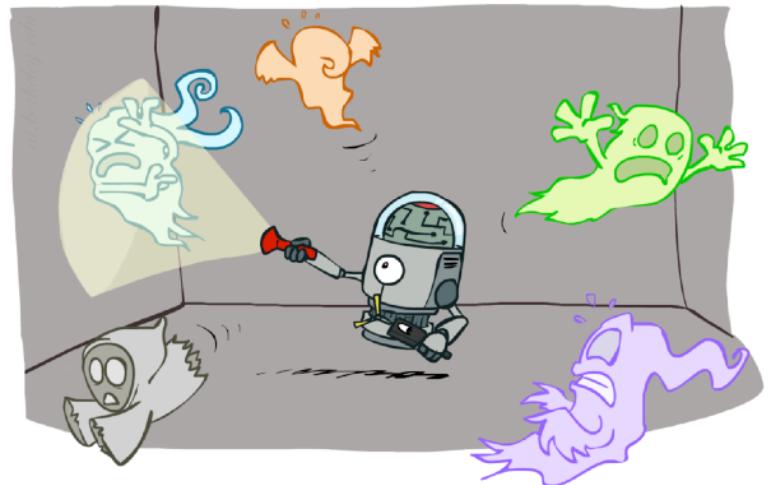
[Demo: pacman sonar ghost DBN model (L15D6)]

Dynamic Bayes Nets (DBNs)

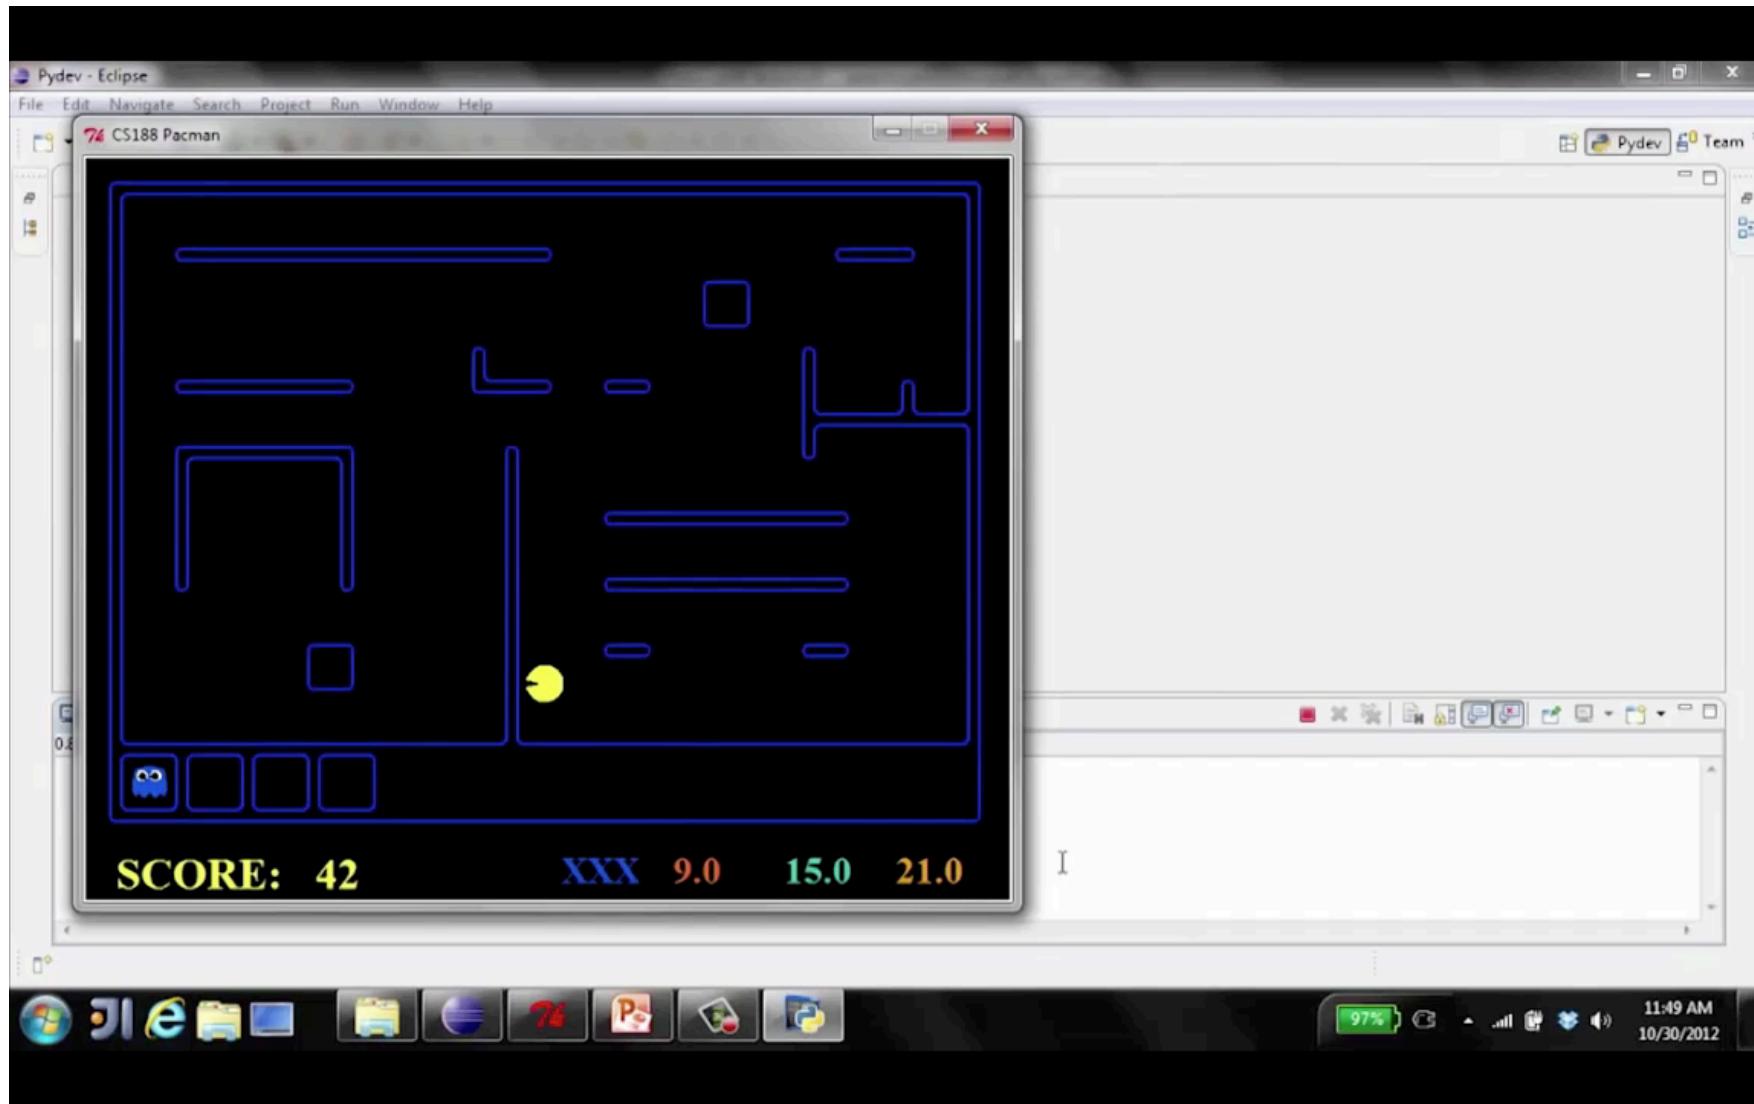
- We want to track multiple variables over time, using multiple sources of evidence
- Idea: Repeat a fixed Bayes net structure at each time
- Variables from time t can condition on those from $t-1$



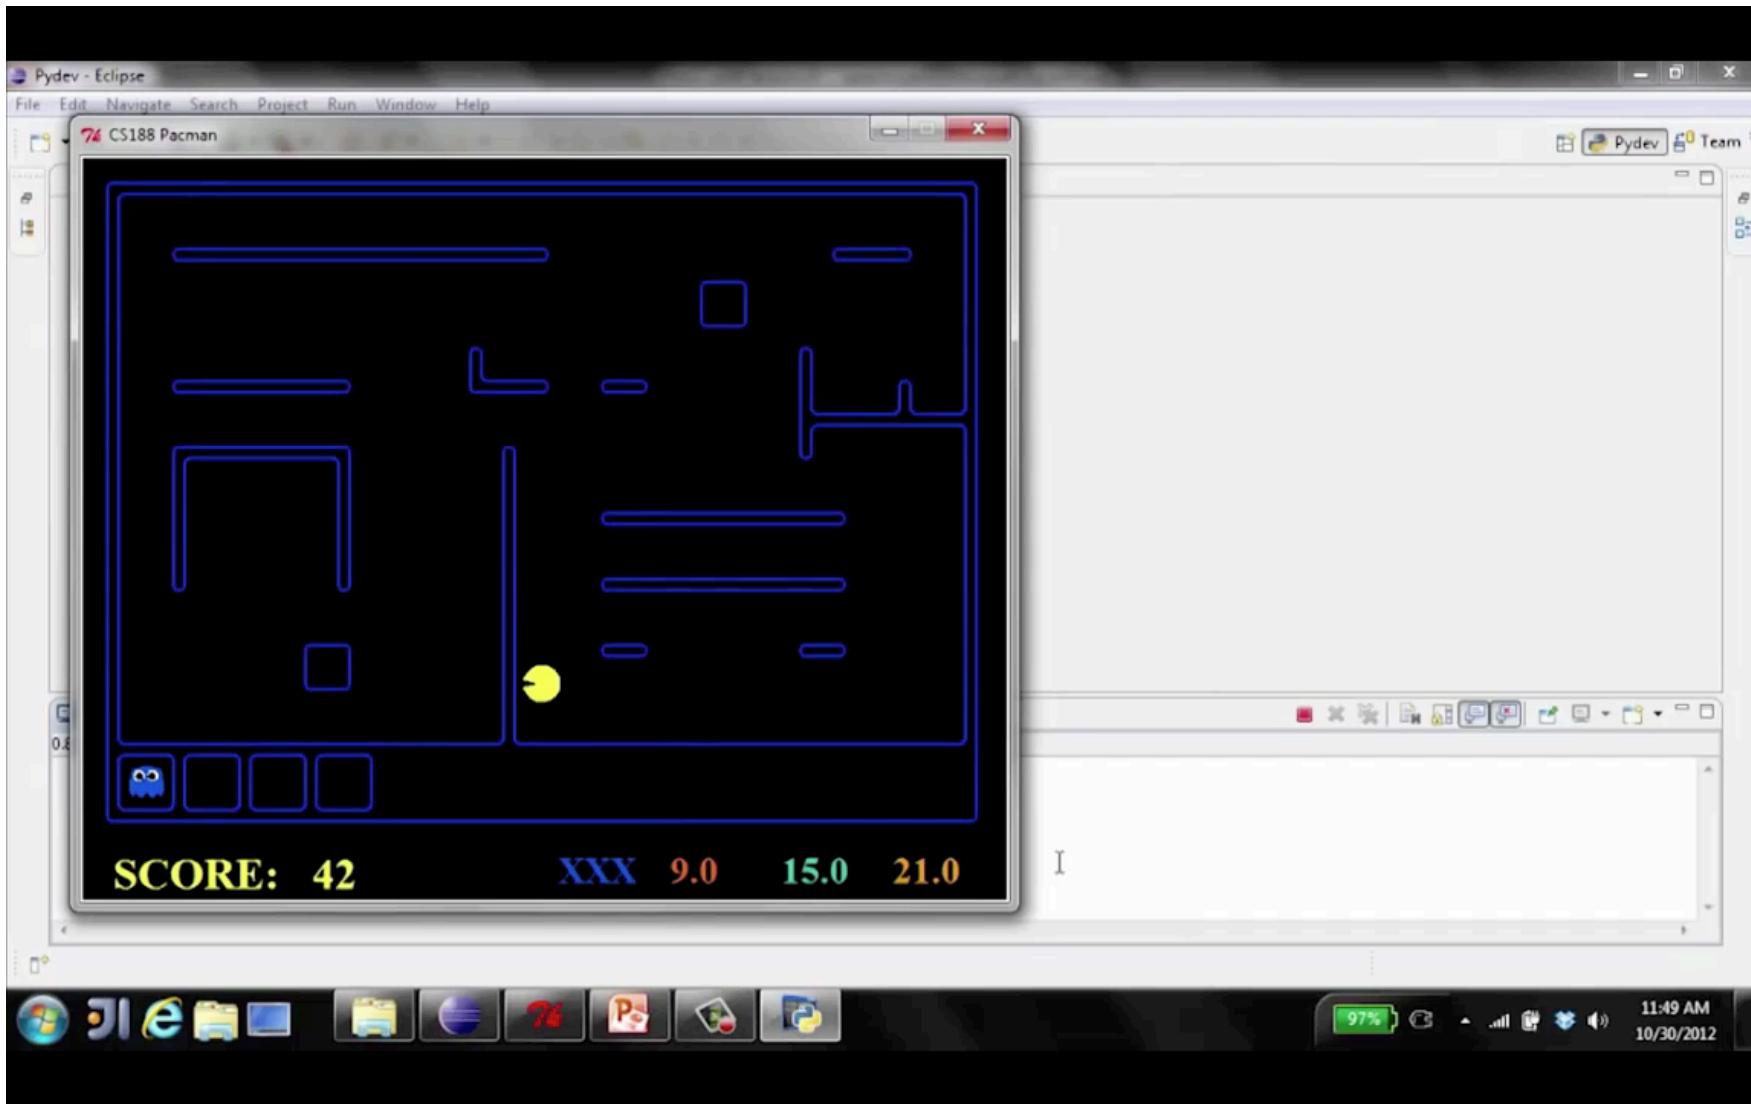
- Dynamic Bayes nets are a generalization of HMMs



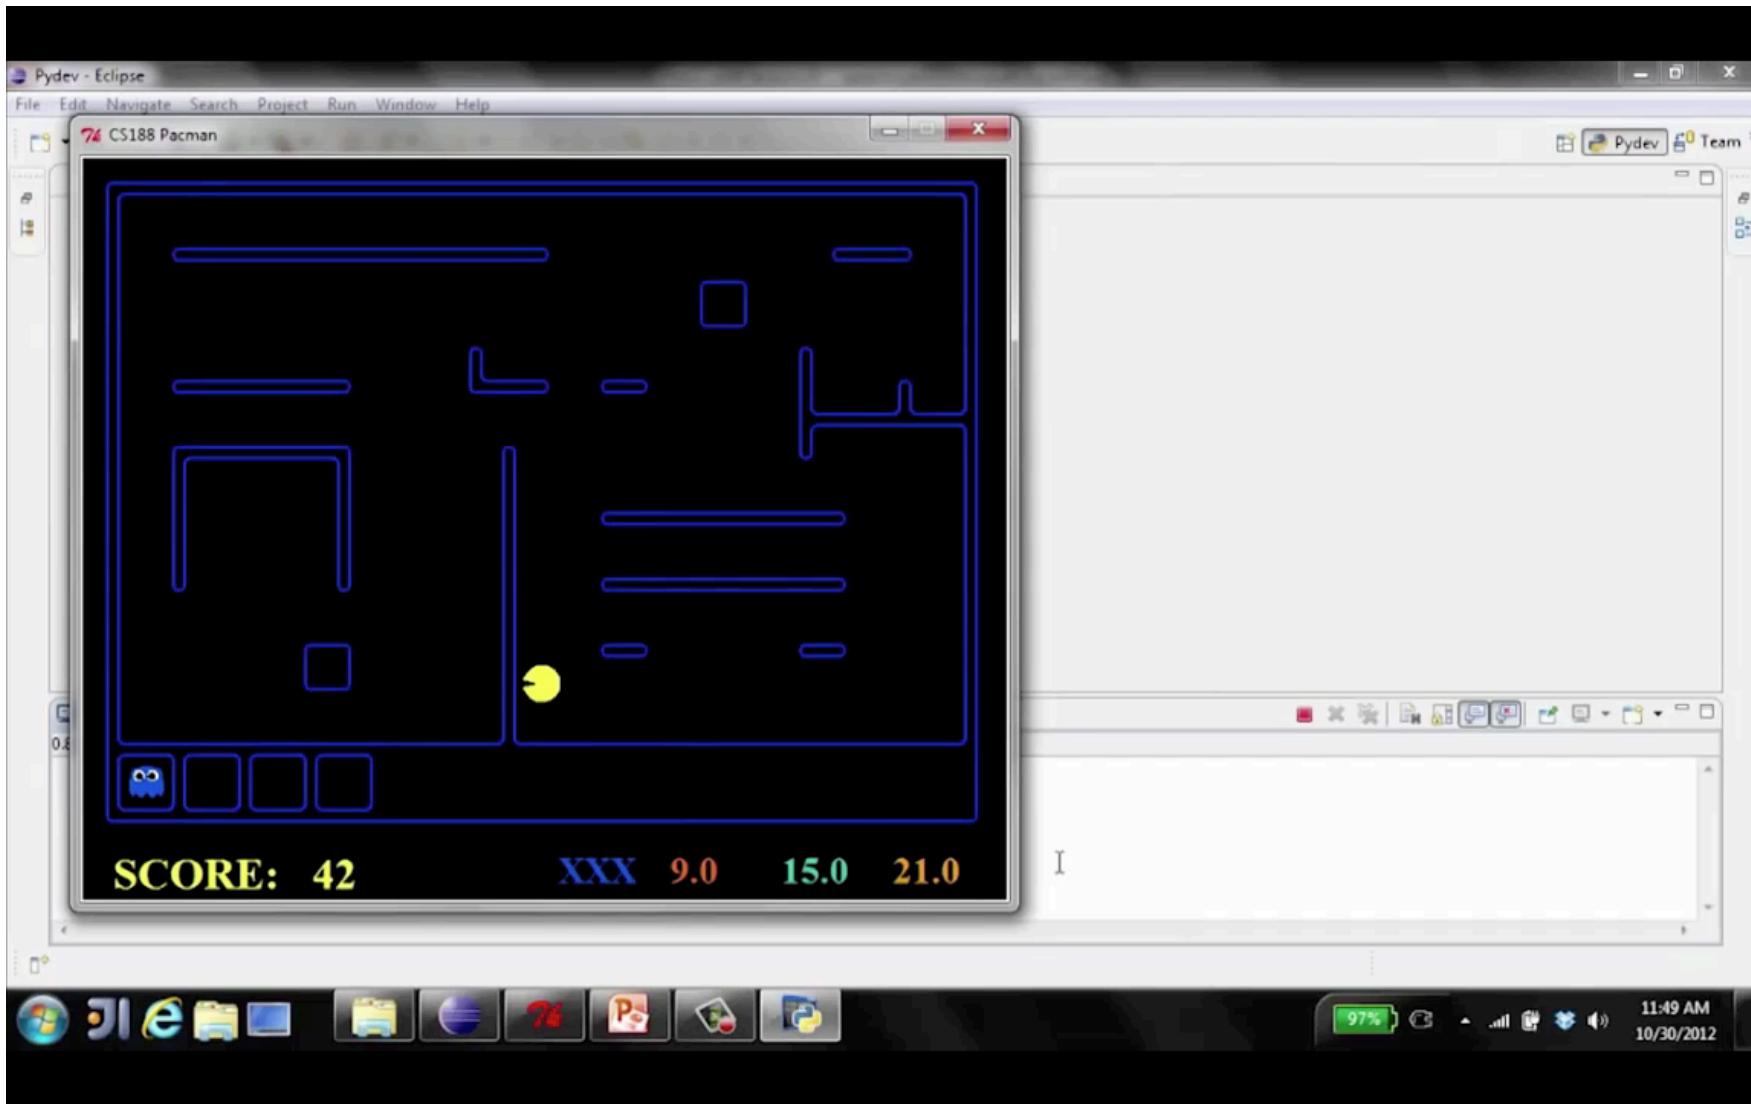
Video of Demo Pacman Sonar Ghost DBN Model



Video of Demo Pacman Sonar Ghost DBN Model



Video of Demo Pacman Sonar Ghost DBN Model



DBN Particle Filters

- A particle is a complete sample for a time step
- **Initialize:** Generate prior samples for the t=1 Bayes net
 - Example particle: $G_1^a = (3,3)$ $G_1^b = (5,3)$
- **Elapse time:** Sample a successor for each particle
 - Example successor: $G_2^a = (2,3)$ $G_2^b = (6,3)$
- **Observe:** Weight each entire sample by the likelihood of the evidence conditioned on the sample
 - Likelihood: $P(E_1^a | G_1^a) * P(E_1^b | G_1^b)$
- **Resample:** Select prior samples (tuples of values) in proportion to their likelihood

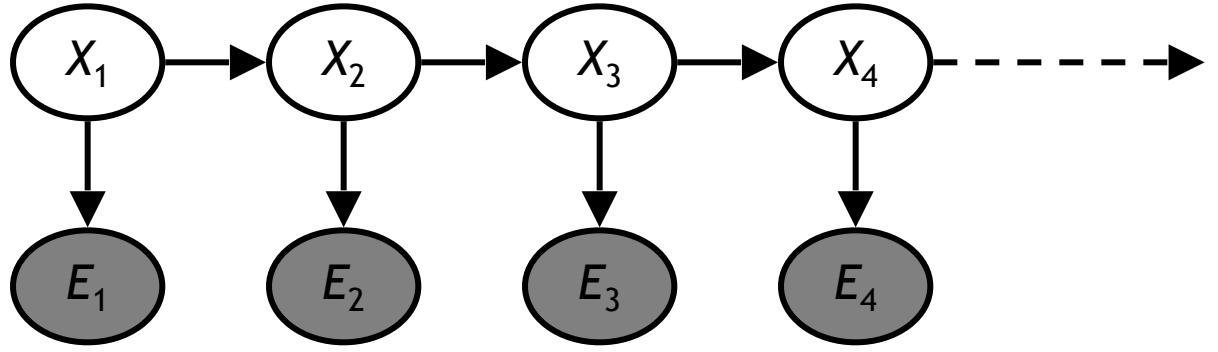
Most Likely Explanation



HMMs: MLE Queries

- HMMs defined by
 - States X
 - Observations E
 - Initial distribution:
 - Transitions:
 - Emissions:
- New query: most likely explanation:

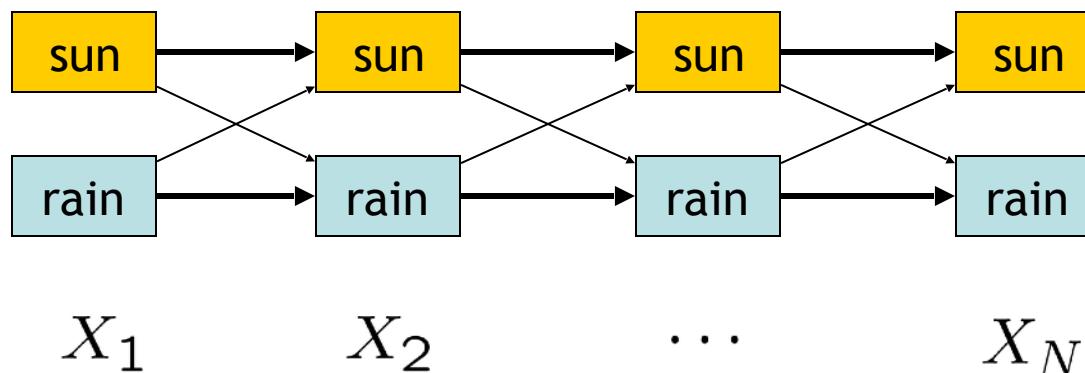
$$\begin{aligned} & P(X_1) \\ & P(X|X_{-1}) \\ & P(E|X) \end{aligned}$$



$$\arg \max_{x_{1:t}} P(x_{1:t}|e_{1:t})$$

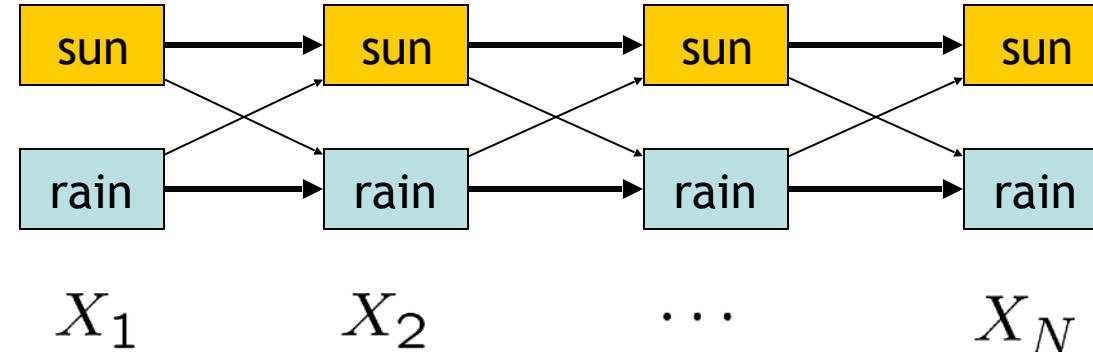
State Trellis

- State trellis: graph of states and transitions over time



- Each arc represents some transition $x_{t-1} \rightarrow x_t$
- Each arc has weight $P(x_t|x_{t-1})P(e_t|x_t)$
- Each path is a sequence of states
- The product of weights on a path is that sequence's probability along with the evidence
- Forward algorithm computes sums of paths, Viterbi computes best paths

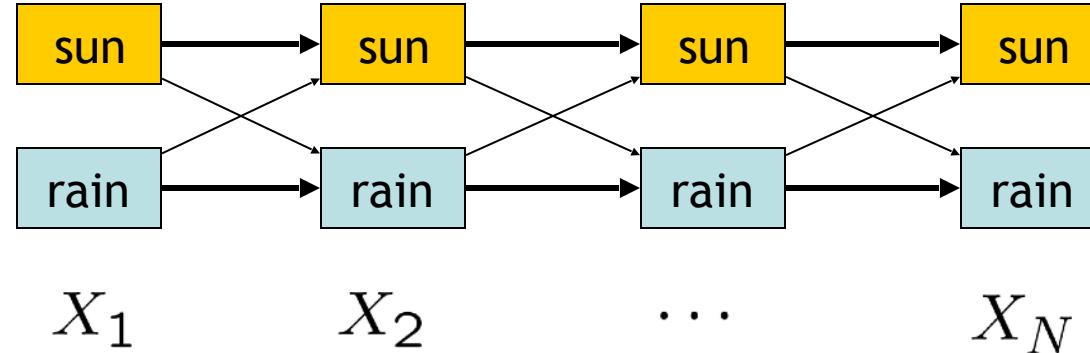
Forward / Viterbi Algorithms



Forward Algorithm (Sum)

$$f_t[x_t] = P(x_t, e_{1:t})$$

Forward / Viterbi Algorithms

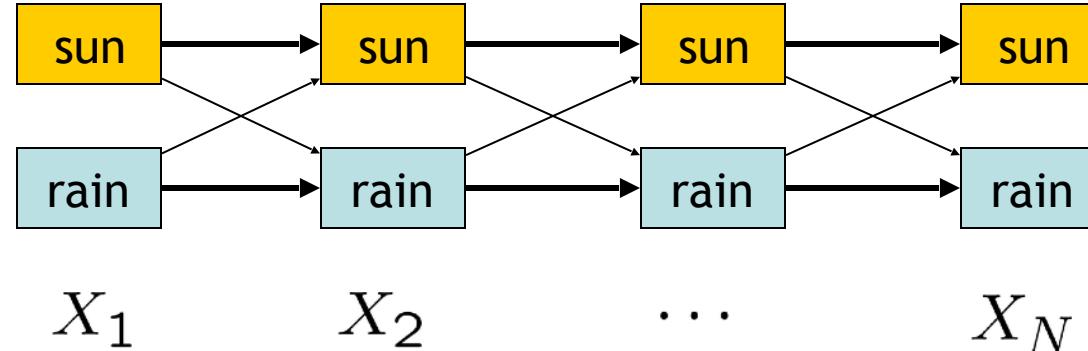


Forward Algorithm (Sum)

$$f_t[x_t] = P(x_t, e_{1:t})$$

$$= P(e_t|x_t) \sum_{x_{t-1}} P(x_t|x_{t-1}) f_{t-1}[x_{t-1}]$$

Forward / Viterbi Algorithms



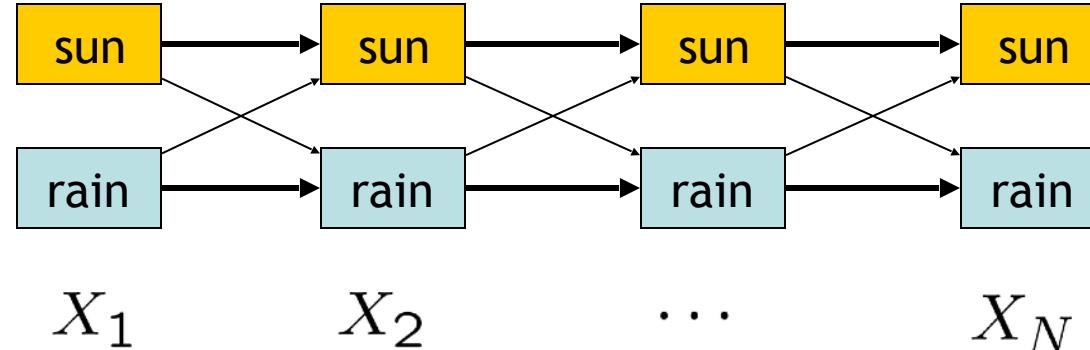
Forward Algorithm (Sum)

$$f_t[x_t] = P(x_t, e_{1:t})$$

Viterbi Algorithm (Max)

$$= P(e_t|x_t) \sum_{x_{t-1}} P(x_t|x_{t-1}) f_{t-1}[x_{t-1}]$$

Forward / Viterbi Algorithms



Forward Algorithm (Sum)

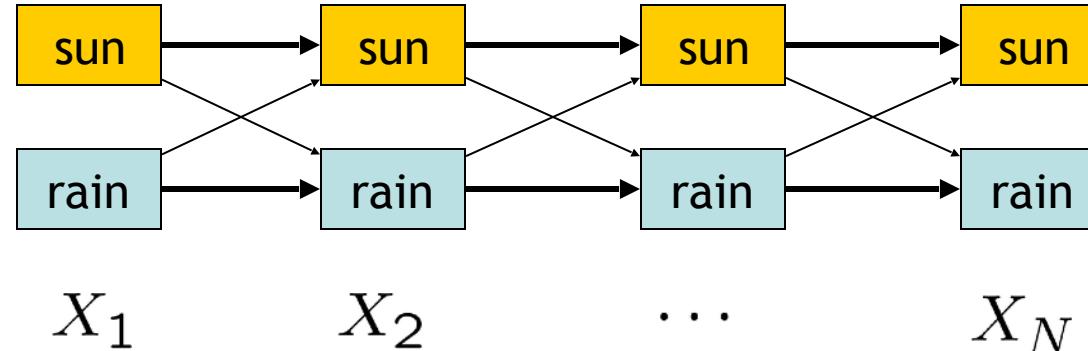
$$f_t[x_t] = P(x_t, e_{1:t})$$

$$= P(e_t|x_t) \sum_{x_{t-1}} P(x_t|x_{t-1}) f_{t-1}[x_{t-1}]$$

Viterbi Algorithm (Max)

$$m_t[x_t] = \max_{x_{1:t-1}} P(x_{1:t-1}, x_t, e_{1:t})$$

Forward / Viterbi Algorithms



Forward Algorithm (Sum)

$$f_t[x_t] = P(x_t, e_{1:t})$$

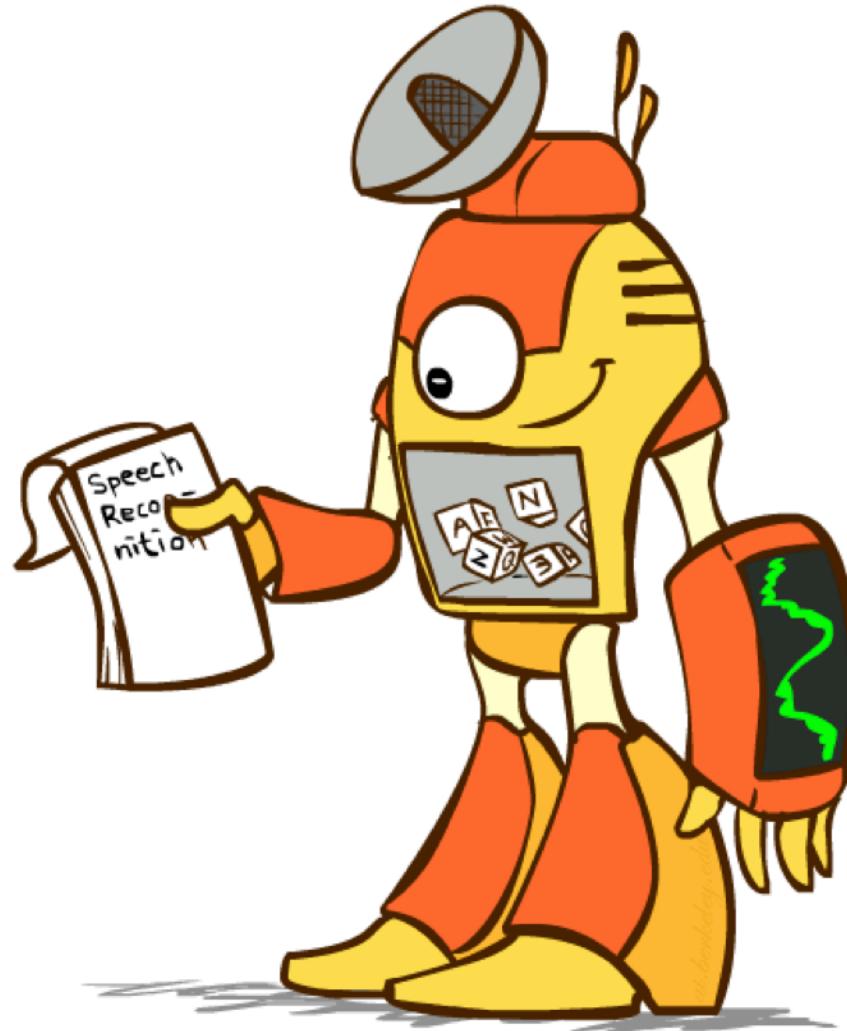
$$= P(e_t|x_t) \sum_{x_{t-1}} P(x_t|x_{t-1}) f_{t-1}[x_{t-1}]$$

Viterbi Algorithm (Max)

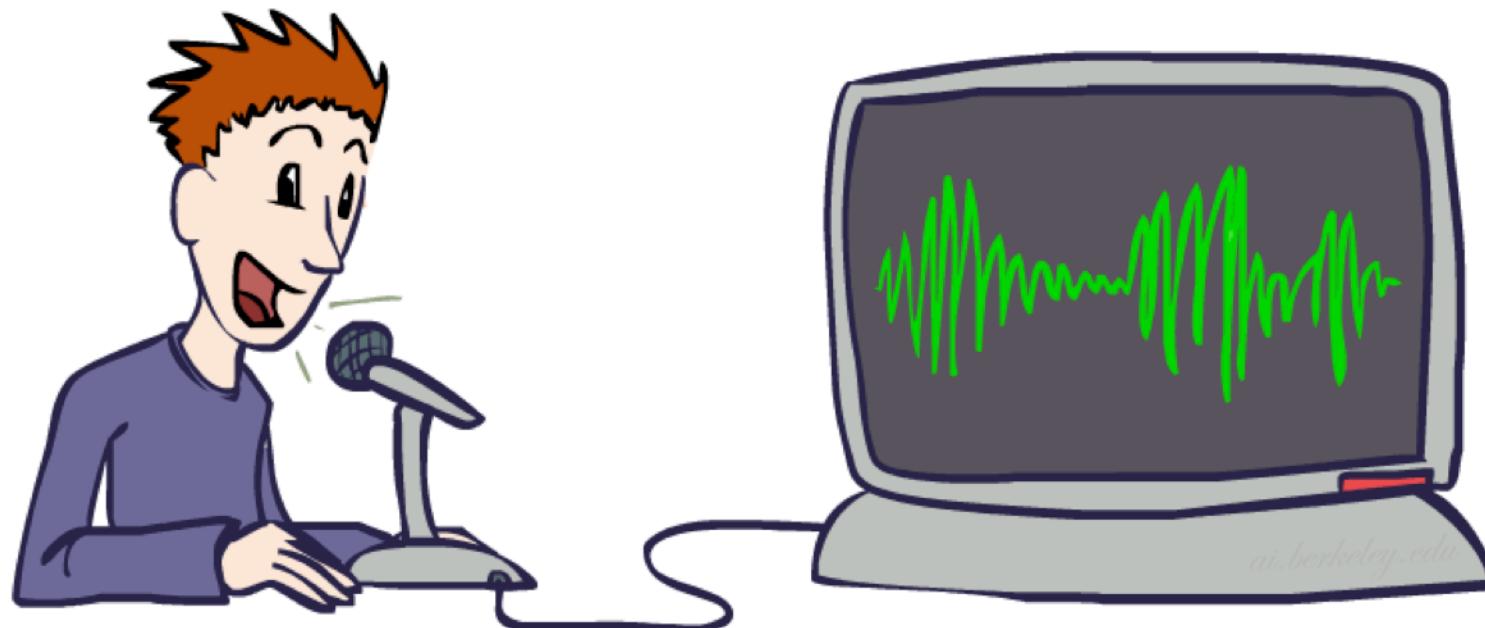
$$m_t[x_t] = \max_{x_{1:t-1}} P(x_{1:t-1}, x_t, e_{1:t})$$

$$= P(e_t|x_t) \max_{x_{t-1}} P(x_t|x_{t-1}) m_{t-1}[x_{t-1}]$$

Speech Recognition

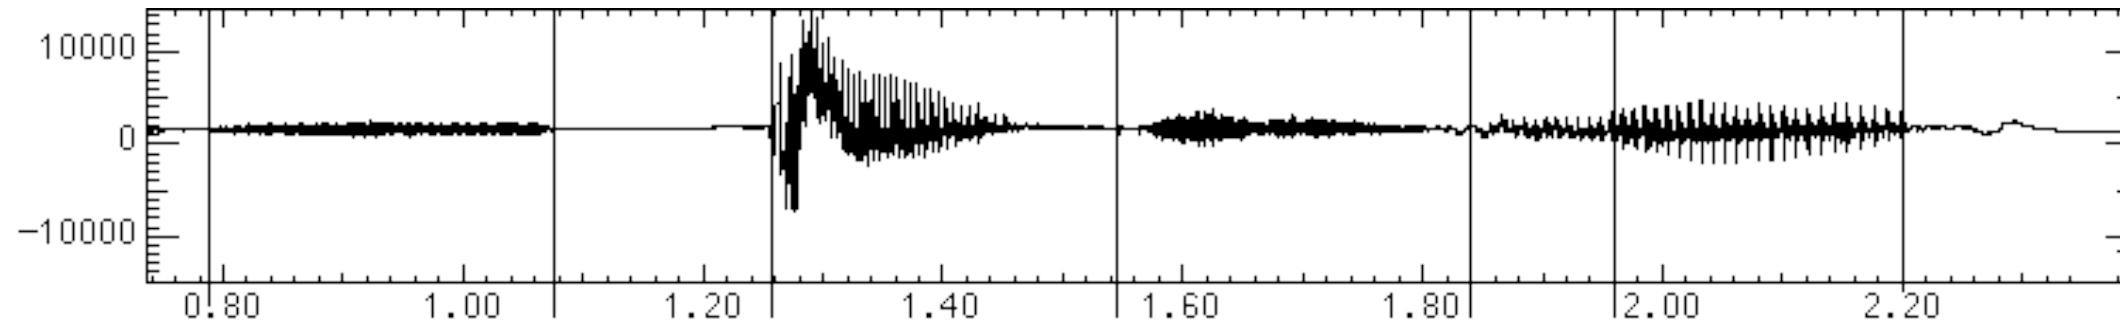


Digitizing Speech



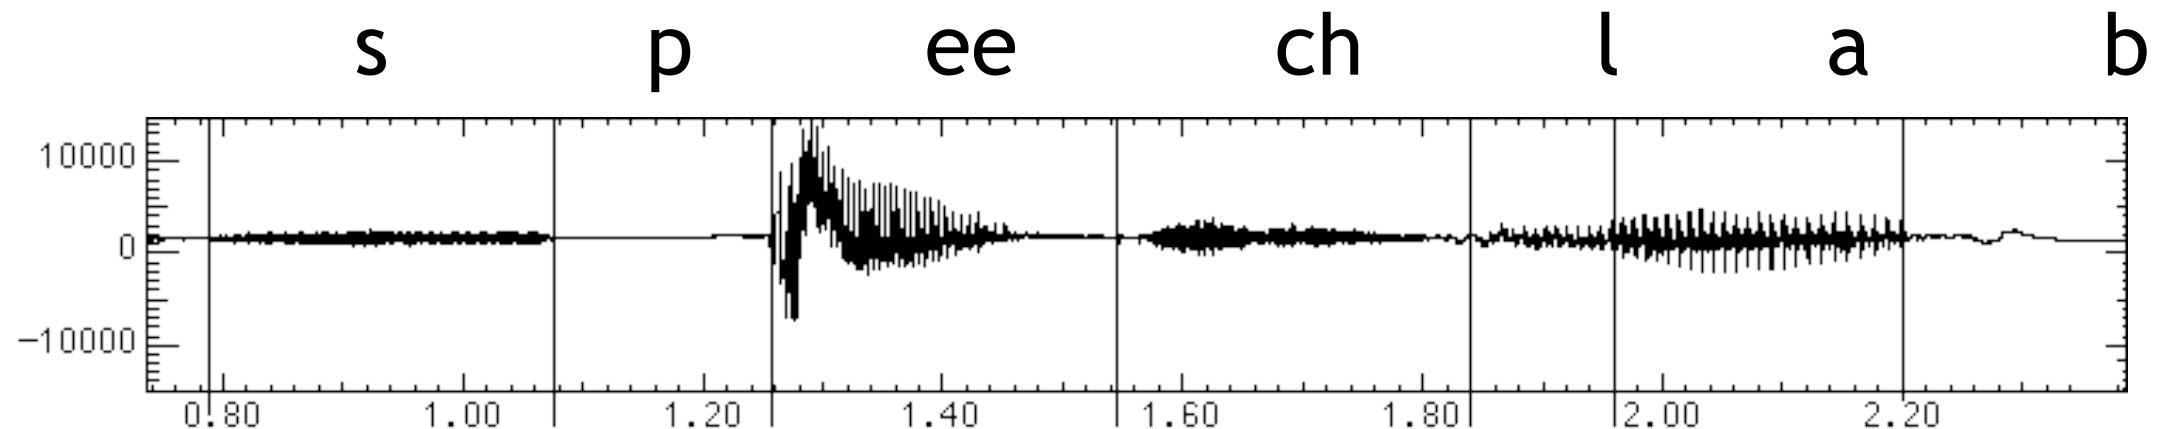
Speech in an Hour

- Speech input is an acoustic waveform



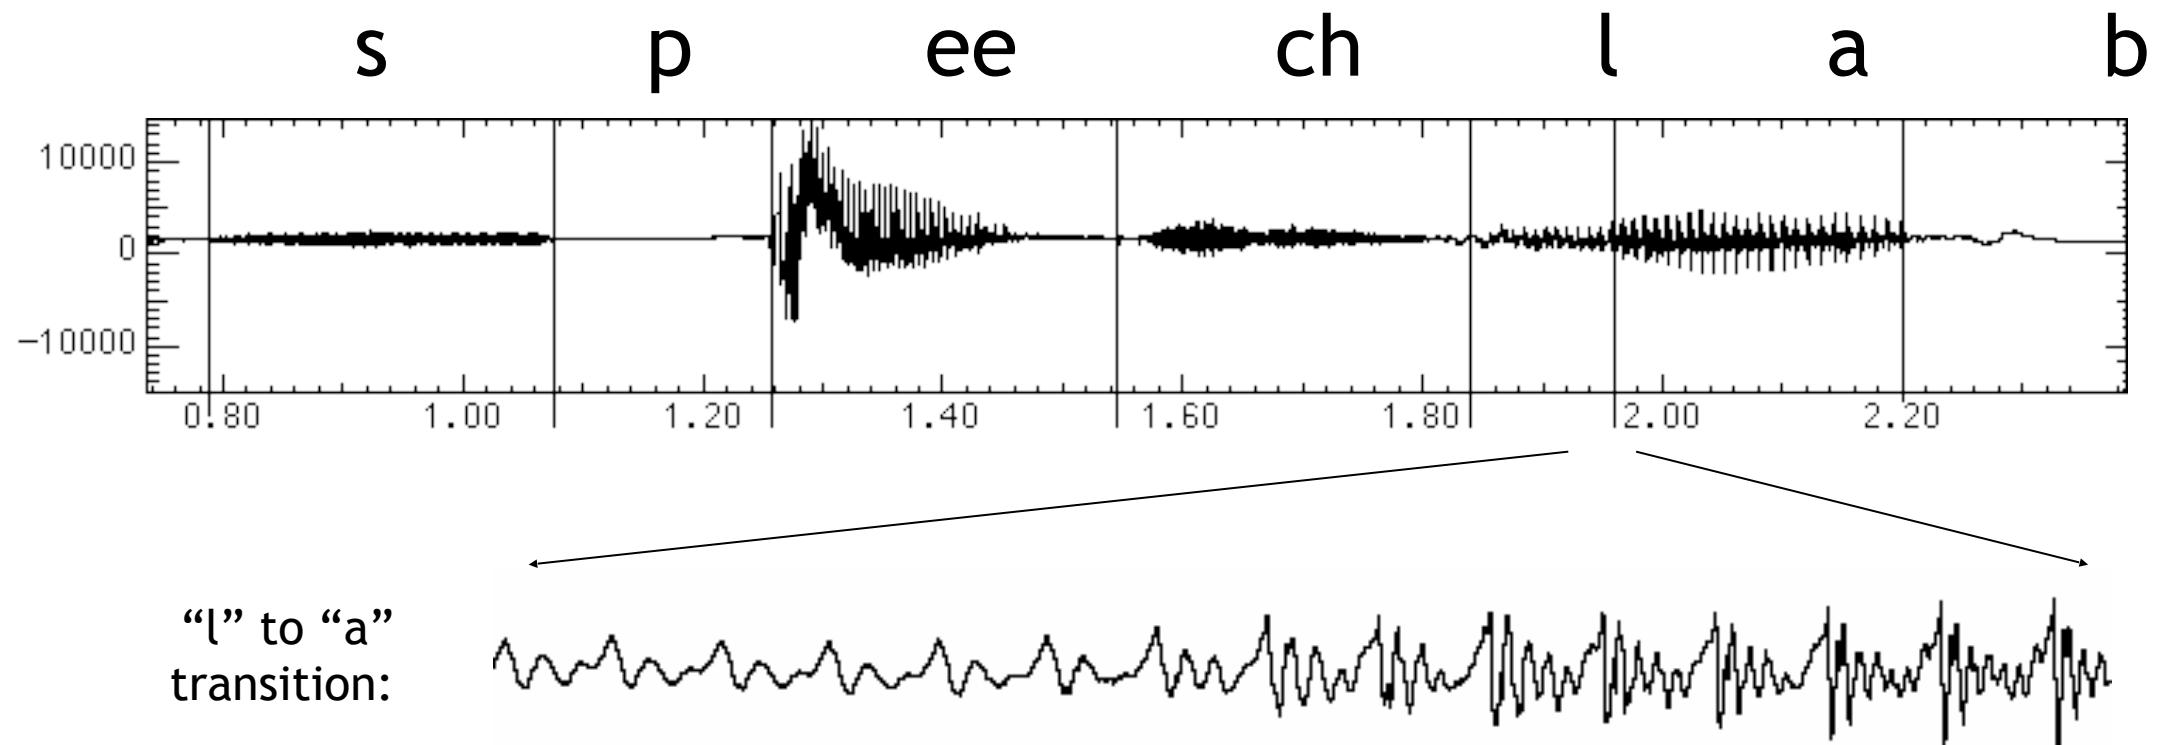
Speech in an Hour

- Speech input is an acoustic waveform



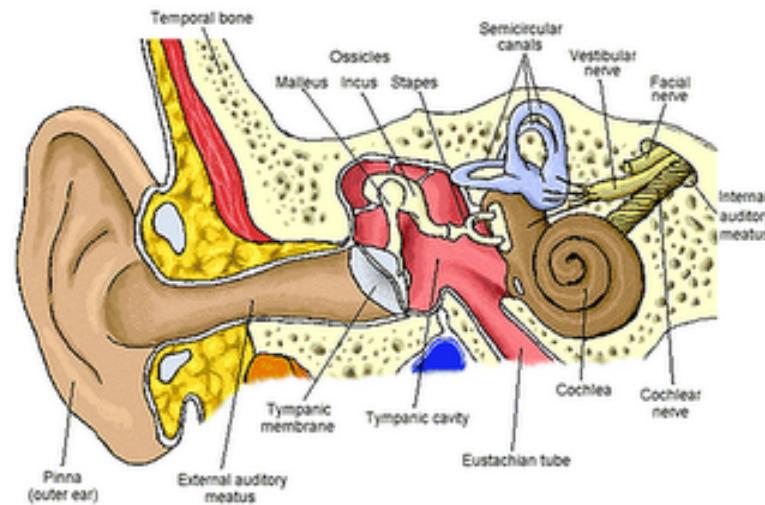
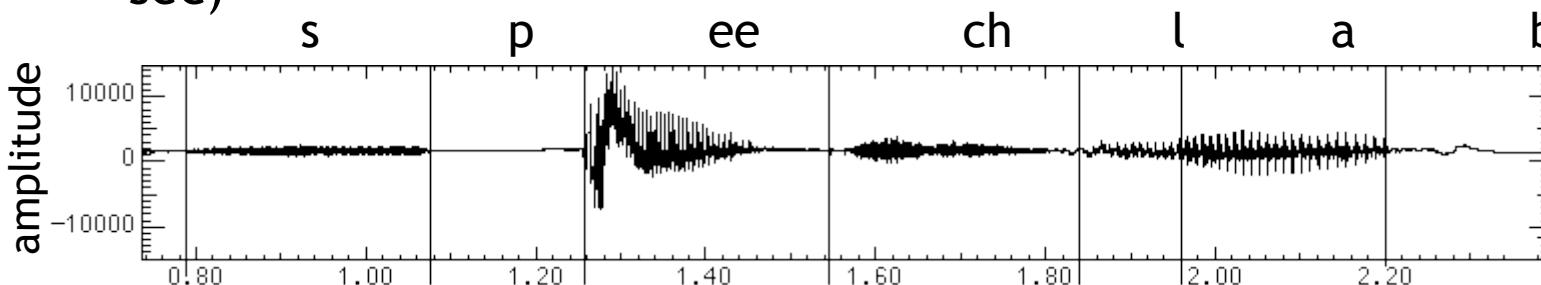
Speech in an Hour

- Speech input is an acoustic waveform

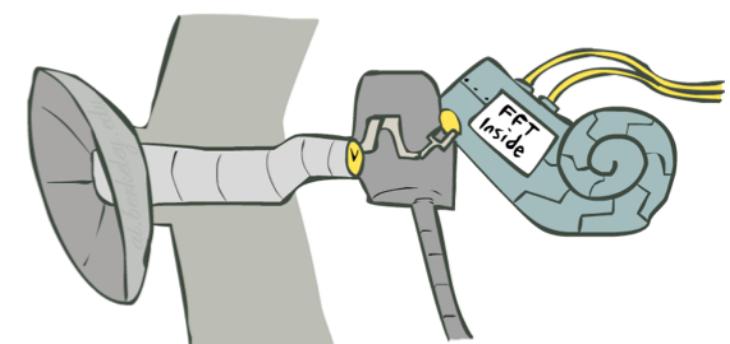
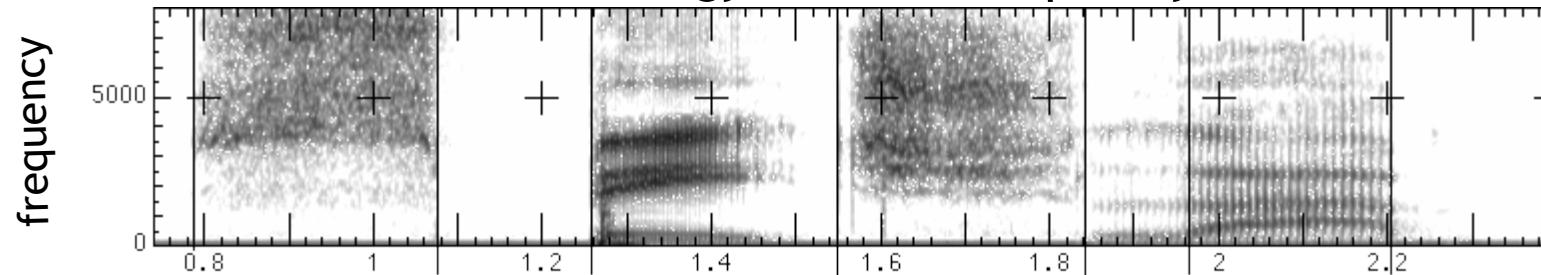


Spectral Analysis

- Frequency gives pitch; amplitude gives volume
 - Sampling at ~8 kHz (phone), ~16 kHz (mic) (kHz=1000 cycles/sec)

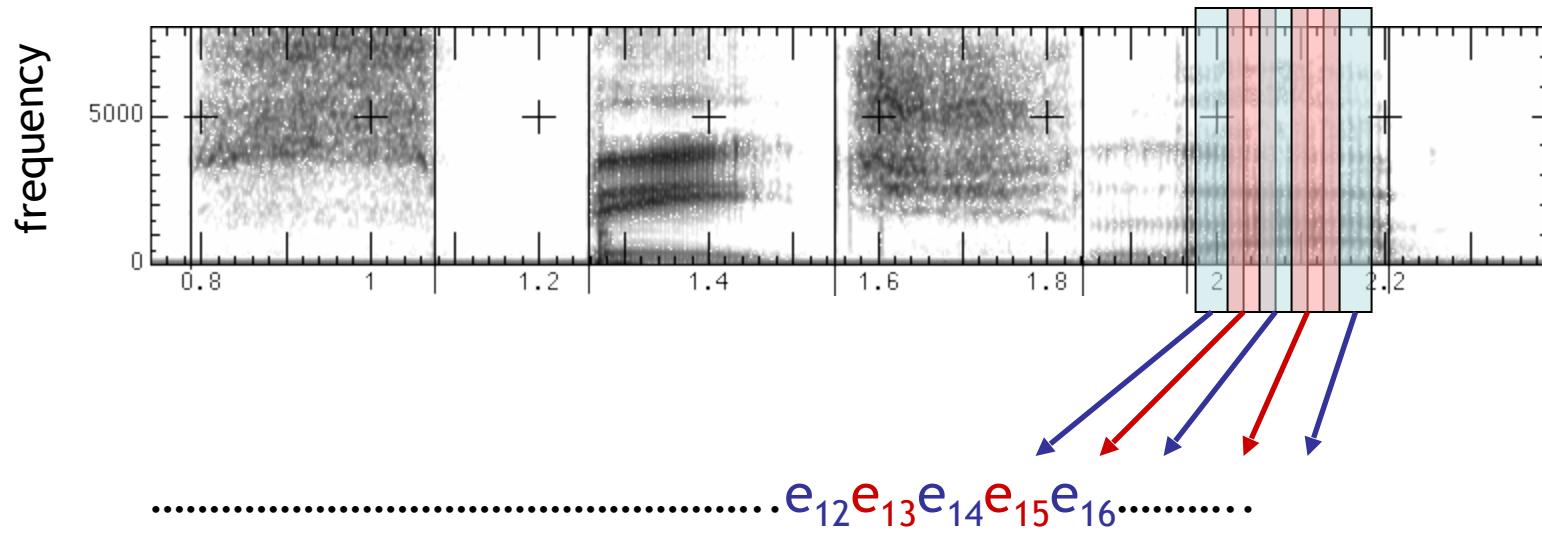


- Fourier transform of wave displayed as a spectrogram
 - Darkness indicates energy at each frequency



Acoustic Feature Sequence

- Time slices are translated into acoustic feature vectors (~39 real numbers per slice)

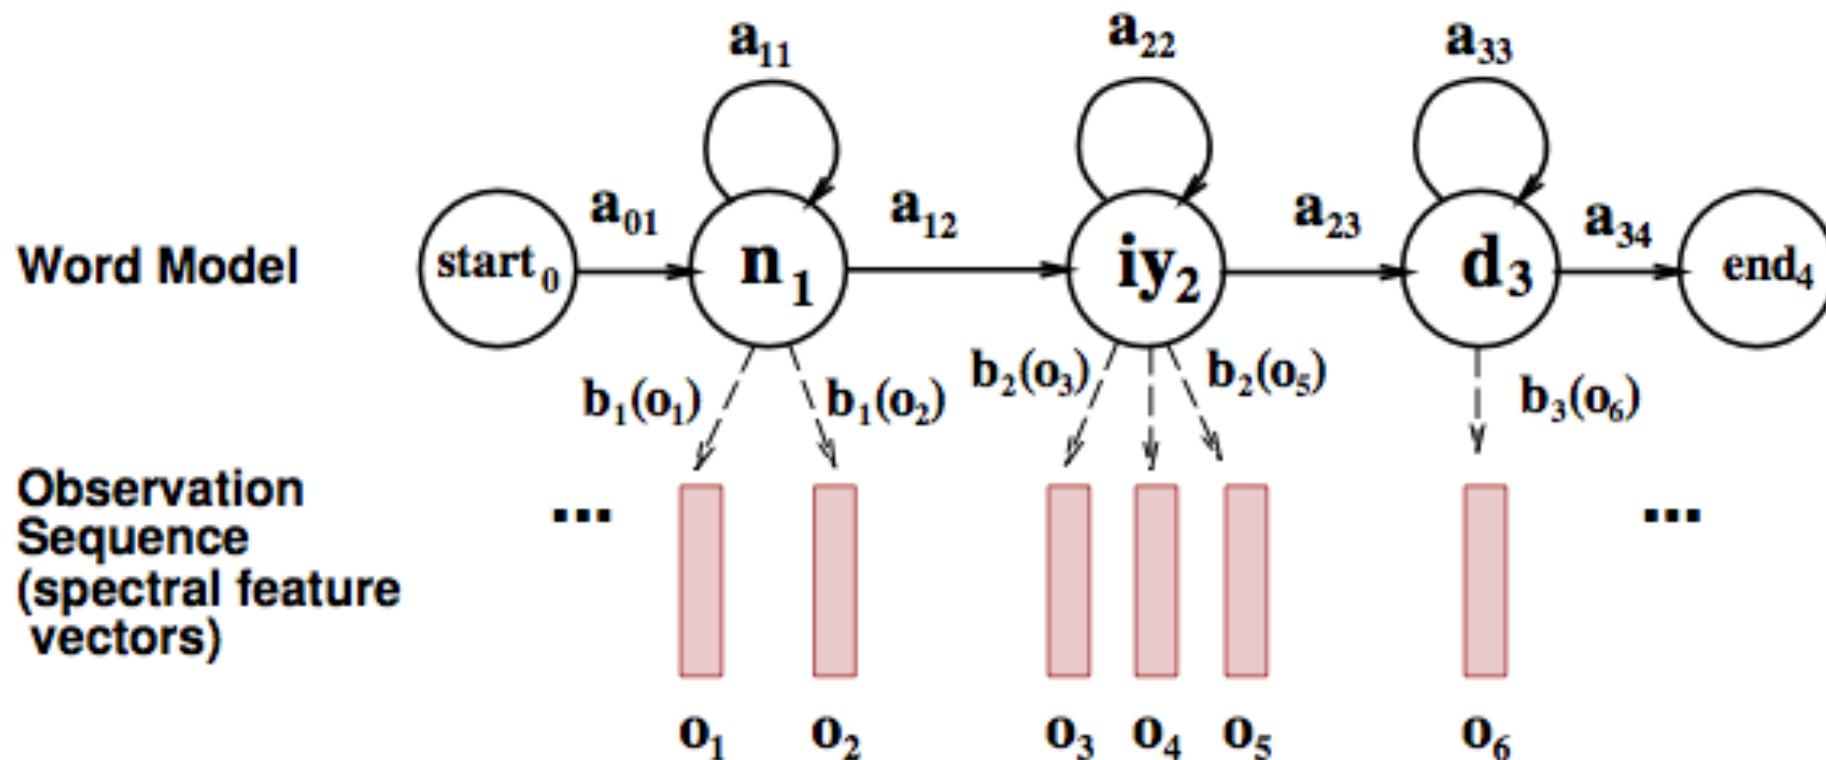


- These are the observations E , now we need the hidden states X

Speech State Space

- **HMM Specification**
 - $P(E|X)$ encodes which acoustic vectors are appropriate for each phoneme (each kind of sound)
 - $P(X|X')$ encodes how sounds can be strung together
- **State Space**
 - We will have one state for each sound in each word
 - Mostly, states advance sound by sound
 - Build a little state graph for each word and chain them together to form the state space X

States in a Word



Transitions with a Bigram Model

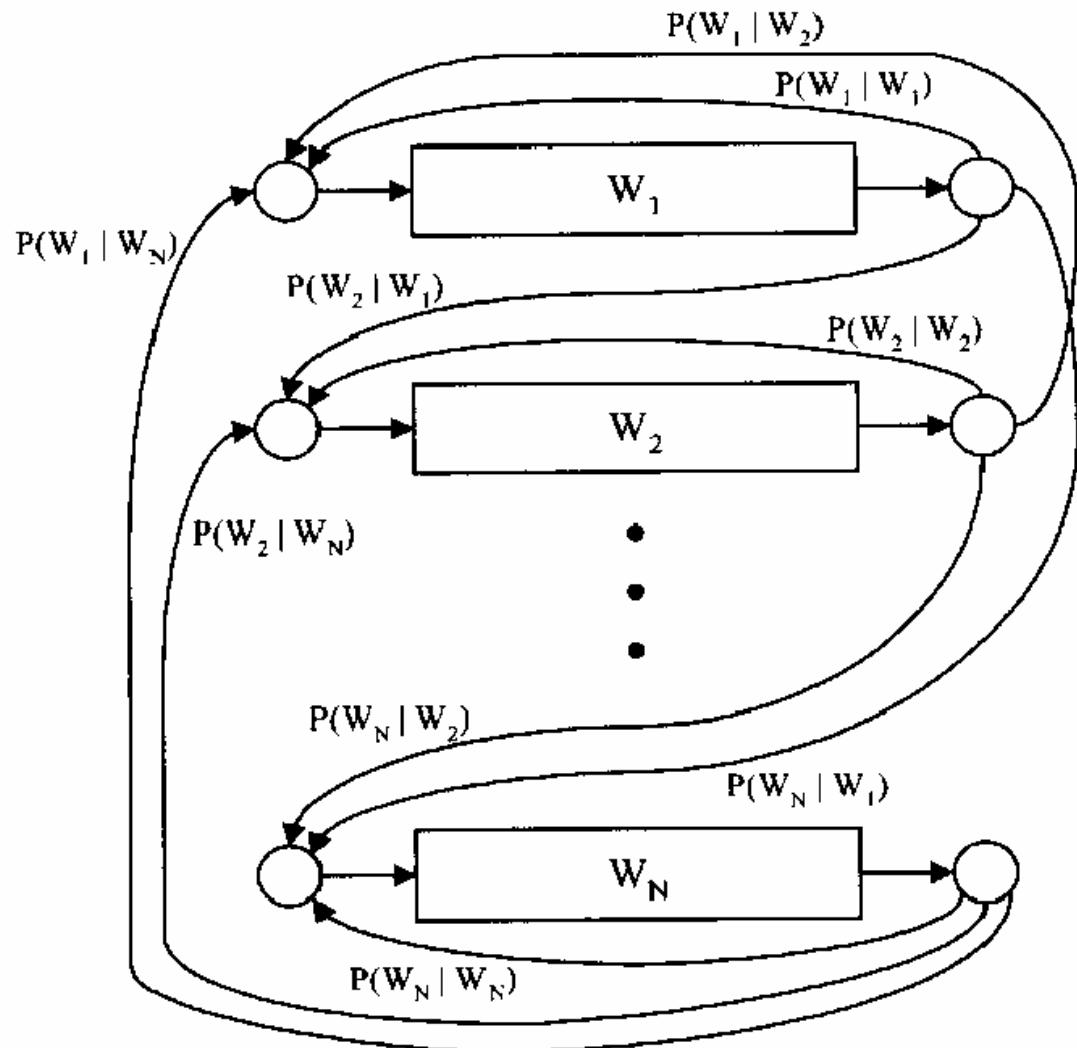
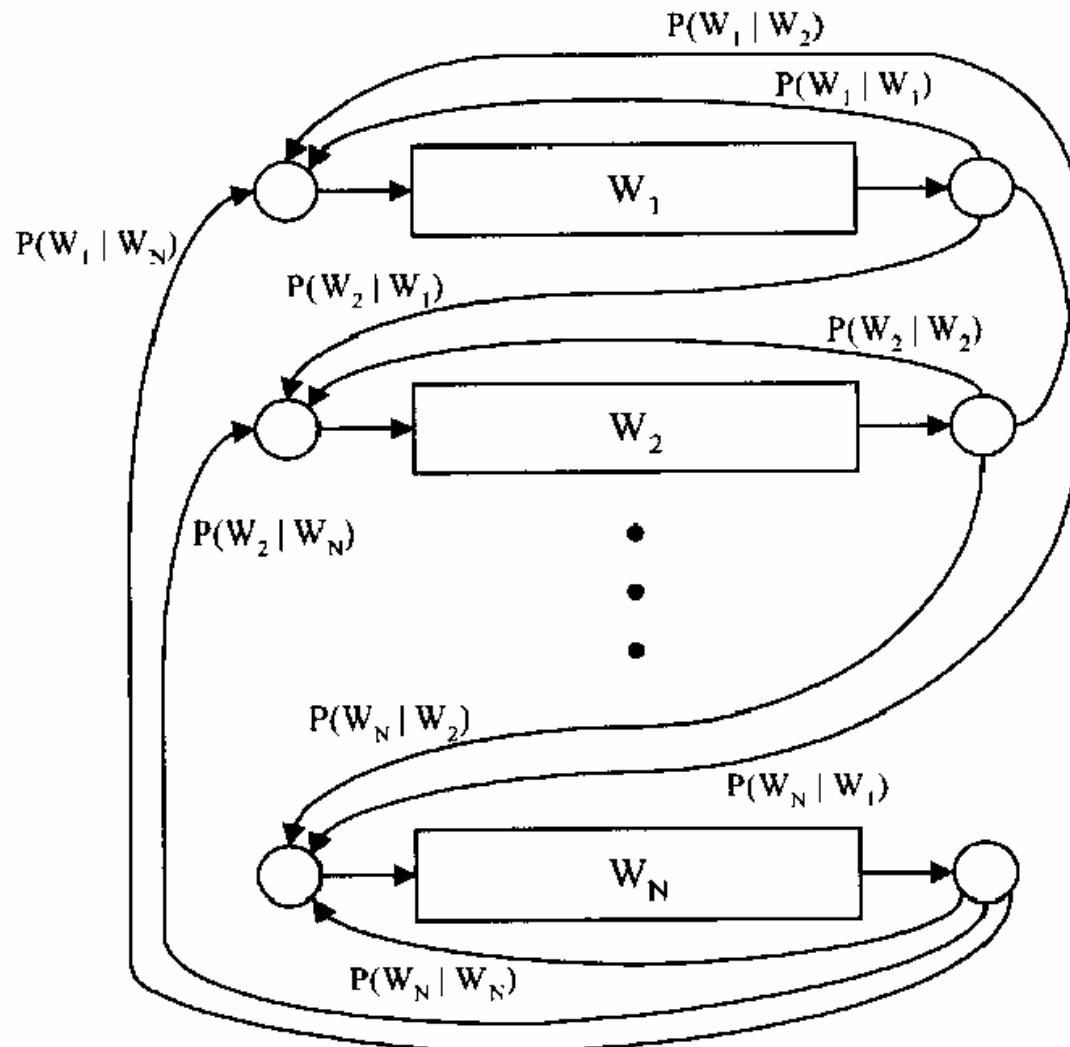


Figure: Huang et al, p. 61

Transitions with a Bigram Model

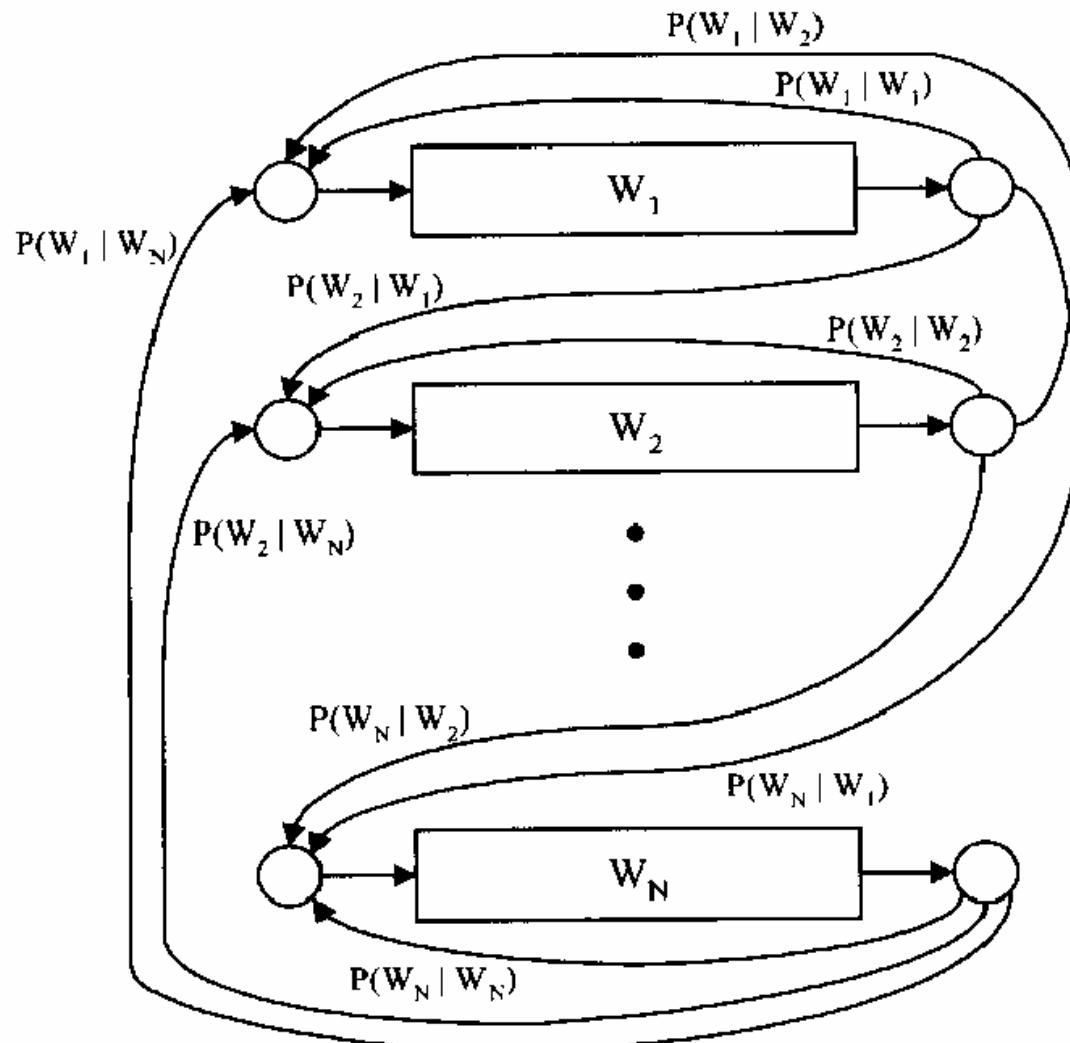


Training Counts

198015222	the first
194623024	the same
168504105	the following
158562063	the world
...	
14112454	the door

..	
23135851162	the *

Transitions with a Bigram Model



Training Counts

198015222	the first
194623024	the same
168504105	the following
158562063	the world
..	
14112454	the door

..	
23135851162	the *

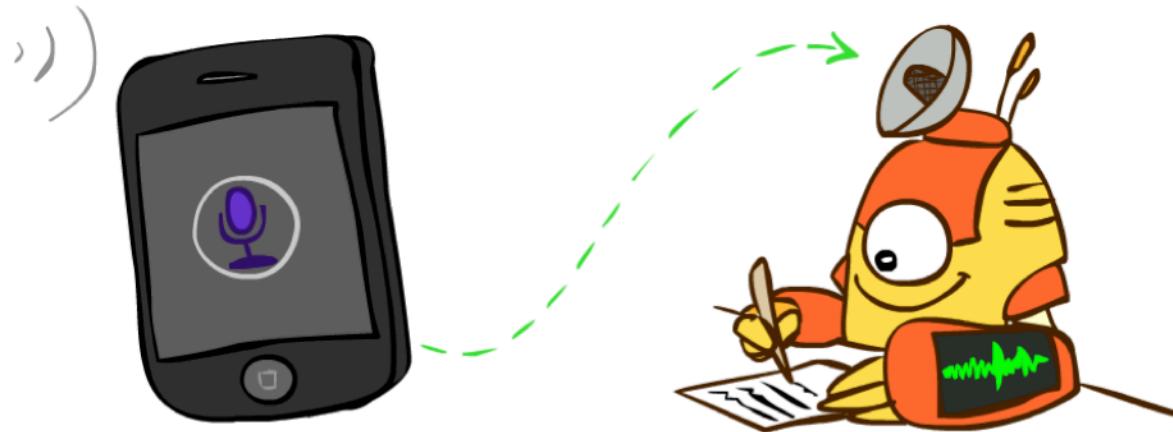
$$\hat{P}(\text{door}|\text{the}) = \frac{14112454}{23135851162}$$
$$= 0.0006$$

Decoding

- Finding the words given the acoustics is an HMM inference problem
- Which state sequence $x_{1:T}$ is most likely given the evidence $e_{1:T}$?

$$x_{1:T}^* = \arg \max_{x_{1:T}} P(x_{1:T}|e_{1:T}) = \arg \max_{x_{1:T}} P(x_{1:T}, e_{1:T})$$

- From the sequence x , we can simply read off the words



Next Time: Bayes' Nets!
