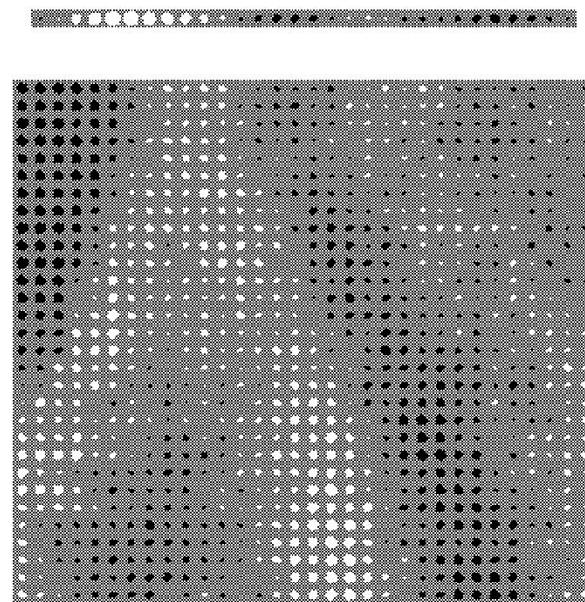
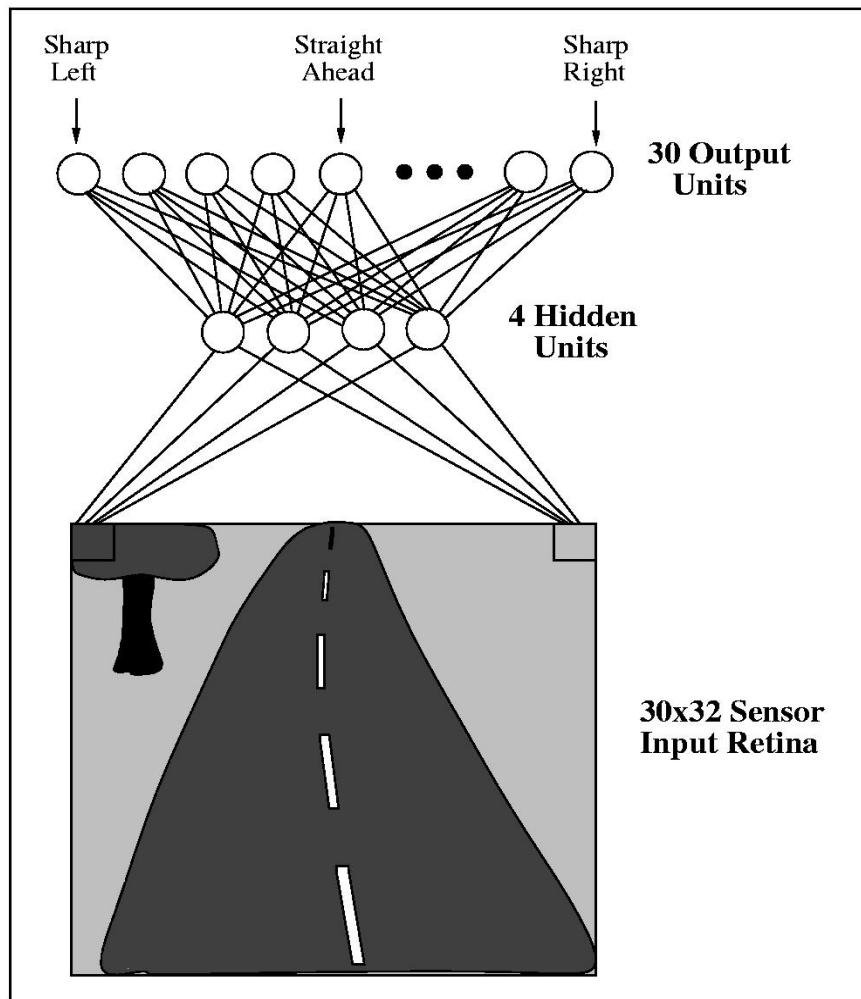


Neural Networks

Instructor: Alan Ritter

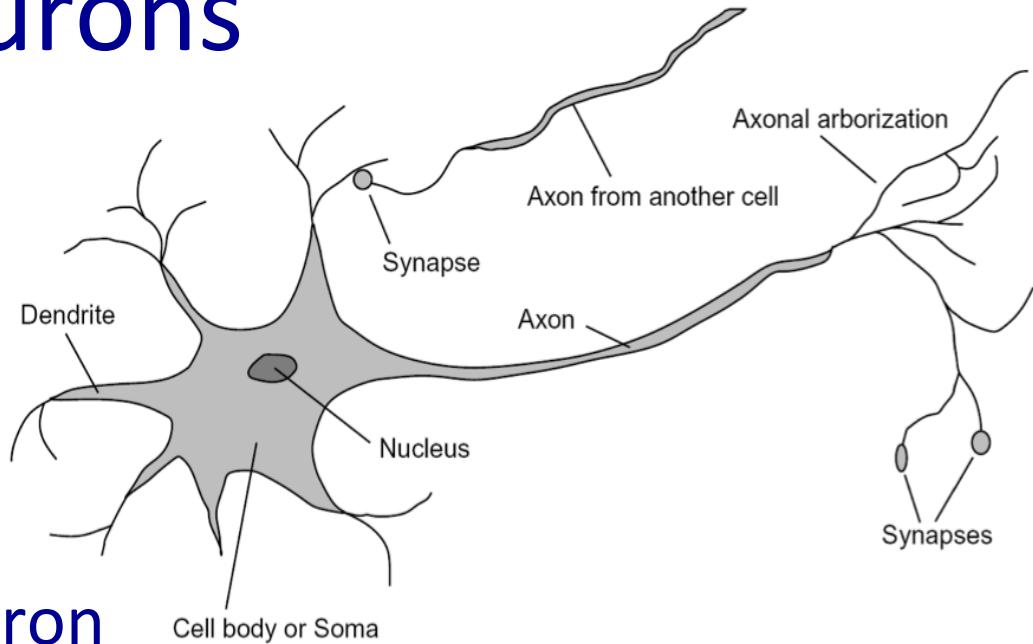
Slides Adapted from Carlos Guestrin and Richard Socher





Human Neurons

- Switching time
 - ~ 0.001 second
- Number of neurons
 - 10^{10}
- Connections per neuron
 - 10^{4-5}
- Scene recognition time
 - 0.1 seconds
- Number of cycles per scene recognition?
 - 100 → much parallel computation!



How to Extend Linear Models to Learn Nonlinear Functions?

- » Transform the input (e.g. kernel methods)

$$x \rightarrow \phi(x)$$

- » Q: how do we decide on the right representation?
 - Very high-dimensional representation?
 - Neural Networks: Let's learn the representation!

Representation Learning

$$y = f(x; \theta, w) = \phi(x; \theta) \cdot w$$

» 2 sets of parameters to learn:

- Linear model: w
- Data representation: θ

Training Data

$$(x_1, y_1)$$

$$(x_2, y_2)$$

$$(x_3, y_3)$$

:

Neural Networks

» Just function composition...

$$f(x) = f^{(3)}(f^{(2)}(f^{(1)}(x)))$$

Depth=3 in this case

Third Layer

Second Layer

First Layer

The diagram illustrates the depth of a neural network. The equation $f(x) = f^{(3)}(f^{(2)}(f^{(1)}(x)))$ shows three layers of function composition. Three light blue callout boxes point to each term: 'First Layer' points to $f^{(1)}(x)$, 'Second Layer' points to $f^{(2)}(f^{(1)}(x))$, and 'Third Layer' points to $f^{(3)}(f^{(2)}(f^{(1)}(x)))$. A purple box labeled 'Depth=3 in this case' is positioned above the equation.

Needs Some Nonlinearity

$$f(x) = f^{(2)}(f^{(1)}(x))$$

» Q: what if both functions are all linear?

$$\begin{aligned} h &= f^{(1)}(x) = W \cdot x \\ y &= f^{(2)}(h) = w \cdot h \end{aligned}$$

$$y = f(x) = w \cdot f^{(1)}(x) = w \cdot h = w^T \cdot W \cdot x = w' \cdot x$$

Needs Some Nonlinearity

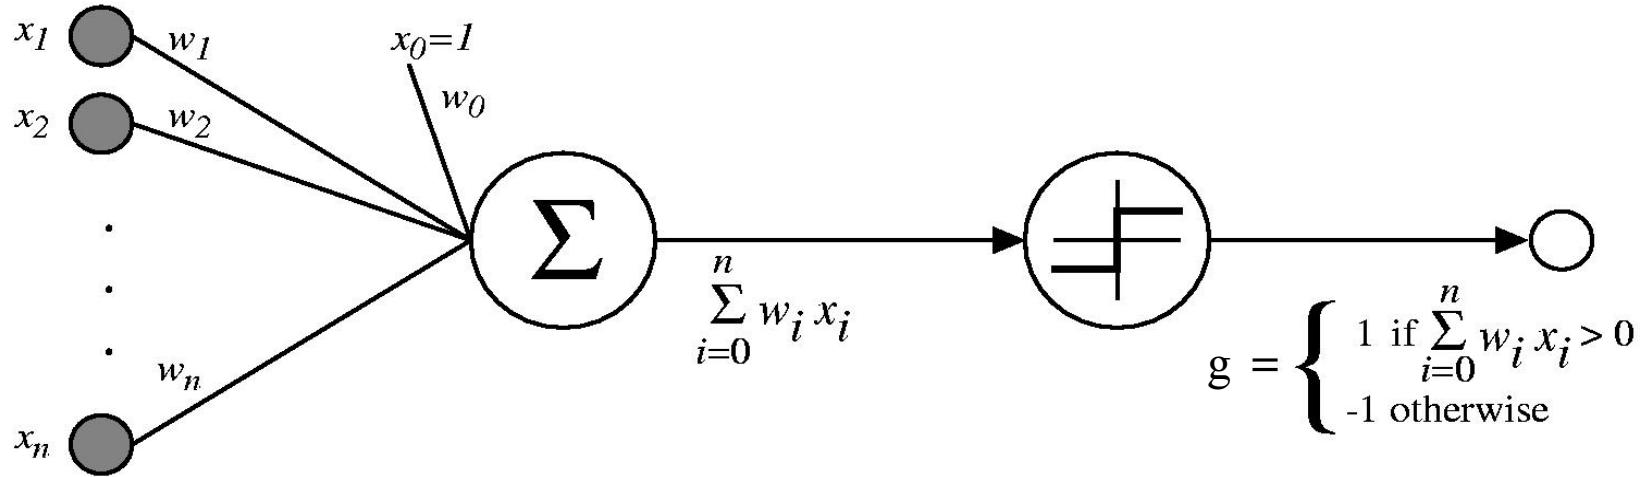
$$f(x) = f^{(1)}(f^2(x))$$

- » Typically use some non-linear transformation (e.g. element-wise sigmoid)

$$h = f^{(1)}(x) = \textcolor{red}{g}(W \cdot x)$$

$$y = f^{(2)}(h) = w \cdot h$$

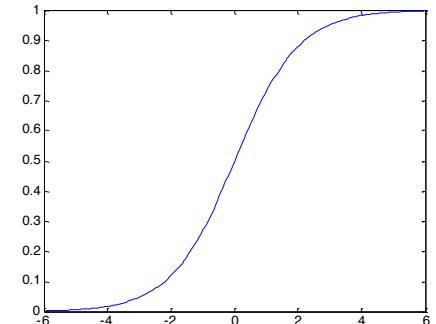
Perceptron as a Neural Network



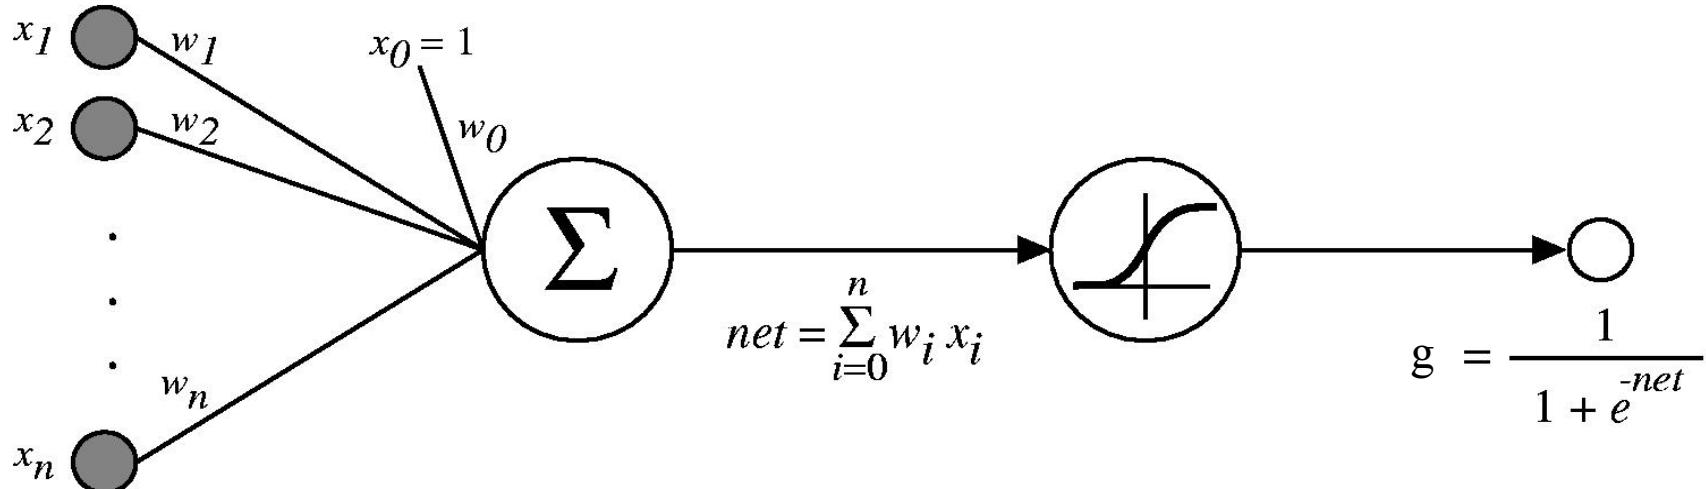
This is one neuron:

- Input edges $x_1 \dots x_n$, along with basis
- The sum is represented graphically
- Sum passed through an activation function g

Sigmoid Neuron



$$g(w_0 + \sum_i w_i x_i) = \frac{1}{1 + e^{-(w_0 + \sum_i w_i x_i)}}$$



Just change g !

- Why would we want to do this?
- Notice new output range $[0, 1]$. What was it before?
- Look familiar?

Optimizing a neuron

$$\frac{\partial}{\partial x} f(g(x)) = f'(g(x))g'(x)$$

We train to minimize sum-squared error

$$\ell(W) = \frac{1}{2} \sum_j [y^j - g(w_0 + \sum_i w_i x_i^j)]^2$$

$$\frac{\partial l}{\partial w_i} = - \sum_j [y_j - g(w_0 + \sum_i w_i x_i^j)] \frac{\partial}{\partial w_i} g(w_0 + \sum_i w_i x_i^j)$$

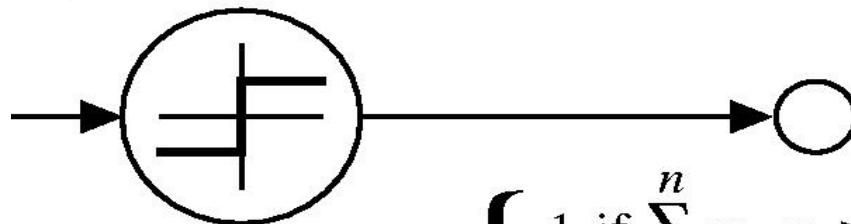
$$\frac{\partial}{\partial w_i} g(w_0 + \sum_i w_i x_i^j) = x_i^j g'(w_0 + \sum_i w_i x_i^j)$$

$$\frac{\partial \ell(W)}{\partial w_i} = - \sum_j [y^j - g(w_0 + \sum_i w_i x_i^j)] x_i^j g'(w_0 + \sum_i w_i x_i^j)$$

Solution just depends on g' : derivative of activation function!

Re-deriving the perceptron update

$$\frac{\partial \ell(W)}{\partial w_i} = - \sum_j [y^j - g(w_0 + \sum_i w_i x_i^j)] x_i^j g'(w_0 + \sum_i w_i x_i^j)$$



=0
everywhere :(

$$g = \begin{cases} 1 & \text{if } \sum_{i=0}^n w_i x_i > 0 \\ -1 & \text{otherwise} \end{cases}$$

$$\frac{\partial \ell(W)}{\partial w_i} = - \sum_j [y^j - g(w_0 + \sum_i w_i x_i^j)] x_i^j$$

For a specific, incorrect example:

- $w = w + y^*x$ (our familiar update!)

Sigmoid units: have to differentiate g

$$\frac{\partial \ell(W)}{\partial w_i} = - \sum_j [y^j - g(w_0 + \sum_i w_i x_i^j)] x_i^j g'(w_0 + \sum_i w_i x_i^j)$$

$$g(x) = \frac{1}{1 + e^{-x}} \quad g'(x) = g(x)(1 - g(x))$$

$$w_i \leftarrow w_i + \eta \sum_j x_i^j \delta^j$$

$$\delta^j = [y^j - g(w_0 + \sum_i w_i x_i^j)] g^j (1 - g^j)$$

$$g^j = g(w_0 + \sum_i w_i x_i^j)$$

Aside: Comparison to logistic regression

- $P(Y|X)$ represented by:

$$\begin{aligned} P(Y = 1 \mid x, W) &= \frac{1}{1 + e^{-(w_0 + \sum_i w_i x_i)}} \\ &= g(w_0 + \sum_i w_i x_i) \end{aligned}$$

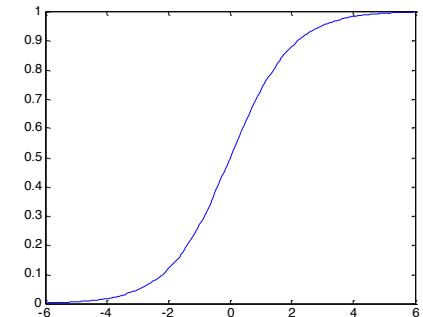
- Learning rule – MLE:

$$\frac{\partial \ell(W)}{\partial w_i} = \sum_j x_i^j [y^j - P(Y^j = 1 \mid x^j, W)]$$

$$= \sum_j x_i^j [y^j - g(w_0 + \sum_i w_i x_i^j)]$$

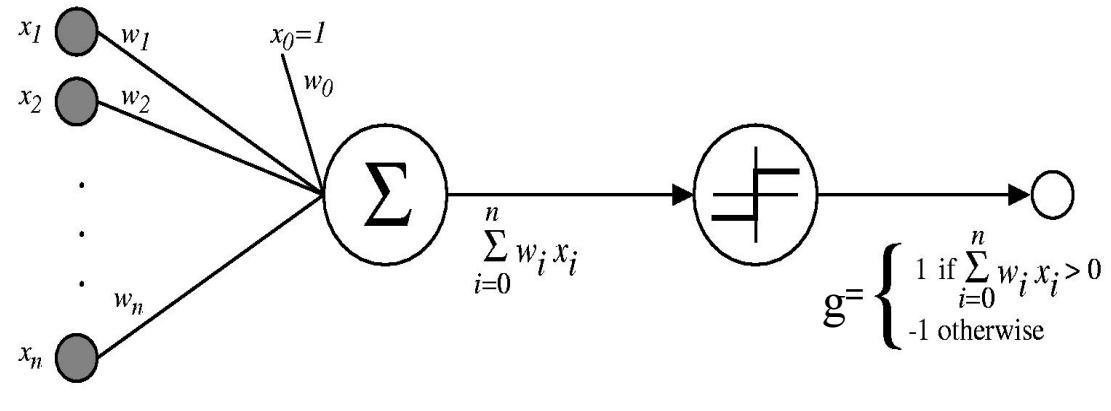
$$w_i \leftarrow w_i + \eta \sum_j x_i^j \delta^j$$

$$\delta^j = y^j - g(w_0 + \sum_i w_i x_i^j)$$



Perceptron, linear classification, Boolean functions: $x_i \in \{0,1\}$

- Can learn $x_1 \vee x_2$
 - $0.5 + x_1 + x_2$
- Can learn $x_1 \wedge x_2$
 - $-1.5 + x_1 + x_2$
- Can learn any conjunction or disjunction?
 - $0.5 + x_1 + \dots + x_n$
 - $(n-0.5) + x_1 + \dots + x_n$
- Can learn majority?
 - $(-0.5*n) + x_1 + \dots + x_n$
- What are we missing? The dreaded XOR!, etc.



Going beyond linear classification

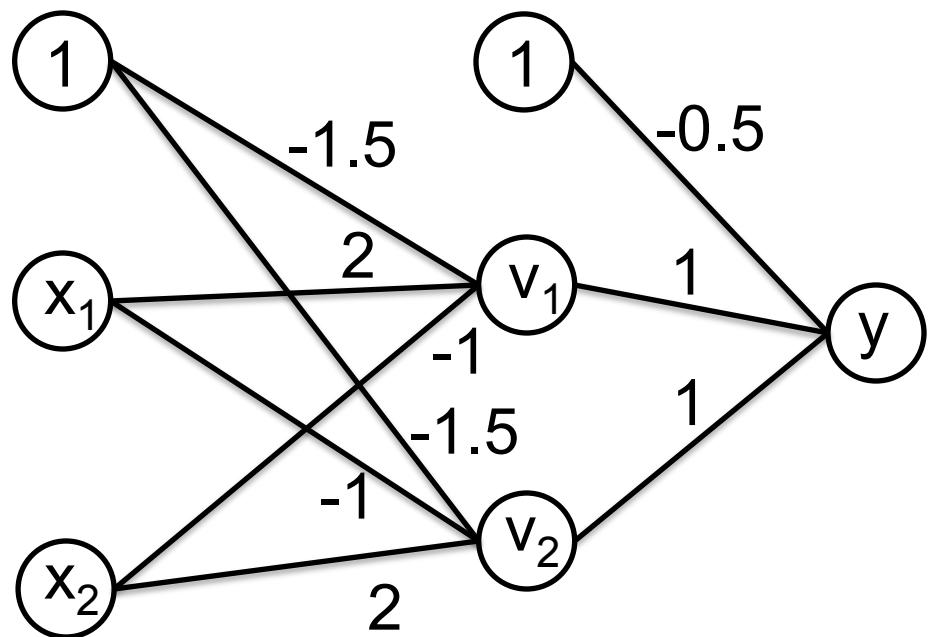
Solving the XOR problem

$$y = x_1 \text{ XOR } x_2 = (x_1 \wedge \neg x_2) \vee (x_2 \wedge \neg x_1)$$

$$\begin{aligned}v_1 &= (x_1 \wedge \neg x_2) \\&= -1.5 + 2x_1 - x_2\end{aligned}$$

$$\begin{aligned}v_2 &= (x_2 \wedge \neg x_1) \\&= -1.5 + 2x_2 - x_1\end{aligned}$$

$$\begin{aligned}y &= v_1 \vee v_2 \\&= -0.5 + v_1 + v_2\end{aligned}$$



Hidden layer

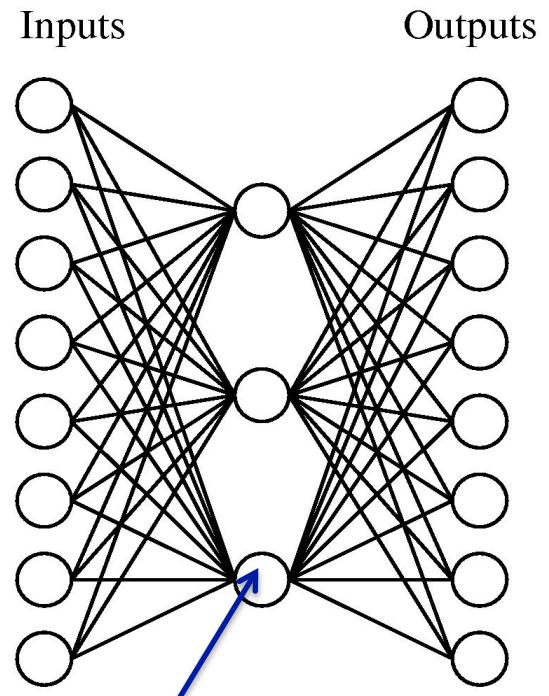
- Single unit:

$$out(\mathbf{x}) = g(w_0 + \sum_i w_i x_i)$$

- 1-hidden layer:

$$out(\mathbf{x}) = g \left(w_0 + \sum_k w_k g(w_0^k + \sum_i w_i^k x_i) \right)$$

- No longer convex function!

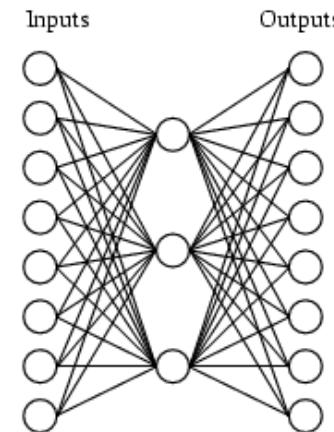


Example data for NN with hidden layer

Autoencoder
(kind of unsupervised learning)

Dimensionality Reduction

A target function:

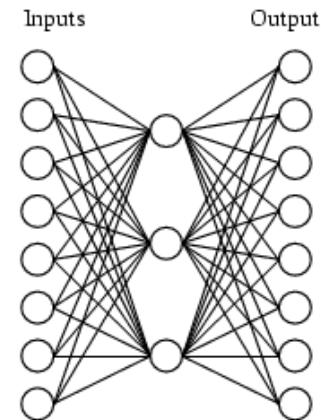


Input	Output
10000000	→ 10000000
01000000	→ 01000000
00100000	→ 00100000
00010000	→ 00010000
00001000	→ 00001000
00000100	→ 00000100
00000010	→ 00000010
00000001	→ 00000001

Can this be learned??

Learned weights for hidden layer

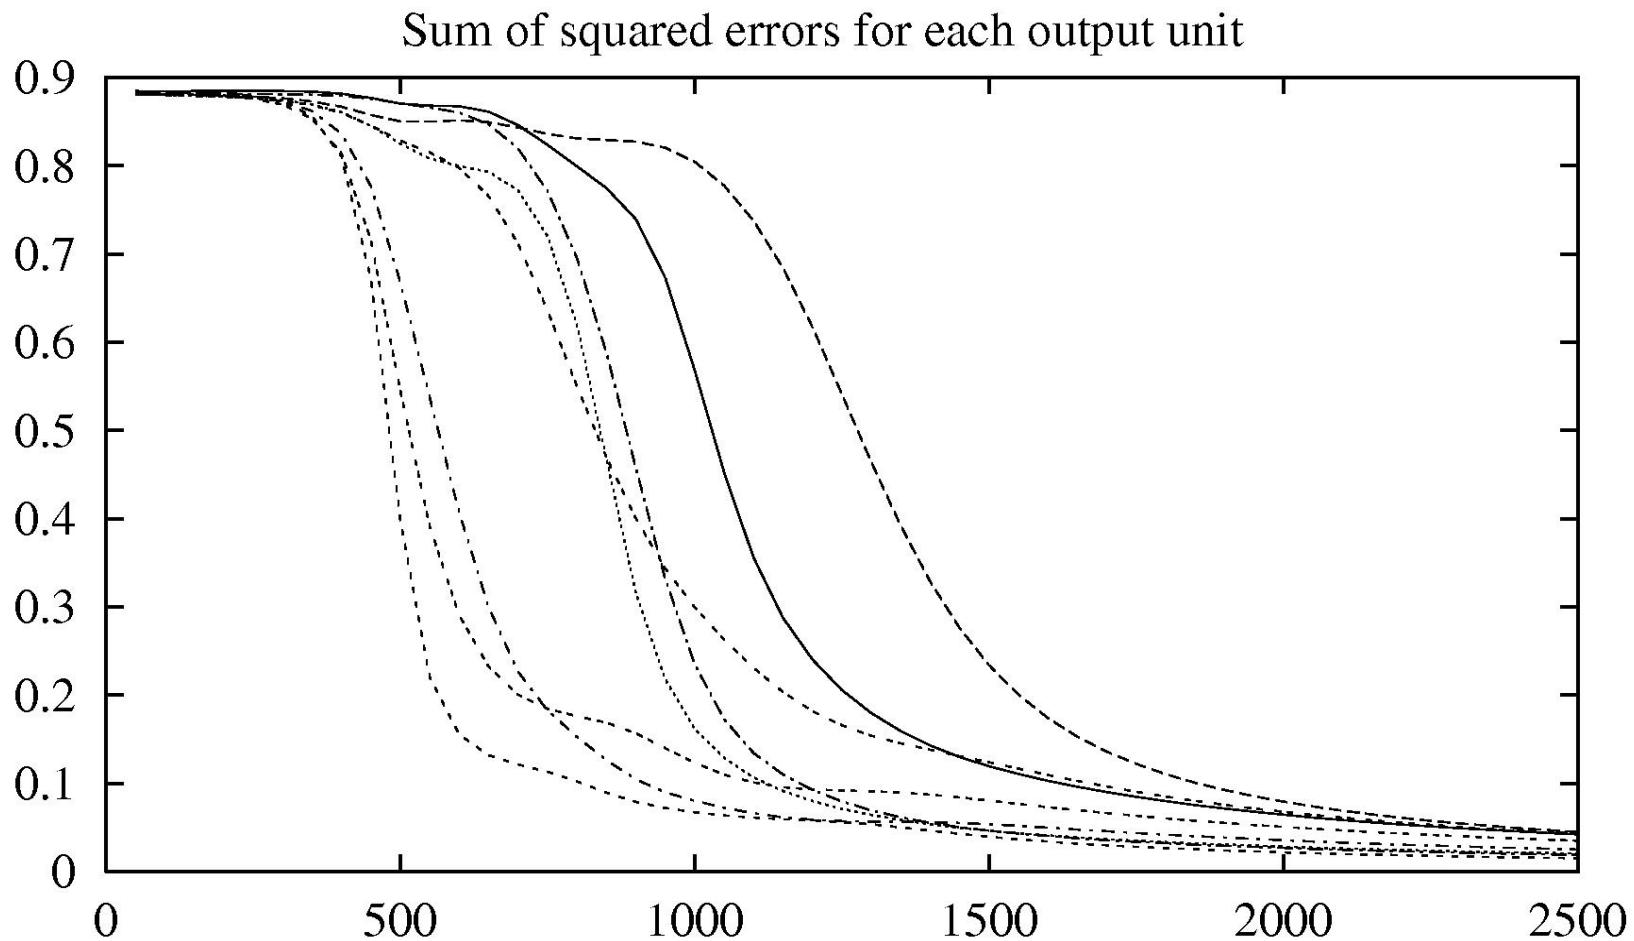
A network:



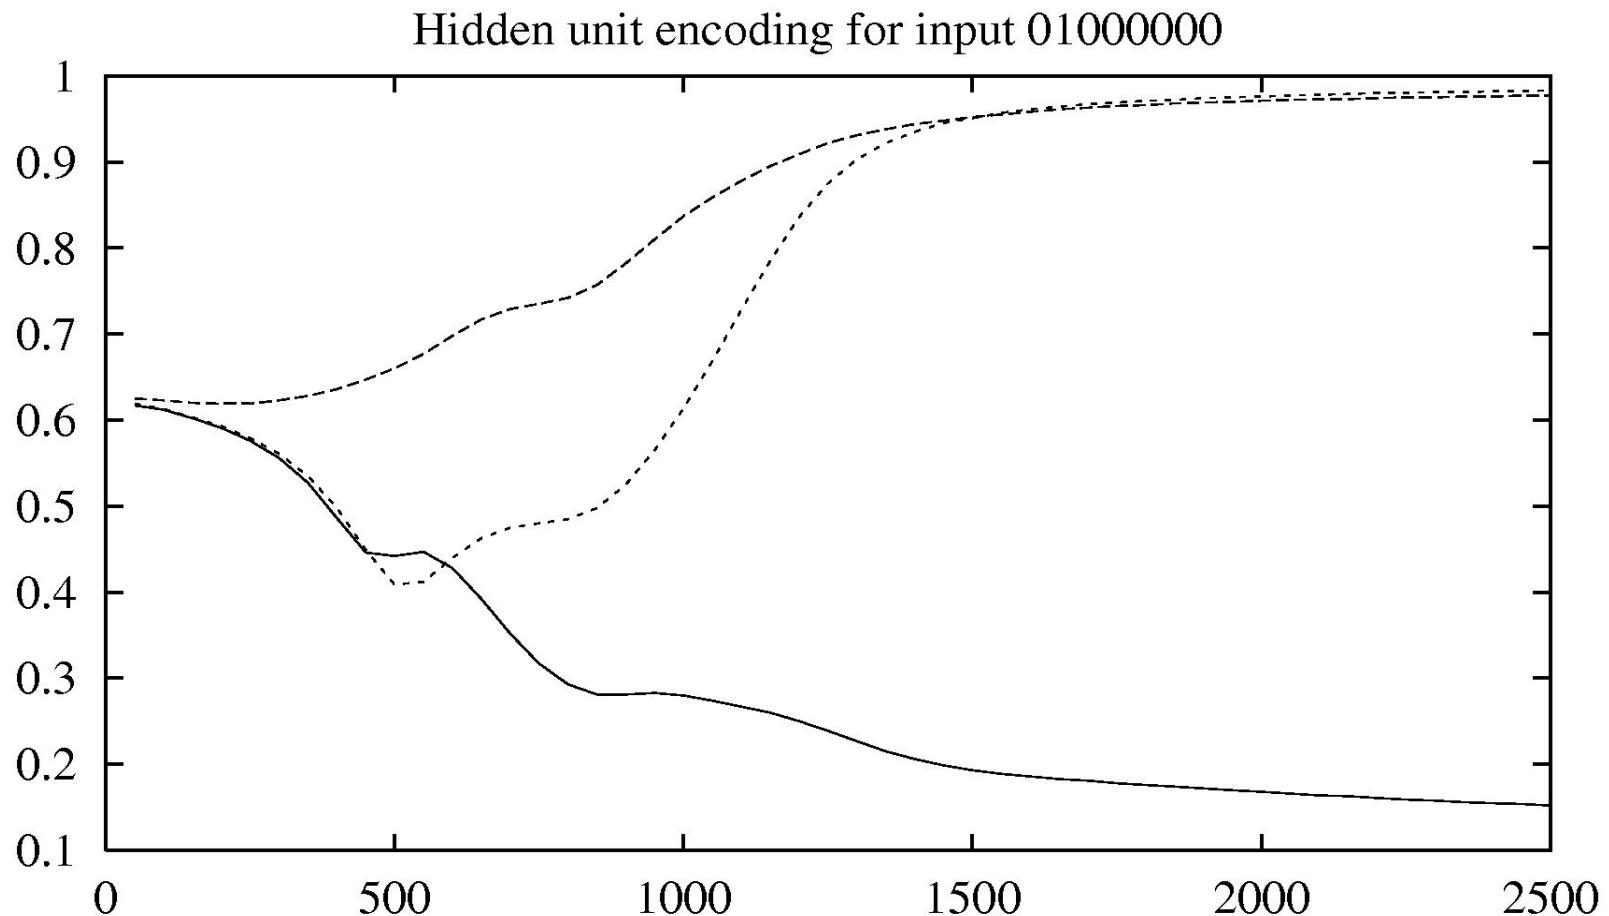
Learned hidden layer representation:

Input	Hidden Values	Output
10000000	→ .89 .04 .08	→ 10000000
01000000	→ .01 .11 .88	→ 01000000
00100000	→ .01 .97 .27	→ 00100000
00010000	→ .99 .97 .71	→ 00010000
00001000	→ .03 .05 .02	→ 00001000
00000100	→ .22 .99 .99	→ 00000100
00000010	→ .80 .01 .98	→ 00000010
00000001	→ .60 .94 .01	→ 00000001

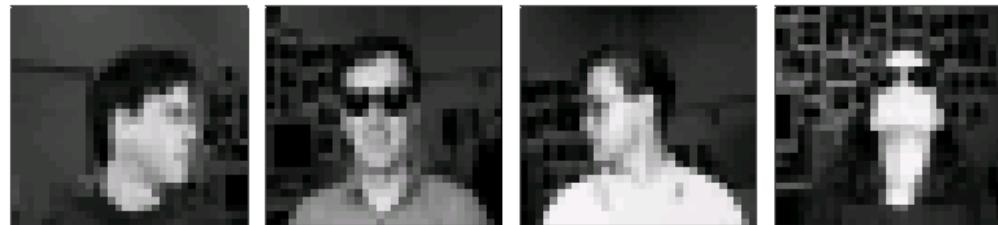
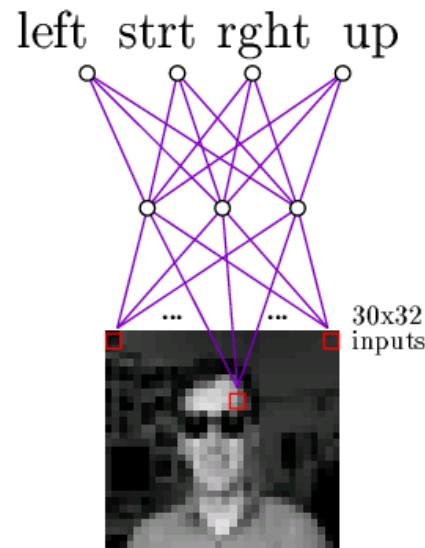
Learning the weights



Learning an encoding



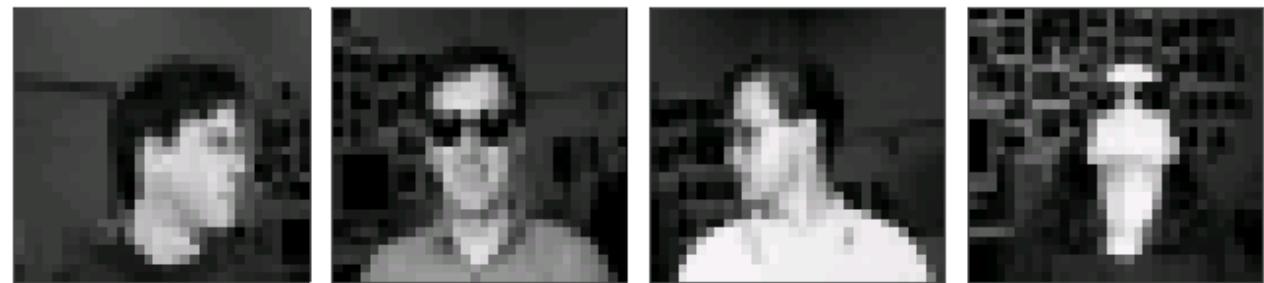
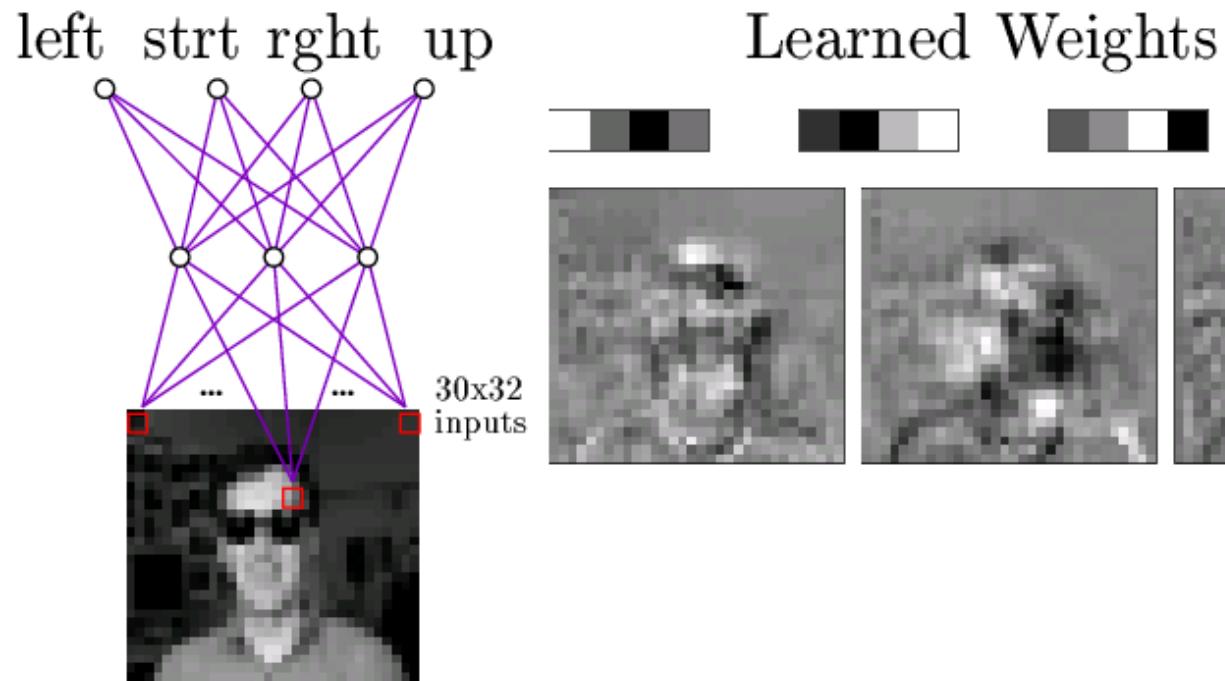
NN for images



Typical input images

90% accurate learning head pose, and recognizing 1-of-20 faces

Weights in NN for images



Typical input images

Forward propagation

1-hidden layer:

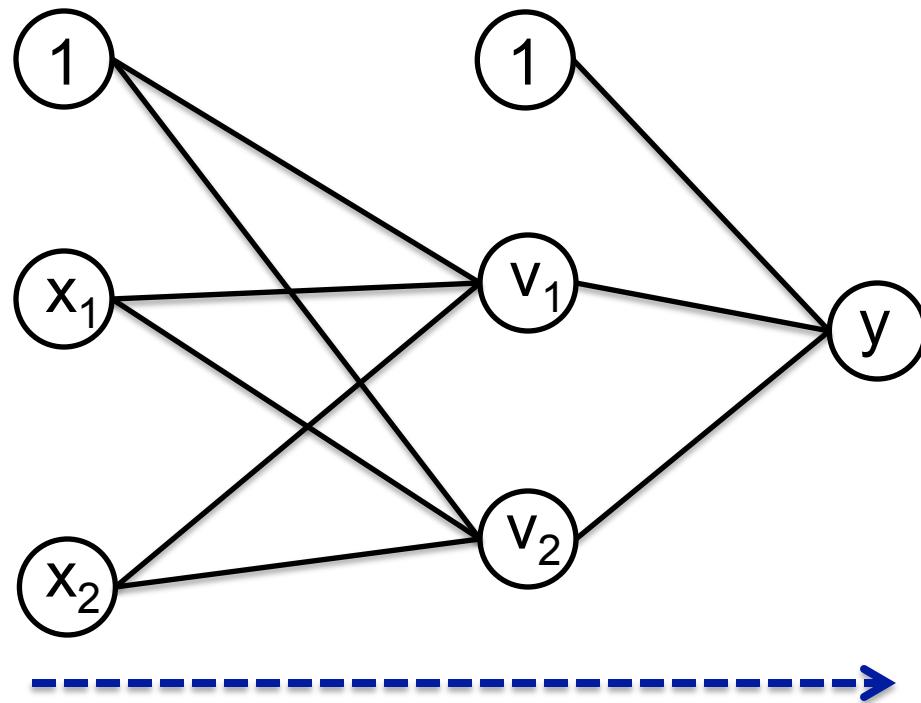
$$out(\mathbf{x}) = g \left(w_0 + \sum_k w_k g(w_0^k + \sum_i w_i^k x_i) \right)$$

k sums over hidden layer

i sums over input layer

Compute values left to right

1. Inputs: x_1, \dots, x_n
2. Hidden: v_1, \dots, v_n
3. Output: y



Gradient descent for 1-hidden layer

$$\frac{\partial \ell(W)}{\partial w_k}$$

Dropped w_0 to make derivation simpler

$$\ell(W) = \frac{1}{2} \sum_j [y^j - out(\mathbf{x}^j)]^2$$

$$out(\mathbf{x}) = g \left(\sum_{k'} w_{k'} g \left(\sum_{i'} w_{i'}^{k'} x_{i'} \right) \right)$$

$$v_k^j = g \left(\sum_{i'} w_{i'}^{k'} x_{i'} \right)$$

$$\frac{\partial \ell(W)}{\partial w_k} = \sum_{j=1}^m -[y^j - out(\mathbf{x}^j)] \frac{\partial out(\mathbf{x}^j)}{\partial w_k}$$

$$out(x) = g \left(\sum_{k'} w_{k'} v_k^j \right)$$

$$\frac{\partial out(\mathbf{x})}{\partial w_k} = v_k^j g' \left(\sum_{k'} w_{k'} v_k^j \right)$$



Gradient for last layer same as the single node case, but with hidden nodes v as input!

Gradient descent for 1-hidden layer

$$\frac{\partial \ell(W)}{\partial w_i^k}$$

$$\ell(W) = \frac{1}{2} \sum_j [y^j - \text{out}(\mathbf{x}^j)]^2$$

$$\text{out}(\mathbf{x}) = g \left(\sum_{k'} w_{k'} g \left(\sum_{i'} w_{i'}^{k'} x_{i'} \right) \right)$$

Dropped w_0 to make derivation simpler

$$\frac{\partial}{\partial x} f(g(x)) = f'(g(x))g'(x)$$

$$\frac{\partial \ell(W)}{\partial w_i^k} = \sum_{j=1}^m -[y - \text{out}(\mathbf{x}^j)] \frac{\partial \text{out}(\mathbf{x}^j)}{\partial w_i^k}$$

$$\frac{\partial \text{out}(\mathbf{x})}{\partial w_i^k} = g' \left(\sum_{k'} w_{k'} g \left(\sum_{i'} w_{i'}^{k'} x_{i'} \right) \right) \frac{\partial}{\partial w_i^k} g \left(\sum_{i'} w_{i'}^{k'} x_{i'} \right) w_k$$

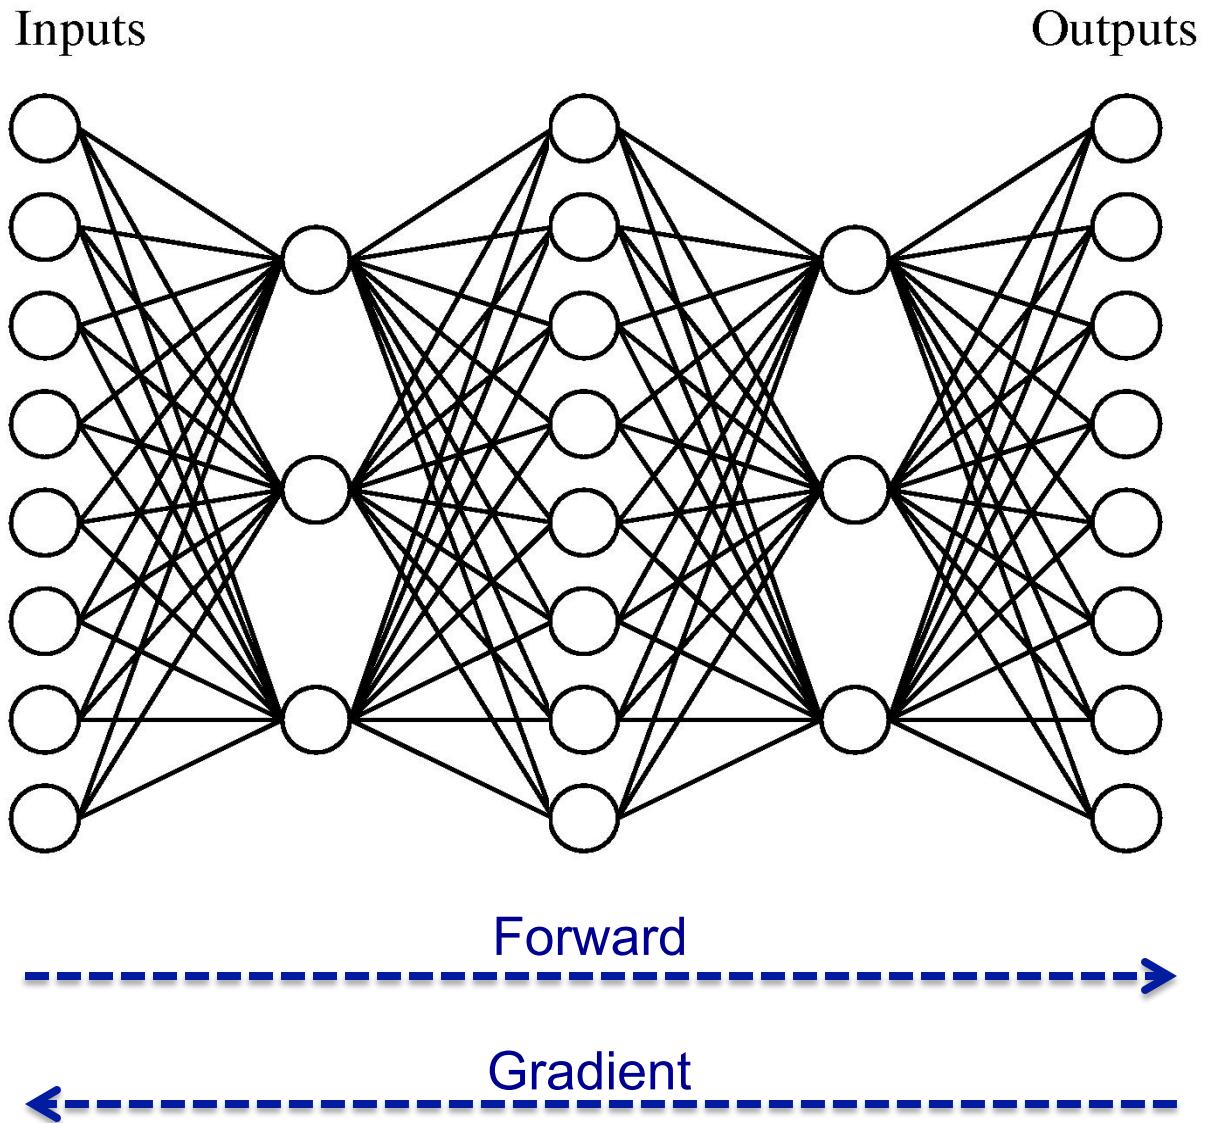

For hidden layer,
two parts:

- Normal update for single neuron
- Recursive computation of gradient on output layer

Multilayer neural networks

Inference and Learning:

- Forward pass: left to right, each hidden layer in turn
- Gradient computation: right to left, propagating gradient for each node



Forward propagation – prediction

- Recursive algorithm
- Start from input layer
- Output of node V_k with parents U_1, U_2, \dots :

$$V_k = g \left(\sum_i w_i^k U_i \right)$$

Back-propagation – learning

- Just gradient descent!!!
- Recursive algorithm for computing gradient
- For each example
 - Perform forward propagation
 - Start from output layer
 - Compute gradient of node V_k with parents U_1, U_2, \dots
 - Update weight w_i^k
 - Repeat (move to preceding layer)

Back-propagation – pseudocode

Initialize all weights to small random numbers

- Until convergence, do:
 - For each training example x, y :
 1. Forward propagation, compute node values V_k
 2. For each output unit o (with labeled output y):
$$\delta_o = V_o(1-V_o)(y-V_o)$$
 3. For each hidden unit h :
$$\delta_h = V_h(1-V_h) \sum_{k \text{ in output}(h)} w_{h,k} \delta_k$$
 4. Update each network weight $w_{i,j}$ from node i to node j
$$w_{i,j} = w_{i,j} + \eta \delta_j x_{i,j}$$

Convergence of backprop

- Perceptron leads to convex optimization
 - Gradient descent reaches **global minima**
- Multilayer neural nets **not convex**
 - Gradient descent gets stuck in local minima
 - Selecting number of hidden units and layers = fuzzy process
 - NNs have made a HUGE comeback in the last few years!!!
 - Neural nets are back with a new name!!!!
 - Deep belief networks
 - Huge error reduction when trained with lots of data on GPUs

Weight Initialization

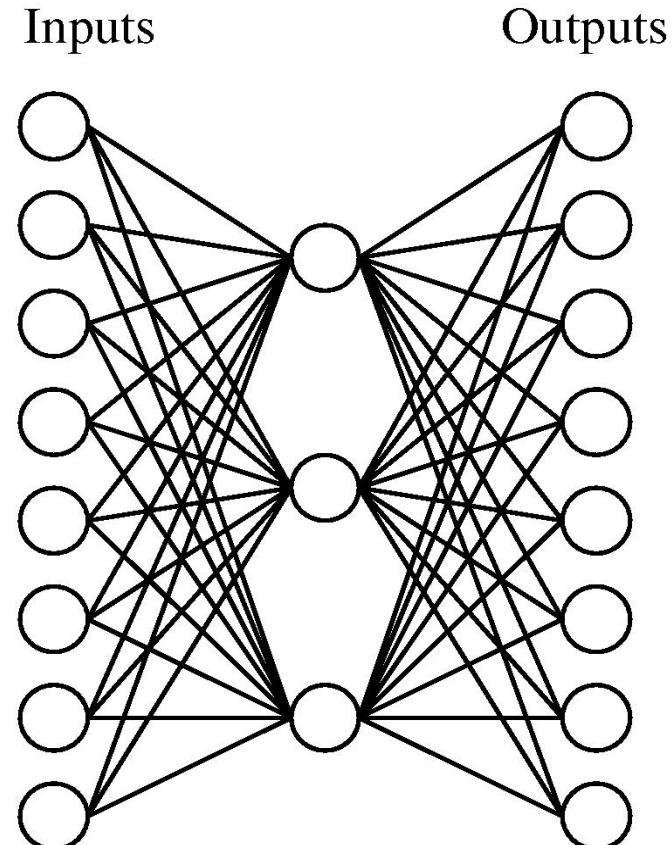
- » Don't just initialize weights to 0!
 - (like we did with linear models)
 - Bad local optima
- » Initial weights = 0
 - => all hidden units = 0
 - => Gradient on output layer will be 0
 - => Gradients of weights on each hidden unit will be the same.
 - => Values of hidden units always the same

Breadth vs. Depth

- » 2 layer networks can represent any function
- » But, can require exponentially many hidden units
- » Analogy to circuit complexity
 - Parity function
 - Easy to define $O(\log(D))$ depth circuit
 - Exponential # of gates required for constant depth

Overfitting in NNs

- Are NNs likely to overfit?
 - Yes, they can represent arbitrary functions!!!
- Avoiding overfitting?
 - More training data
 - Fewer hidden nodes / better topology
 - Regularization
 - Early stopping



Disadvantages of Neural Networks

- » Lots of hyperparameters!
 - Learning rate
 - Early stopping
 - How many hidden units?
 - How many layers?
 - ...

Disadvantages of Neural Networks

» Vanishing Gradients

- gradients shrink during backprop
- At beginning of deep network, changing one weight doesn't have much effect on the output
- Derivatives are small

Computation Graphs

- » Arbitrary graphs of function composition
- » Choose whatever activation functions
- » Choose any loss function
- » Computing gradients is all very mechanical
- » Also automatically parallelize on GPUs



theano

Keras == simple
10 lines of code



The PyTorch logo features the letters 'PYTORCH' in a white, rounded font, with a blue lightning bolt symbol integrated into the letter 'T'.

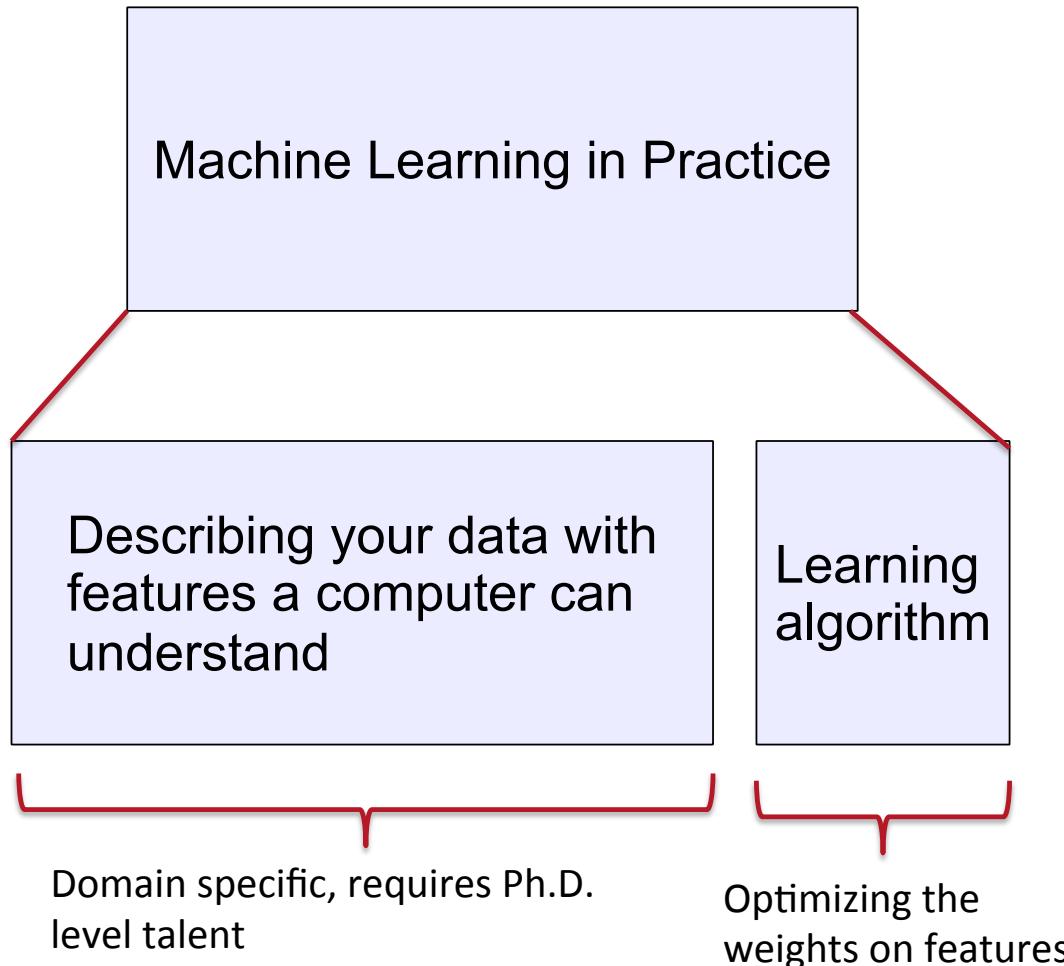
th>

```
input_encoder = Sequential()
input_encoder.add.Embedding(input_dim=vocab_size, output_dim=64)
input_encoder.add(LSTM(64))

initial()
embedding(input_dim=vocab_size, output_dim=64)
lstm(64)

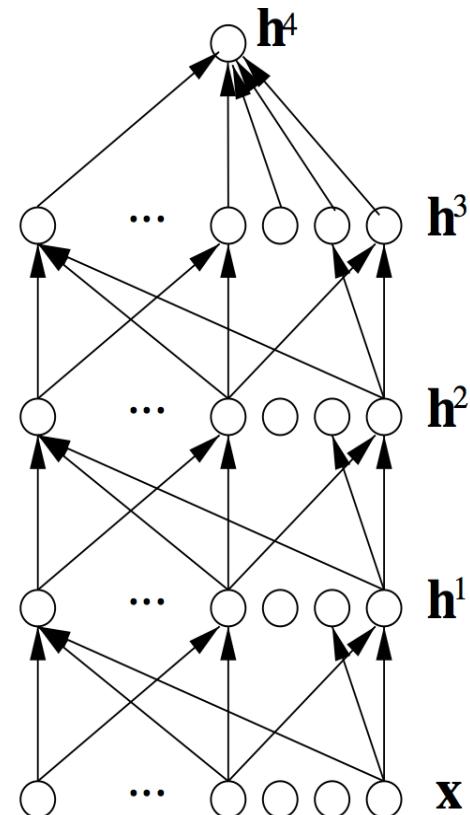
encoder, question_encoder], mode='concat'))
softmax()
```

Machine Learning vs Deep Learning



What's Deep Learning (DL)?

- Representation learning attempts to automatically learn good features or representations
- Deep learning algorithms attempt to learn (multiple levels of) representation and an output
- From “raw” inputs \mathbf{x} (e.g. words)

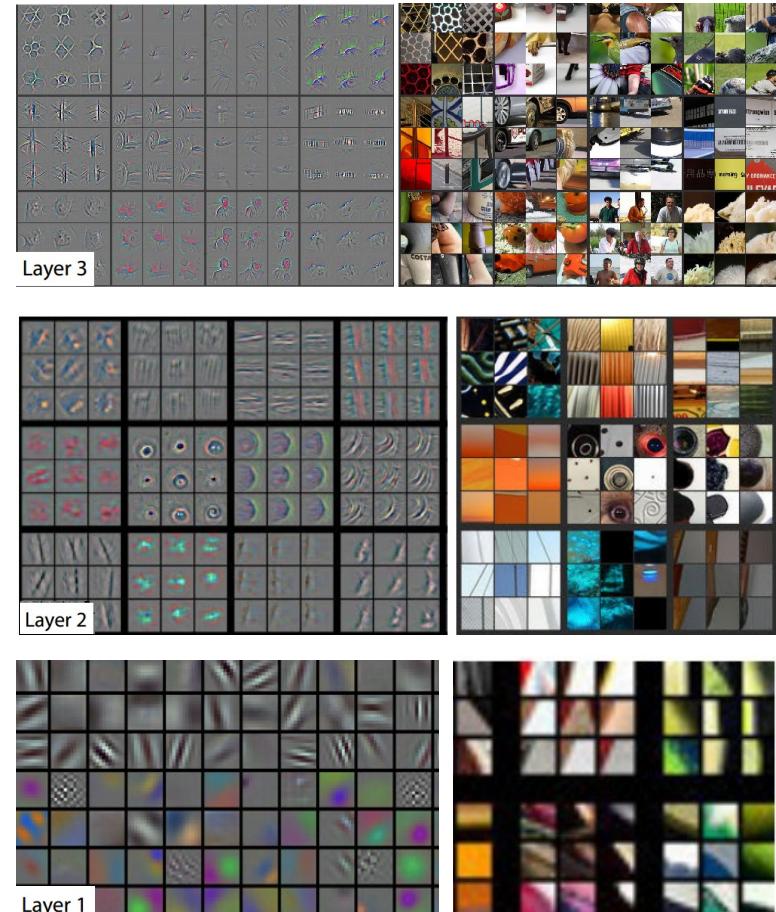


Reasons for Exploring Deep Learning

- In 2006 **deep** learning techniques started outperforming other machine learning techniques. Why now?
 - DL techniques benefit more from a lot of data
 - Faster machines and multicore CPU/GPU help DL
 - New models, algorithms, ideas
- **Improved performance** (first in speech and vision, then NLP)

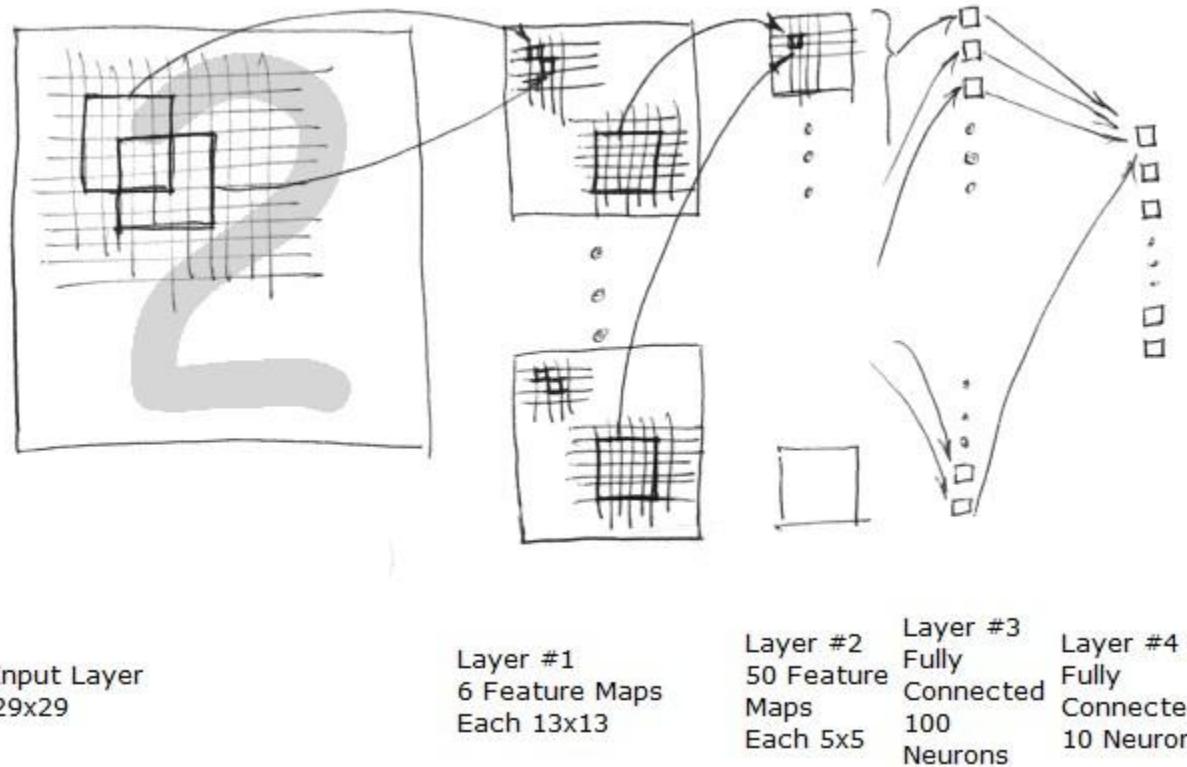
Deep Learning for Computer Vision

- Most deep learning groups have (until recently) largely focused on computer vision
- Break through paper:
ImageNet Classification with Deep Convolutional Neural Networks by Krizhevsky et al. 2012



Zeiler and Fergus (2013)

Convolutional Neural Networks



AlexNet

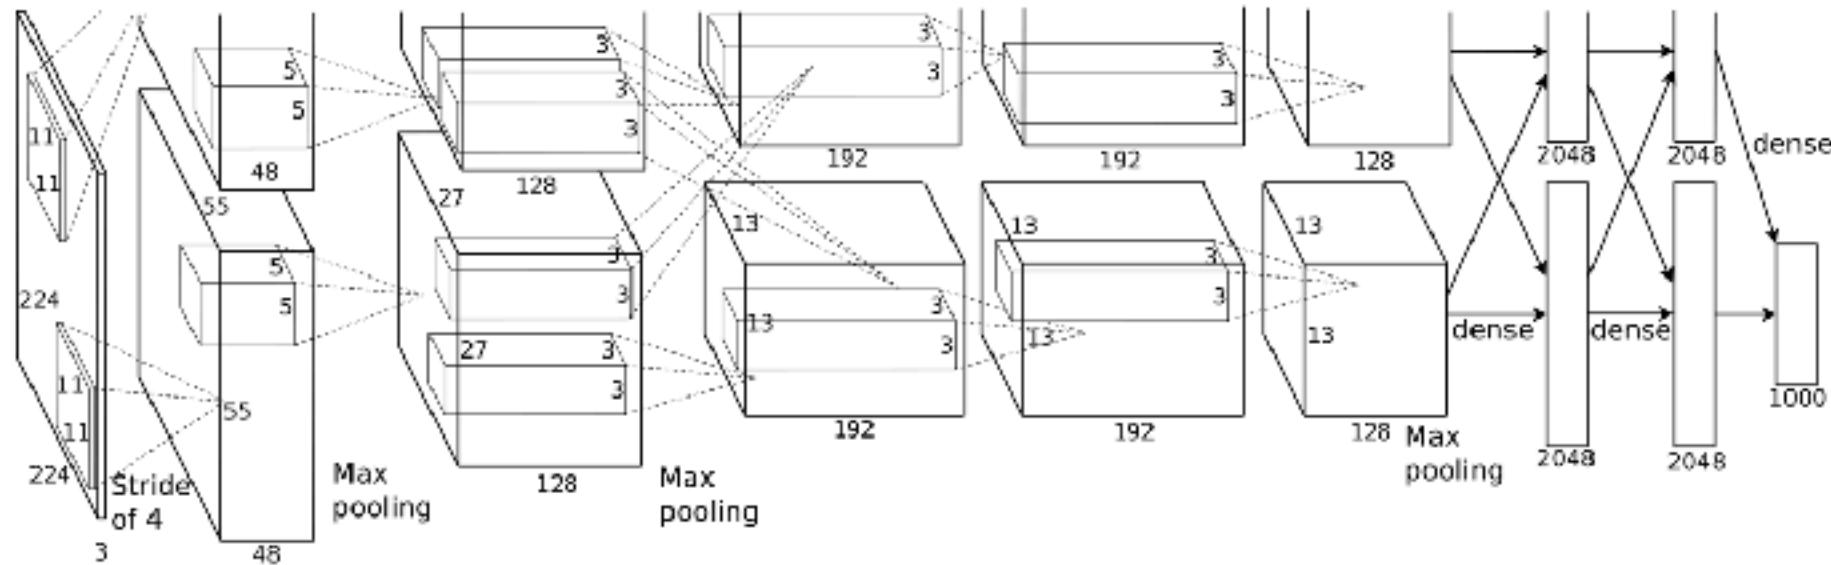


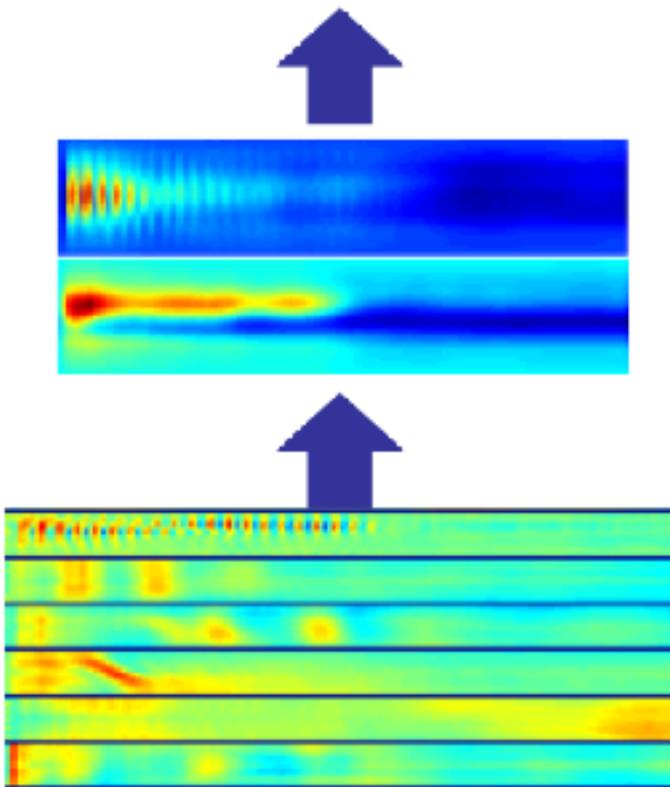
Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

Deep Learning for Speech

- The first breakthrough results of “deep learning” on large datasets happened in speech recognition
- Context-Dependent Pre-trained Deep Neural Networks for Large Vocabulary Speech Recognition
Dahl et al. (2010)

Acoustic model	Recog \\ WER	RT03S FSH	Hub5 SWB
Traditional features	1-pass -adapt	27.4	23.6
Deep Learning	1-pass -adapt	18.5 (-33%)	16.1 (-32%)

Phonemes/Words



Machine Translation

- Many levels of translation have been tried in the past:
- Traditional MT systems are very large complex systems

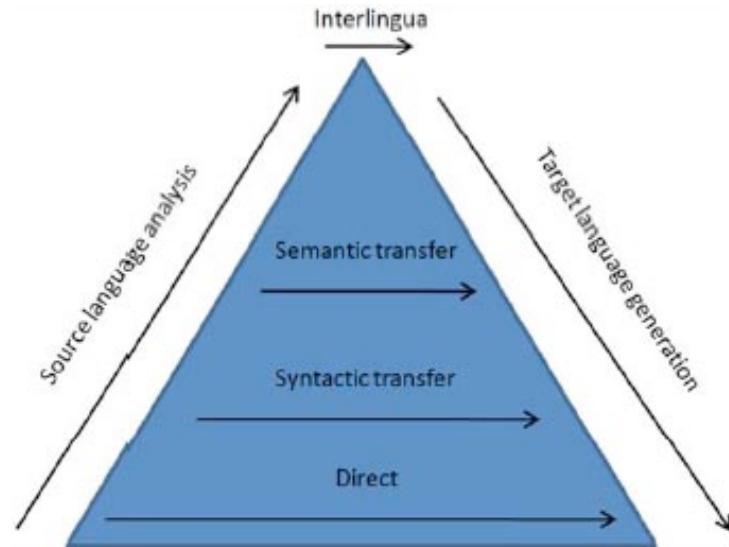
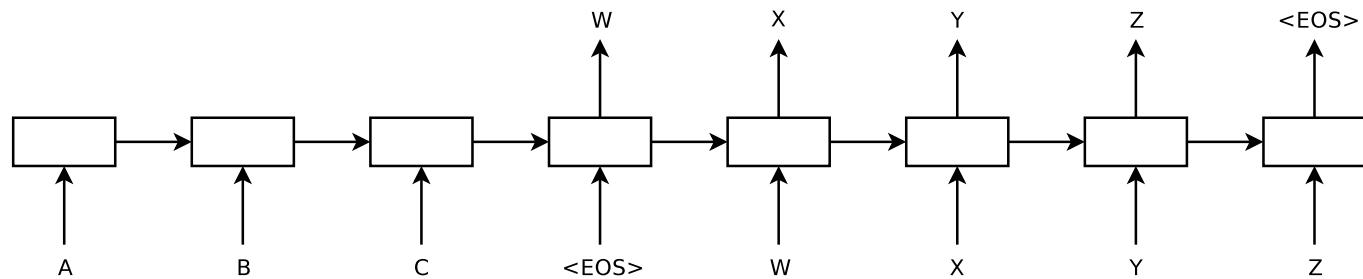


Figure 1: The Vauquois triangle

- What do you think is the interlingua for the DL approach to translation?

Machine Translation

- Source sentence mapped to vector, then output sentence generated.



- Sequence to Sequence Learning with Neural Networks by Sutskever et al. 2014

What you need to know about neural networks

- Perceptron:
 - Relationship to general neurons
- Multilayer neural nets
 - Representation
 - Derivation of backprop
 - Learning rule
- Overfitting