

# Introduction to Machine Learning

<http://bootcamp.lif.univ-mrs.fr:8080/mainpage>

PASCAL Bootcamp 2010

$$P(\theta|\mathcal{D}, \mathcal{M}) = \frac{P(\mathcal{D}|\theta, \mathcal{M}) P(\theta|\mathcal{M})}{P(\mathcal{D}|\mathcal{M})}$$

Sponsored links

## Web-based Text Mining

Natural language processing API:  
entity extraction, text categ, etc.  
[www.alchemyapi.com/](http://www.alchemyapi.com/)

## Machine Learning Jobs

Google is looking for **machine learning** experts  
[www.google.com/jobs](http://www.google.com/jobs)

## Seven Interesting Things

**Machine learning** + Trends + News  
7 tasty **machine**-picked things a day  
[7interestingthings.blogspot.com](http://7interestingthings.blogspot.com)



$$R(T_i) \leq R_{\text{emp}}(T_i) + \frac{\ln N - \ln \eta}{\ell} \left( 1 + \sqrt{1 + \frac{2R_{\text{emp}}(T_i)\ell}{\ln N - \ln \eta}} \right)$$

Iain Murray

<http://homepages.inf.ed.ac.uk/imurray2/>

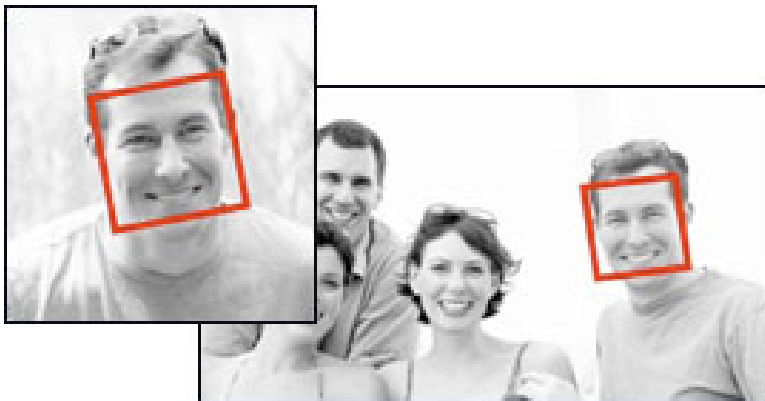
# Face detection

---

How would you detect a face?



(R. Vaillant, C. Monrocq and Y. LeCun, 1994)



How does album software  
tag your friends?

# What do we do?

---



<http://brain.mada.org.il/upside-down-e.html>





# Response surface optimization

---

Consult on making a new:  
concrete, weld, . . . widget?

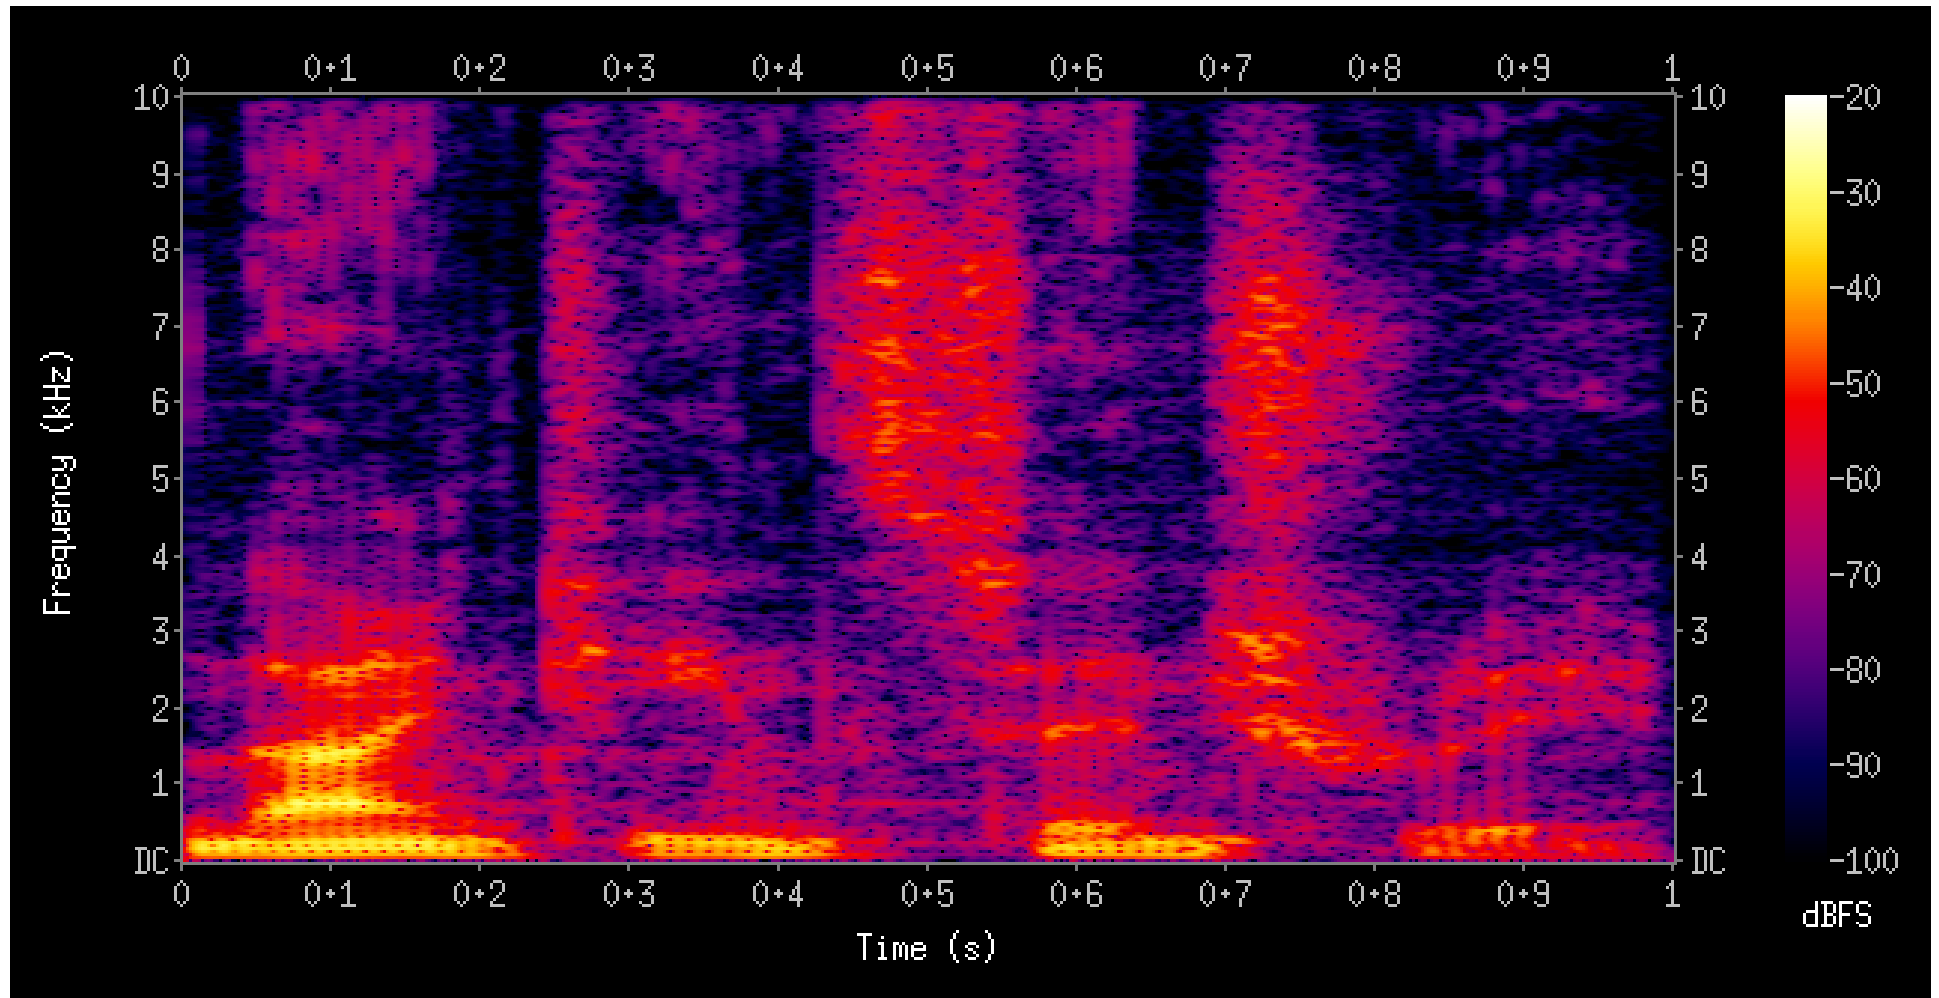


(antmoose on Flickr, cc-by-2.0)



(Bhadeshia et al.)

# Speech recognition



“nineteenth century”

# Recommender systems

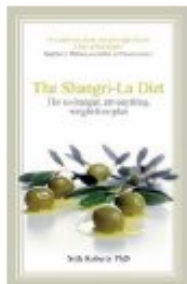
---

Amazon, Netflix, tell you what you might like

**Collaborative filtering** useful in many settings?

## Today's Recommendations For You

Here's a daily sample of items recommended for you. Click here to [see all recommendations](#).



[The Shangri-la Diet](#)  
(Paperback) by Seth Roberts

★★★★★ (3) £5.81


[Fix this recommendation](#)



[C++ Design Patterns and Derivativ...](#) (Paperback) by M. S. Joshi

★★★★★ (7) £22.78

[Fix this recommendation](#)



[What the Dog Saw: and other...](#) (Paperback) by Malcolm Gladwell

★★★★☆ (17) £5.00

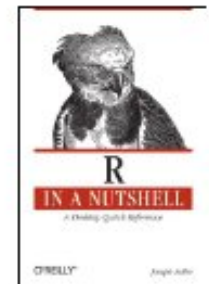
[Fix this recommendation](#)



[Garden State \[DVD\] \[2004\]](#)  
DVD ~ Zach Braff

★★★★☆ (98) £3.99

[Fix this recommendation](#)



[R in a Nutshell \(In a Nutshell ...\)](#)  
(Paperback) by Joseph Adler

£20.40

[Fix this recommendation](#)

# Roadmap

---

## **Binary classification**

— Parametric and non-parametric prediction

## **Other ‘Supervised’ settings**

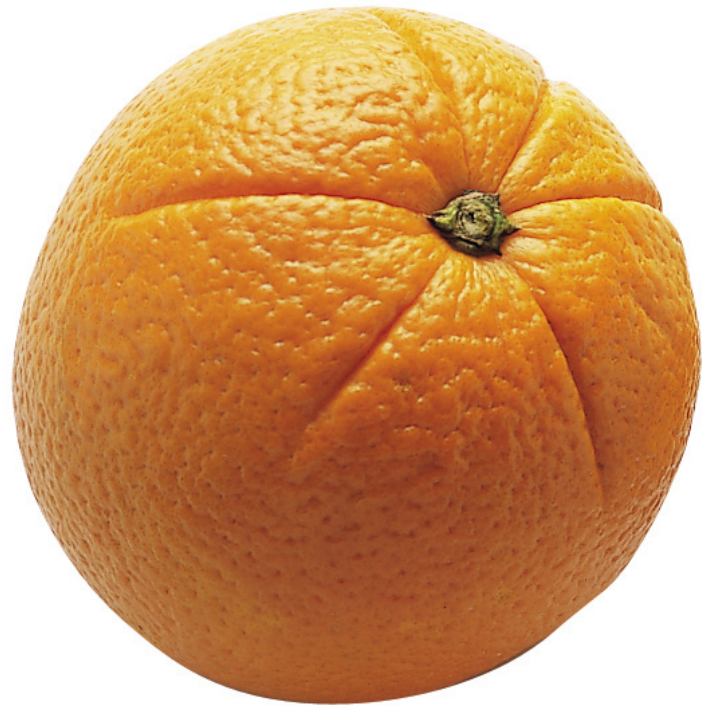
## **Principles for learning**

## **More machine learning paradigms**





# Oranges and Lemons

---




# A two-dimensional space

---

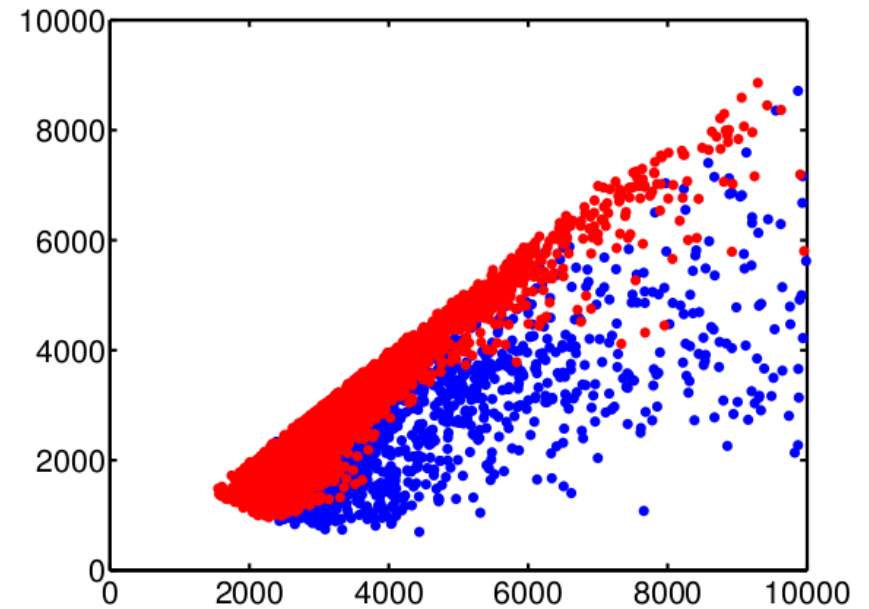


# Stars and Galaxies

---




<http://www-wfau.roe.ac.uk/ss/>



Minor elliptical axis (y) against Major elliptical axis (x) for stars (red) and galaxies (blue). (Amos Storkey)

# Linear classifier

---



Straight cut dividing input space into two

# The weight vector

---


Define positive class region:  $\mathbf{w}^\top \mathbf{x} + b > 0$

$$w_1x_1 + w_2x_2 + w_3x_3 + \dots b > 0$$

$$\sum_d w_d x_d + b > 0$$

We will set  $b = 0$ :

$$\mathbf{w}^\top \mathbf{x} > 0$$



# The weight vector

---

Positive class region  $\mathbf{w}^\top \mathbf{x} > 0$

**Matlab / Octave:**

```
%      ww = Dx1 weights  
% Xstar = MxD test cases  
y_prediction = sign(Xstar*ww); % Mx1
```

Learn a numerical package: Python+NumPy, R, Matlab/Octave, Eigen+C++,...



# Learning the weights





---

## A ‘perceptron’ learning rule:

$$\hat{y} \leftarrow \text{sgn}(\mathbf{w}^\top \mathbf{x})$$


$$\mathbf{w} \leftarrow \mathbf{w} + (y - \hat{y})\mathbf{x}$$

```
% Matlab/Octave code
old_ww = []; ww = zeros(size(xx,2), 1);
while ~isequal(ww, old_ww)
    old_ww = ww;
    for ii = 1:N
        pred = sign(xx(ii, :)*ww);
        ww = ww + (yy(ii) - pred)*xx(ii, :);
    end
end
```



# Implementing the bias

---




$\text{sgn}(\mathbf{w}^\top \mathbf{x})$  model is inflexible

Need a bias parameter

Or add constant feature to  $\mathbf{x}$ :

$$\mathbf{x}' = \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix}, \quad \mathbf{w}' = \begin{bmatrix} \mathbf{w} \\ b \end{bmatrix}$$



# THE PERCEPTRON: A PROBABILISTIC MODEL FOR INFORMATION STORAGE AND ORGANIZATION IN THE BRAIN<sup>1</sup>

F. ROSENBLATT

*Cornell Aeronautical Laboratory*

If we are eventually to understand the capability of higher organisms for perceptual recognition, generalization, recall, and thinking, we must first have answers to three fundamental questions:

1. How is information about the physical world sensed, or detected, by the biological system?
2. In what form is information stored, or remembered?
3. How does information contained in storage, or in memory, influence recognition and behavior?





FIG. 1. Organization of a perceptron.

# K-Nearest neighbours classifier

---



Fix and Hodges (1951)



# The odd-ball orange


---



# Decision boundaries

---

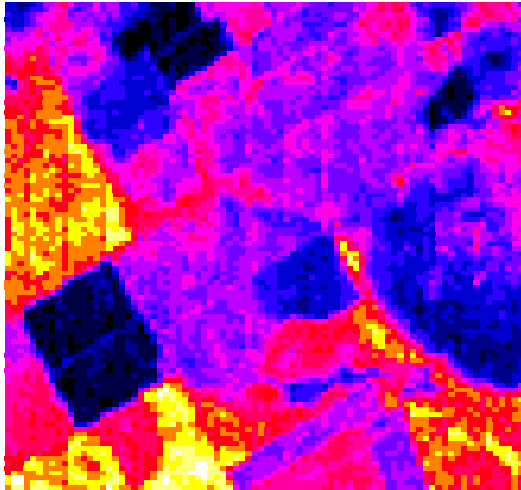
Classifying oil flow



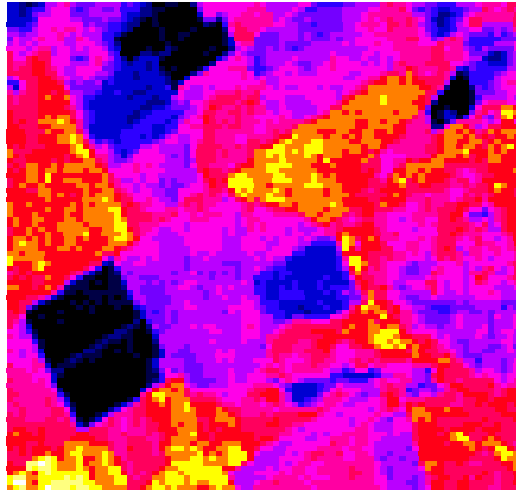
# LANDSAT application

---

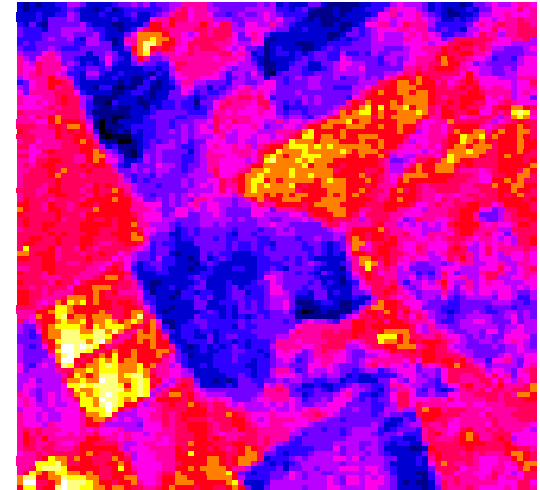
Spectral Band 1



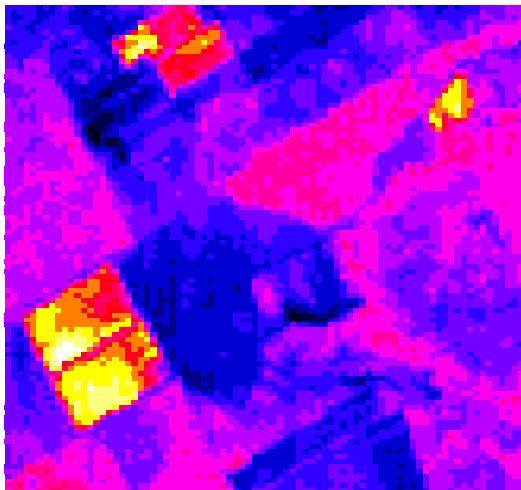
Spectral Band 2




Spectral Band 3




Spectral Band 4



Land Usage



Predicted Land Usage



# Parametric vs. non-parametric

---


Started assuming decision boundary is a plane

Non-parametric KNN has no fixed assumption:  
boundary gets more complicated with more data

Non-parametric methods may need more data  
and can be computationally intensive

# Linear classifier revisited

---



$\text{sgn}(\mathbf{w}^\top \mathbf{x})$  model is inflexible

Need a bias parameter

Or add constant feature to  $\mathbf{x}$ :



$$\mathbf{x}' = \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix}, \quad \mathbf{w}' = \begin{bmatrix} \mathbf{w} \\ b \end{bmatrix}$$

If not linear separable must extend model. . .

. . . or add features.

# Nonlinear basis functions

Create new feature vector:  $\mathbf{x}' = \begin{bmatrix} \phi_1(\mathbf{x}) \\ \phi_2(\mathbf{x}) \\ 1 \end{bmatrix}$





# Batch supervised learning

---

**Given:** example inputs and targets (training set)


**Task:** predict targets for new inputs (test set)

## Examples:

- classification (binary and multi-class)
- (real-valued) regression
- ordinal regression
- Poisson regression
- ranking(?)

# Linear regression

---




Find linear function  
 $\hat{\mathbf{y}} = X\mathbf{w}$   
that minimizes sum  
of squared residuals  
from  $\mathbf{y}$ .

Matlab/Octave:  
 $\mathbf{w} = X \setminus \mathbf{y}$

# Linear regression (with features)


---



```
X = [ones(N, 1), xx, xx.^2, xx.^3, xx.^4, xx.^5, xx.^6];  
Xnew = [ones(N, 1), xnew, xnew.^2, xnew.^3, xnew.^4, xnew.^5, xnew.^6];  
ww = X \ yy;  
ynew = Xnew * ww;
```

# Neighbour-based regression

---




Take height from nearest input

# Simplest kernel smoothing

---


Weight points in proportion to a *kernel*:



$$\hat{y}_* = \sum_i w_i y_i, \quad \text{where } w_i = k(x_*, x_i) / \sum_j k(x_*, x_j)$$

# Simplest kernel smoothing

---



$$\hat{y}_* = \sum_i w_i y_i, \quad \text{where } w_i = k(x_*, x_i) / \sum_j k(x_*, x_j)$$




# Least squares classifier

---

Why not run  
regressors on  
binary targets?

```
% fit  $yy = xx * ww$   
% by least squares  
 $ww = xx \setminus yy;$ 
```




Line shows  $xx * ww == 0$ .  $ww(3)$  is a bias:  $xx(:,3) = 1$

# Least squares classifier (2)

---

Why not run  
regressors on  
binary targets?

```
% fit  $yy = xx * ww$   
% by least squares  
 $ww = xx \setminus yy;$ 
```



PRML C. M. Bishop (2006)

Magenta line shows  $xx * ww == 0$ .  $ww(3)$  is a bias:  $xx(:,3) = 1$

# Roadmap

---

## **Supervised learning:**

Many ways of mapping inputs  $\rightarrow$  outputs

How do we choose what to do?

How do we know if we are doing well?

**Later:** a little on practical issues,  
other ideas in machine learning.

# Algorithm's Objective/Cost

---

## **Formal objective for algorithms:**

- Minimize a cost function
- Or maximize an objective function

## **Proving convergence:**

- does objective monotonically improve?

## **Considering alternatives:**

- does another algorithm score better?

# Loss functions

---

We want to specify the objective of an algorithm

One idea: consider *loss function*  $L(\hat{y}(\mathbf{x}_*); y_*)$

Would be good to minimize loss at test time

Minimizing empirical loss might be a reasonable proxy:  $\sum_i L(\hat{y}(\mathbf{x}_i); y_i)$

# Choosing loss functions

---

## **Motivated by the application:**

0–1 error, achieving a tolerance, business cost



## **Computational convenience:**

Differentiability, convexity

## **Beware of loss being dominated by artifacts:**

- Outliers
- Unbalanced classes


# Squared error



$$z = \mathbf{w}^\top \mathbf{x}, \quad \hat{y} = \text{sgn}(z)$$

# Comparing loss functions

---






# Logistic regression

---

Probabilistic model for labels:

$$P(y | \mathbf{x}, \mathbf{w}) = \frac{1}{1 + e^{-y(\mathbf{w}^\top \mathbf{x})}}, \quad \text{activation } z = \mathbf{w}^\top \mathbf{x}$$



# Logistic regression

---


Maximize the **likelihood** of  $\mathbf{w}$ :

$$P(\mathbf{y} | X, \mathbf{w}) = \prod_i \frac{1}{1 + e^{-y_i(\mathbf{w}^\top \mathbf{x}_i)}}$$

$$\log P(\mathbf{y} | X, \mathbf{w}) = - \sum_i \log(1 + e^{-y_i(\mathbf{w}^\top \mathbf{x}_i)})$$

Pay loss:



$$\log(1 + e^{-y_i(\mathbf{w}^\top \mathbf{x}_i)})$$



# Over fitting

---

We can make the empirical loss zero:




# Generalization

---

**Want to do well on future, unknown data**

There is no one right way to proceed




Multiple theoretical points of view

First, a simple practical procedure

# Validation set


---



$$X = [x.^0, x.^1, x.^2, \dots, x.^p];$$

# Learning curves

---



# Using validation sets

---

## **Can overfit the validation set:**

- Always hold out a true “test set” for final evaluation

## **Validation set is not used (much) for fitting:**

- make training set as large as you can
- consider multiple folds/rounds

# Early stopping

---

Validation set used to estimate test error for fitted models.

**Tracking a validation set is also used during fitting of a single model:**

- *Ad hoc*, that is, a hack
- depends on optimizer
- can work annoyingly well
- sometimes fast compared to alternatives




# Regularization

---

We don't need the validation set to know some fits are silly. Regularization discourages solutions we don't like

**Shrinkage or weight decay:**  
make parameters closer to zero.

Early stopping can be similar:  
stop before weights grow big



# Regularization

---

## Maximize a penalized objective:

Penalize  $\mathcal{F}$  = negative training error, or training probability

$$\mathcal{F}(\mathbf{w}; X, \mathbf{y}) - \lambda \sum_i |w_i|, \quad L_1, \text{ (e.g., 'lasso')}$$

$$\mathcal{F}(\mathbf{w}; X, \mathbf{y}) - \lambda \sum_i w_i^2, \quad L_2, \text{ (e.g., 'ridge regression')}$$

## Equivalently, constrain parameters:


Maximize  $\mathcal{F}(\mathbf{w}; X, \mathbf{y})$  such that  $\sum_i |w_i| = t$  or  $\sum_i w_i^2 = t$

One-one correspondence between  $t$  and  $\lambda$

$\lambda$  is a Lagrange multiplier for the constrained problem.

# Different regularizers

---



# Regularization constants

---

## How do we pick $\lambda$ or $t$ ?

- based on validation
- based on bounds on generalization error
- based on “empirical Bayes”

Or given a well-defined procedure, reinterpret  $\lambda$  and do something else.  
Fully Bayesian methods “integrate out”  $\lambda$ .

# Learning theory

---

Compute bounds on generalization error

**Based on held-out validation set:**

- easy to construct
- easy to interpret theoretically
- harder to extend to  $K$ -fold validation

**Based on training set performance:**

- philosophically pleasing and interesting
- bounds can be very loose

# The Bayesian view

---

## **Why do we use training data?**

Because there are things we don't know.

There are also things we do know (regularization)

Under certain assumptions, probabilities and probability calculus (e.g., Bayes rule) are *the* way to combine prior beliefs and observations.

# Bayesian procedure

---

Specify possible worlds by putting a *prior* distribution over parameters of a model  $p(\mathbf{w})$

Given data  $\mathcal{D} = \{X, \mathbf{y}\}$  we have a *posterior*:

$$p(\mathbf{w} | \mathcal{D}) = \frac{p(\mathcal{D} | \mathbf{w}) P(\mathbf{w})}{P(\mathcal{D})} \propto p(\mathcal{D} | \mathbf{w}) P(\mathbf{w})$$

Minimize expected loss:

$$\begin{aligned} \hat{y} &= \arg \min_y \sum_{y_*} L(y; y_*) p(y_* | \mathbf{x}_*, \mathcal{D}), \quad \text{cf "Bayes classifier"} \\ &= \arg \min_y \sum_{y_*} L(y; y_*) \int p(y_* | \mathbf{w}) p(\mathbf{w} | \mathcal{D}) d\mathbf{w} \end{aligned}$$

# Bayesian connections

---

Approximating the integral can involve maximizing:

$$\log p(\mathbf{w} | \mathcal{D}) = \log p(\mathcal{D} | \mathbf{w}) + \log p(\mathbf{w}) + \text{const.}$$

$$\text{where } \log p(\mathbf{w}) = -\lambda \sum_i w_i^2, \quad \text{if } p(\mathbf{w}) = \mathcal{N}(0, 2\lambda)$$

$\Rightarrow$  Optimize L2-regularized log-likelihood.

However, *not* tied to  $-ve$  log-prob loss.

The loss we care about is optimized later.



# Bayesian advantages

---

In principle the solution to any problem

(e.g., outliers, data corruption, missing inputs)

is solved by modelling it and probability calculus.




Don't know  $\lambda$ ?

$$\begin{aligned}\hat{y} &= \arg \min_y \sum_{y_*} L(y; y_*) \int p(y_* | \mathbf{w}) p(\mathbf{w} | \mathcal{D}) d\mathbf{w} \\ &= \arg \min_y \sum_{y_*} L(y; y_*) \int p(y_* | \mathbf{w}) \int p(\mathbf{w} | \mathcal{D}, \lambda) p(\lambda | \mathcal{D}) d\lambda d\mathbf{w}\end{aligned}$$

Might maximize  $p(\lambda | \mathcal{D})$  instead of integrating over  $\lambda$ .

“Empirical Bayes” / Type-II maximum likelihood maximizes  $p(\mathcal{D} | \lambda)$

# Bayesian logistic regression



Information theory, inference, and learning algorithms, David J. C. MacKay (2003)  
Logistic regression is referred to as a single-neuron classifier

# Bayesian computations

---

Given a model, the Bayesian solution is a purely mechanical manipulation of probabilities.

(given a little practice)

The required integrals can be formidable.

Huge research effort in **approximate inference**:

- variational objectives
- Monte Carlo sampling

# Bayesian models

---

We can fit arbitrary rules for  $\mathbf{x} \rightarrow y$  and estimate future losses.

**To take posterior beliefs seriously,  
need to start with sensible prior beliefs.**

Need flexible distributions over models and priors that are “sensible enough”.

Graphical models and non-parametric Bayesian methods offer useful machinery.

# Supervised learning

---

## Overview:

- identify task and loss function
- build flexible model of mapping  $\mathbf{x} \rightarrow y$
- use data to set all the free parameters

## **Some form of regularization is crucial**

Large, complicated models are useful in applications.

- penalize empirical score
- Bayesian methods

# Pre-processing

---

## **Back to thinking about applications**

Given a dataset of example inputs and outputs,  
will what we have seen so far work?

What's  $x$ ?

# Rescaling and centering

---

## Centering:

- Might have all features at  $1000 \pm 10$ .
- Hard to predict ball park of bias.
- Subtract mean from all input features.

## Rescaling:


- I might measure heights in cm or m.
- Rescale inputs to have unit variance  
... or interquartile range.

Similarly center and scale real-valued outputs

**Care at test time:** apply same scaling to inputs and reverse scaling to predictions.

# Log-transform +ve inputs

---



(Wisconsin breast cancer data from the UCI repository)

Positive quantities are often highly skewed  
The log-domain is often much more natural



# Encoding attributes

---

## **Categorical variables:**

- A study has three individuals.
- Three different colours
- Possible encoding: 100, 010, 001

## **Ordinal variables:**

- Movie rating, 1–3 stars
- Tissue anomaly rating, expert scores in 1–3
- Possible encoding: 00, 10, 11

# Basis functions features

---

Toy examples used polynomials:  $x, x^2, x^3 \dots$   
Often a bad choice

Polynomials of sparse binary features make sense:

$$x_1x_2, x_1x_3, \dots, x_1x_2x_3$$

## Other options:

- Radial basis functions:  $\exp(-|\mathbf{x} - \boldsymbol{\mu}|^2/h^2)$
- Sigmoids:  $1/(1 + \exp(-\mathbf{v}^\top \mathbf{x}))$
- Fourier, wavelets,  $\dots$

# Feature engineering

---

The difference in performance in an application can depend more on the original features than the algorithm.

Working out clever ways of making features from complex objects like images can be worthwhile, is hard, and isn't always respected. . .

# The SIFT story

---

Taken from: <http://www.lecun.com/ex/pamphlets/publishing-models.html>

Many of us have horror stories about how some of our best papers have been rejected by conferences. Perhaps the best case in point of the last few years is David Lowe's work on the SIFT method. After years of being rejected from conference[s] starting in 1997, the journal version published in 2004 went on to become the most highly cited paper in all of engineering sciences in 2005.

David Lowe relates the story:

*I did submit papers on earlier versions of SIFT to both ICCV and CVPR (around 1997/98) and both were rejected. I then added more of a systems flavor and the paper was published at ICCV 1999, but just as a poster. By then I had decided the computer vision community was not interested, so I applied for a patent and intended to promote it just for industrial applications.*

# Creating extra data


---

**A dirty trick:** create more training 'data' by corrupting examples in the real training set.

Changes could respect invariances that would be difficult or burdensome to model directly.

# PCA — Principal Components

Can reduce dimensionality of inputs with PCA



```
[N,D] = size(X);  
[V,S] = svds(cov(X,1), K);  
% Project X onto K principal eigenvectors  
Z = (X - repmat(mean(X),N,1))*V; % NxK  
% Project back into D-dims (for figure)  
Xproj = Z*V' + repmat(mean(X),N,1); % NxD
```

K = 1    × = X    • = Xproj

Projects onto a linear subspace minimizing square error such that new features are uncorrelated.

# PCA applied to DNA data

---

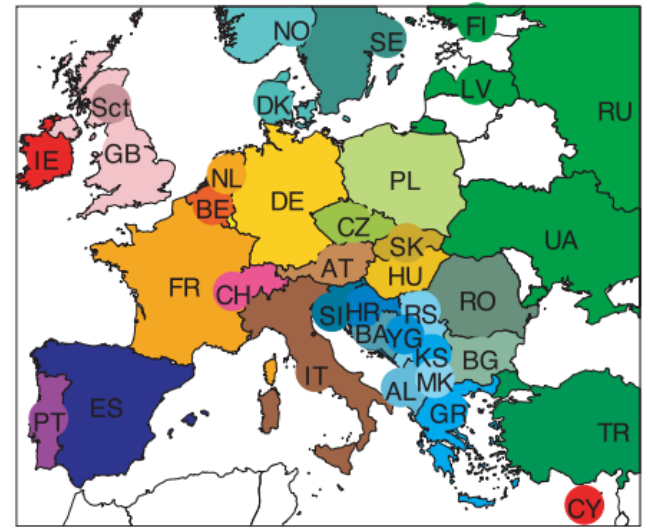
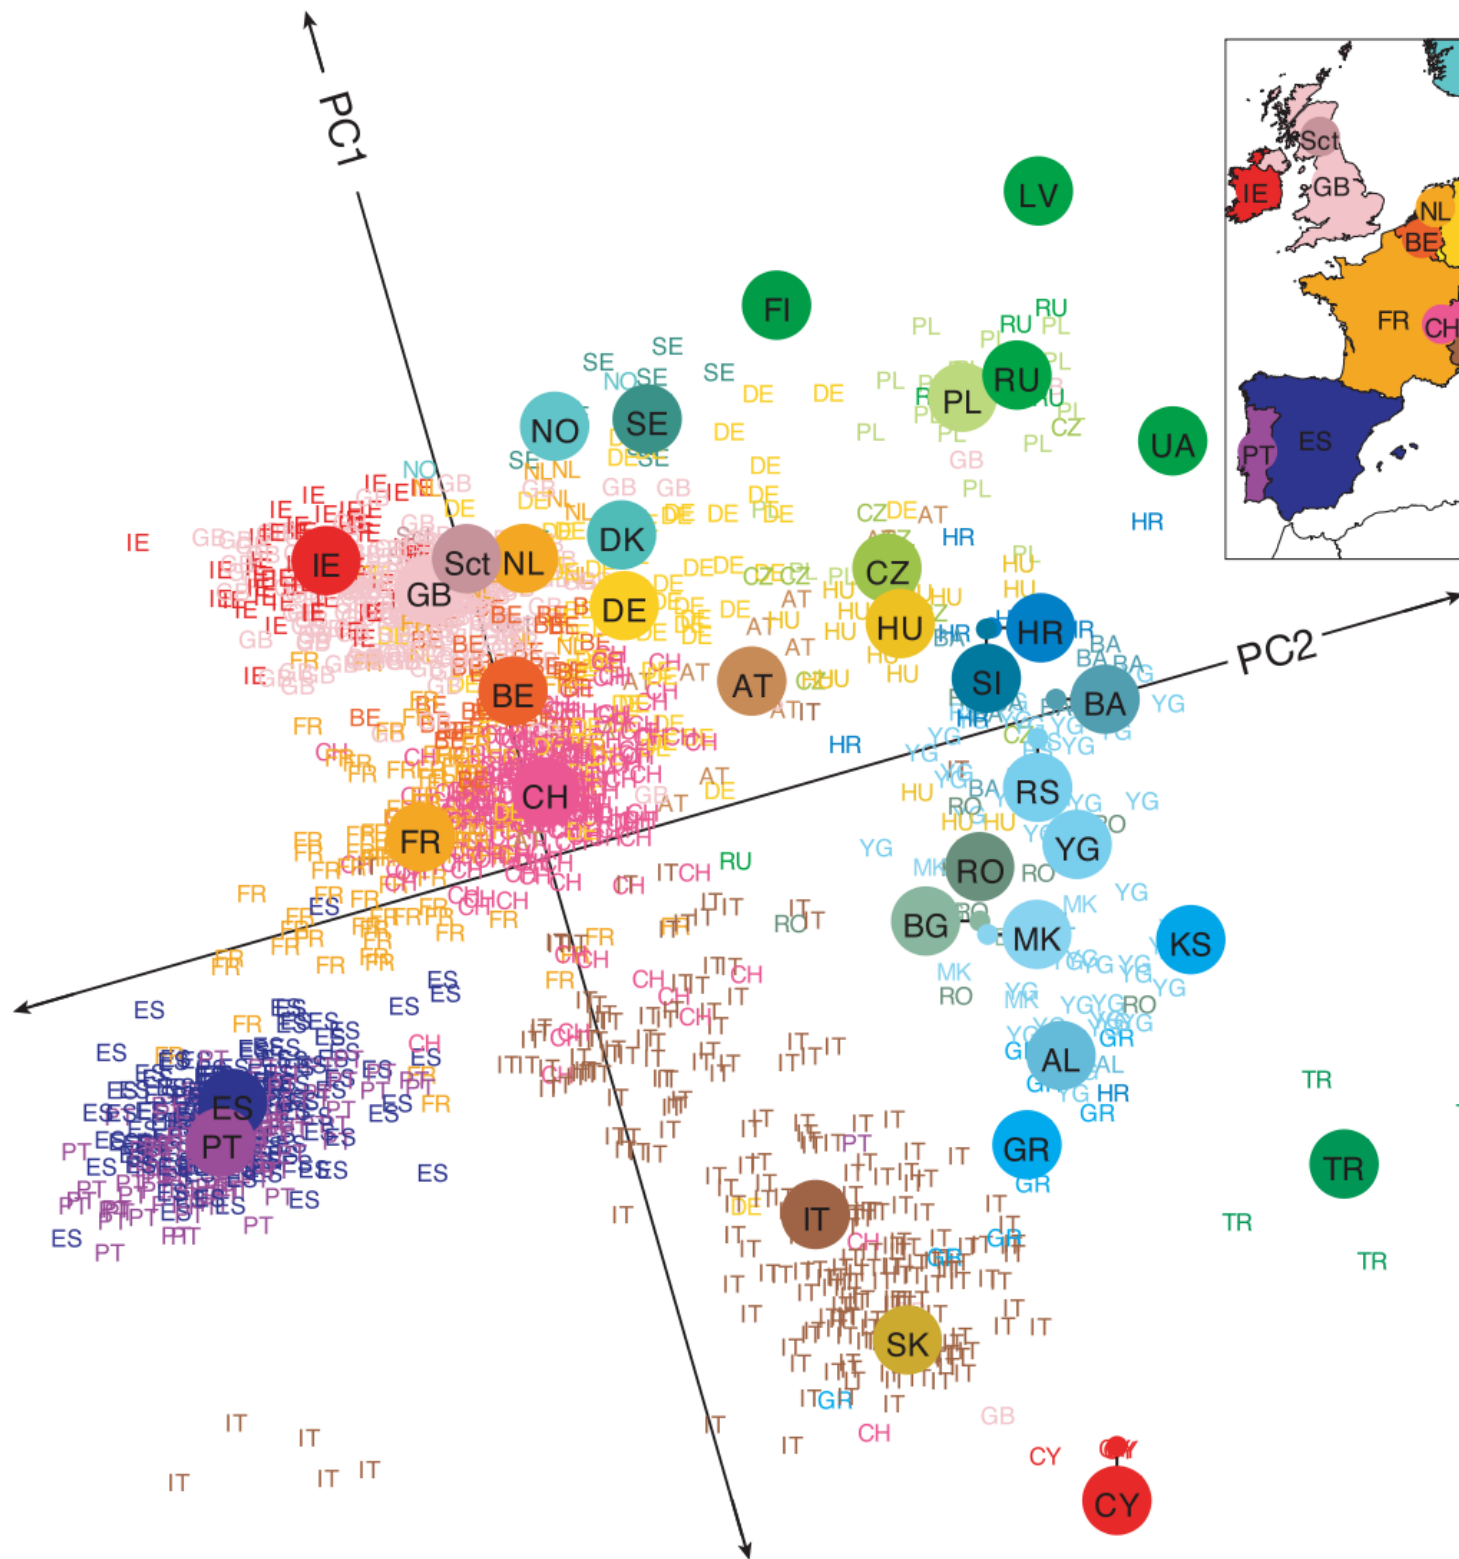
Novembre et al. (2008) — doi:10.1038/nature07331

After selection of both individuals and features:

1,387 individuals

197,146 single nucleotide polymorphisms (SNPs)

Use PCA to reduce features to two(!) dimensions





# Visualization

---

Visualization may suggest issues, transformations, features for a particular problem.

The insights could be results in themselves.

Artifacts that could cause problems:  
outliers, thresholding, kinks, . . .

How noisy is the data?

How relevant are features?

Did you remember to hold out your test set?

# Things to try

---

Histograms of individual features

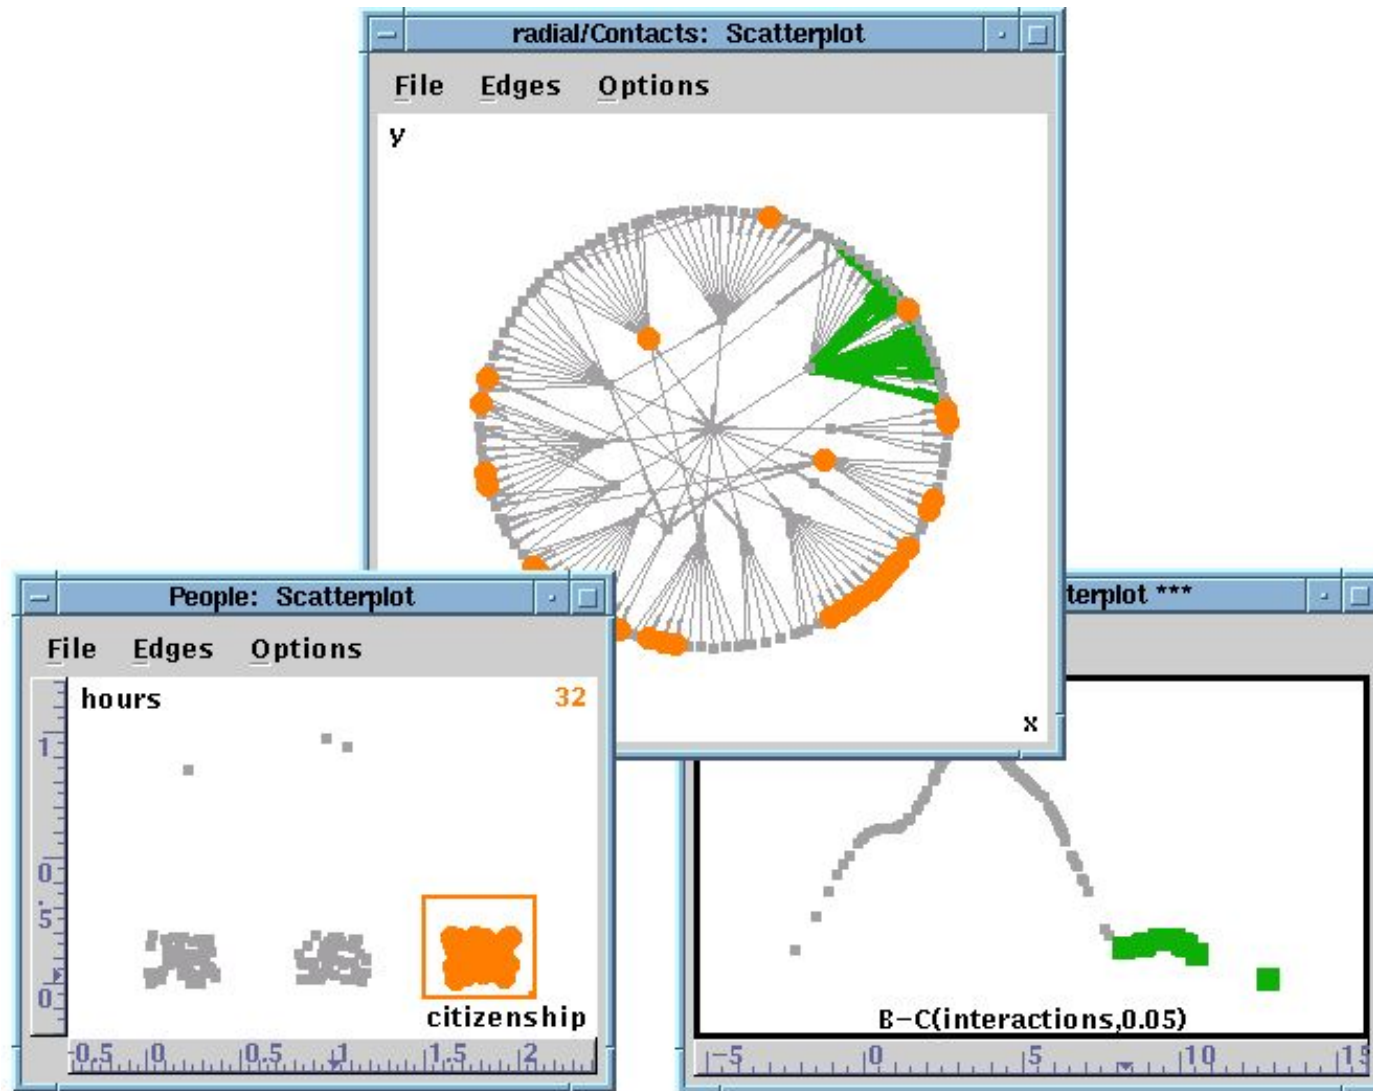
Scatter plots of pairs of features

Scatter plots of PCA, and random projections

Colour plots using  $y$ , or plot against it

# Tours, linked plots, more

---



<http://www.ggobi.org/>

Swayne, Buja and Lang (2004)

# Unsupervised Learning

---


Learn structure in data automatically without a supervisor providing targets

Much less well defined.

Human tweaking often important, but we would always like computers to do more by themselves.


# I see clusters

---



# K-means demo


---





# K-means demo

---






## height/cm




# K-means demo

---




# K-means demo

---




# K-means demo

---




# K-means demo

---




# K-means demo

---




# K-means demo

---



# K-means demo

---





# Mixtures of Gaussians

---

MoG is a probabilistic, fancier K-means

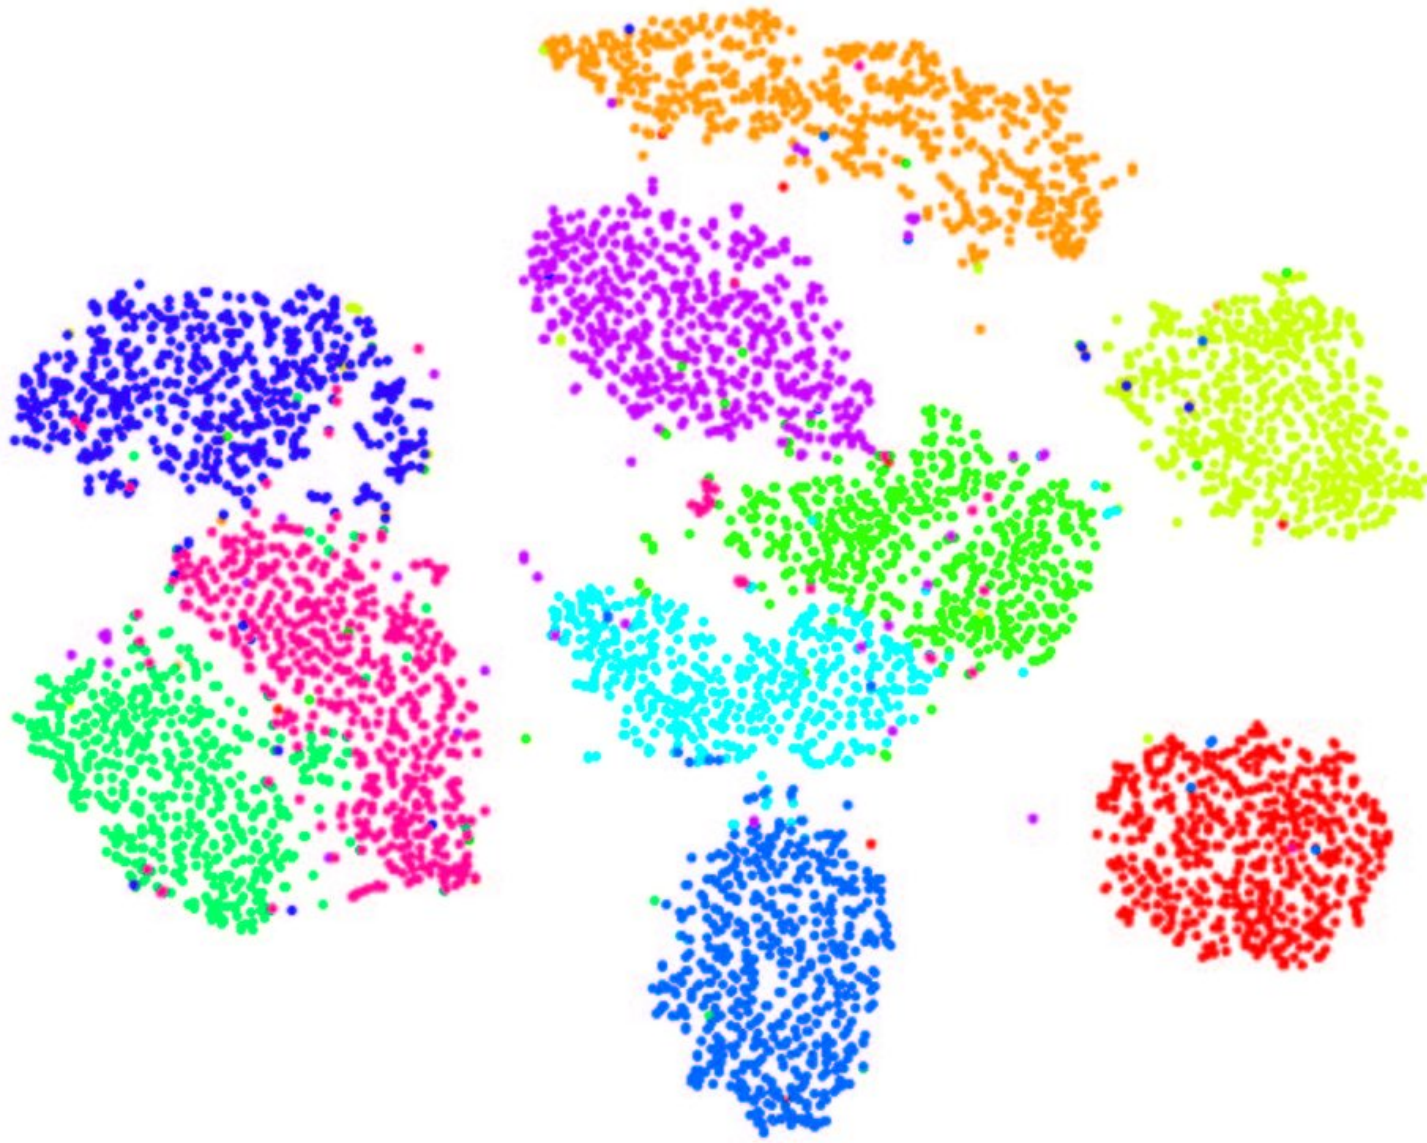


Dahlkamp et al. (2006)

<http://robots.stanford.edu/talks/stanley/>

# Dimensionality reduction

---



t-SNE on MNIST digits, Van der Maaten and Hinton (2008)

See also: <http://www.cs.toronto.edu/~hinton/turian.png>


# Density estimation

---

## One principle for unsupervised learning:

Model joint probability of all data

$$p(X | \mathbf{w}) \quad \text{or} \quad p(\mathbf{y}, X | \mathbf{w}) = p(\mathbf{y} | X, \mathbf{w}_d) p(X | \mathbf{w}_g)$$



Jointly modelling/regularizing  $\mathbf{w}_d$ ,  $\mathbf{w}_g$  uses data more fully

# “Missing data”

---

Having a fully generative model of  $(\mathbf{x}, y)$  allows probabilistic/Bayesian inference to deal with any type of unknown


## **The varieties of “missing data”:**

- features missing at random
- data censored for a reason
- truncated data
- unlabelled data

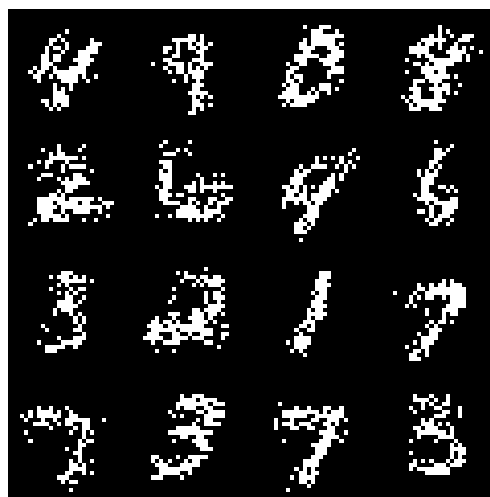
# Underfitting

---


The main problem with unsupervised or part-supervised learning of structure is **underfitting**.



Data samples



MoB samples



RBM samples

# More underfitting

---

Even non-parametric methods can underfit

## **Recommender systems:**

- Amazon
- Netflix

Netflix prize dataset:

100,480,507 ratings; 480,189 users; 17,770 movies

# Building richer models

---

**Learn features and relationships between variables:** neural networks, graphical models.

**Model combination:** boosting, cascading, mixtures of experts, products of experts, stacking models (deep learning)

Note: Bayesian model averaging, does inference for the model you have (simple or otherwise)

# Summary

---

Machine learning uses data to refine flexible computer programs

We seek formalize learning objectives that allow for fitting the richness present in real data

There are many useful tools out there, but huge practical and theoretical challenges to enjoy



# What I haven't covered

---

## **Lots more**

Other settings: transductive learning,  
semi-supervised learning, online learning

Reinforcement learning, games

Time series, structured outputs, relational  
learning and many other specialist settings

# A rant about least squares

---

Whenever a person eagerly inquires if my computer can solve a set of 300 equations in 300 unknowns. . . The odds are all too high that our inquiring friend. . . has collected a set of experimental data and is now attempting to fit a 300-parameter model to it—by Least Squares! The sooner this guy can be eased out of your office, the sooner you will be able to get back to useful work—but these chaps are persistent. . . you end up by getting angry and throwing the guy out of your office.

There is usually a reasonable procedure. Unfortunately, it is undramatic, laborious, and requires thought—which most of these charlatans avoid like the plague. They should merely fit a five-parameter model, then a six-parameter one. . . Somewhere along the line—and it will be much closer to 15 parameters than to 300—the significant improvement will cease and the fitting operation is over. There is no system of 300 equations, no 300 parameters, and no glamor.

The computer center's director must prevent the looting of valuable computer time by these would-be fitters of many parameters. The task is not a pleasant one, but the legitimate computer users have rights, too. . . the impasse finally has to be broken by violence—which therefore might as well be used in the very beginning.