

# Lecture 9: CNNs, Neural CRFs


Alan Ritter

(many slides from Greg Durrett)

# Recall: RNNs


---

- ▶ Cell that takes some input  $x$ , has some hidden state  $h$ , and updates that hidden state and produces output  $y$  (all vector-valued)



# Recall: RNN Abstraction

---



the movie was great

- ▶ **Encoding of the sentence** – can pass this a decoder or make a classification decision about the sentence
- ▶ **Encoding of each word** – can pass this to another layer to make a prediction (can also pool these to get a different sentence encoding)
- ▶ RNN can be viewed as a transformation of a sequence of vectors into a sequence of context-dependent vectors

# This Lecture

---

- ▶ CNNs
- ▶ CNNs for Sentiment
- ▶ Neural CRFs

# CNNs

# Convolutional Layer

---


- ▶ Applies a *filter* over patches of the input and returns that filter's activations
- ▶ Convolution: take dot product of filter with a patch of the input

# Convolutional Layer

---

- ▶ Applies a *filter* over patches of the input and returns that filter's activations
- ▶ Convolution: take dot product of filter with a patch of the input

image:  $n \times n \times k$




Each of these cells is a vector with multiple values  
Images: RGB values (3 dim)

# Convolutional Layer

---

- ▶ Applies a *filter* over patches of the input and returns that filter's activations
- ▶ Convolution: take dot product of filter with a patch of the input

image:  $n \times n \times k$       filter:  $m \times m \times k$





Each of these cells is a vector with multiple values  
Images: RGB values (3 dim)

# Convolutional Layer

- ▶ Applies a *filter* over patches of the input and returns that filter's activations
- ▶ Convolution: take dot product of filter with a patch of the input

image:  $n \times n \times k$     filter:  $m \times m \times k$



sum over dot products

$$\text{activation}_{ij} = \sum_{i_o=0}^{k-1} \sum_{j_o=0}^{k-1} \text{image}(i + i_o, j + j_o) \cdot \text{filter}(i_o, j_o)$$

↑  
offsets



Each of these cells is a vector with multiple values  
Images: RGB values (3 dim)

# Convolutional Layer

---

- ▶ Applies a *filter* over patches of the input and returns that filter's activations
- ▶ Convolution: take dot product of filter with a patch of the input

image:  $n \times n \times k$       filter:  $m \times m \times k$




# Convolutional Layer

---

- ▶ Applies a *filter* over patches of the input and returns that filter's activations
- ▶ Convolution: take dot product of filter with a patch of the input

image:  $n \times n \times k$       filter:  $m \times m \times k$




# Convolutional Layer

---

- ▶ Applies a *filter* over patches of the input and returns that filter's activations
- ▶ Convolution: take dot product of filter with a patch of the input

image:  $n \times n \times k$     filter:  $m \times m \times k$     activations:  $(n - m + 1) \times (n - m + 1) \times 1$



# Convolutions for NLP

---

- ▶ Input and filter are 2-dimensional instead of 3-dimensional


# Convolutions for NLP

---

- ▶ Input and filter are 2-dimensional instead of 3-dimensional

sentence:  $n$  words  $\times k$  vec dim

the movie was good



vector for each word


# Convolutions for NLP

---

- ▶ Input and filter are 2-dimensional instead of 3-dimensional

sentence:  $n$  words  $\times k$  vec dim    filter:  $m \times k$

the movie was good




vector for each word

# Convolutions for NLP

---

- ▶ Input and filter are 2-dimensional instead of 3-dimensional

sentence:  $n$  words  $\times k$  vec dim    filter:  $m \times k$




# Convolutions for NLP

---

- ▶ Input and filter are 2-dimensional instead of 3-dimensional

sentence:  $n$  words  $\times k$  vec dim    filter:  $m \times k$

activations:  $(n - m + 1) \times 1$



# Convolutions for NLP

---


- ▶ Input and filter are 2-dimensional instead of 3-dimensional

sentence:  $n$  words  $\times k$  vec dim

filter:  $m \times k$

activations:  $(n - m + 1) \times 1$

the movie was good



- ▶ Combines evidence locally in a sentence and produces a new (but still variable-length) representation

# Compare: CNNs vs. LSTMs

---




the movie was good

# Compare: CNNs vs. LSTMs

---



c filters,  
 $m \times k$  each



$n \times k$

the movie was good

# Compare: CNNs vs. LSTMs

---



$O(n) \times c$



c filters,  
 $m \times k$  each



$n \times k$

the movie was good

# Compare: CNNs vs. LSTMs


---



$O(n) \times c$



c filters,  
 $m \times k$  each



$n \times k$

the movie was good



$n \times k$

the movie was good

# Compare: CNNs vs. LSTMs

---



$O(n) \times c$




c filters,  
 $m \times k$  each



$n \times k$

the movie was good




BiLSTM with  
hidden size c

$n \times k$

the movie was good

# Compare: CNNs vs. LSTMs

---



$O(n) \times c$




c filters,  
 $m \times k$  each



$n \times k$

the movie was good




BiLSTM with  
hidden size c

$n \times k$

the movie was good

# Compare: CNNs vs. LSTMs


---



$O(n) \times c$




c filters,  
 $m \times k$  each



$n \times k$

the movie was good



$n \times 2c$


BiLSTM with  
hidden size  $c$

$n \times k$


the movie was good

# Compare: CNNs vs. LSTMs

---




$O(n) \times c$



$n \times k$

the movie was good



$n \times 2c$

BiLSTM with  
hidden size  $c$


$n \times k$

the movie was good

- ▶ Both LSTMs and convolutional layers transform the input using context

# Compare: CNNs vs. LSTMs


---



$O(n) \times c$




c filters,  
 $m \times k$  each



$n \times k$

the movie was good



$n \times 2c$

BiLSTM with  
hidden size  $c$


$n \times k$

the movie was good

- ▶ Both LSTMs and convolutional layers transform the input using context
- ▶ LSTM: “globally” looks at the entire sentence (but local for many problems)

# Compare: CNNs vs. LSTMs

---



$O(n) \times c$




c filters,  
 $m \times k$  each



$n \times k$

the movie was good



$n \times 2c$

BiLSTM with  
hidden size  $c$

$n \times k$


the movie was good

- ▶ Both LSTMs and convolutional layers transform the input using context
- ▶ LSTM: “globally” looks at the entire sentence (but local for many problems)
- ▶ CNN: local depending on filter width + number of layers

# CNNs for Sentiment

# CNNs for Sentiment Analysis

---




$n \times c$   
c filters,  
 $m \times k$  each  
 $n \times k$

the movie was good

# CNNs for Sentiment Analysis

---



$c$ -dimensional vector

max pooling over the sentence

$n \times c$


$c$  filters,  
 $m \times k$  each

$n \times k$

- ▶ Max pooling: return the max activation of a given filter over the entire sentence; like a logical OR (sum pooling is like logical AND)

the movie was good

# CNNs for Sentiment Analysis



$$P(y|x)$$

projection + softmax

$c$ -dimensional vector

max pooling over the sentence

$$n \times c$$

$c$  filters,  
 $m \times k$  each


$$n \times k$$


the movie was good

- ▶ Max pooling: return the max activation of a given filter over the entire sentence; like a logical OR (sum pooling is like logical AND)

# Understanding CNNs for Sentiment


---

*the* 

*movie* 

*was* 

*good* 

. 

- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment


---



- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment


---



- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment


---



- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment


---



- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment



---



- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment


---

<i>the</i>		0.03
<i>movie</i>		0.02
<i>was</i>		
<i>good</i>		
.		

- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment


---



- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment



---



- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment


---

<i>the</i>		0.03
<i>movie</i>		0.02
<i>was</i>		0.1
<i>good</i>		
.		

- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment


---



- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment



---



- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment




---

<i>the</i>		0.03
<i>movie</i>		0.02
<i>was</i>		0.1
<i>good</i>		1.1
.		

- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment


---

<i>the</i>		0.03
<i>movie</i>		0.02
<i>was</i>		0.1
<i>good</i>		1.1
.		

- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment



---



- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment


---

<i>the</i>		0.03
<i>movie</i>		0.02
<i>was</i>		0.1
<i>good</i>		1.1
.		0.0

- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment


---



- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment


---




- ▶ Filter “looks like” the things that will cause it to have high activation

# Understanding CNNs for Sentiment


---




# Understanding CNNs for Sentiment




# Understanding CNNs for Sentiment




# Understanding CNNs for Sentiment




# Understanding CNNs for Sentiment



# Understanding CNNs for Sentiment




# Understanding CNNs for Sentiment




- ▶ Takes variable-length input and turns it into fixed-length output

# Understanding CNNs for Sentiment



- ▶ Takes variable-length input and turns it into fixed-length output
- ▶ Filters are initialized randomly and then learned


# Understanding CNNs for Sentiment





- ▶ Word vectors for similar words are similar, so convolutional filters will have similar outputs

# Understanding CNNs for Sentiment

---

*the* 


*movie* 

*was* 

*not* 


*good* 

.




# Understanding CNNs for Sentiment

---




# Understanding CNNs for Sentiment

---




# Understanding CNNs for Sentiment

---




# Understanding CNNs for Sentiment

---




# Understanding CNNs for Sentiment

---




# Understanding CNNs for Sentiment


---




# Understanding CNNs for Sentiment



# Understanding CNNs for Sentiment




# Understanding CNNs for Sentiment



- Analogous to bigram features in bag-of-words models

# Understanding CNNs for Sentiment



- ▶ Analogous to bigram features in bag-of-words models
- ▶ Indicator feature of text containing bigram  $\leftrightarrow$  max pooling of a filter that matches that bigram

# What can CNNs learn?

---

*the movie was not good*



*the movie was not really all that good*

*the cinematography was good, the music great, but the movie was bad*

*I entered the theater in the bloom of youth and left as an old man*

# Deep Convolutional Networks



- ▶ Low-level filters: extract low-level features from the data



Zeiler and Fergus (2014)

# Deep Convolutional Networks

- ▶ High-level filters: match larger and more “semantic patterns”



Zeiler and Fergus (2014)

# CNNs: Implementation

---

- ▶ Input is  $\text{batch\_size} \times n \times k$  matrix, filters are  $c \times m \times k$  matrix ( $c$  filters)

# CNNs: Implementation

---

- ▶ Input is  $\text{batch\_size} \times n \times k$  matrix, filters are  $c \times m \times k$  matrix ( $c$  filters)
- ▶ Typically use filters with  $m$  ranging from 1 to 5 or so (multiple filter widths in a single convnet)


# CNNs: Implementation

---

- ▶ Input is  $\text{batch\_size} \times n \times k$  matrix, filters are  $c \times m \times k$  matrix ( $c$  filters)
- ▶ Typically use filters with  $m$  ranging from 1 to 5 or so (multiple filter widths in a single convnet)
- ▶ All computation graph libraries support efficient convolution operations

# CNNs for Sentence Classification

- ▶ Question classification, sentiment, etc.
- ▶ Conv+pool, then use feedforward layers to classify
- ▶ Can use multiple types of input vectors (fixed initializer and learned)



the movie was good

Kim (2014)

# Sentence Classification


---

Model	MR	SST-1	SST-2	Subj	TREC	CR	MPQA
CNN-multichannel	81.1	47.4	88.1	93.2	92.2	85.0	89.4
NBSVM (Wang and Manning, 2012)	79.4	–	–	93.2	–	81.8	86.3

# Sentence Classification

---

movie review  
sentiment



Model	MR	SST-1	SST-2	Subj	TREC	CR	MPQA
CNN-multichannel	81.1	47.4	88.1	93.2	92.2	85.0	89.4
NBSVM (Wang and Manning, 2012)	79.4	-	-	93.2	-	81.8	86.3

# Sentence Classification

---

movie review  
sentiment

subjectivity/objectivity  
detection

Model	MR	SST-1	SST-2	Subj	TREC	CR	MPQA
CNN-multichannel	81.1	47.4	88.1	93.2	92.2	85.0	89.4
NBSVM (Wang and Manning, 2012)	79.4	-	-	93.2	-	81.8	86.3

# Sentence Classification

Model	MR	SST-1	SST-2	Subj	TREC	CR	MPQA
CNN-multichannel	81.1	47.4	88.1	93.2	92.2	85.0	89.4
NBSVM (Wang and Manning, 2012)	79.4	-	-	93.2	-	81.8	86.3

movie review sentiment

subjectivity/objectivity detection

question type classification

# Sentence Classification

Model	MR	SST-1	SST-2	Subj	TREC	CR	MPQA
CNN-multichannel	81.1	47.4	88.1	93.2	92.2	85.0	89.4
NBSVM (Wang and Manning, 2012)	79.4	-	-	93.2	-	81.8	86.3

movie review  
sentiment

subjectivity/objectivity  
detection

product  
reviews

question type  
classification

# Sentence Classification

Model	MR	SST-1	SST-2	Subj	TREC	CR	MPQA
CNN-multichannel	81.1	47.4	88.1	93.2	92.2	85.0	89.4
NBSVM (Wang and Manning, 2012)	79.4	-	-	93.2	-	81.8	86.3

movie review  
sentiment

subjectivity/objectivity  
detection

product  
reviews

question type  
classification

```
graph TD; A[Movie review sentiment] --> B[Model]; C[Subjectivity/objectivity detection] --> B; D[Product reviews] --> B; E[Question type classification] --> B; B --- R1[NBSVM]; B --- R2[CNN];
```

- ▶ Also effective at document-level text classification

# Neural CRF Basics

# NER Revisited

---

B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

*Barack Obama will travel to Hangzhou today for the G20 meeting .*

PERSON

LOC

ORG

# NER Revisited

- Features in CRFs:  $I[\text{tag}=\text{B-LOC} \& \text{curr\_word}=\text{Hangzhou}]$ ,  
 $I[\text{tag}=\text{B-LOC} \& \text{prev\_word}=to]$ ,  $I[\text{tag}=\text{B-LOC} \& \text{curr\_prefix}=\text{Han}]$

# NER Revisited

- ▶ Features in CRFs:  $I[\text{tag}=\text{B-LOC} \ \& \ \text{curr\_word}=\text{Hangzhou}]$ ,  
 $I[\text{tag}=\text{B-LOC} \ \& \ \text{prev\_word}=to]$ ,  $I[\text{tag}=\text{B-LOC} \ \& \ \text{curr\_prefix}=\text{Han}]$
  - ▶ Linear model over features

# NER Revisited

B-PER I-PER O O O B-LOC : O O O B-ORG O O

*Barack Obama* will travel to *Hangzhou* today for the *G20* meeting .

PERSON LOC ORG

- ▶ Features in CRFs:  $I[\text{tag}=\text{B-LOC} \& \text{curr\_word}=\text{Hangzhou}]$ ,  
 $I[\text{tag}=\text{B-LOC} \& \text{prev\_word}=\text{to}]$ ,  $I[\text{tag}=\text{B-LOC} \& \text{curr\_prefix}=\text{Han}]$
  - ▶ Linear model over features
  - ▶ Downsides:

# NER Revisited

- ▶ Features in CRFs:  $I[\text{tag}=\text{B-LOC} \ \& \ \text{curr\_word}=\text{Hangzhou}]$ ,  
 $I[\text{tag}=\text{B-LOC} \ \& \ \text{prev\_word}=\text{to}]$ ,  $I[\text{tag}=\text{B-LOC} \ \& \ \text{curr\_prefix}=\text{Han}]$
  - ▶ Linear model over features
  - ▶ Downsides:
    - ▶ Lexical features mean that words need to be seen in the training data

# NER Revisited

- ▶ Features in CRFs:  $I[\text{tag}=\text{B-LOC} \ \& \ \text{curr\_word}=\text{Hangzhou}]$ ,  
 $I[\text{tag}=\text{B-LOC} \ \& \ \text{prev\_word}=to]$ ,  $I[\text{tag}=\text{B-LOC} \ \& \ \text{curr\_prefix}=\text{Han}]$
  - ▶ Linear model over features
  - ▶ Downsides:
    - ▶ Lexical features mean that words need to be seen in the training data
    - ▶ Linear model can't capture feature conjunctions as effectively (doesn't work well to look at more than 2 words with a single feature)

# LSTMs for NER

B-PER I-PER O O O B-LOC O O O B-ORG O O


*Barack Obama will travel to Hangzhou today for the G20 meeting .*

PERSON

LOC

ORG

B-PER I-PER O O O B-LOC



Barack Obama will travel to Hangzhou

- ▶ Transducer (LM-like model)
- ▶ What are the strengths and weaknesses of this model compared to CRFs?

# LSTMs for NER

B-PER I-PER O O O B-LOC O O O B-ORG O O


*Barack Obama will travel to Hangzhou today for the G20 meeting.*

PERSON

LOC

ORG

B-PER I-PER O O O B-LOC



Barack Obama will travel to Hangzhou

- ▶ Bidirectional transducer model
- ▶ What are the strengths and weaknesses of this model compared to CRFs?

# Neural CRFs


B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

*Barack Obama will travel to Hangzhou today for the G20 meeting.*

PERSON

LOC

ORG



Barack Obama will travel to Hangzhou

# Neural CRFs


B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

*Barack Obama will travel to Hangzhou today for the G20 meeting.*

PERSON

LOC

ORG



Barack Obama will travel to Hangzhou

# Neural CRFs


B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

*Barack Obama will travel to Hangzhou today for the G20 meeting.*

PERSON

LOC

ORG




Barack Obama will travel to Hangzhou

- ▶ Neural CRFs: bidirectional LSTMs (or some NN) compute emission potentials, capture structural constraints in transition potentials

# Neural CRFs

---

$$P(\mathbf{y}|\mathbf{x}) = \frac{1}{Z} \prod_{i=2}^n \exp(\phi_t(y_{i-1}, y_i)) \prod_{i=1}^n \exp(\phi_e(y_i, i, \mathbf{x}))$$




- ▶ Conventional:  $\phi_e(y_i, i, \mathbf{x}) = w^\top f_e(y_i, i, \mathbf{x})$

# Neural CRFs

---

$$P(\mathbf{y}|\mathbf{x}) = \frac{1}{Z} \prod_{i=2}^n \exp(\phi_t(y_{i-1}, y_i)) \prod_{i=1}^n \exp(\phi_e(y_i, i, \mathbf{x}))$$




- ▶ Conventional:  $\phi_e(y_i, i, \mathbf{x}) = w^\top f_e(y_i, i, \mathbf{x})$
- ▶ Neural:  $\phi_e(y_i, i, \mathbf{x}) = W_{y_i}^\top f(i, \mathbf{x})$     $W$  is a `num_tags x len(f)` matrix

# Neural CRFs

---

$$P(\mathbf{y}|\mathbf{x}) = \frac{1}{Z} \prod_{i=2}^n \exp(\phi_t(y_{i-1}, y_i)) \prod_{i=1}^n \exp(\phi_e(y_i, i, \mathbf{x}))$$




- ▶ Conventional:  $\phi_e(y_i, i, \mathbf{x}) = w^\top f_e(y_i, i, \mathbf{x})$
- ▶ Neural:  $\phi_e(y_i, i, \mathbf{x}) = W_{y_i}^\top f(i, \mathbf{x})$     $W$  is a `num_tags x len(f)` matrix
- ▶  $f(i, \mathbf{x})$  could be the output of a feedforward neural network looking at the words around position  $i$ , or the  $i$ th output of an LSTM, ...

# Neural CRFs

---

$$P(\mathbf{y}|\mathbf{x}) = \frac{1}{Z} \prod_{i=2}^n \exp(\phi_t(y_{i-1}, y_i)) \prod_{i=1}^n \exp(\phi_e(y_i, i, \mathbf{x}))$$




- ▶ Conventional:  $\phi_e(y_i, i, \mathbf{x}) = w^\top f_e(y_i, i, \mathbf{x})$
- ▶ Neural:  $\phi_e(y_i, i, \mathbf{x}) = W_{y_i}^\top f(i, \mathbf{x})$     $W$  is a `num_tags` x `len(f)` matrix
- ▶  $f(i, \mathbf{x})$  could be the output of a feedforward neural network looking at the words around position  $i$ , or the  $i$ th output of an LSTM, ...
- ▶ Neural network computes unnormalized potentials that are consumed and “normalized” by a structured model

# Neural CRFs

---


$$P(\mathbf{y}|\mathbf{x}) = \frac{1}{Z} \prod_{i=2}^n \exp(\phi_t(y_{i-1}, y_i)) \prod_{i=1}^n \exp(\phi_e(y_i, i, \mathbf{x}))$$



- ▶ Conventional:  $\phi_e(y_i, i, \mathbf{x}) = w^\top f_e(y_i, i, \mathbf{x})$
- ▶ Neural:  $\phi_e(y_i, i, \mathbf{x}) = W_{y_i}^\top f(i, \mathbf{x})$     $W$  is a `num_tags x len(f)` matrix
- ▶  $f(i, \mathbf{x})$  could be the output of a feedforward neural network looking at the words around position  $i$ , or the  $i$ th output of an LSTM, ...
- ▶ Neural network computes unnormalized potentials that are consumed and “normalized” by a structured model
- ▶ Inference: compute  $f$ , use Viterbi

# Computing Gradients

---


$$P(\mathbf{y}|\mathbf{x}) = \frac{1}{Z} \prod_{i=2}^n \exp(\phi_t(y_{i-1}, y_i)) \prod_{i=1}^n \exp(\phi_e(y_i, i, \mathbf{x}))$$


- ▶ Conventional:  $\phi_e(y_i, i, \mathbf{x}) = w^\top f_e(y_i, i, \mathbf{x})$
- ▶ Neural:  $\phi_e(y_i, i, \mathbf{x}) = W_{y_i}^\top f(i, \mathbf{x})$

# Computing Gradients

---

$$P(\mathbf{y}|\mathbf{x}) = \frac{1}{Z} \prod_{i=2}^n \exp(\phi_t(y_{i-1}, y_i)) \prod_{i=1}^n \exp(\phi_e(y_i, i, \mathbf{x}))$$




- ▶ Conventional:  $\phi_e(y_i, i, \mathbf{x}) = w^\top f_e(y_i, i, \mathbf{x})$
  - ▶ Neural:  $\phi_e(y_i, i, \mathbf{x}) = W_{y_i}^\top f(i, \mathbf{x})$
- $$\frac{\partial \mathcal{L}}{\partial \phi_{e,i}} = -P(y_i = s | \mathbf{x}) + I[s \text{ is gold}] \text{ “error signal”, compute with F-B}$$

# Computing Gradients

---


$$P(\mathbf{y}|\mathbf{x}) = \frac{1}{Z} \prod_{i=2}^n \exp(\phi_t(y_{i-1}, y_i)) \prod_{i=1}^n \exp(\phi_e(y_i, i, \mathbf{x}))$$



- ▶ Conventional:  $\phi_e(y_i, i, \mathbf{x}) = w^\top f_e(y_i, i, \mathbf{x})$
- ▶ Neural:  $\phi_e(y_i, i, \mathbf{x}) = W_{y_i}^\top f(i, \mathbf{x})$
- $$\frac{\partial \mathcal{L}}{\partial \phi_{e,i}} = -P(y_i = s | \mathbf{x}) + I[s \text{ is gold}]$$
 “error signal”, compute with F-B
- ▶ For linear model: 
$$\frac{\partial \phi_{e,i}}{w_i} = f_{e,i}(y_i, i, \mathbf{x})$$

# Computing Gradients

$$P(\mathbf{y}|\mathbf{x}) = \frac{1}{Z} \prod_{i=2}^n \exp(\phi_t(y_{i-1}, y_i)) \prod_{i=1}^n \exp(\phi_e(y_i, i, \mathbf{x}))$$



► Conventional:  $\phi_e(y_i, i, \mathbf{x}) = w^\top f_e(y_i, i, \mathbf{x})$

► Neural:  $\phi_e(y_i, i, \mathbf{x}) = W_{y_i}^\top f(i, \mathbf{x})$

$$\frac{\partial \mathcal{L}}{\partial \phi_{e,i}} = -P(y_i = s | \mathbf{x}) + I[s \text{ is gold}]$$


“error signal”, compute with F-B

► For linear model:  $\frac{\partial \phi_{e,i}}{w_i} = f_{e,i}(y_i, i, \mathbf{x})$

chain rule say to multiply together, gives our update

# Computing Gradients

$$P(\mathbf{y}|\mathbf{x}) = \frac{1}{Z} \prod_{i=2}^n \exp(\phi_t(y_{i-1}, y_i)) \prod_{i=1}^n \exp(\phi_e(y_i, i, \mathbf{x}))$$



- ▶ Conventional:  $\phi_e(y_i, i, \mathbf{x}) = w^\top f_e(y_i, i, \mathbf{x})$
- ▶ Neural:  $\phi_e(y_i, i, \mathbf{x}) = W_{y_i}^\top f(i, \mathbf{x})$
- ▶ For linear model:  $\frac{\partial \mathcal{L}}{\partial \phi_{e,i}} = -P(y_i = s | \mathbf{x}) + I[s \text{ is gold}]$  “error signal”, compute with F-B
  - ▶ chain rule say to multiply together, gives our update
- ▶ For neural model: compute gradient of phi w.r.t. parameters of neural net

# Neural CRFs


B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

*Barack Obama will travel to Hangzhou today for the G20 meeting .*

PERSON

LOC

ORG



Barack Obama will travel to Hangzhou

# Neural CRFs


B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

*Barack Obama will travel to Hangzhou today for the G20 meeting.*

PERSON

LOC

ORG



1) Compute  $f(x)$

Barack Obama will travel to Hangzhou

# Neural CRFs


B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

*Barack Obama will travel to Hangzhou today for the G20 meeting.*

PERSON

LOC

ORG



Barack Obama will travel to Hangzhou

2) Run forward-backward

1) Compute  $f(x)$

# Neural CRFs


B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

*Barack Obama will travel to Hangzhou today for the G20 meeting.*

PERSON

LOC

ORG



Barack Obama will travel to Hangzhou

2) Run forward-backward

3) Compute error signal

1) Compute  $f(x)$

# Neural CRFs


B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

*Barack Obama will travel to Hangzhou today for the G20 meeting .*

PERSON

LOC

ORG



Barack Obama will travel to Hangzhou

- 1) Compute  $f(x)$
- 2) Run forward-backward
- 3) Compute error signal
- 4) Backprop (no knowledge of sequential structure required)

# FFNN Neural CRF for NER

---

B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

*Barack Obama will travel to Hangzhou today for the G20 meeting .*

PERSON

LOC

ORG

# FFNN Neural CRF for NER


B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

*Barack Obama will travel to Hangzhou today for the G20 meeting .*

PERSON

LOC

ORG



*to Hangzhou today*

# FFNN Neural CRF for NER


B-PER I-PER O O O B-LOC O O O B-ORG O O

*Barack Obama will travel to Hangzhou today for the G20 meeting.*

PERSON

LOC

ORG



$$f(\mathbf{x}, i) = [\text{emb}(\mathbf{x}_{i-1}), \text{emb}(\mathbf{x}_i), \text{emb}(\mathbf{x}_{i+1})]$$



previous word curr word next word

*to Hangzhou today*

# FFNN Neural CRF for NER


B-PER I-PER 0 0 0 B-LOC 0 0 0 B-ORG 0 0

*Barack Obama will travel to Hangzhou today for the G20 meeting.*

PERSON

LOC

ORG



$$f(\mathbf{x}, i) = [\text{emb}(\mathbf{x}_{i-1}), \text{emb}(\mathbf{x}_i), \text{emb}(\mathbf{x}_{i+1})]$$

previous word curr word next word

*to Hangzhou today*

# FFNN Neural CRF for NER


B-PER I-PER O O O B-LOC O O O B-ORG O O

*Barack Obama will travel to Hangzhou today for the G20 meeting.*

PERSON

LOC

ORG



$$\phi_e = Wg(Vf(\mathbf{x}, i))$$

$$f(\mathbf{x}, i) = [\text{emb}(\mathbf{x}_{i-1}), \text{emb}(\mathbf{x}_i), \text{emb}(\mathbf{x}_{i+1})]$$

previous word curr word next word

*to Hangzhou today*

# LSTM Neural CRFs


B-PER	I-PER	0	0	0	B-LOC	0	0	0	B-ORG	0	0
-------	-------	---	---	---	-------	---	---	---	-------	---	---

*Barack Obama will travel to Hangzhou today for the G20 meeting .*

PERSON

LOC

ORG



Barack Obama will travel to Hangzhou

- Bidirectional LSTMs compute emission (or transition) potentials

# LSTMs for NER

B-PER I-PER O O O B-LOC O O O B-ORG O O


*Barack Obama will travel to Hangzhou today for the G20 meeting.*

PERSON

LOC

ORG

B-PER I-PER O O O B-LOC



Barack Obama will travel to Hangzhou

- ▶ How does this compare to neural CRF?


# “NLP (Almost) From Scratch”



Collobert, Weston, et al. 2008, 2011

# “NLP (Almost) From Scratch”


<b>Approach</b>	<b>POS</b> (PWA)	<b>CHUNK</b> (F1)	<b>NER</b> (F1)	<b>SRL</b> (F1)
<b>Benchmark Systems</b>	97.24	94.29	89.31	77.92
NN+WLL	96.31	89.13	79.53	55.40
NN+SLL	96.37	90.33	81.47	70.99
NN+WLL+LM1	97.05	91.91	85.68	58.18
NN+SLL+LM1	97.10	93.65	87.58	73.84
NN+WLL+LM2	97.14	92.04	86.96	58.34
NN+SLL+LM2	97.20	93.63	88.67	74.15



# “NLP (Almost) From Scratch”

<b>Approach</b>	<b>POS</b> (PWA)	<b>CHUNK</b> (F1)	<b>NER</b> (F1)	<b>SRL</b> (F1)
<b>Benchmark Systems</b>	97.24	94.29	89.31	77.92
NN+WLL	96.31	89.13	79.53	55.40
NN+SLL	96.37	90.33	81.47	70.99
NN+WLL+LM1	97.05	91.91	85.68	58.18
NN+SLL+LM1	97.10	93.65	87.58	73.84
NN+WLL+LM2	97.14	92.04	86.96	58.34
NN+SLL+LM2	97.20	93.63	88.67	74.15


- WLL: independent classification; SLL: neural CRF



# “NLP (Almost) From Scratch”

<b>Approach</b>	<b>POS</b> (PWA)	<b>CHUNK</b> (F1)	<b>NER</b> (F1)	<b>SRL</b> (F1)
<b>Benchmark Systems</b>	97.24	94.29	89.31	77.92
NN+WLL	96.31	89.13	79.53	55.40
NN+SLL	96.37	90.33	81.47	70.99
NN+WLL+LM1	97.05	91.91	85.68	58.18
NN+SLL+LM1	97.10	93.65	87.58	73.84
NN+WLL+LM2	97.14	92.04	86.96	58.34
NN+SLL+LM2	97.20	93.63	88.67	74.15

- ▶ WLL: independent classification; SLL: neural CRF
- ▶ LM2: word vectors learned from a precursor to word2vec/GloVe, trained for 2 weeks (!) on Wikipedia



# CNN Neural CRFs

---

*travel to Hangzhou today for*

# CNN Neural CRFs

---


- ▶ Append to each word vector an *embedding of the relative position* of that word

*travel to Hangzhou today for*

# CNN Neural CRFs


---

- ▶ Append to each word vector an *embedding of the relative position* of that word



# CNN Neural CRFs


---



- ▶ Append to each word vector an *embedding of the relative position* of that word

# CNN Neural CRFs

---



- ▶ Append to each word vector an *embedding of the relative position* of that word
- ▶ Convolution over the sentence produces a position-dependent representation

# CNN NCRFs vs. FFNN NCRFs

---

<b>Approach</b>	<b>POS</b> (PWA)	<b>CHUNK</b> (F1)	<b>NER</b> (F1)	<b>SRL</b> (F1)
<b>Benchmark Systems</b>	97.24	94.29	89.31	77.92
<i>Window Approach</i>				
NN+SLL+LM2	97.20	93.63	88.67	-
<i>Sentence Approach</i>				
NN+SLL+LM2	97.12	93.37	88.78	74.15

- ▶ Sentence approach (CNNs) is comparable to window approach (FFNNs) except for SRL where they claim it works much better


# Neural CRFs with LSTMs

---

- ▶ Neural CRF using character LSTMs to compute word representations



# Neural CRFs with LSTMs

- ▶ Neural CRF using character LSTMs to compute word representations



# Neural CRFs with LSTMs

- ▶ Neural CRF using character LSTMs to compute word representations



# Neural CRFs with LSTMs

- ▶ Chiu+Nichols: character CNNs instead of LSTMs
- ▶ Lin/Passos/Luo: use external resources like Wikipedia
- ▶ LSTM-CRF captures the important aspects of NER: word context (LSTM), sub-word features (character LSTMs), outside knowledge (word embeddings)

Model	F <sub>1</sub>
Collobert et al. (2011)*	89.59
Lin and Wu (2009)	83.78
Lin and Wu (2009)*	90.90
Huang et al. (2015)*	90.10
Passos et al. (2014)	90.05
Passos et al. (2014)*	90.90
Luo et al. (2015)* + gaz	89.9
Luo et al. (2015)* + gaz + linking	<b>91.2</b>
Chiu and Nichols (2015)	90.69
Chiu and Nichols (2015)*	90.77
LSTM-CRF (no char)	90.20
LSTM-CRF	<b>90.94</b>

# Takeaways

---

- ▶ CNNs are a flexible way of extracting features analogous to bag of n-grams, can also encode positional information
- ▶ All kinds of NNs can be integrated into CRFs for structured inference. Can be applied to NER, other tagging, parsing, ...