

Lecture 17: Unsupervised Learning

Alan Ritter


(many slides from Greg Durrett)

Administrivia

- ▶ Wei Xu will present on Friday
- ▶ No Class on December 4
- ▶ Final Project Presentations are during the final exam time scheduled on December 12

What data do we learn from?

- ▶ Supervised settings:
 - ▶ Tagging: POS, NER
 - ▶ Parsing: constituency, dependency, semantic parsing
 - ▶ IE, MT, QA, ...
- ▶ Semi-supervised models
 - ▶ Word embeddings / word clusters (helpful for nearly all tasks)
 - ▶ Language models for machine translation
 - ▶ Learn linguistic structure from unlabeled data and use it?




This Lecture

- ▶ Discrete linguistic structure from generative models: unsupervised POS induction
 - ▶ Expectation maximization for learning HMMs
- ▶ Continuous structure with generative models: variational autoencoders
- ▶ Continuous structure with “discriminative” models: transfer learning

EM for HMMs

Recall: Hidden Markov Models

- ▶ Input $\mathbf{x} = (x_1, \dots, x_n)$ Output $\mathbf{y} = (y_1, \dots, y_n)$



$$P(\mathbf{y}, \mathbf{x}) = P(y_1) \underbrace{\prod_{i=2}^n P(y_i | y_{i-1})}_{\text{Initial distribution probabilities}} \underbrace{\prod_{i=1}^n P(x_i | y_i)}_{\text{Transition probabilities Emission probabilities}}$$

- ▶ Observation (x) depends only on current state (y)
- ▶ Multinomials: tag x tag transitions, tag x word emissions
- ▶ $P(x|y)$ is a distribution over all words in the vocabulary
 - not a distribution over features (but could be!)

Unsupervised Learning

- ▶ Can we induce linguistic structure? Thought experiment...

a b a c c c c

b a c c c

- ▶ What's a two-state HMM that could produce this?
- ▶ What if I show you this sequence?

a a b c c a a

- ▶ What did you do? Use current model parameters + data to refine your model. This is what EM will do

Part-of-Speech Induction

- ▶ Input $\mathbf{x} = (x_1, \dots, x_n)$ Output $\mathbf{y} = (y_1, \dots, y_n)$
- ▶ Assume we don't have access to labeled examples — how can we learn a POS tagger?
- ▶ Key idea: optimize $P(\mathbf{x}) = \sum_{\mathbf{y}} P(\mathbf{y}, \mathbf{x}) \leftarrow$ Generative model explains the data \mathbf{x} ; the right HMM makes it look likely
- ▶ Optimizing marginal log-likelihood with no labels \mathbf{y} :

$$\mathcal{L}(\mathbf{x}_1, \dots, \mathbf{x}_D) = \sum_{i=1}^D \log \sum_{\mathbf{y}} P(\mathbf{y}, \mathbf{x}_i)$$

▶ non-convex optimization problem

Part-of-Speech Induction

- ▶ Input $\mathbf{x} = (x_1, \dots, x_n)$ Output $\mathbf{y} = (y_1, \dots, y_n)$
- ▶ Optimizing marginal log-likelihood with no labels \mathbf{y} :

$$\mathcal{L}(\mathbf{x}_1, \dots, \mathbf{x}_D) = \sum_{i=1}^D \log \sum_{\mathbf{y}} P(\mathbf{y}, \mathbf{x}_i)$$

- ▶ Can't use a discriminative model; $\sum_{\mathbf{y}} P(\mathbf{y}|\mathbf{x}) = 1$, doesn't model \mathbf{x}
- ▶ What's the point of this? Model has inductive bias and so should learn some useful latent structure \mathbf{y} (clustering effect)
- ▶ EM is just one procedure for optimizing this kind of objective

Expectation Maximization

$$\begin{aligned} & \log \sum_{\mathbf{y}} P(\mathbf{x}, \mathbf{y} | \theta) \\ &= \log \sum_{\mathbf{y}} q(\mathbf{y}) \frac{P(\mathbf{x}, \mathbf{y} | \theta)}{q(\mathbf{y})} \\ &\geq \sum_{\mathbf{y}} q(\mathbf{y}) \log \frac{P(\mathbf{x}, \mathbf{y} | \theta)}{q(\mathbf{y})} \\ &= \mathbb{E}_{q(\mathbf{y})} \log P(\mathbf{x}, \mathbf{y} | \theta) + \text{Entropy}[q(\mathbf{y})] \end{aligned}$$

- ▶ Condition on parameters θ
- ▶ Variational approximation q — this is a trick we'll return to later!
- ▶ Jensen's inequality (uses concavity of \log)

- ▶ Can optimize this lower-bound on log likelihood instead of log-likelihood

Expectation Maximization

$$\log \sum_{\mathbf{y}} P(\mathbf{x}, \mathbf{y} | \theta) \geq \mathbb{E}_{q(\mathbf{y})} \log P(\mathbf{x}, \mathbf{y} | \theta) + \text{Entropy}[q(\mathbf{y})]$$


- ▶ If $q(\mathbf{y}) = P(\mathbf{y} | \mathbf{x}, \theta)$, this bound ends up being tight
- ▶ Expectation-maximization: alternating maximization of the lower bound over q and θ
 - ▶ Current timestep = t , have parameters θ^{t-1}
 - ▶ E-step: maximize w.r.t. q ; that is, $q^t = P(\mathbf{y} | \mathbf{x}, \theta^{t-1})$
 - ▶ M-step: maximize w.r.t. θ ; that is, $\theta^t = \operatorname{argmax}_{\theta} \mathbb{E}_{q^t} \log P(\mathbf{x}, \mathbf{y} | \theta)$

EM for HMMs


- ▶ Expectation-maximization: alternating maximization
 - ▶ E-step: maximize w.r.t. q ; that is, $q^t = P(\mathbf{y}|\mathbf{x}, \theta^{t-1})$
 - ▶ M-step: maximize w.r.t. θ ; that is, $\theta^t = \operatorname{argmax}_{\theta} \mathbb{E}_{q^t} \log P(\mathbf{x}, \mathbf{y}|\theta)$
- ▶ E-step: for an HMM: run forward-backward with the given parameters
- ▶ Compute $P(y_i = s|\mathbf{x}, \theta^{t-1})$, $P(y_i = s_1, y_{i+1} = s_2|\mathbf{x}, \theta^{t-1})$
 - tag marginals at each position
 - tag pair marginals at each position
- ▶ M-step: set parameters to optimize the crazy argmax term

M-Step

- Recall how we maximized $\log P(x,y)$: read counts off data




- Same procedure, but maximizing $P(x,y)$ in expectation under q means that q specifies *fractional counts*



M-Step

- ▶ Same for transition probabilities



How does EM learn things?

- ▶ Initialize (M-step 0):

- ▶ Emissions

$$P(\text{the} \mid \text{DT}) = 0.9$$

$$P(\text{dog} \mid \text{DT}) = 0.05$$

$$P(\text{marsupial} \mid \text{DT}) = 0.05$$

$$P(\text{the} \mid \text{NN}) = 0.05$$

$$P(\text{dog} \mid \text{NN}) = 0.9$$

$$P(\text{marsupial} \mid \text{NN}) = 0.05$$

- ▶ Transition probabilities: uniform

- ▶ E-step 1: (all values are approximate)

DT: 0.95 DT: 0.05

NN: 0.05 NN: 0.95

the

dog

DT: 0.95 DT: 0.5

NN: 0.05 NN: 0.5

the

marsupial

▶ uniform

How does EM learn things?

► E-step 1:

DT: **0.95** DT: 0.05

NN: 0.05 **NN: 0.95**

the *dog*

DT: **0.95** DT: 0.5

NN: 0.05 **NN: 0.5**

the **marsupial**

► M-step 1:

► Emissions aren't so different

► Transition probabilities (approx): $P(\text{NN} | \text{DT}) = 3/4$, $P(\text{DT} | \text{DT}) = 1/4$

How does EM learn things?

► E-step 2:

DT: 0.95 DT: 0.05

NN: 0.05 NN: 0.95

the *dog*

DT: 0.95 DT: 0.30

NN: 0.05 NN: 0.70

the *marsupial*

► M-step 1:

► Emissions aren't so different

► Transition probabilities (approx): $P(\text{NN} | \text{DT}) = 3/4$, $P(\text{DT} | \text{DT}) = 1/4$

How does EM learn things?

► E-step 2:

DT: 0.95 DT: 0.05

NN: 0.05 NN: 0.95

the *dog*

DT: 0.95 DT: 0.30

NN: 0.05 NN: 0.70

the *marsupial*

► M-step 2:

► Emission $P(\text{marsupial} | \text{NN}) > P(\text{marsupial} | \text{DT})$

► Remember to tag marsupial as NN in the future!


► Context constrained what we learned! That's how data helped us

How does EM learn things?

- ▶ Can think of q as a kind of “fractional annotation”
- ▶ E-step: compute annotations (posterior under current model)
- ▶ M-step: supervised learning with those fractional annotations
- ▶ Initialize with some reasonable weights, alternate E and M until convergence

EM's Lower Bound


$$\mathcal{L}(\mathbf{x}_{1,\dots,D}) = \sum_{i=1}^D \log \sum_{\mathbf{y}} P(\mathbf{y}, \mathbf{x}_i)$$



Initialize probabilities θ
repeat
 ● Compute expected counts \mathbf{e}
 ● Fit parameters θ
until convergence

EM's Lower Bound

$$\mathcal{L}(\mathbf{x}_{1,\dots,D}) = \sum_{i=1}^D \log \sum_{\mathbf{y}} P(\mathbf{y}, \mathbf{x}_i)$$




Initialize probabilities θ
repeat
 ● Compute expected counts \mathbf{e}
 ● Fit parameters θ
until convergence

► E-step: compute q which gives this lower bound

EM's Lower Bound


$$\mathcal{L}(\mathbf{x}_{1,\dots,D}) = \sum_{i=1}^D \log \sum_{\mathbf{y}} P(\mathbf{y}, \mathbf{x}_i)$$



Initialize probabilities θ
repeat
 ● Compute expected counts e
 ● Fit parameters θ
until convergence

EM's Lower Bound

$$\mathcal{L}(\mathbf{x}_{1,\dots,D}) = \sum_{i=1}^D \log \sum_{\mathbf{y}} P(\mathbf{y}, \mathbf{x}_i)$$




Initialize probabilities θ
repeat
 ● Compute expected counts e
 ● Fit parameters θ
until convergence

► E-step 2: re-estimate q

EM's Lower Bound

$$\mathcal{L}(\mathbf{x}_{1,\dots,D}) = \sum_{i=1}^D \log \sum_{\mathbf{y}} P(\mathbf{y}, \mathbf{x}_i)$$




Initialize probabilities θ
repeat
 ● Compute expected counts e
 ● Fit parameters θ
until convergence

► E-step 2: re-estimate q

EM's Lower Bound


$$\mathcal{L}(\mathbf{x}_{1,\dots,D}) = \sum_{i=1}^D \log \sum_{\mathbf{y}} P(\mathbf{y}, \mathbf{x}_i)$$



Initialize probabilities θ
repeat
 ● Compute expected counts \mathbf{e}
 ● Fit parameters θ
until convergence

EM's Lower Bound

$$\mathcal{L}(\mathbf{x}_{1,\dots,D}) = \sum_{i=1}^D \log \sum_{\mathbf{y}} P(\mathbf{y}, \mathbf{x}_i)$$



Initialize probabilities θ

repeat


- Compute expected counts e

- Fit parameters θ

until convergence

EM's Lower Bound

$$\mathcal{L}(\mathbf{x}_{1,\dots,D}) = \sum_{i=1}^D \log \sum_{\mathbf{y}} P(\mathbf{y}, \mathbf{x}_i)$$



Initialize probabilities θ
repeat
 ● Compute expected counts e
 ● Fit parameters θ
until convergence

Part-of-speech Induction

- ▶ Merialdo (1994): you have a whitelist of tags for each word
- ▶ Learn parameters on k examples to start, use those to initialize EM, run on 1 million words of unlabeled data
- ▶ Tag dictionary + data should get us started in the right direction...

Part-of-speech Induction

| Number of tagged sentences used for the initial model | | | | | | | |
|---|---|------|------|------|-------|-------|------|
| | 0 | 100 | 2000 | 5000 | 10000 | 20000 | all |
| Iter | Correct tags (% words) after ML on 1M words | | | | | | |
| 0 | 77.0 | 90.0 | 95.4 | 96.2 | 96.6 | 96.9 | 97.0 |
| 1 | 80.5 | 92.6 | 95.8 | 96.3 | 96.6 | 96.7 | 96.8 |
| 2 | 81.8 | 93.0 | 95.7 | 96.1 | 96.3 | 96.4 | 96.4 |
| 3 | 83.0 | 93.1 | 95.4 | 95.8 | 96.1 | 96.2 | 96.2 |
| 4 | 84.0 | 93.0 | 95.2 | 95.5 | 95.8 | 96.0 | 96.0 |
| 5 | 84.8 | 92.9 | 95.1 | 95.4 | 95.6 | 95.8 | 95.8 |
| 6 | 85.3 | 92.8 | 94.9 | 95.2 | 95.5 | 95.6 | 95.7 |
| 7 | 85.8 | 92.8 | 94.7 | 95.1 | 95.3 | 95.5 | 95.5 |
| 8 | 86.1 | 92.7 | 94.6 | 95.0 | 95.2 | 95.4 | 95.4 |
| 9 | 86.3 | 92.6 | 94.5 | 94.9 | 95.1 | 95.3 | 95.3 |
| 10 | 86.6 | 92.6 | 94.4 | 94.8 | 95.0 | 95.2 | 95.2 |

- ▶ Small amounts of data > large amounts of unlabeled data
- ▶ Running EM *hurts* performance once you have labeled data

Two Hours of Annotation

| Human Annotations | 0. No EM | | | 1. EM only | | | 2. With LP | | |
|-------------------|----------|----|----|------------|----|----|------------|----|----|
| | T | K | U | T | K | U | T | K | U |
| Initial data | | | | | | | | | |
| KIN tokens A | 72 | 90 | 58 | 55 | 82 | 32 | 71 | 86 | 58 |
| KIN types A | | | | 63 | 77 | 32 | 78 | 83 | 69 |
| MLG tokens B | 74 | 89 | 49 | 68 | 87 | 39 | 74 | 89 | 49 |
| MLG types B | | | | 71 | 87 | 46 | 72 | 81 | 57 |
| ENG tokens A | 63 | 83 | 38 | 62 | 83 | 37 | 72 | 85 | 55 |
| ENG types A | | | | 66 | 76 | 37 | 75 | 81 | 56 |
| ENG tokens B | 70 | 87 | 44 | 70 | 87 | 43 | 78 | 90 | 60 |
| ENG types B | | | | 69 | 83 | 38 | 75 | 82 | 61 |

- ▶ Kinyarwanda and Malagasy (two actual low-resource languages)
- ▶ Label propagation (technique for using dictionary labels) helps a lot, with data that was collected in two hours

Variational Autoencoders

Continuous Latent Variables

- ▶ For discrete latent variables y , we optimized: $P(\mathbf{x}) = \sum_y P(y, \mathbf{x})$
- ▶ What if we want to use continuous latent variables?


$$P(z, \mathbf{x}) = P(z)P(\mathbf{x}|z)$$

$$P(\mathbf{x}) = \int P(z)P(\mathbf{x}|z)\partial z$$

- ▶ Can use EM here when $P(z)$ and $P(\mathbf{x}|z)$ are Gaussians
- ▶ What if we want $P(\mathbf{x}|z)$ to be something more complicated, like an LSTM with z as the initial state?

Deep Generative Models

$$P(z, \mathbf{x}) = P(z)P(\mathbf{x}|z)$$



- ▶ z is a latent variable which should control the generation of the sentence, maybe capture something about its topic

Deep Generative Models

$$\log \int_z P(\mathbf{x}, z|\theta) = \log \int_z q(z) \frac{P(\mathbf{x}, z|\theta)}{q(z)} \geq \int_z q(z) \log \frac{P(\mathbf{x}, z|\theta)}{q(z)}$$

Jensen

$$= \mathbb{E}_{q(z|\mathbf{x})}[-\log q(z|\mathbf{x}) + \log P(\mathbf{x}, z|\theta)]$$

$$= \mathbb{E}_{q(z|\mathbf{x})}[\log P(\mathbf{x}|z, \theta)] - \text{KL}(q(z|\mathbf{x}) \| P(z))$$


“make the data likely under q” “make q close to the prior”
(discriminative)

- ▶ KL divergence: distance metric over distributions (more dissimilar \Leftrightarrow higher KL)


Variational Autoencoders

$$\mathbb{E}_{q(z|x)}[\log P(x|z, \theta)] - \text{KL}(q(z|x) \| P(z))$$

Generative model (test):



Autoencoder (training):



Training VAEs

$$\mathbb{E}_{q(z|x)}[\log P(x|z, \theta)] - \text{KL}(q(z|x) \| P(z))$$


- ▶ Choose q to be Gaussian with parameters that are computed from x

$$q = N(\mu(x), \text{diag}(\sigma^2(x)))$$

- ▶ mu and sigma are computed from an LSTM over x , call their parameters ϕ

- ▶ How to handle the expectation?
Sampling

Autoencoder (training):



Training VAEs

For each example \mathbf{x}

Compute q (run forward pass to
compute mu and sigma)


For some number of samples

Sample $z \sim q$


Compute $P(\mathbf{x}|z)$ and compute loss

Backpropagate to update phi, theta

Autoencoder (training):



Autoencoders




- ▶ Another interpretation: train an autoencoder and add Gaussian noise
- ▶ Same computation graph as VAE, add KL divergence term to make the objective the same
- ▶ Inference network (q) is the encoder and generator is the decoder

Visualization

$$\mathbb{E}_{q(z|\mathbf{x})}[\log P(\mathbf{x}|z, \theta)] + \text{KL}(q(z|\mathbf{x})||P(z))$$

- ▶ What does gradient encourage latent space to do?



What do VAEs do?

- ▶ Let us encode a sentence and generate similar sentences:

| | | | |
|---------|--|-------------------------------------|-----------------------------|
| INPUT | we looked out at the setting sun . | i went to the kitchen . | how are you doing ? |
| MEAN | <i>they were laughing at the same time .</i> | <i>i went to the kitchen .</i> | <i>what are you doing ?</i> |
| SAMP. 1 | <i>ill see you in the early morning .</i> | <i>i went to my apartment .</i> | <i>“ are you sure ?</i> |
| SAMP. 2 | <i>i looked up at the blue sky .</i> | <i>i looked around the room .</i> | <i>what are you doing ?</i> |
| SAMP. 3 | <i>it was down on the dance floor .</i> | <i>i turned back to the table .</i> | <i>what are you doing ?</i> |


- ▶ Style transfer: also condition on sentiment, change sentiment
 - Positive
⇒ ARAE
⇒ Cross-AE
 - great indoor mall .
no smoking mall .
terrible outdoor urine .
- ▶ ...or use the latent representations for semi-supervised learning
 - Positive
⇒ ARAE
⇒ Cross-AE
 - it has a great atmosphere , with wonderful service .
it has no taste , with a complete jerk .
it has a great horrible food and run out service .

BERT

Goals of Unsupervised Learning

- ▶ We want to use unlabeled data, but EM “requires” generative models.
Are models like this really necessary?
- ▶ word2vec: predict nearby word given context. This wasn’t generative, but the supervision is free...
- ▶ Language modeling is a “more contextualized” form of word2vec

ELMo



learn a linear classifier on top of
this vector to get a POS tagger
with 97.3% accuracy (~SOTA)

$$P(x_i|x_1, \dots, x_{i-1}) = \text{LSTM}(x_1, \dots, x_{i-1})$$


- ▶ Generative model of the data!
- ▶ Train one model in each direction on 1B words, use the LSTM hidden states as context-aware token representations

Recall: Self-Attention

- ▶ Each word forms a “query” which then computes attention over each word

$$\alpha_{i,j} = \text{softmax}(x_i^\top x_j) \quad \text{scalar}$$



$$x'_i = \sum_{j=1}^n \alpha_{i,j} x_j \quad \text{vector} = \text{sum of scalar} * \text{vector}$$



- ▶ Multiple “heads” analogous to different convolutional filters. Use parameters W_k and V_k to get different attention values + transform vectors

$$\alpha_{k,i,j} = \text{softmax}(x_i^\top W_k x_j) \quad x'_{k,i} = \sum_{j=1}^n \alpha_{k,i,j} V_k x_j$$

Recall: Transformers




- ▶ Augment word embedding with position embeddings, each dim is a sine/cosine wave of a different frequency. Closer points = higher dot products
- ▶ Works essentially as well as just encoding position as a one-hot vector

BERT


- ▶ AI2 made ELMo in spring 2018, GPT was released in summer 2018, BERT came out October 2018
- ▶ Three major changes compared to ELMo:
 - ▶ Transformers instead of LSTMs (transformers in GPT as well)
 - ▶ Bidirectional <=> Masked LM objective instead of standard LM
 - ▶ Fine-tune instead of freeze at test time

BERT

- ▶ ELMo is a unidirectional model (as is GPT): we can concatenate two unidirectional models, but is this the right thing to do?
- ▶ ELMo reprs look at each direction in isolation; BERT looks at them jointly





A stunning ballet dancer, Copeland is one of the best performers to see live.




BERT

- ▶ How to learn a “deeply bidirectional” model? What happens if we just replace an LSTM with a transformer?




- John visited Madagascar yesterday*
- ▶ Transformer LMs have to be “one-sided” (only attend to previous tokens), not what we want

Masked Language Modeling

- ▶ How to prevent cheating? Next word prediction fundamentally doesn't work for bidirectional models, instead do *masked language modeling*
 - ▶ BERT formula: take a chunk of text, predict 15% of the tokens
 - ▶ For 80% (of the 15%), replace the input token with [MASK]
 - ▶ For 10%, replace w/random
 - ▶ For 10%, keep same
- 
- John visited [MASK] yesterday
- John visited of yesterday
- John visited Madagascar yesterday


Next “Sentence” Prediction

- ▶ Input: [CLS] Text chunk 1 [SEP] Text chunk 2
- ▶ 50% of the time, take the true next chunk of text, 50% of the time take a random other chunk. Predict whether the next chunk is the “true” next
- ▶ BERT objective: masked LM + next sentence prediction




BERT Architecture


- ▶ BERT Base: 12 layers, 768-dim per wordpiece token, 12 heads.
Total params = 110M
- ▶ BERT Large: 24 layers, 1024-dim per wordpiece token, 16 heads.
Total params = 340M



- ▶ Positional embeddings and segment embeddings, 30k word pieces
- ▶ This is the model that gets pre-trained on a large corpus



What can BERT do?




- ▶ CLS token is used to provide classification decisions
- ▶ Sentence pair tasks (entailment): feed both sentences into BERT
- ▶ BERT can also do tagging by predicting tags at each word piece

Devlin et al. (2019)


What can BERT do?

Entails



[CLS] A boy plays in the snow [SEP] A boy is outside

- ▶ How does BERT model this sentence pair stuff?
- ▶ Transformers can capture interactions between the two sentences, even though the NSP objective doesn't really cause this to happen




(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG

What can BERT NOT do?

- ▶ BERT **cannot** generate text (at least not in an obvious way)
- ▶ Not an autoregressive model, can do weird things like stick a [MASK] at the end of a string, fill in the mask, and repeat
- ▶ Masked language models are intended to be used primarily for “analysis” tasks

Fine-tuning BERT

- ▶ Fine-tune for 1-3 epochs, batch size 2-32, learning rate 2e-5 - 5e-5



(b) Single Sentence Classification Tasks:
SST-2, CoLA

- ▶ Large changes to weights up here (particularly in last layer to route the right information to [CLS])
- ▶ Smaller changes to weights lower down in the transformer
- ▶ Small LR and short fine-tuning schedule mean weights don't change much
- ▶ More complex “triangular learning rate” schemes exist

Fine-tuning BERT

| Pretraining | Adaptation | NER CoNLL 2003 | SA SST-2 | Nat. lang. inference MNLI | SICK-E | Semantic textual similarity SICK-R | MRPC | STS-B |
|---------------|------------|-------------------|-------------|------------------------------|-------------|---------------------------------------|-------------|-------------|
| Skip-thoughts | ❄️ | - | 81.8 | 62.9 | - | 86.6 | 75.8 | 71.8 |
| ELMo | ❄️ | 91.7 | 91.8 | 79.6 | 86.3 | 86.1 | 76.0 | 75.9 |
| | 🔥 | 91.9 | 91.2 | 76.4 | 83.3 | 83.3 | 74.7 | 75.5 |
| | Δ=🔥-❄️ | 0.2 | -0.6 | -3.2 | -3.3 | -2.8 | -1.3 | -0.4 |
| BERT-base | ❄️ | 92.2 | 93.0 | 84.6 | 84.8 | 86.4 | 78.1 | 82.9 |
| | 🔥 | 92.4 | 93.5 | 84.6 | 85.8 | 88.7 | 84.8 | 87.1 |
| | Δ=🔥-❄️ | 0.2 | 0.5 | 0.0 | 1.0 | 2.3 | 6.7 | 4.2 |

- ▶ BERT is typically better if the whole network is fine-tuned, unlike ELMo

Evaluation: GLUE

| Corpus | Train | Test | Task | Metrics | Domain |
|---------------------------------|-------|-------------|---------------------|------------------------------|---------------------|
| Single-Sentence Tasks | | | | | |
| CoLA | 8.5k | 1k | acceptability | Matthews corr. | misc. |
| SST-2 | 67k | 1.8k | sentiment | acc. | movie reviews |
| Similarity and Paraphrase Tasks | | | | | |
| MRPC | 3.7k | 1.7k | paraphrase | acc./F1 | news |
| STS-B | 7k | 1.4k | sentence similarity | Pearson/Spearman corr. | misc. |
| QQP | 364k | 391k | paraphrase | acc./F1 | social QA questions |
| Inference Tasks | | | | | |
| MNLI | 393k | 20k | NLI | matched acc./mismatched acc. | misc. |
| QNLI | 105k | 5.4k | QA/NLI | acc. | Wikipedia |
| RTE | 2.5k | 3k | NLI | acc. | news, Wikipedia |
| WNLI | 634 | 146 | coreference/NLI | acc. | fiction books |

Results

| System | MNLI-(m/mm) 392k | QQP 363k | QNLI 108k | SST-2 67k | CoLA 8.5k | STS-B 5.7k | MRPC 3.5k | RTE 2.5k | Average |
|-----------------------|---------------------|-------------|--------------|--------------|--------------|---------------|--------------|-------------|-------------|
| Pre-OpenAI SOTA | 80.6/80.1 | 66.1 | 82.3 | 93.2 | 35.0 | 81.0 | 86.0 | 61.7 | 74.0 |
| BiLSTM+ELMo+Attn | 76.4/76.1 | 64.8 | 79.9 | 90.4 | 36.0 | 73.3 | 84.9 | 56.8 | 71.0 |
| OpenAI GPT | 82.1/81.4 | 70.3 | 88.1 | 91.3 | 45.4 | 80.0 | 82.3 | 56.0 | 75.2 |
| BERT _{BASE} | 84.6/83.4 | 71.2 | 90.1 | 93.5 | 52.1 | 85.8 | 88.9 | 66.4 | 79.6 |
| BERT _{LARGE} | 86.7/85.9 | 72.1 | 91.1 | 94.9 | 60.5 | 86.5 | 89.3 | 70.1 | 81.9 |

- ▶ Huge improvements over prior work (even compared to ELMo)
- ▶ Effective at “sentence pair” tasks: textual entailment (does sentence A imply sentence B), paraphrase detection

RoBERTa

- ▶ “Robustly optimized BERT”
- ▶ 160GB of data instead of 16 GB
- ▶ Dynamic masking: standard BERT uses the same MASK scheme for every epoch, RoBERTa recomputes them
- ▶ New training + more data = better performance

| Model | data | bsz | steps | SQuAD (v1.1/2.0) | MNLI-m | SST-2 |
|--------------------------|-------|-----|-------|---------------------|-------------|-------------|
| RoBERTa | | | | | | |
| with BOOKS + WIKI | 16GB | 8K | 100K | 93.6/87.3 | 89.0 | 95.3 |
| + additional data (§3.2) | 160GB | 8K | 100K | 94.0/87.7 | 89.3 | 95.6 |
| + pretrain longer | 160GB | 8K | 300K | 94.4/88.7 | 90.0 | 96.1 |
| + pretrain even longer | 160GB | 8K | 500K | 94.6/89.4 | 90.2 | 96.4 |

| BERT _{LARGE} | with BOOKS + WIKI | 13GB | 256 | 1M | 90.9/81.8 | 86.6 | 93.7 |
|-----------------------|-------------------|------|-----|----|-----------|------|------|
| | | | | | | | |

GPT/GPT2

OpenAI GPT/GPT2

- ▶ “ELMo with transformers” (works better than ELMo)
- ▶ Train a single unidirectional transformer LM on long contexts
- ▶ GPT2: trained on 40GB of text collected from upvoted links from reddit
- ▶ 1.5B parameters — by far the largest of these models trained as of March 2019
- ▶ Because it's a language model, we can **generate** from it

| Parameters | Layers | d_{model} |
|------------|--------|-------------|
| 117M | 12 | 768 |
| 345M | 24 | 1024 |
| 762M | 36 | 1280 |
| 1542M | 48 | 1600 |

OpenAI GPT2

**SYSTEM PROMPT
(HUMAN-WRITTEN)**

Miley Cyrus was caught shoplifting from Abercrombie and Fitch on Hollywood Boulevard today.

**MODEL COMPLETION
(MACHINE-WRITTEN,
SECOND TRY)**

The 19-year-old singer was caught on camera being escorted out of the store by security guards.

The singer was wearing a black hoodie with the label 'Blurred Lines' on the front and 'Fashion Police' on the back.

Scroll down for video

Shoplifting: Miley Cyrus was caught shoplifting from Abercrombie and Fitch on Hollywood Boulevard today (pictured)

The singer was wearing a black hoodie with the label 'Blurred Lines' on the front and 'Fashion Police' on the back

The singer was also wearing a pair of black-rimmed glasses, a black jacket, black jeans and black sandals.

She was carrying a pair of black and white striped gloves and a small black bag.

slide credit:
OpenAI

Open Questions


- 1) How novel is the stuff being generated? (Is it just doing nearest neighbors on a large corpus?)
- 2) How do we understand and distill what is learned in this model?
- 3) How do we harness these priors for conditional generation tasks (summarization, generate a report of a basketball game, etc.)
- 4) Is this technology dangerous? (OpenAI has only released 774M param model, not 1.5B yet)

Pre-Training Cost (with Google/AWS)

- ▶ BERT: Base \$500, Large \$7000
- ▶ Grover-MEGA: \$25,000
- ▶ XLNet (BERT variant): \$30,000 – \$60,000 (unclear)
- ▶ This is for a single pre-training run...developing new pre-training techniques may require many runs
- ▶ *Fine-tuning* these models can typically be done with a single GPU (but may take 1-3 days for medium-sized datasets)

Pushing the Limits

- ▶ NVIDIA: trained 8.3B parameter GPT model (5.6x the size of GPT-2)
- ▶ Arguably these models are still underfit: larger models still get better held-out perplexities




Google T5

| Number of tokens | Repeats | GLUE | CNNDM | SQuAD | SGLUE | EnDe | EnFr | EnRo |
|------------------|---------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| ★ Full dataset | 0 | 83.28 | 19.24 | 80.88 | 71.36 | 26.98 | 39.82 | 27.65 |
| 2^{29} | 64 | 82.87 | 19.19 | 80.97 | 72.03 | 26.83 | 39.74 | 27.63 |
| 2^{27} | 256 | 82.62 | 19.20 | 79.78 | 69.97 | 27.02 | 39.71 | 27.33 |
| 2^{25} | 1,024 | 79.55 | 18.57 | 76.27 | 64.76 | 26.38 | 39.56 | 26.80 |
| 2^{23} | 4,096 | 76.34 | 18.33 | 70.92 | 59.29 | 26.37 | 38.84 | 25.81 |

- ▶ Colossal Cleaned Common Crawl: 750 GB of text
- ▶ We still haven't hit the limit of bigger data being useful


BART

- ▶ Sequence-to-sequence BERT variant: permute/make/delete tokens, then predict full sequence autoregressively
- ▶ For downstream tasks: feed document into both encoder + decoder, use decoder hidden state as output
- ▶ Good results on dialogue, summarization tasks



Analysis

What does BERT learn?




- ▶ Heads on transformers learn interesting and diverse things: content heads (attend based on content), positional heads (based on position), etc.

What does BERT learn?


Head 8-10

- Direct objects attend to their verbs
- 86.8% accuracy at the dobj relation




Head 8-11

- Noun modifiers (e.g., determiners) attend to their noun
- 94.3% accuracy at the det relation



Head 5-4

- Coreferent mentions attend to their antecedents
- 65.1% accuracy at linking the head of a coreferent mention to the head of an antecedent



► Still way worse than what supervised systems can do, but interesting that this is learned organically


Probing BERT

- ▶ Try to predict POS, etc. from each layer.
Learn mixing weights

$$\mathbf{h}_{i,\tau} = \gamma_\tau \sum_{\ell=0}^L s_\tau^{(\ell)} \mathbf{h}_i^{(\ell)}$$

↑
representation of wordpiece i
for task τ


- ▶ Plot shows s weights (blue) and performance deltas when an additional layer is incorporated (purple)
- ▶ BERT “redisCOVERS the classical NLP pipeline”: first syntactic tasks then semantic ones



Tenney et al. (2019)

Compressing BERT

- ▶ Remove 60+% of BERT's heads with minimal drop in performance
- ▶ DistilBERT (Sanh et al., 2019): nearly as good with half the parameters of BERT (via knowledge distillation)



(b) Evolution of accuracy on the MultiNLI-matched validation set when heads are pruned from BERT according to I_h (solid blue) and accuracy difference (dashed green).

Open Questions

- ▶ BERT-based systems are state-of-the-art for nearly every major text analysis task
- ▶ These techniques are here to stay, unclear what form will win out
- ▶ Role of academia vs. industry: no major pretrained model has come purely from academia
- ▶ Cost/carbon footprint: a single model costs \$10,000+ to train (though this cost should come down)